

# Databases

## Active Databases

**Danilo Montesi**

[danilo.montesi@unibo.it](mailto:danilo.montesi@unibo.it)

# Passive Databases

- Reaction strategies in the integrity constraints are the first example of the need to introduce a **reactive** behaviour in the databases:

```
on      < delete | updated >
        < cascade | set null | set
          default | no action >
```

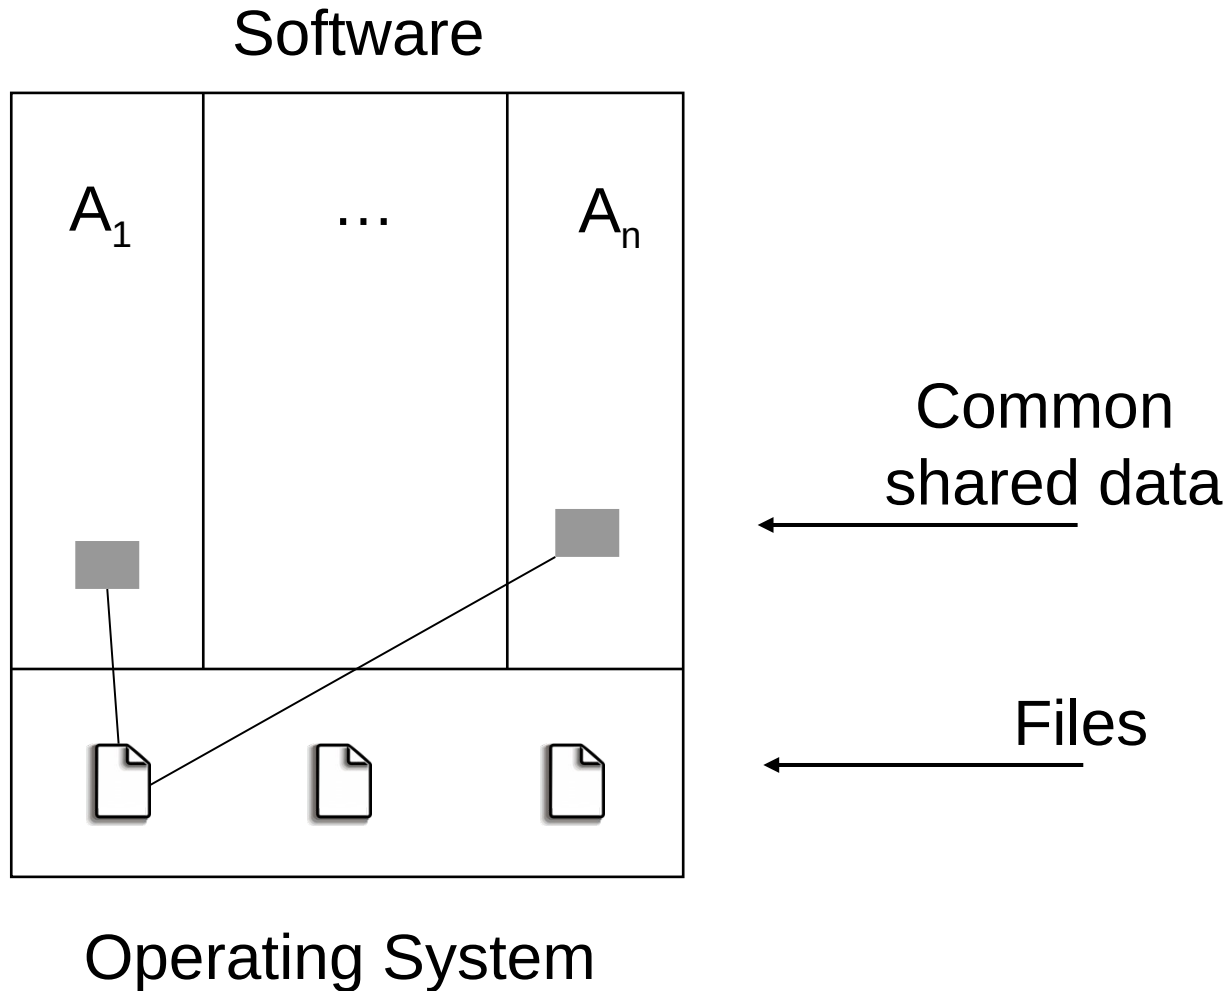
- The idea is to introduce **language constructs** specific for this goal
- These constructs are called (active) **rules** to handle a part of the procedural behaviour of an application
- Being at the database level, this behaviour is therefore “shared” among many applications, achieving data independence

# Active Databases

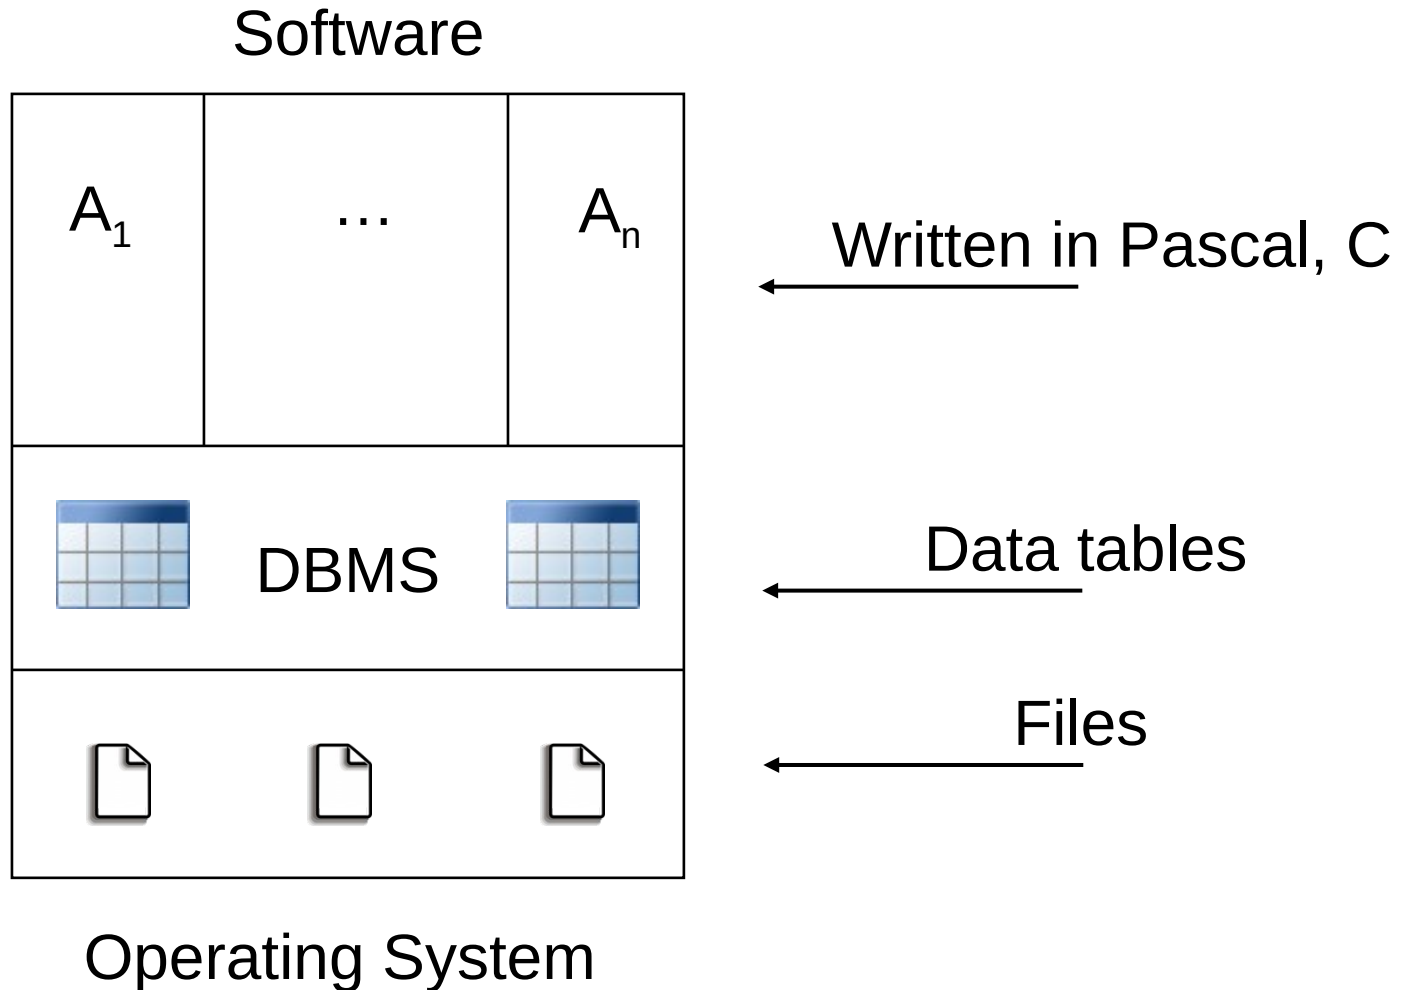
---

- An active database has a component to handle active rules of the kind **Event-Condition-Action**:
  - events are changes in the database
  - a condition is verified, based on a true/false value
  - one or more actions are executed
- These databases have a **reactive** behaviour (opposite to **passive**): they do not just execute the user transactions, but also the rules
- Commercial DBMSs (SQL3) use **triggers** to specify rules, like the ones in the schema definition

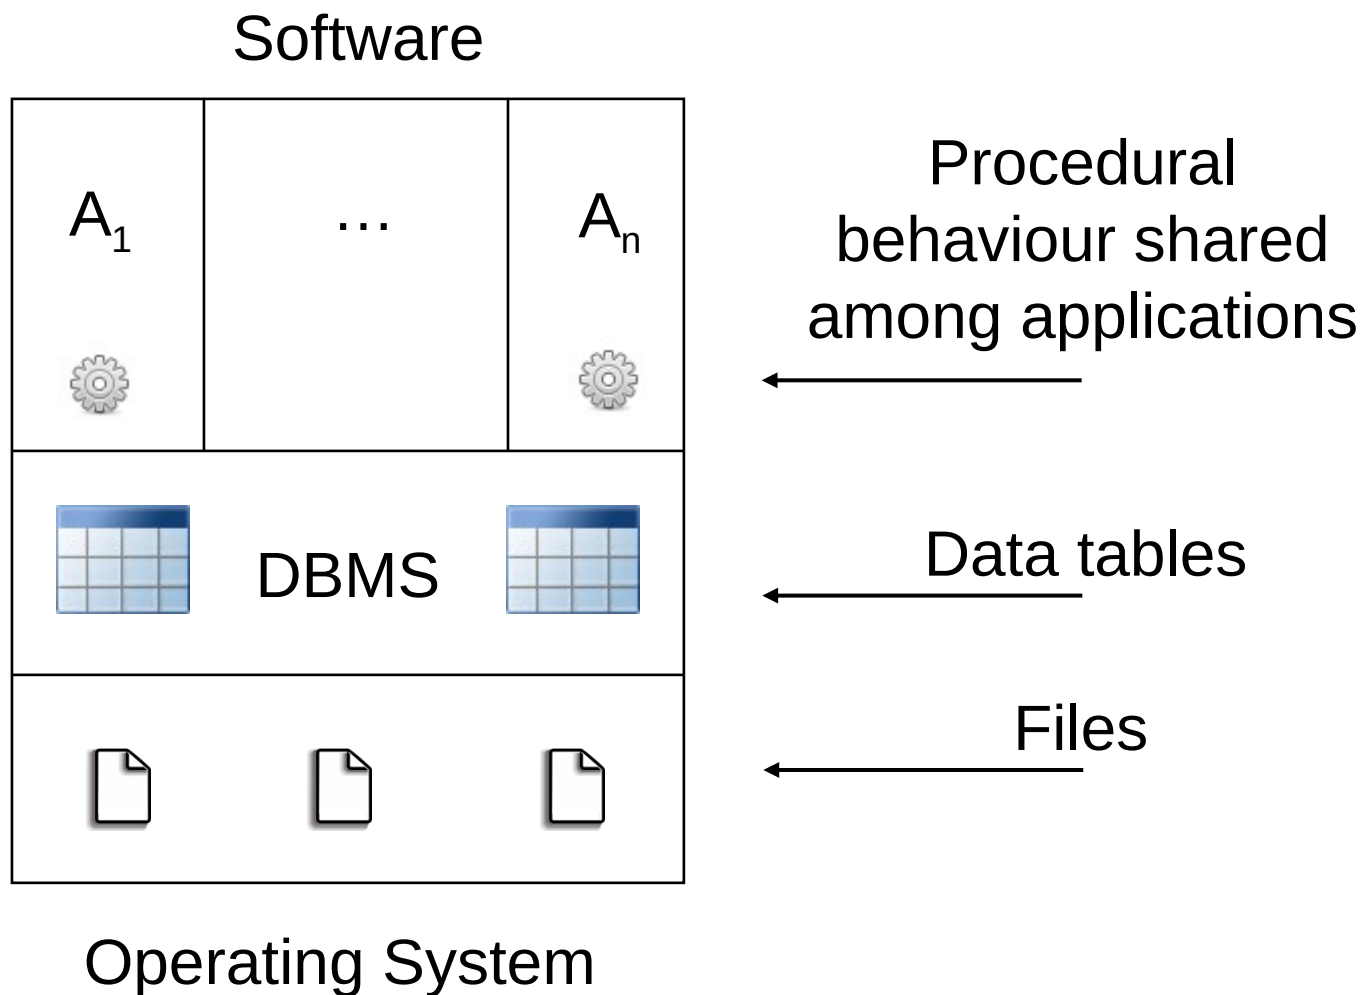
# The 70s: no DBMS



# The 80s: First DBMSs



# The 90s: the Procedural Behaviour



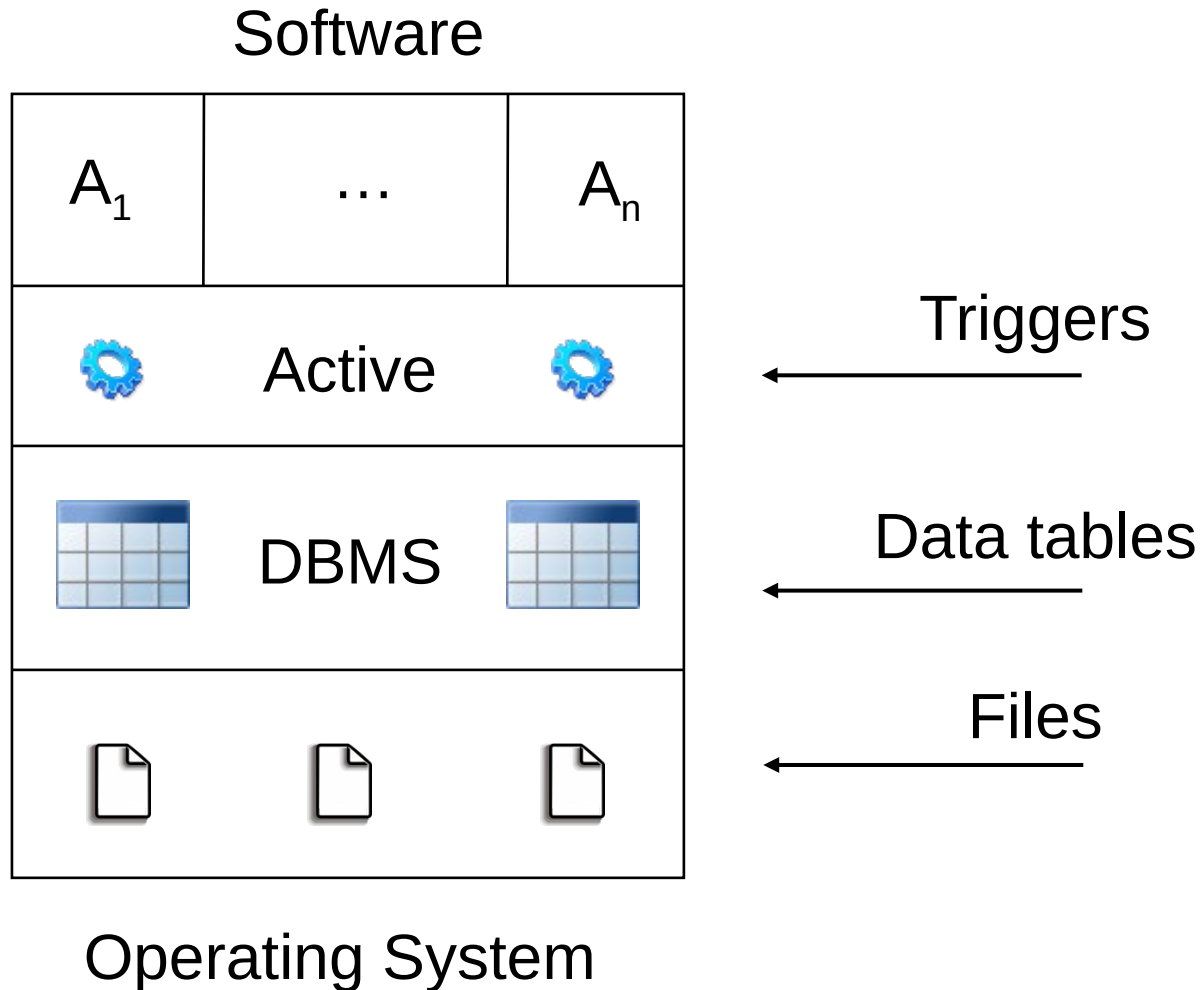


# Advances in DBMSs: Stored Procedures

---

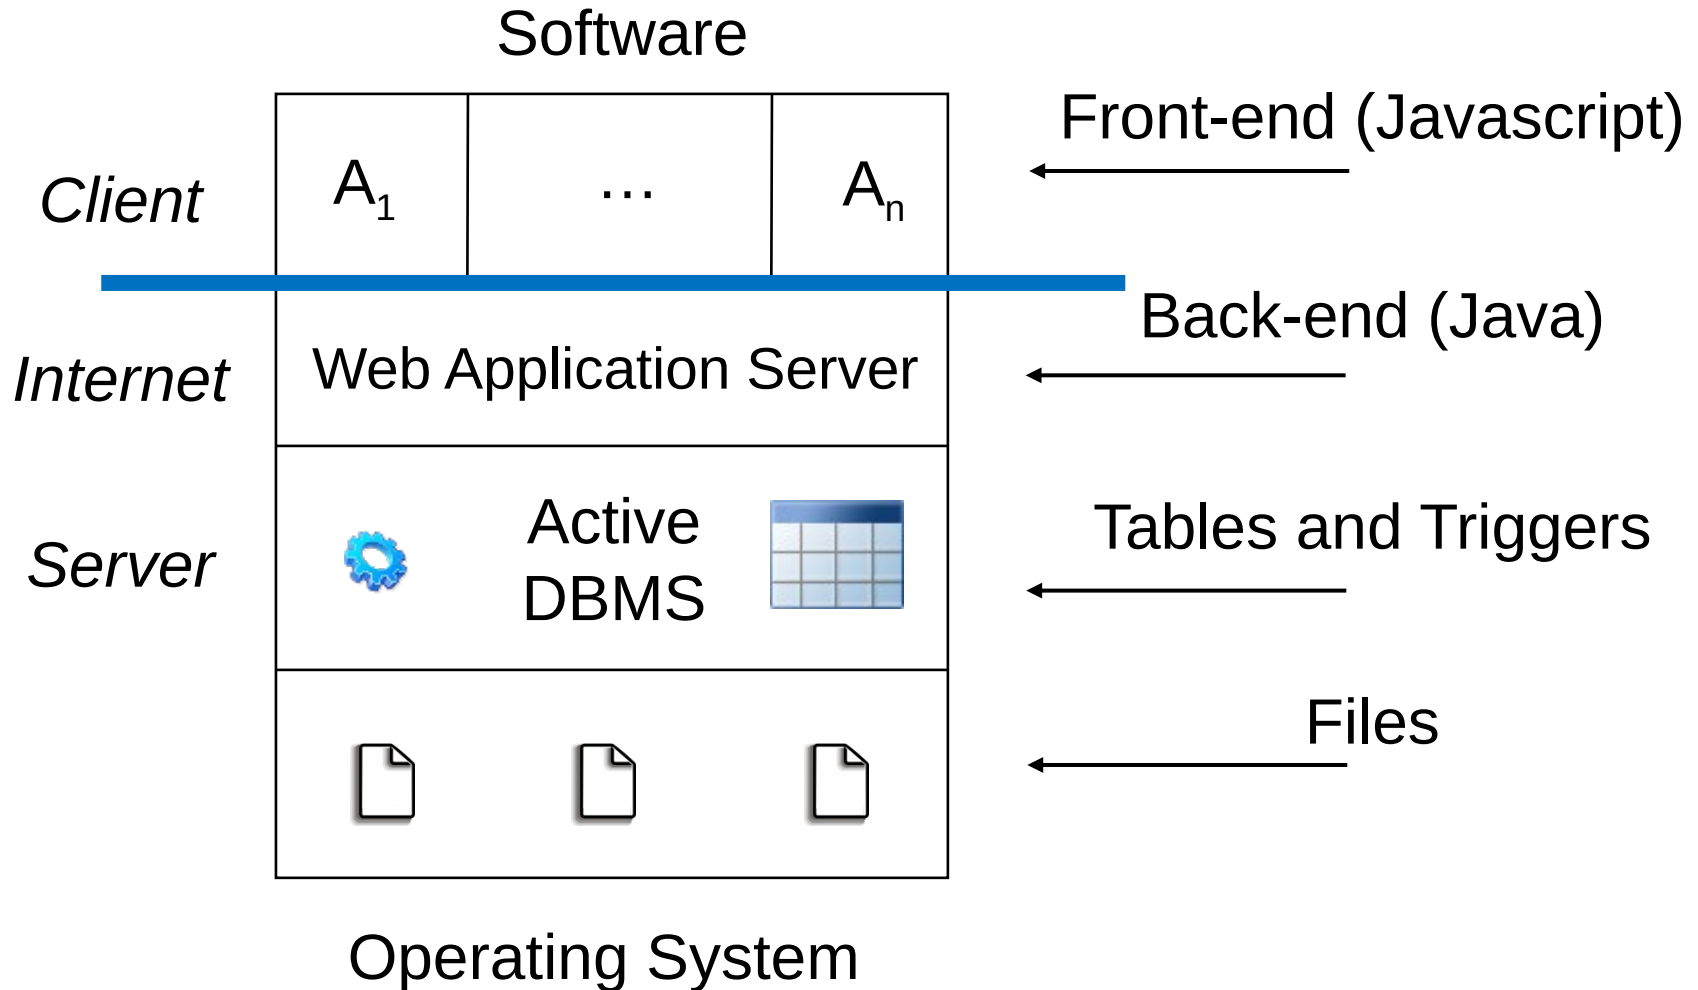
- Stored procedures have been introduced to share the common procedural behaviours between different software
- Stored procedures are not standardized and are affected by the problem of impedance mismatch with the language used to express such procedures
- As a result specific rules (*triggers*) have been introduced to model the procedural behaviour shared among different software and are handled by the DBMS itself

# The 90s: Active DBMS

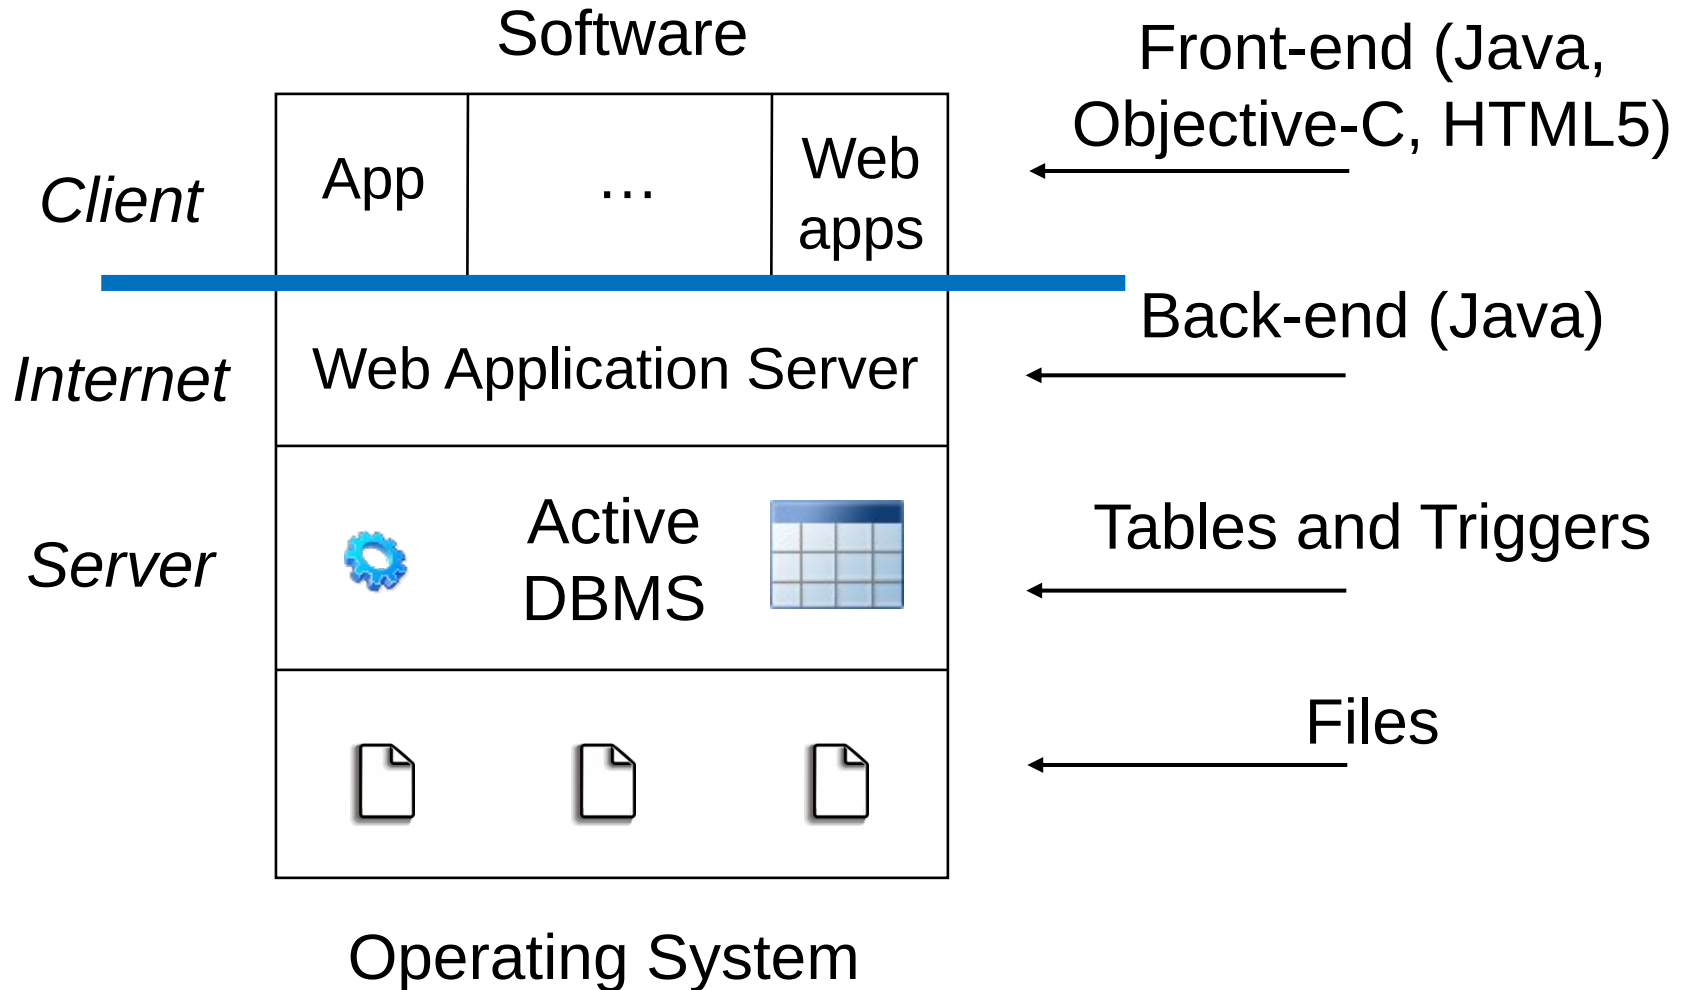




# Advance in the 2000s



# The 2010s: Mobile Apps



# Active Database

---

- A database that contains active rules (called triggers)
  - Definition of triggers in SQL:1999
  - Definition of triggers in Oracle and DB2
  - Design issues for applications based on the use of triggers

# Trigger

---

- Defined with DDL instructions (`create trigger`)
  - based on the event-condition-action (ECA) paradigm:
    - event: change in the data, specified using `insert`, `delete`, `update`
    - condition: (optional) SQL predicate
    - action: sequence of SQL instructions (or extensions, for example PL/SQL in Oracle)
  - intuitively:
    - when there is an event (activation)
    - if the condition is satisfied (verification)
    - then execute the action (execution)
- each trigger refers to a (target) table: it responds to events related to that table



# Trigger: Granularities and Modalities

---

- Granularity
  - of tuple (**row-level**): activation for each tuple involved in the operation
  - of operation (**statement-level**): only one activation for an SQL instruction, with reference to all the tuples involved (“set-oriented”)
- Mode
  - **immediate**: right after (or right before) of the event
  - **deferred**: at the time of commit

# Computational Model

- Let  $T^U = U_1; \dots; U_n$  the user transaction
- If the rules of  $P$  have form  $E, C \rightarrow A$  with  $E$  event,  $C$  condition and  $A$  action, then:
  - **Immediate** semantic generates
$$T^I = U_1; \underline{U_1^P}; \dots; U_n; \underline{U_n^P}$$
  - **Deferred** semantic generates
$$T^D = U_1; \dots; U_n; \underline{U_1^P}; \dots; \underline{U_n^P}$$
- Where  $\underline{U_i^P}$  represents the sequence of actions induced by  $U_i$  on  $P$
- Problems:
  - Termination
  - Confluence
  - Equivalence



# Trigger in SQL:1999: Syntax

---

```
create trigger triggerName
  { before | after }
  { insert | delete | update [of
    column] }
    on targetTable
  [referencing
    { [old table [as] oldTableVar]
      [new table [as] newTableVar] } |
    { [old [row] [as] oldTupleVar]
      [new [row] [as] newTupleVar] } ]
  [for each { row | statement } ]
  [when condition]
SQLProceduralStatement
```

# Types of Events

---

## ■ Before

- The trigger is considered and possibly executed before the event
- It is usually used when you want to verify a change before it occurs and “modify the change”

## ■ After

- The trigger is considered and executed after the event
- This is the most common mode, suitable for most applications



# Referencing Clause

---

- Depends on granularity
  - statement-level mode: two transition tables (**old table** and **new table**) contain the previous and next values of the tuples modified by the statement
  - row-level mode: two transition variables (**old** and **new**) represent the value before or after a tuple is modified
- **old** and **old table** are not present with the insert event
- **new** and **new table** are not present with the delete event



# Trigger in Oracle: Syntax

```
create trigger triggerName
  { before | after } event [, event [,
    event]]
  [[referencing
    [old [row] [as] oldTupleVar]
    [new [row] [as] newTupleVar]]
  for each { row | statement } [when
    condition]]
```

- **PL/SQL block** **before** | **after** }: event type
- *event*: insert, update, delete
- **for each row** specifies the granularity
- *Reference*: it allows to define variable names (usable only for tuple granularity)

**old as** *OldVariable* | **new as** *NewVariable*

# Trigger in Oracle: Semantic

---

- **Immediate** mode, both after and before
- Execution plan
  - before statement trigger
    - for each tuple involved
      - before row trigger
      - execution and constraints check
      - after row trigger
    - after statement trigger
- In the event of an error, everything is discarded
- Priorities between triggers, using timestamp
- Max 32 triggers activated in cascade



# Trigger in Oracle: an Example (1)

```
create trigger Reorder
after update of QtyAvbl on Warehouse
when (new.QtyAvbl < new.QtyLimit)
for each row
  declare X number;
begin
  select count(*) into X
  from PendingOrders
  where Part = new.Part;
  if X = 0
  then
    insert into PendingOrders
    values (new.Part, new.QtyReord,
sysdate);
  end if;
end;
```



# Trigger in Oracle: an Example (2)

Warehouse	Part	QtyAvbl	QtyLimit	QtyReord
	1	200	150	100
	2	780	500	200
	3	450	400	120

T1: update Warehouse  
set QtyAvbl = QtyAvbl -  
70  
where Part = 1

T2: update Warehouse  
set QtyAvbl = QtyAvbl -  
60  
where Part <= 3



# Trigger in DB2: Syntax

```
create trigger triggerName  
  { before | after } event on  
  targetTable  
  [referencing reference]  
  for each level  
  [when (SQLPredicate)]  
  SQLProceduralStatement
```

- { **before** | **after** }: event type
- *event*: insert, update, delete
- *for each level* specifies the granularity
- *Reference*: it allows to define variable names (depending on granularity):

```
old as OldTupleVar | new as NewTupleVar
```

```
old_table as OldTableVar | new_table as NewTableVar
```

# Trigger in DB2: Semantic

---

- **Immediate** mode, both after and before
  - before triggers cannot modify the database, apart from variants on the changes caused by the event (therefore they cannot in general activate other triggers)
- In the event of an error, everything is discarded
- No priority between triggers (the order is defined by the system, *i.e.*, *timestamp*), interaction with compensatory actions on referential integrity constraints
- Max 16 triggers activated in cascade



# Trigger in DB2: an Example

---

```
create trigger checkWage
  after update of Wage on Employee
  for each row
  when (new.Wage < old.Wage * 0.97)
  begin
    update Employee
    set Wage = old.Wage * 0.97
    where EmpCode = new.EmpCode;
  end;
```





# Extensions (not usually available)

---

- Temporal events (also periodical) or “user-defined”
- Boolean combinations of events
- Clause `instead of`: it does not execute the operation that activated the event, but another action instead
- “Detached” execution: an autonomous transaction is activated
- Priorities definition
- Groups rules, can be activated and deactivated
- Rules associated also with queries (not just updates)

# Rules Properties

---

- Termination (essential)
- Confluence
- Determinism of observations

# Applications

---

- Internal features
  - Handling of integrity constraints
  - Replication
  - View management
    - Materialized: propagation
    - Virtual: modification of the queries
- Application features: description of the behaviour of the database