

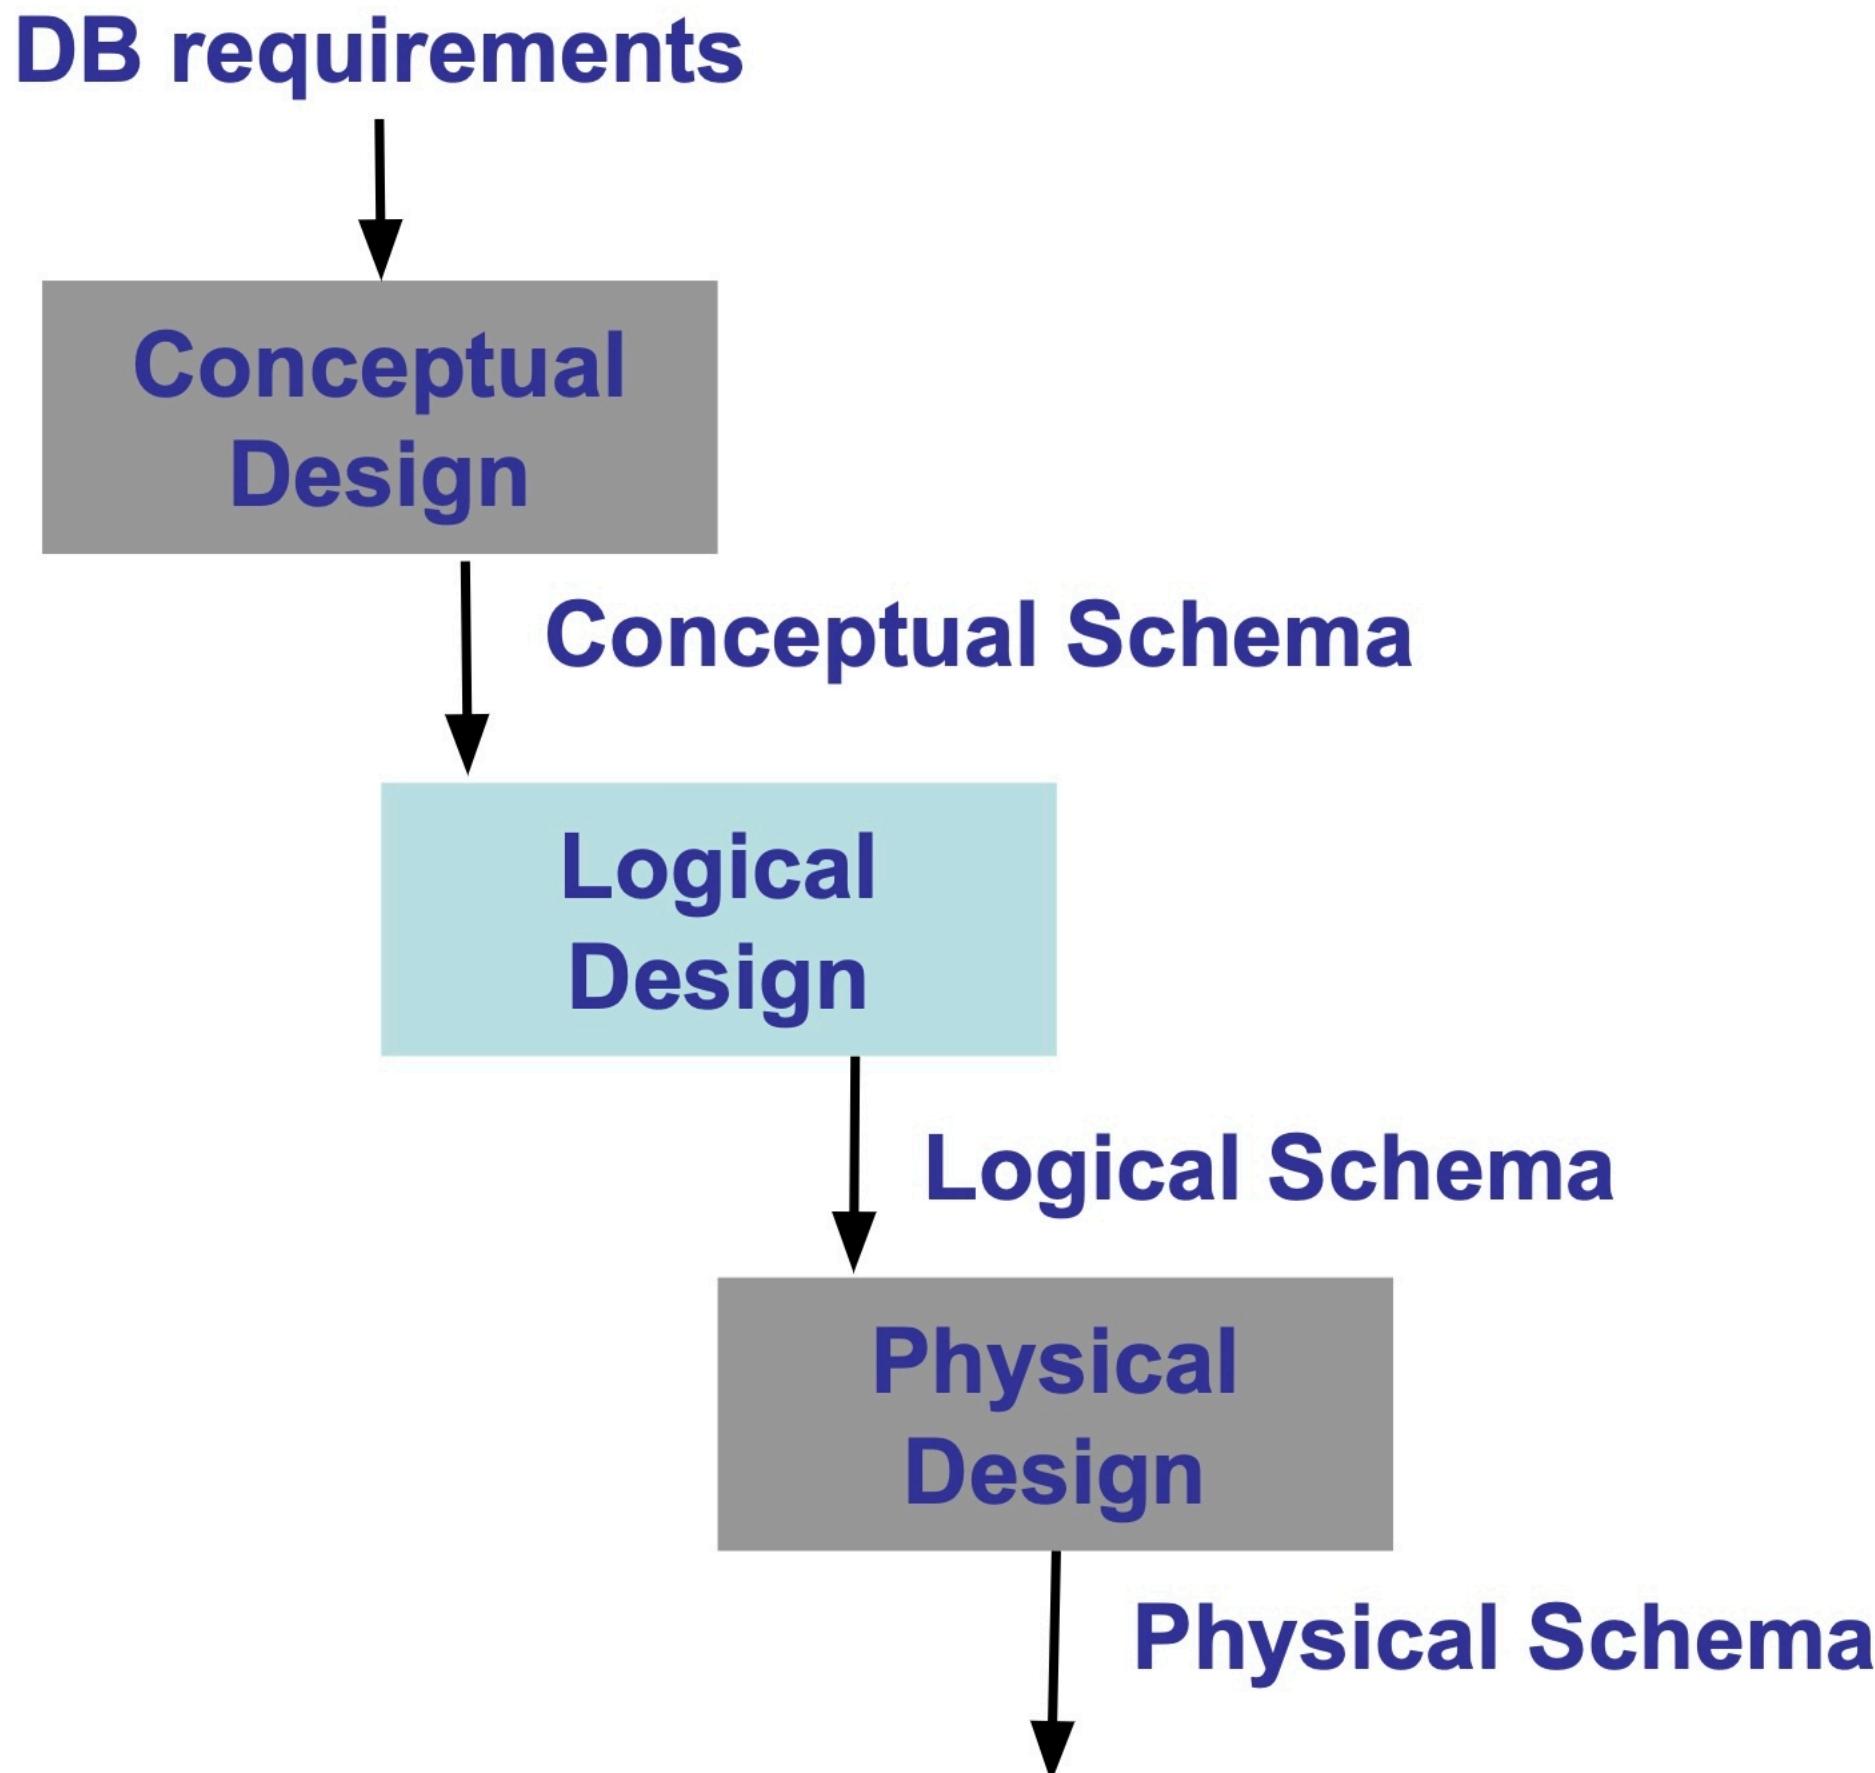


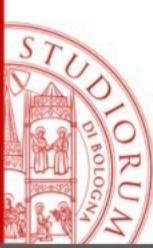
Databases

Logical Design



Requirements & Design





Logical Design: Objective

- The goal is to “translate the conceptual schema into a logical schema representing the same data in a correct and efficient way.



Logical design: input and output data

■ Input:

- Conceptual Schema
- Application workload information
- Logical Model

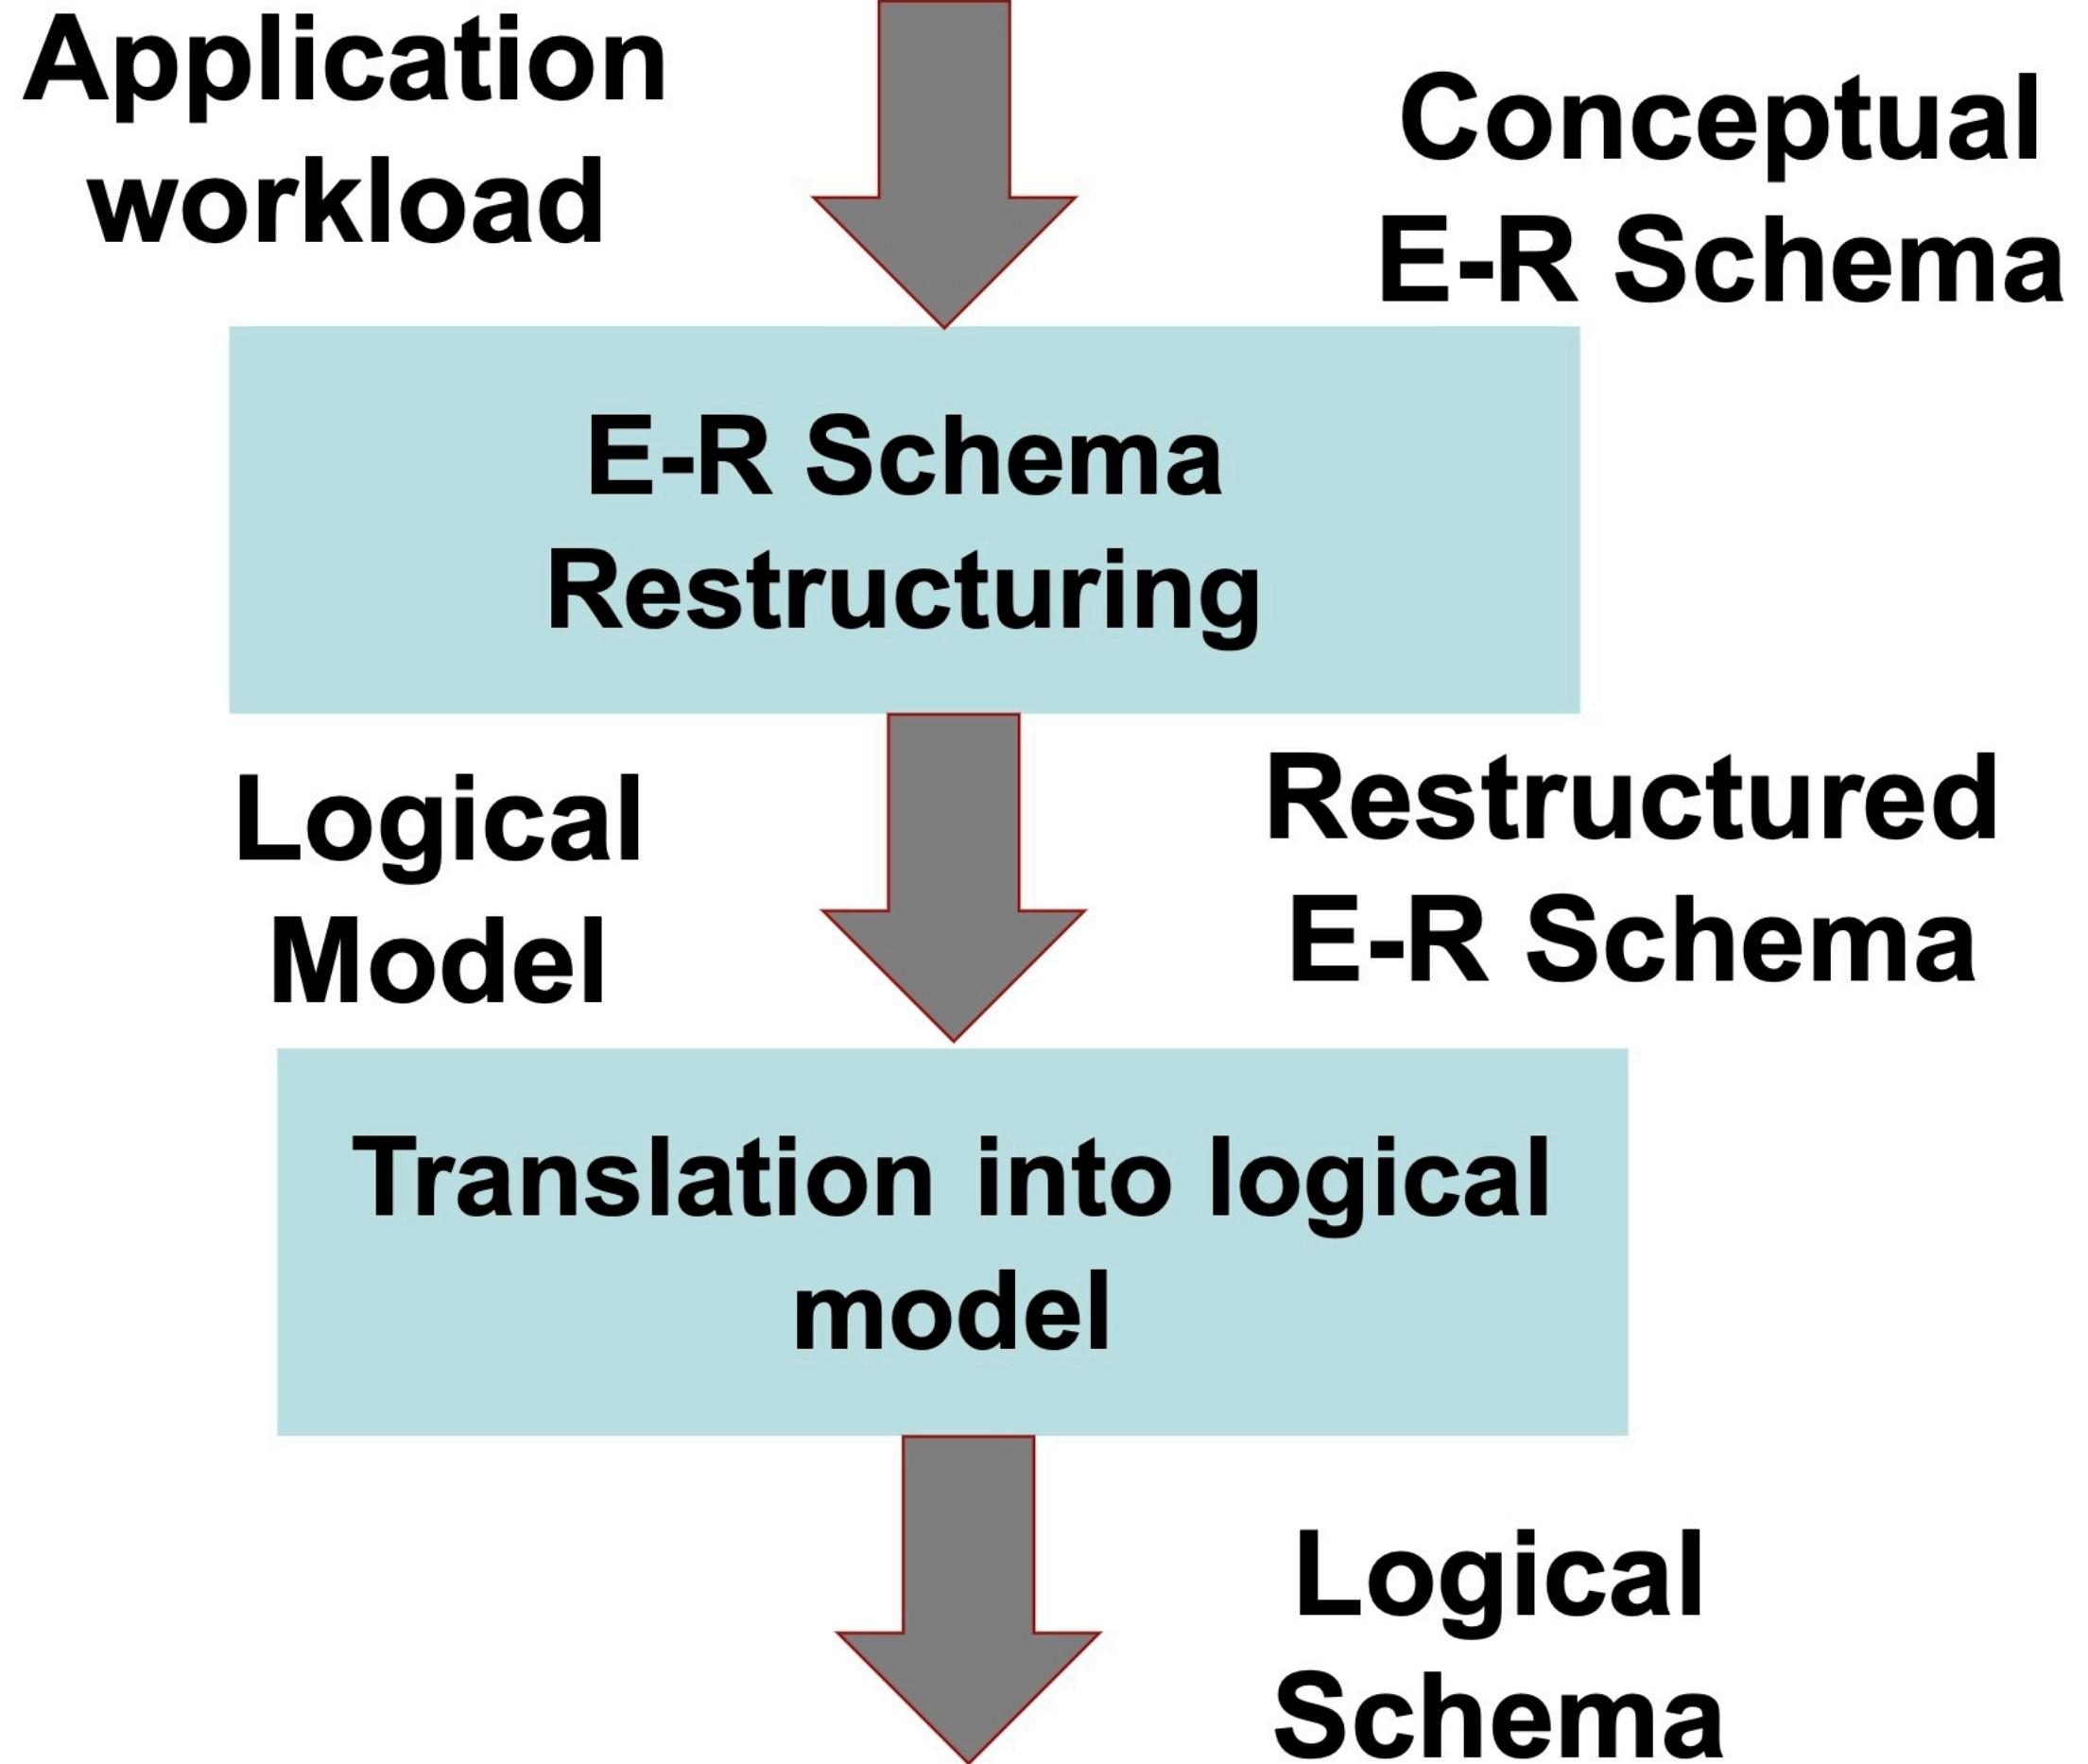
■ Uscita:

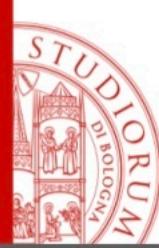
- Logical Schema (either relational, object-oriented, graph ...)
- Associated documentation



It is not a simple translation

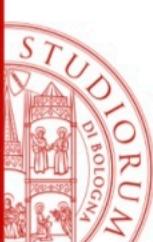
- Some aspects may not be represented directly
- In this phase we must also consider performance (efficiency).





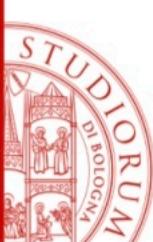
E-R Schema restructuring

- Why?
 - Simplify the translation
 - Optimize the performances
- Please note that:
 - A restructured E-R is no more a “conceptual schema” in the strict sense of the term.



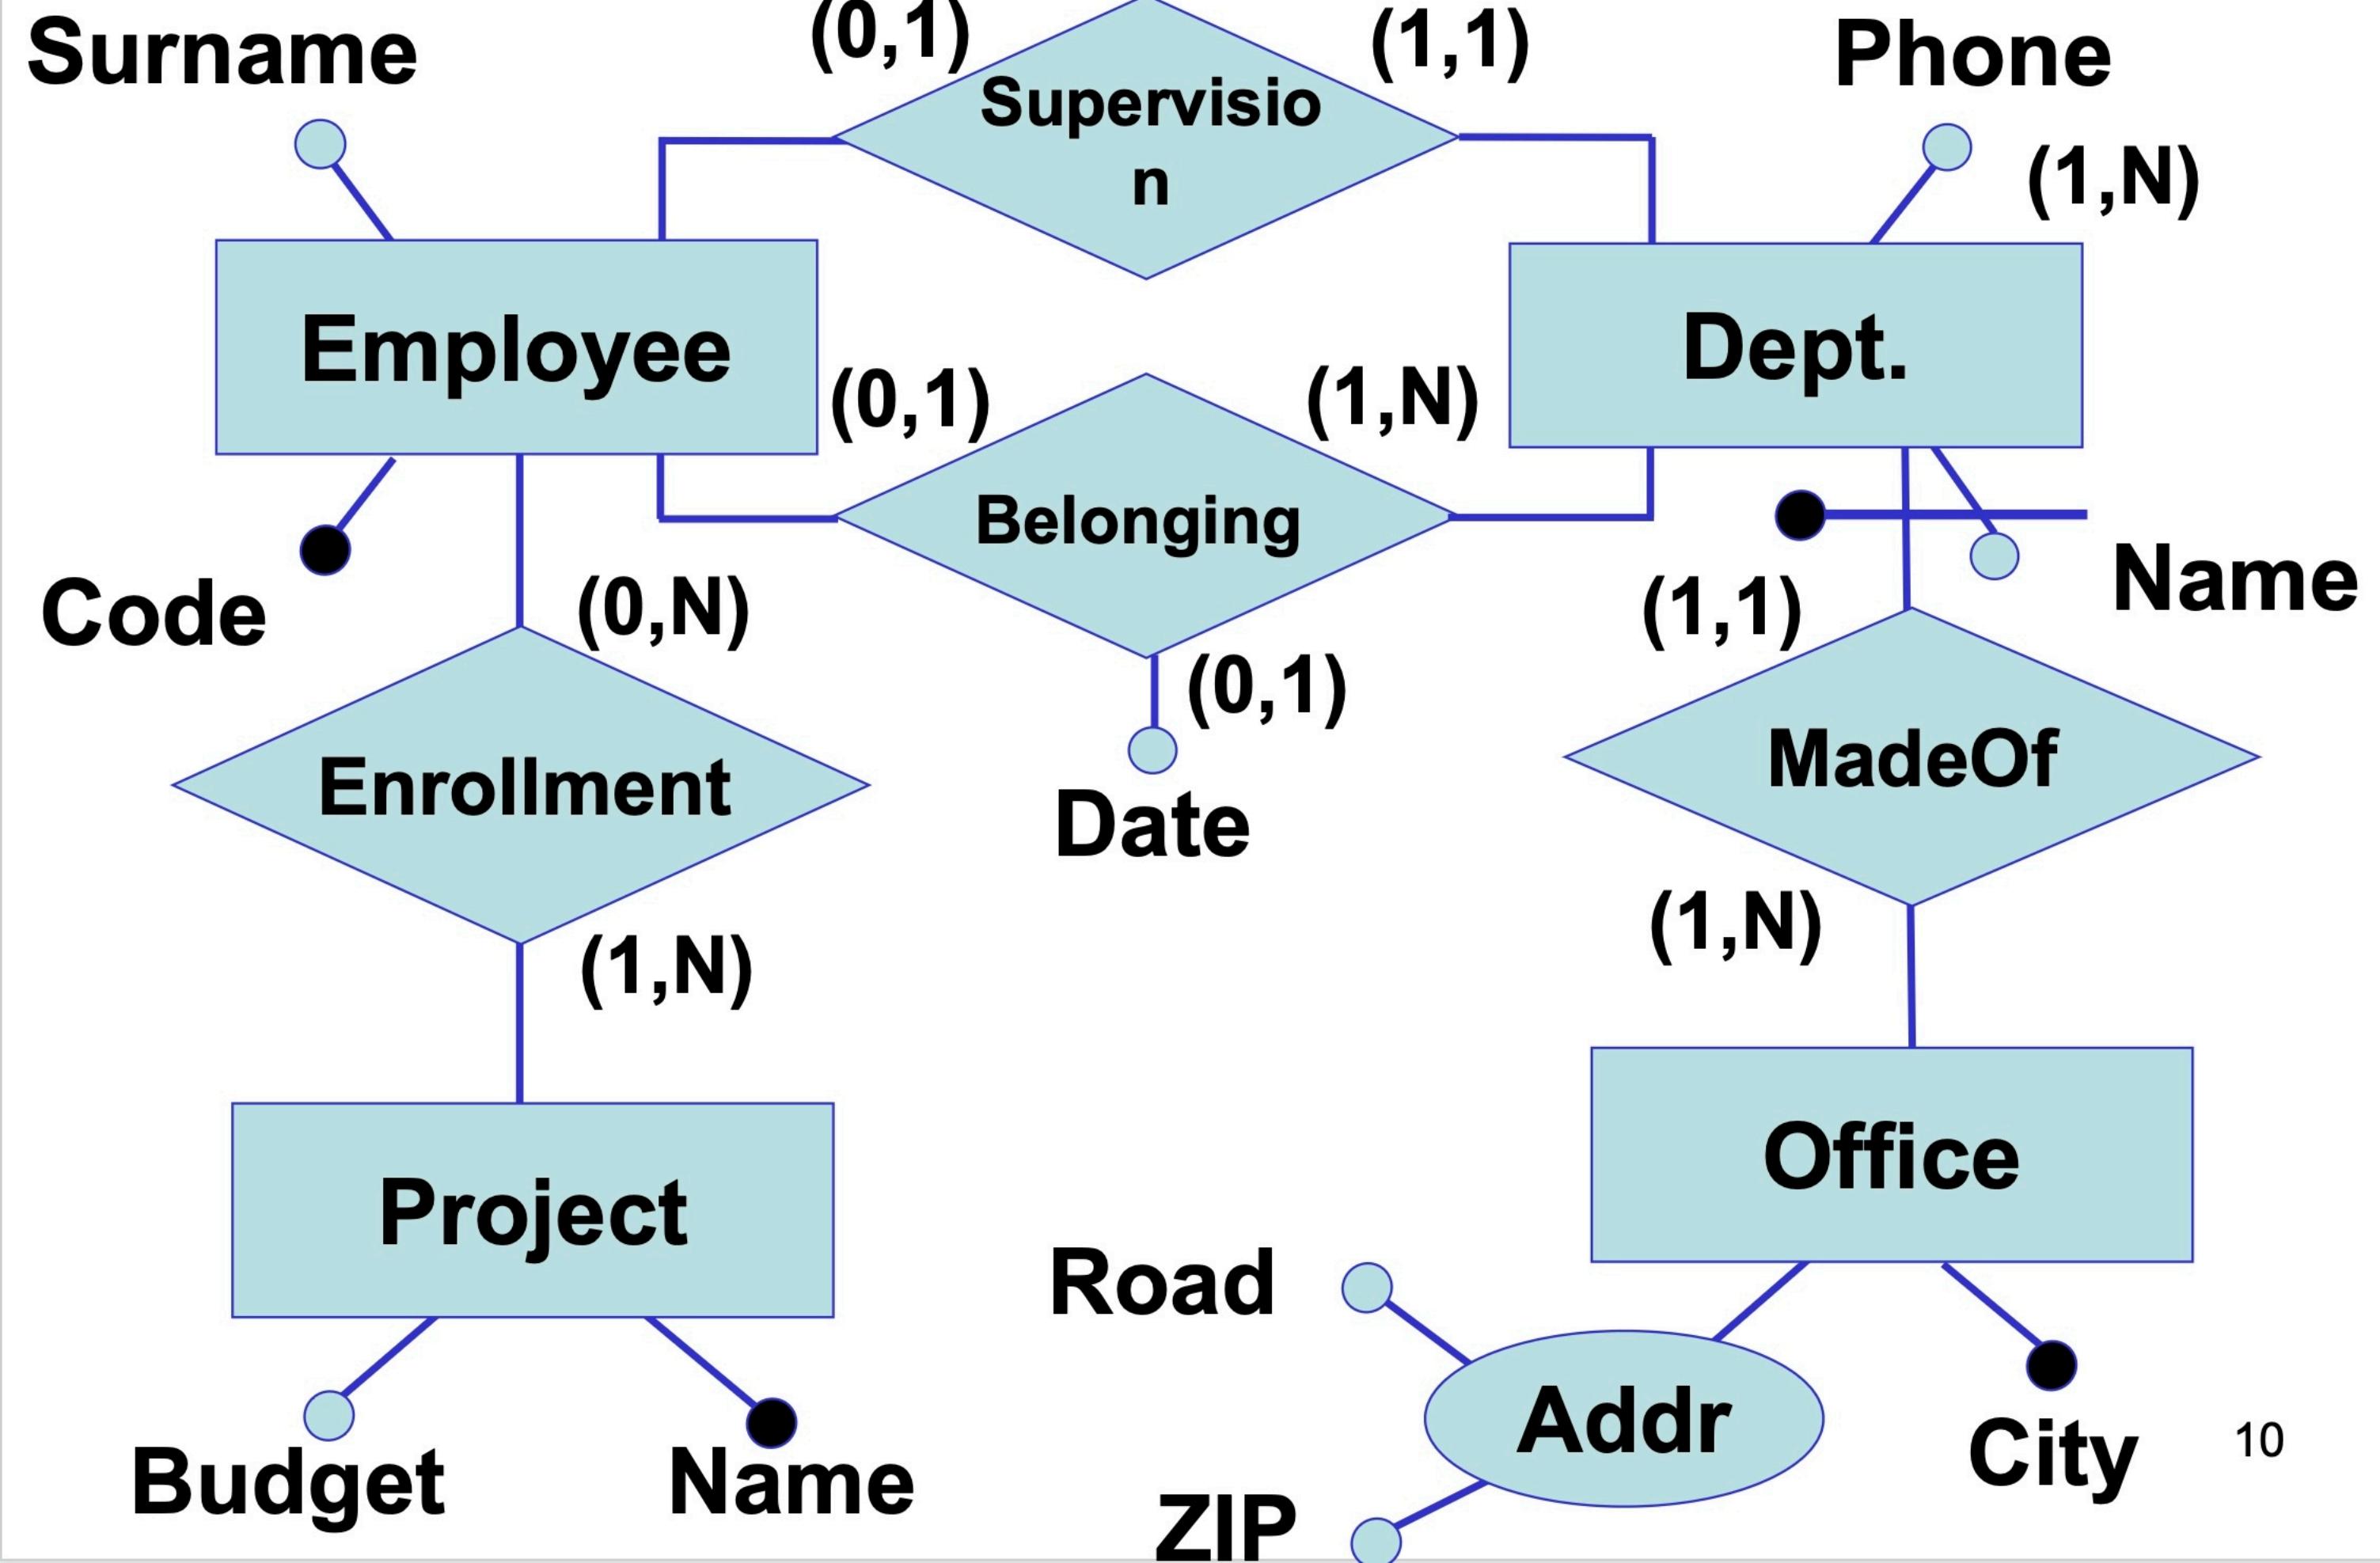
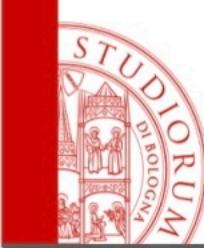
Performances?

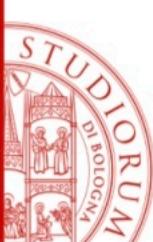
- In order to optimize the results we need to analyze performances at this level
- However:
 - Performances cannot be evaluated precisely on a conceptual schema!



Performances, approximated

- Let us consider:
 - some performance “**indicators**”: parameters on which the performances depend upon
- **space**: number of stored instances expected
- **time**: number of instances (of entities and relationships) visited during an operation





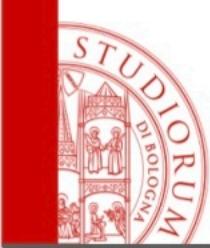
Size Table

Name	Kind	Size
Office	E	10
Dept.	E	80
Employee	E	2000
Project	E	500
MadeOf	R	80
Belonging	R	1900
Supervision	R	80
Enrollment	R	6000

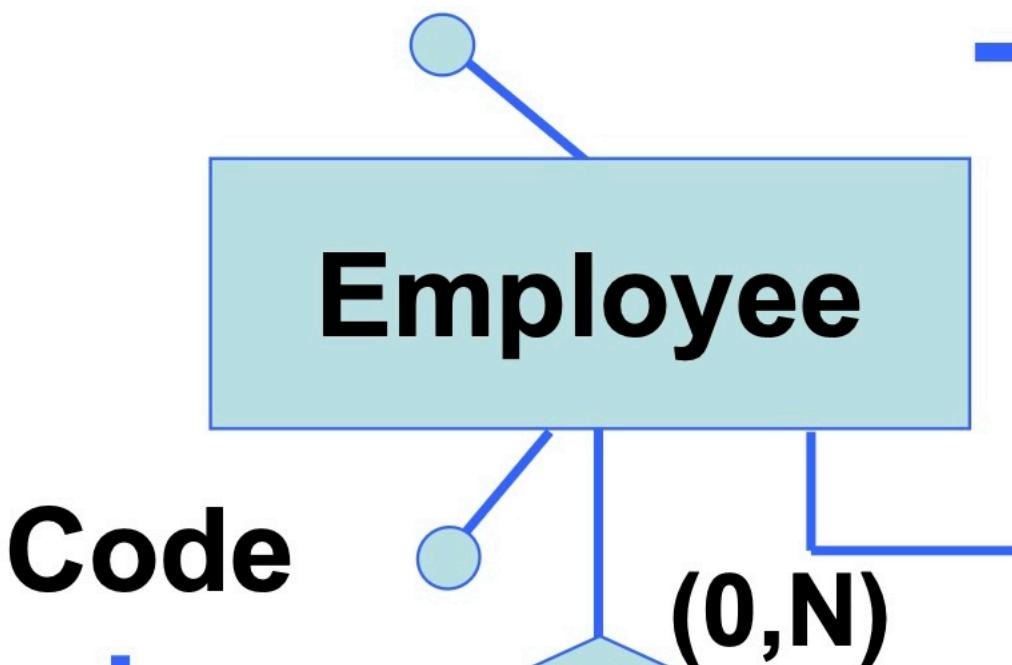


Example of cost evaluation

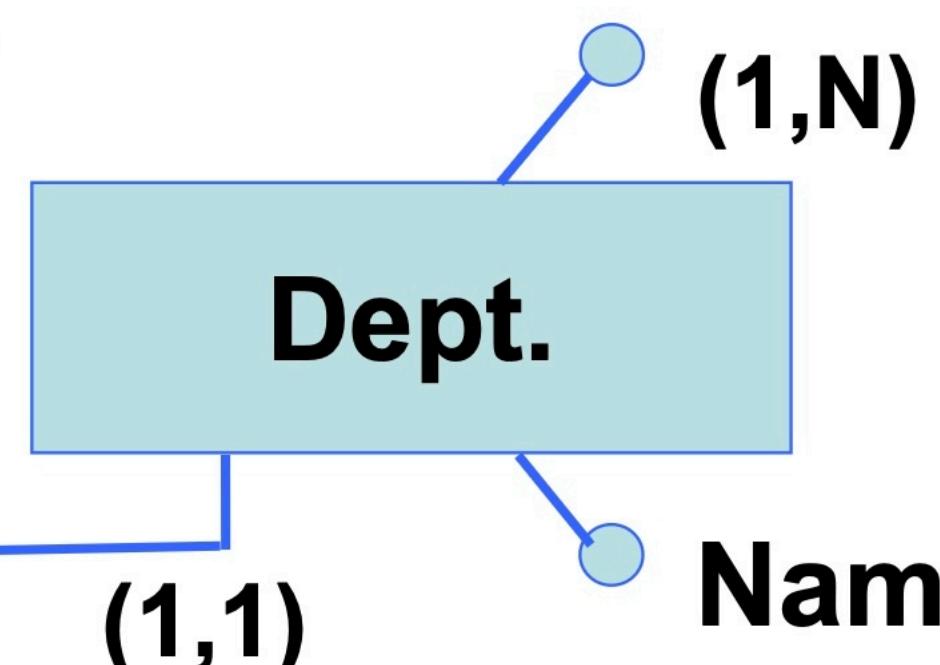
- Operation:
 - Return all the data of an employee, the data of the department where he belongs, and the data of the projects he works on.
- Build an **Access Table** following the **navigation schema**



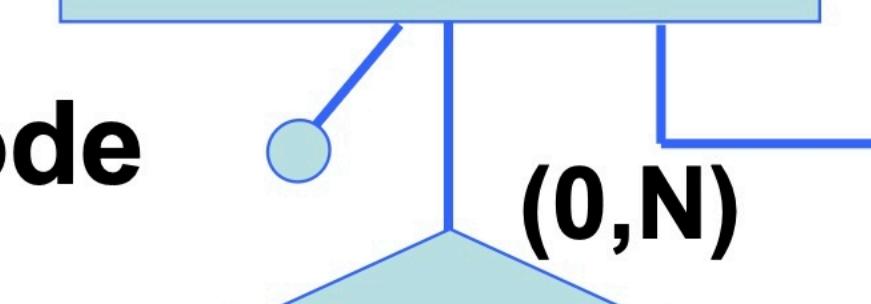
Surname



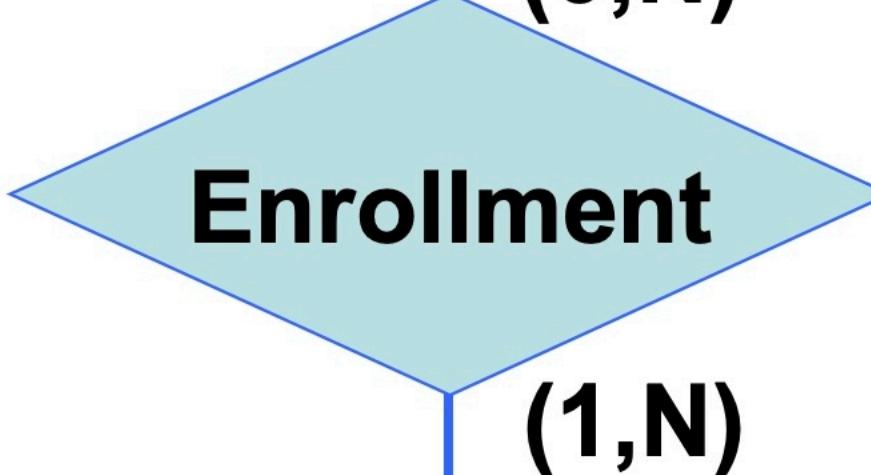
Phone



Code



Enrollment



Project



Budget

(0,1)

(1,N)

(1,1)

Date

(0,1)

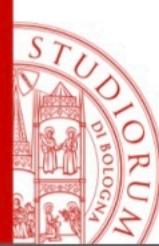
(1,N)

Name

Belonging

(0,1)

Name



Access Table

Name	Kind	#Accesses	Accesses Types	Access order
Dept.	Entity	1	R	3
Employee	Entity	1	R	1
Project	Entity	3	R	5
Belonging	Rel.	1	R	2
Enrollment	Rel.	3	R	4



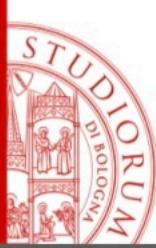
Restructuring activities

- Redundancies Analysis
- Generalizations deletion
- Partitioning/grouping of entities and relationships
- Identifying the primary keys



Redundancy Analysis

- A redundancy in a E-R schema is an information that is relevant but can be derived from others.
- In this phase we have to decide whether we have to keep, remove or create new redundancies.



Redundances

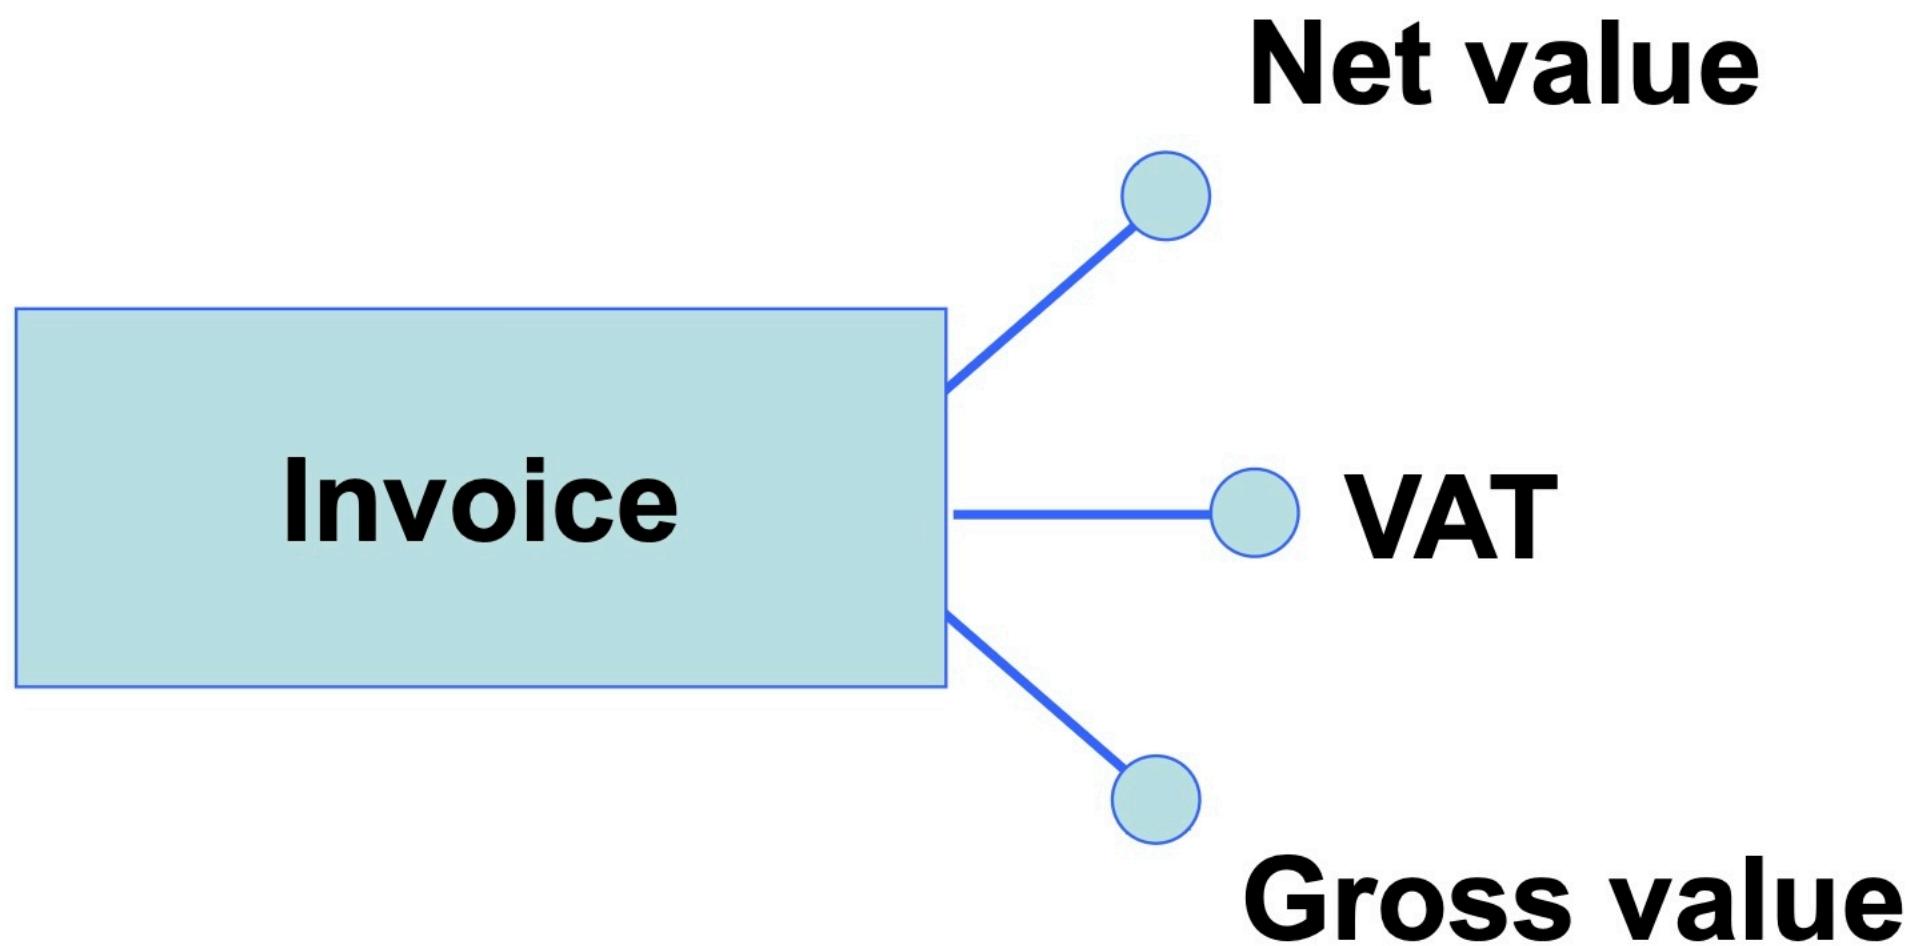
- Pros
 - They simplify queries
- Cons
 - Updates take more time
 - Storage size is increased



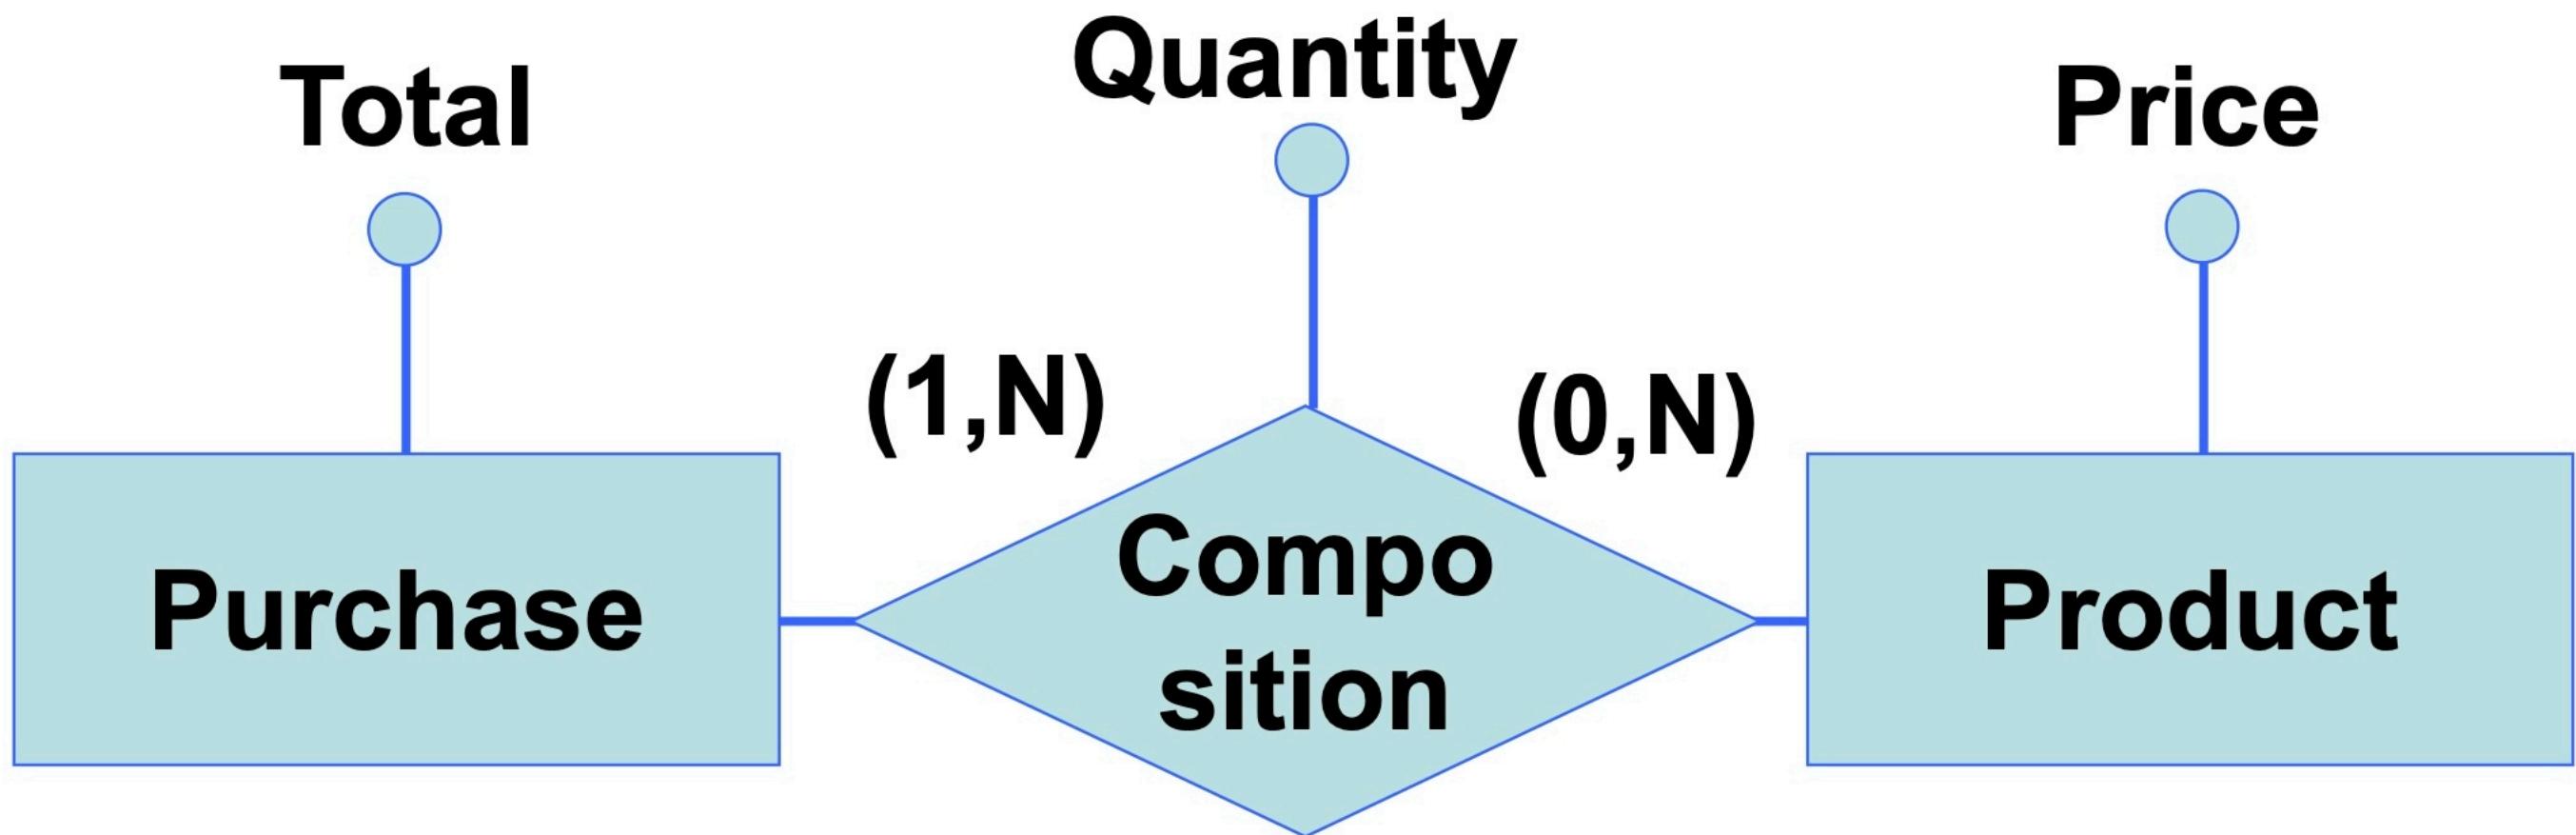
E-R: Types of Redundancies

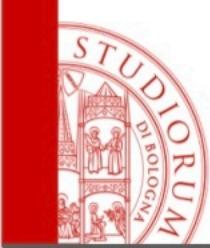
- Derivable attributes:
 - From other attributes within the same entity (or relationship)
 - From attributes of other entities (or relationships)
- Relationships derivable from the composition of different relationships (generally speaking: cycles of relationships)

Derivable attributes

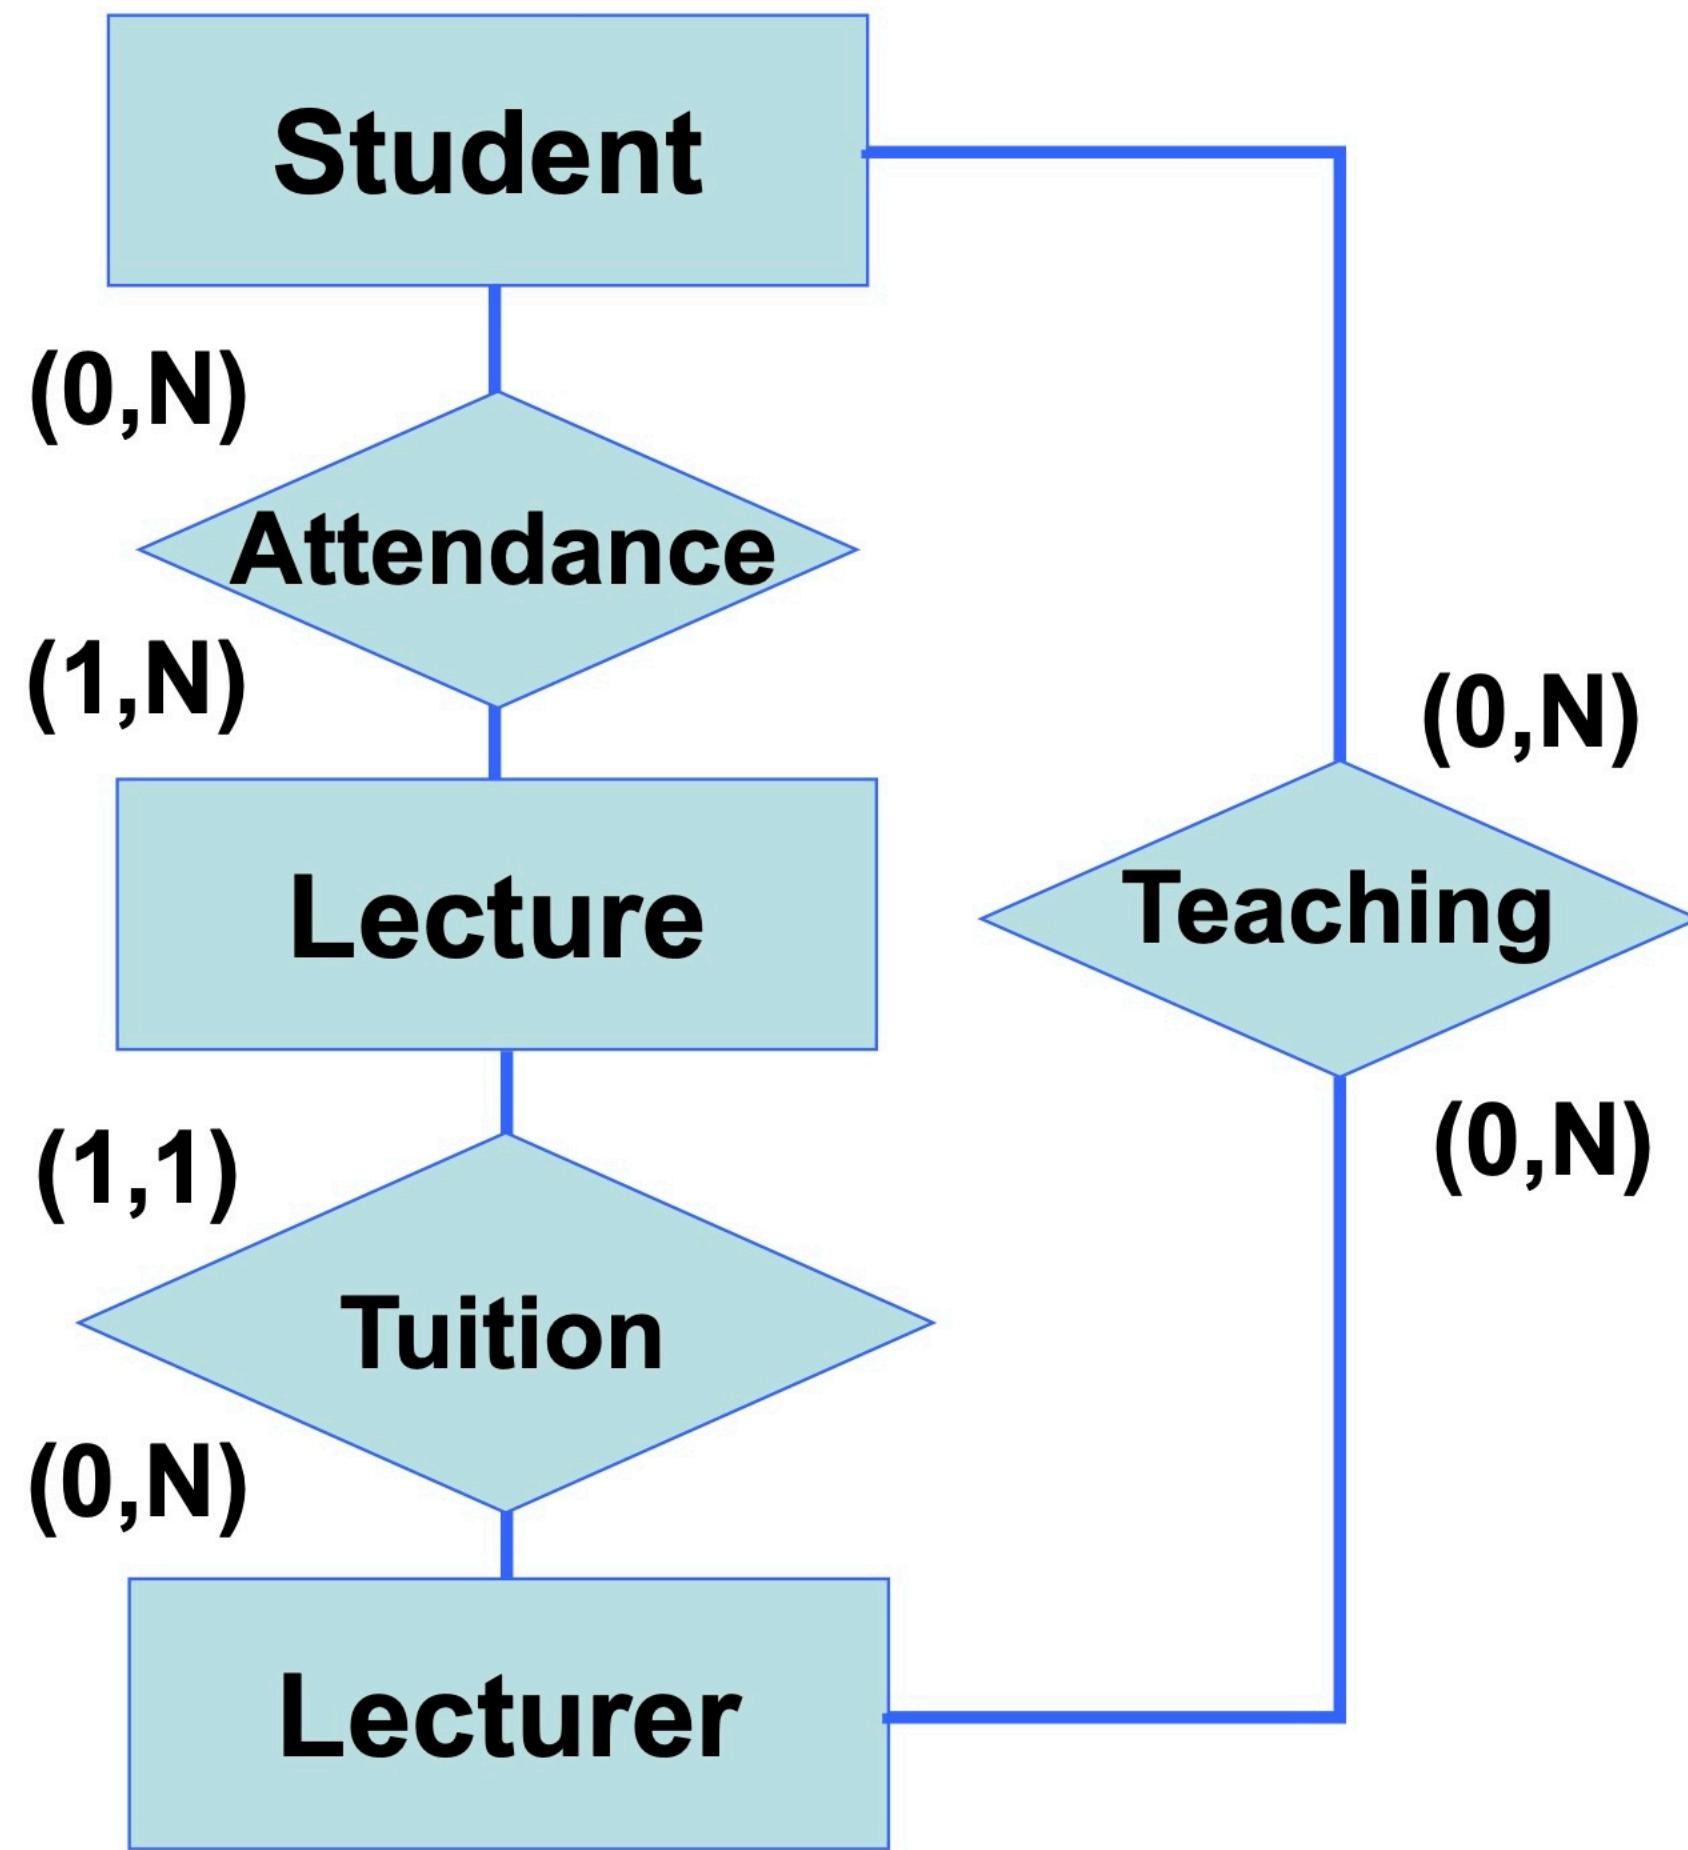


Attribute derivable from other entities

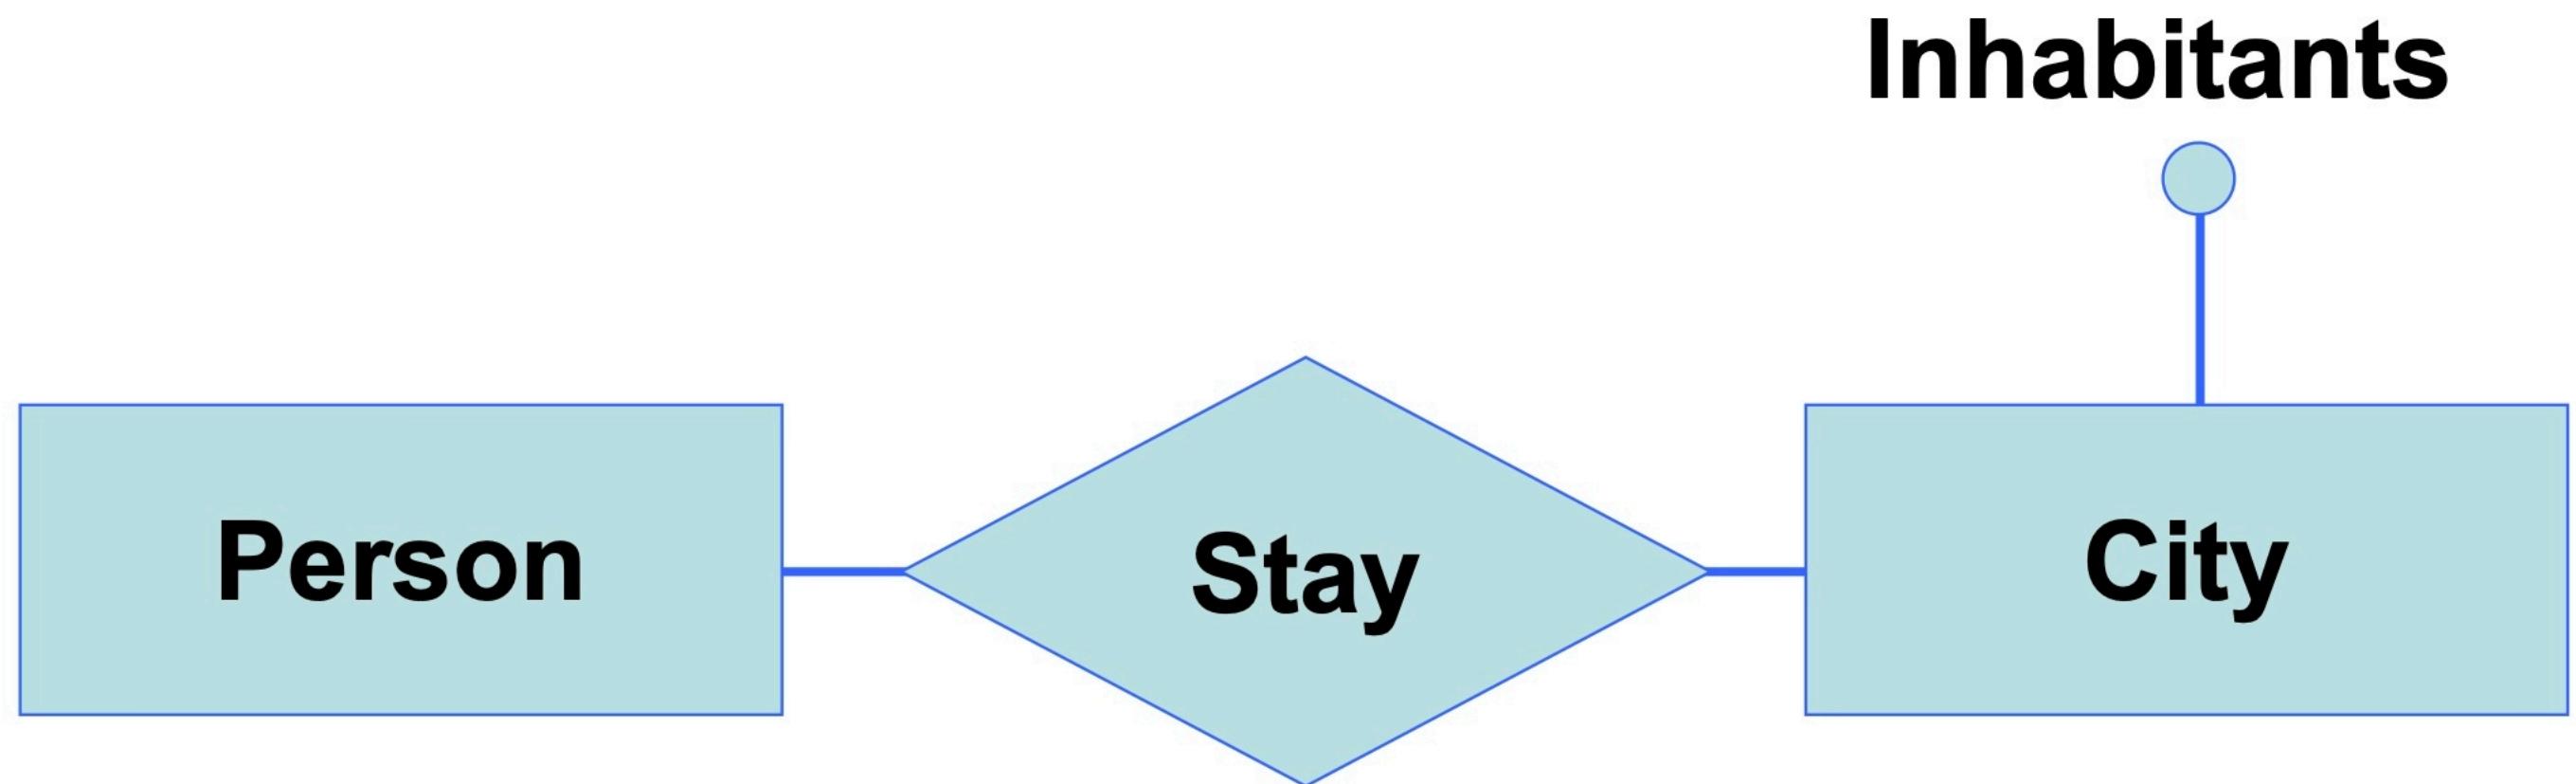




Redundancy from a cycle

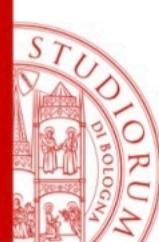


Redundancy Analysis



Name	Kind	Size
City	E	200
Person	E	1000000
Stay	R	1000000

- **Operation #1:** store a new person with his city of staying (500 times a day)
- **Operation #2:** print all the cities' data (including number of inhabitants) (2 times a day)



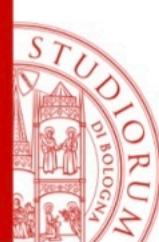
With Redundancy

Operation #1

Name	Kind	#Accesses	Accesses Types
Person	E	1	W
Stay	R	1	W
City	E	1	R
City	E	1	W

Operation #2

Name	Kind	#Accesses	Accesses Types
City	E	1	R



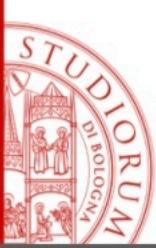
Without Redundancy

Operation #1

Name	Kind	#Accesses	Accesses Types
Person	E	1	W
Stay	R	1	W

Operation #2

Name	Kind	#Accesses	Accesses Types
City	E	1	R
Stay	R	5000	R



With Redundancy

- Costs:
 - Operation #1: 1500 writes and 500 reads per day
 - Operation #2: negligible.
- We count write accesses as doubles
 - A total amount of 3500 operations per day



Without Redundancy

- Costs:
 - Operation #1: 1000 writes
 - Operation #2: 10000 reads per day
- Writings costs double
 - A total amount of 12000 operations per day



Restructuring activities

- Redundancies analysis
- Generalizations deletion
- Partitioning/grouping of entities and relationships
- Identifying the primary keys



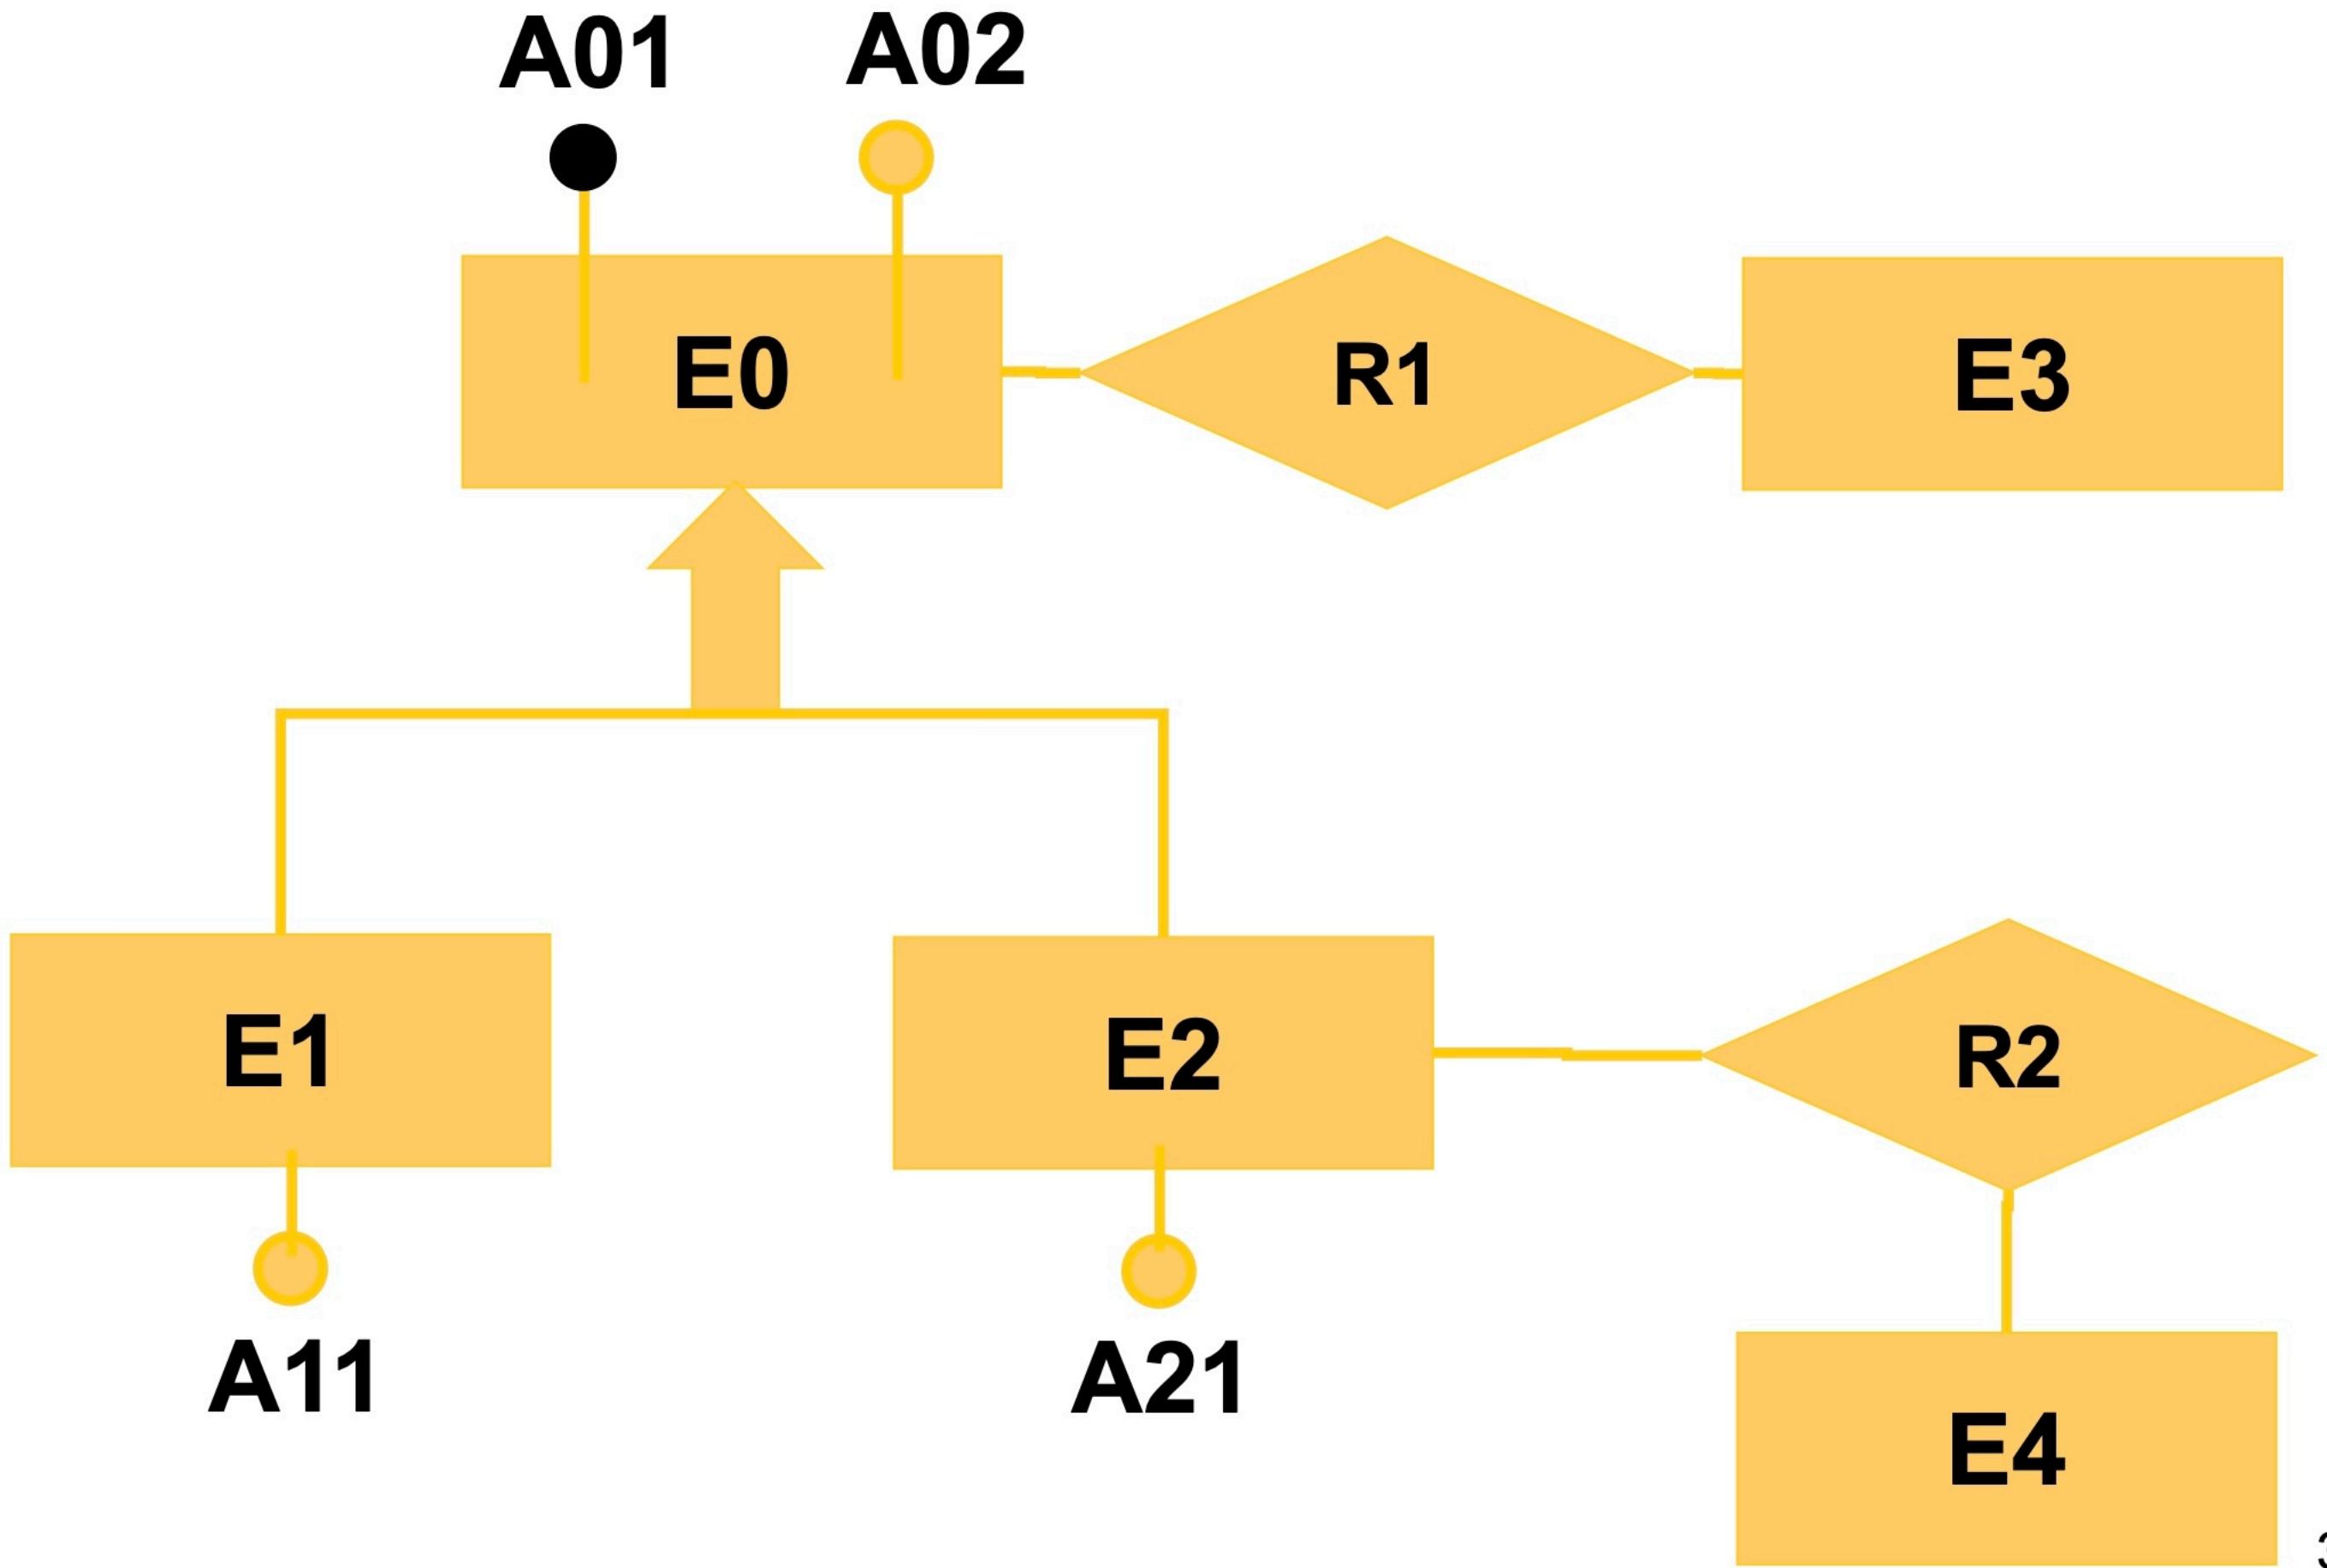
Hierarchies deletion

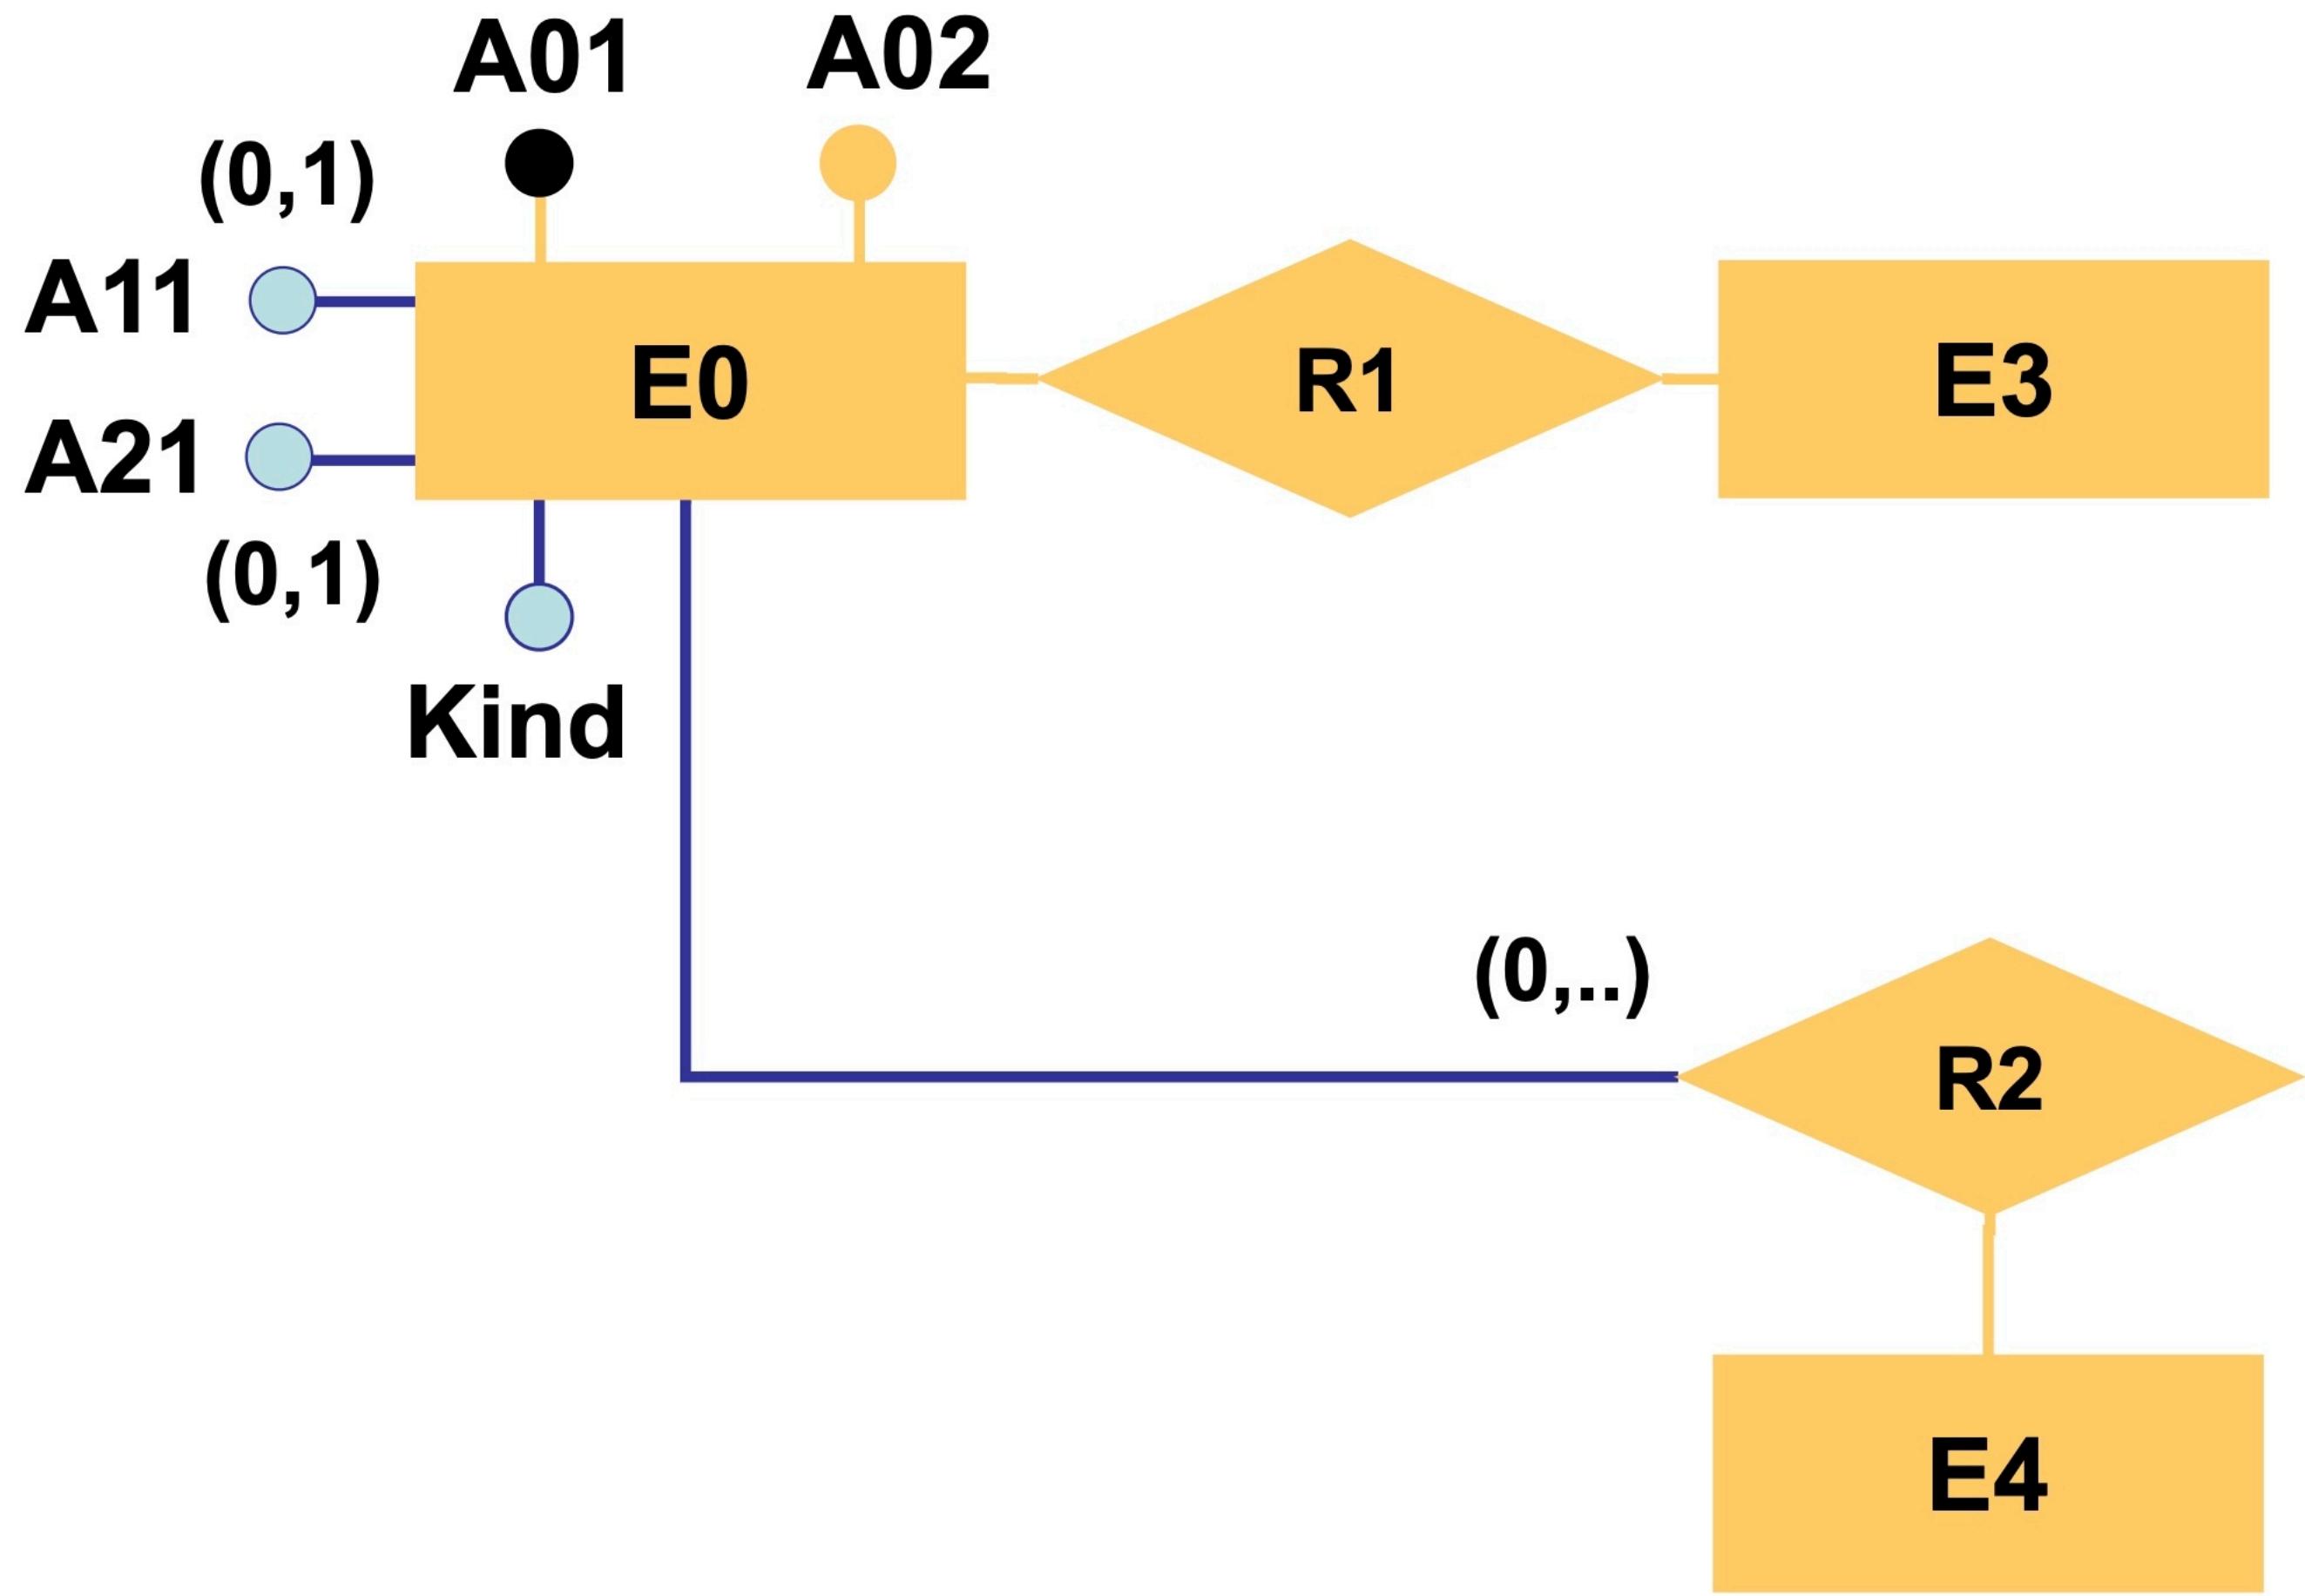
- The relational data model does not directly support generalizations, they cannot be directly represented
- ...while entities and relationships are directly representable
- We therefore remove hierarchies replacing them with entities and relationships

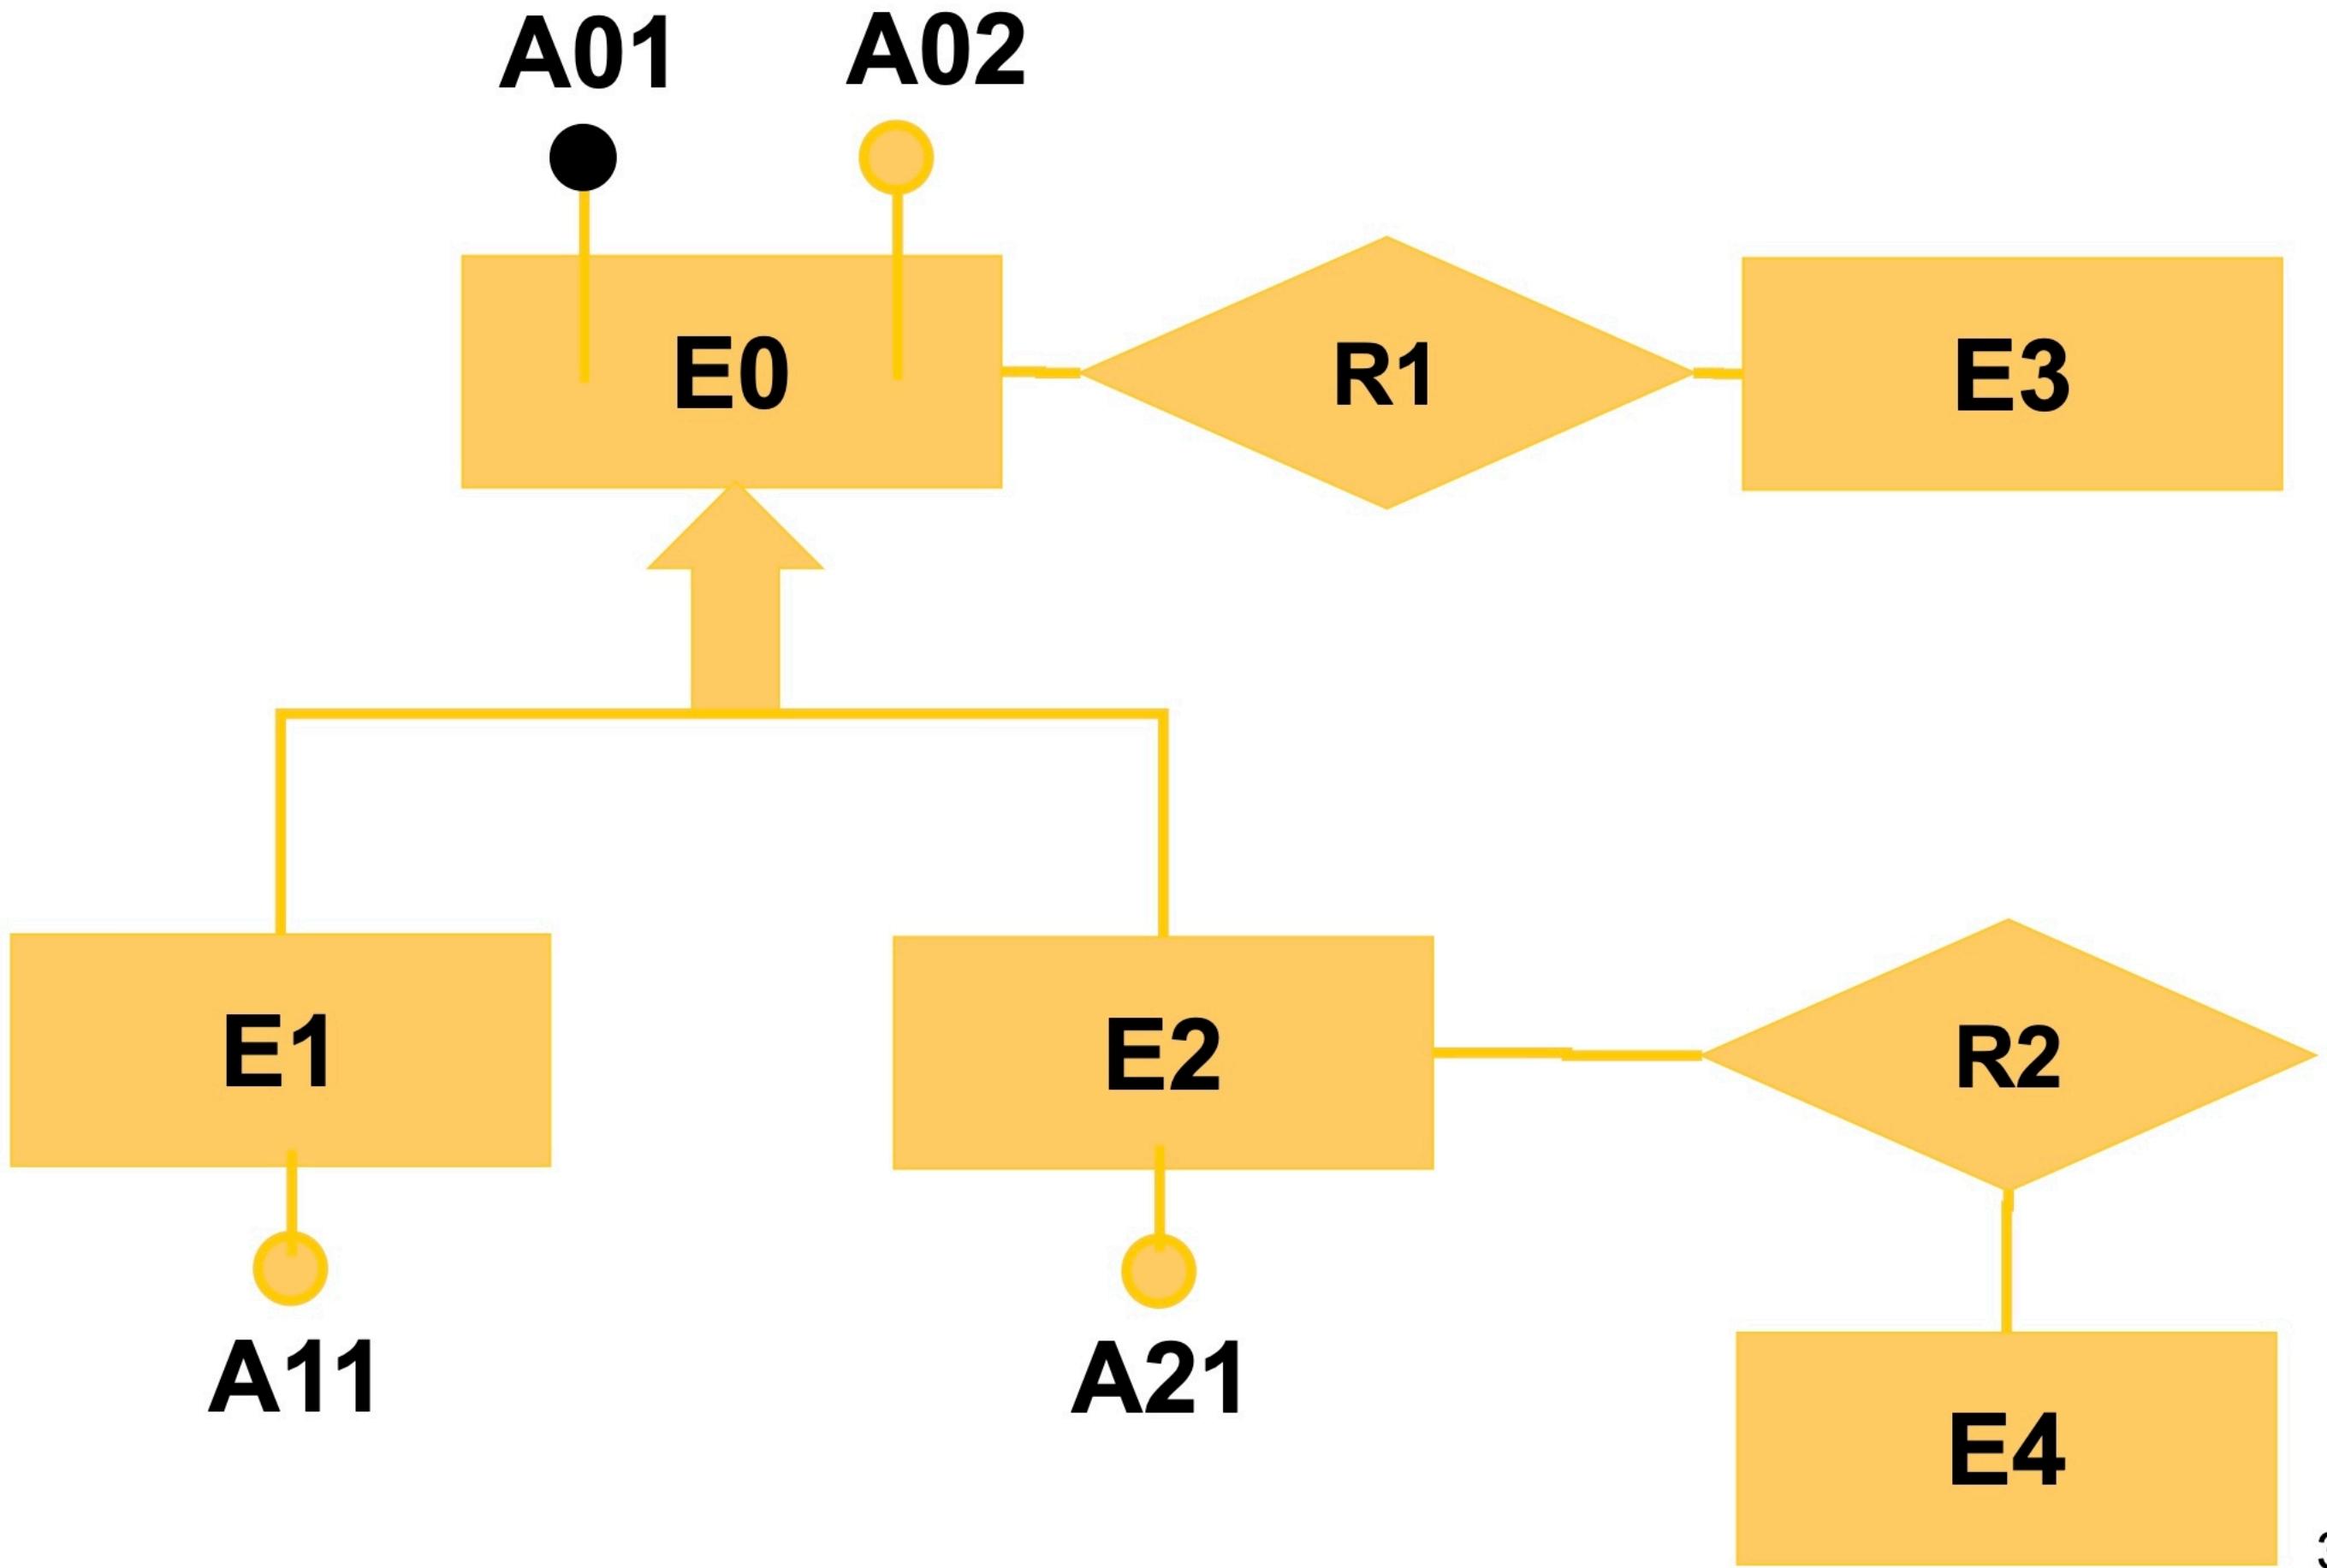


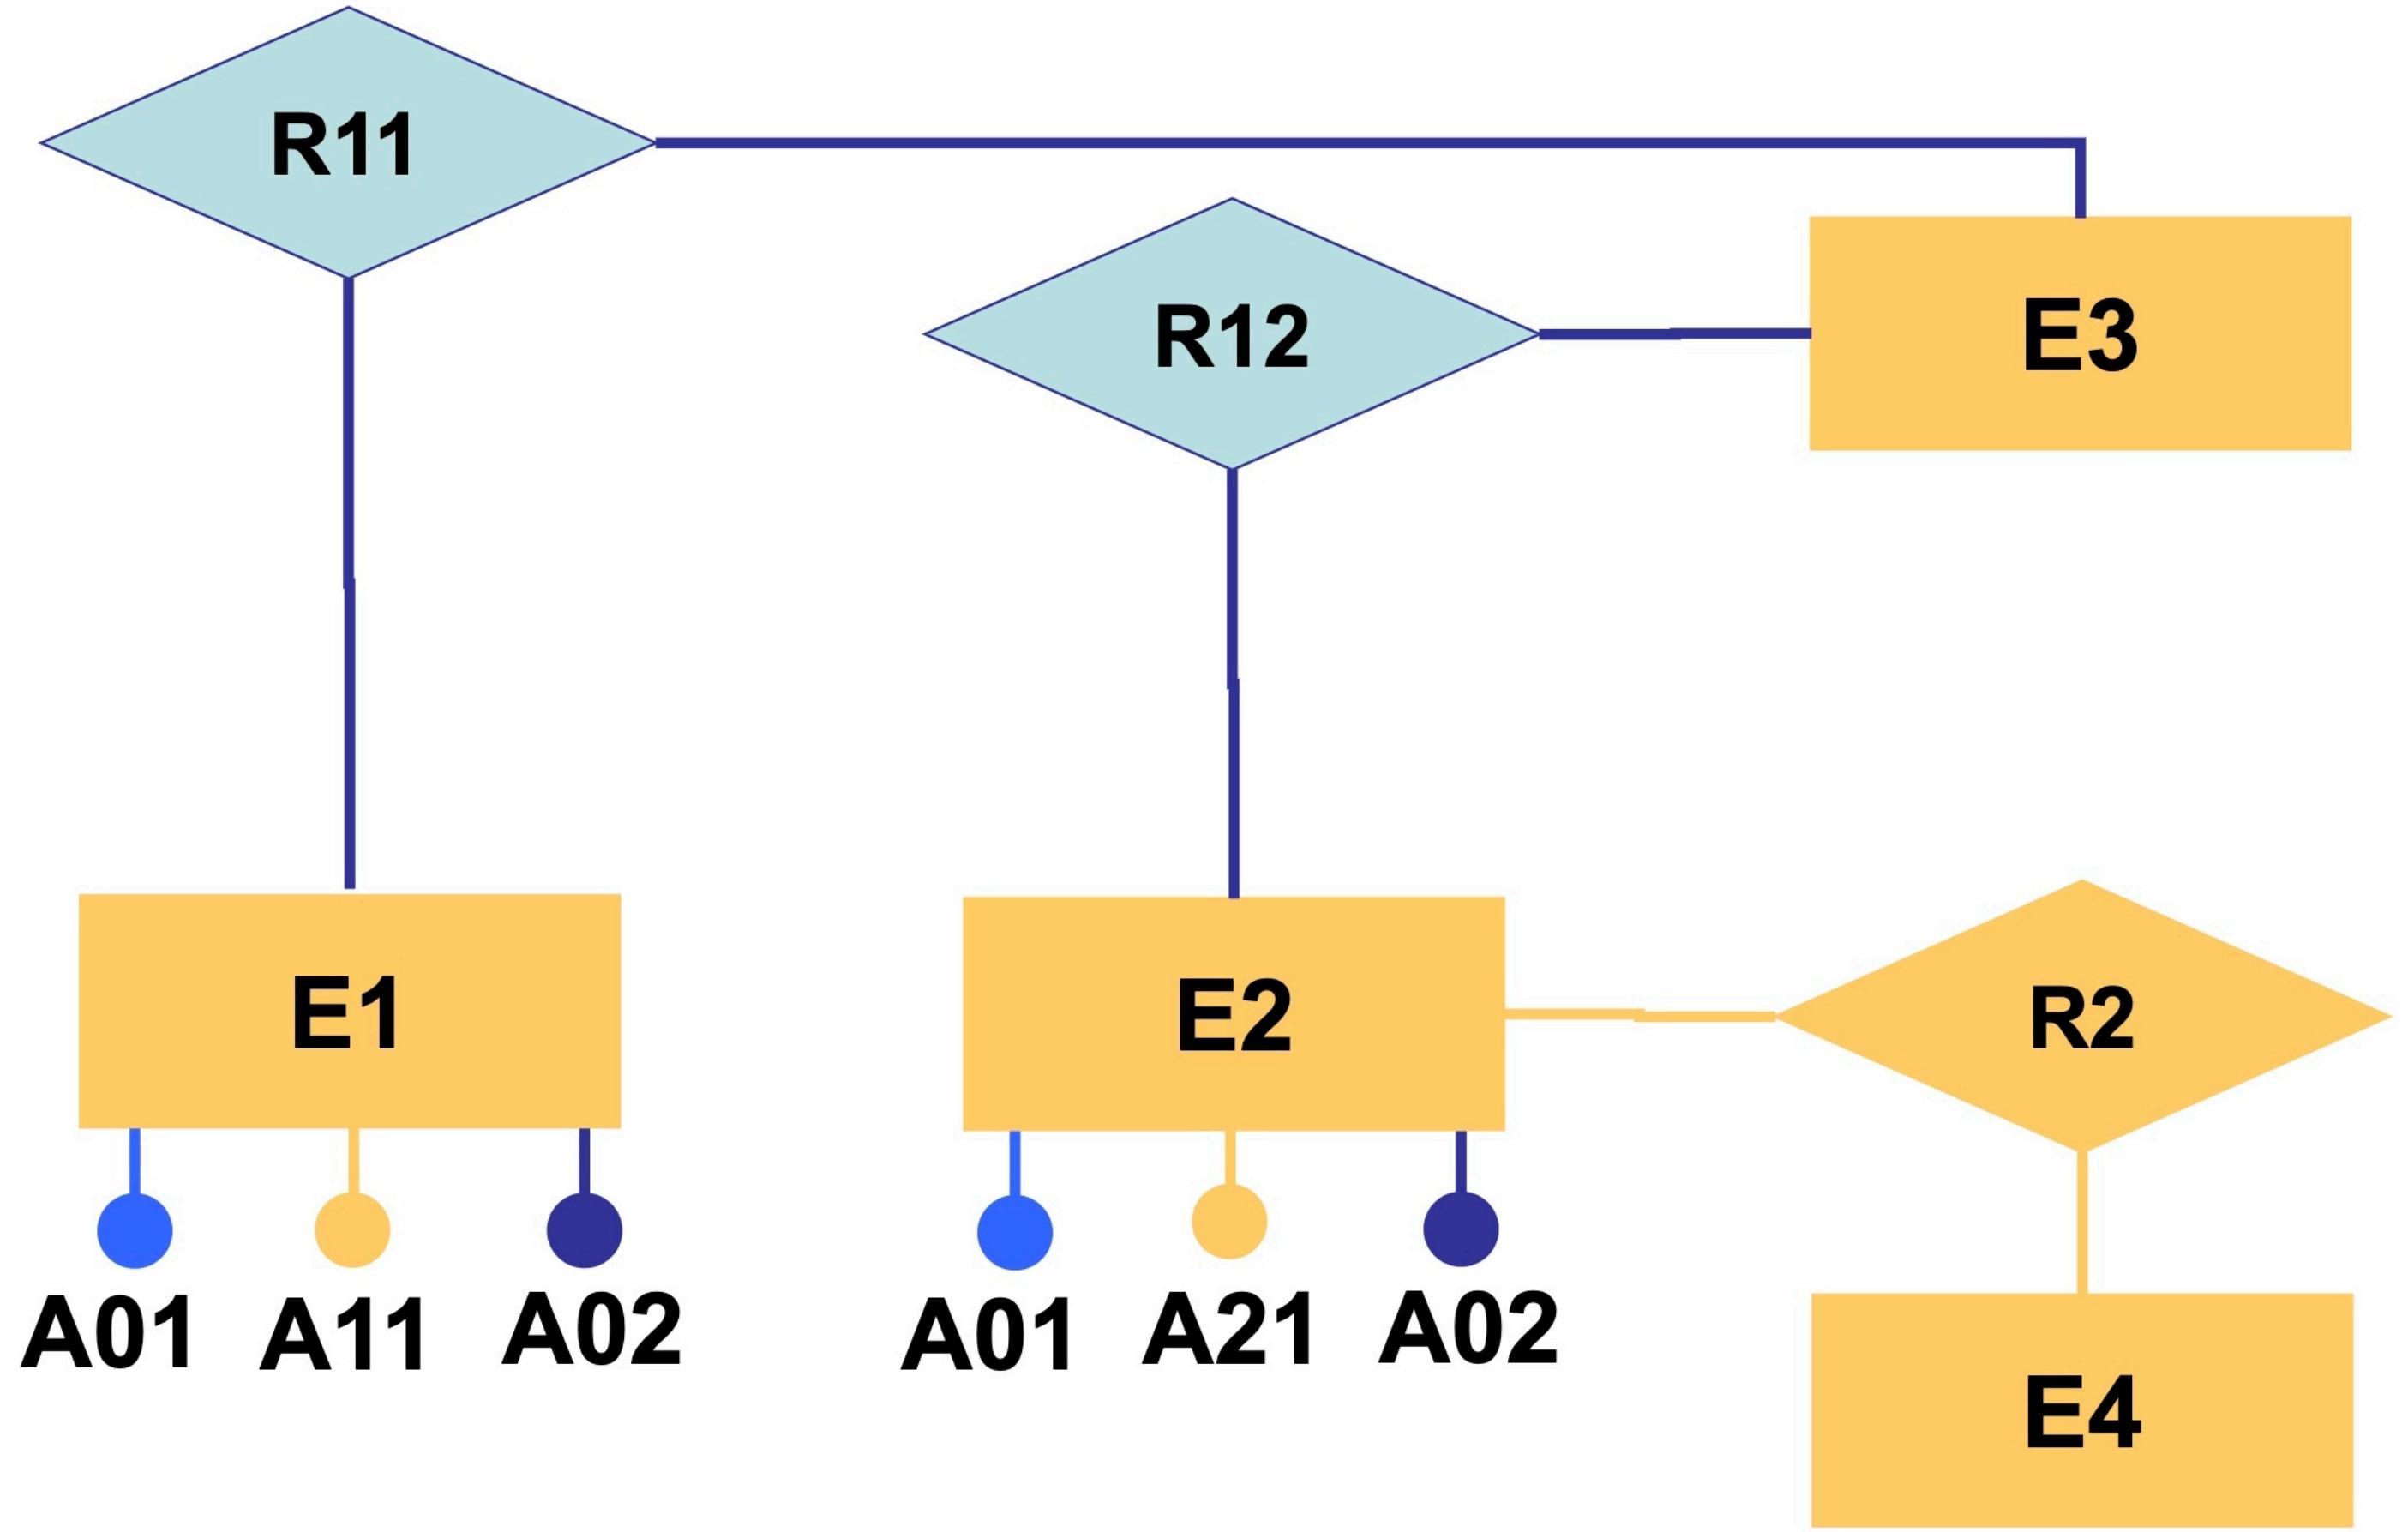
Three possible solutions

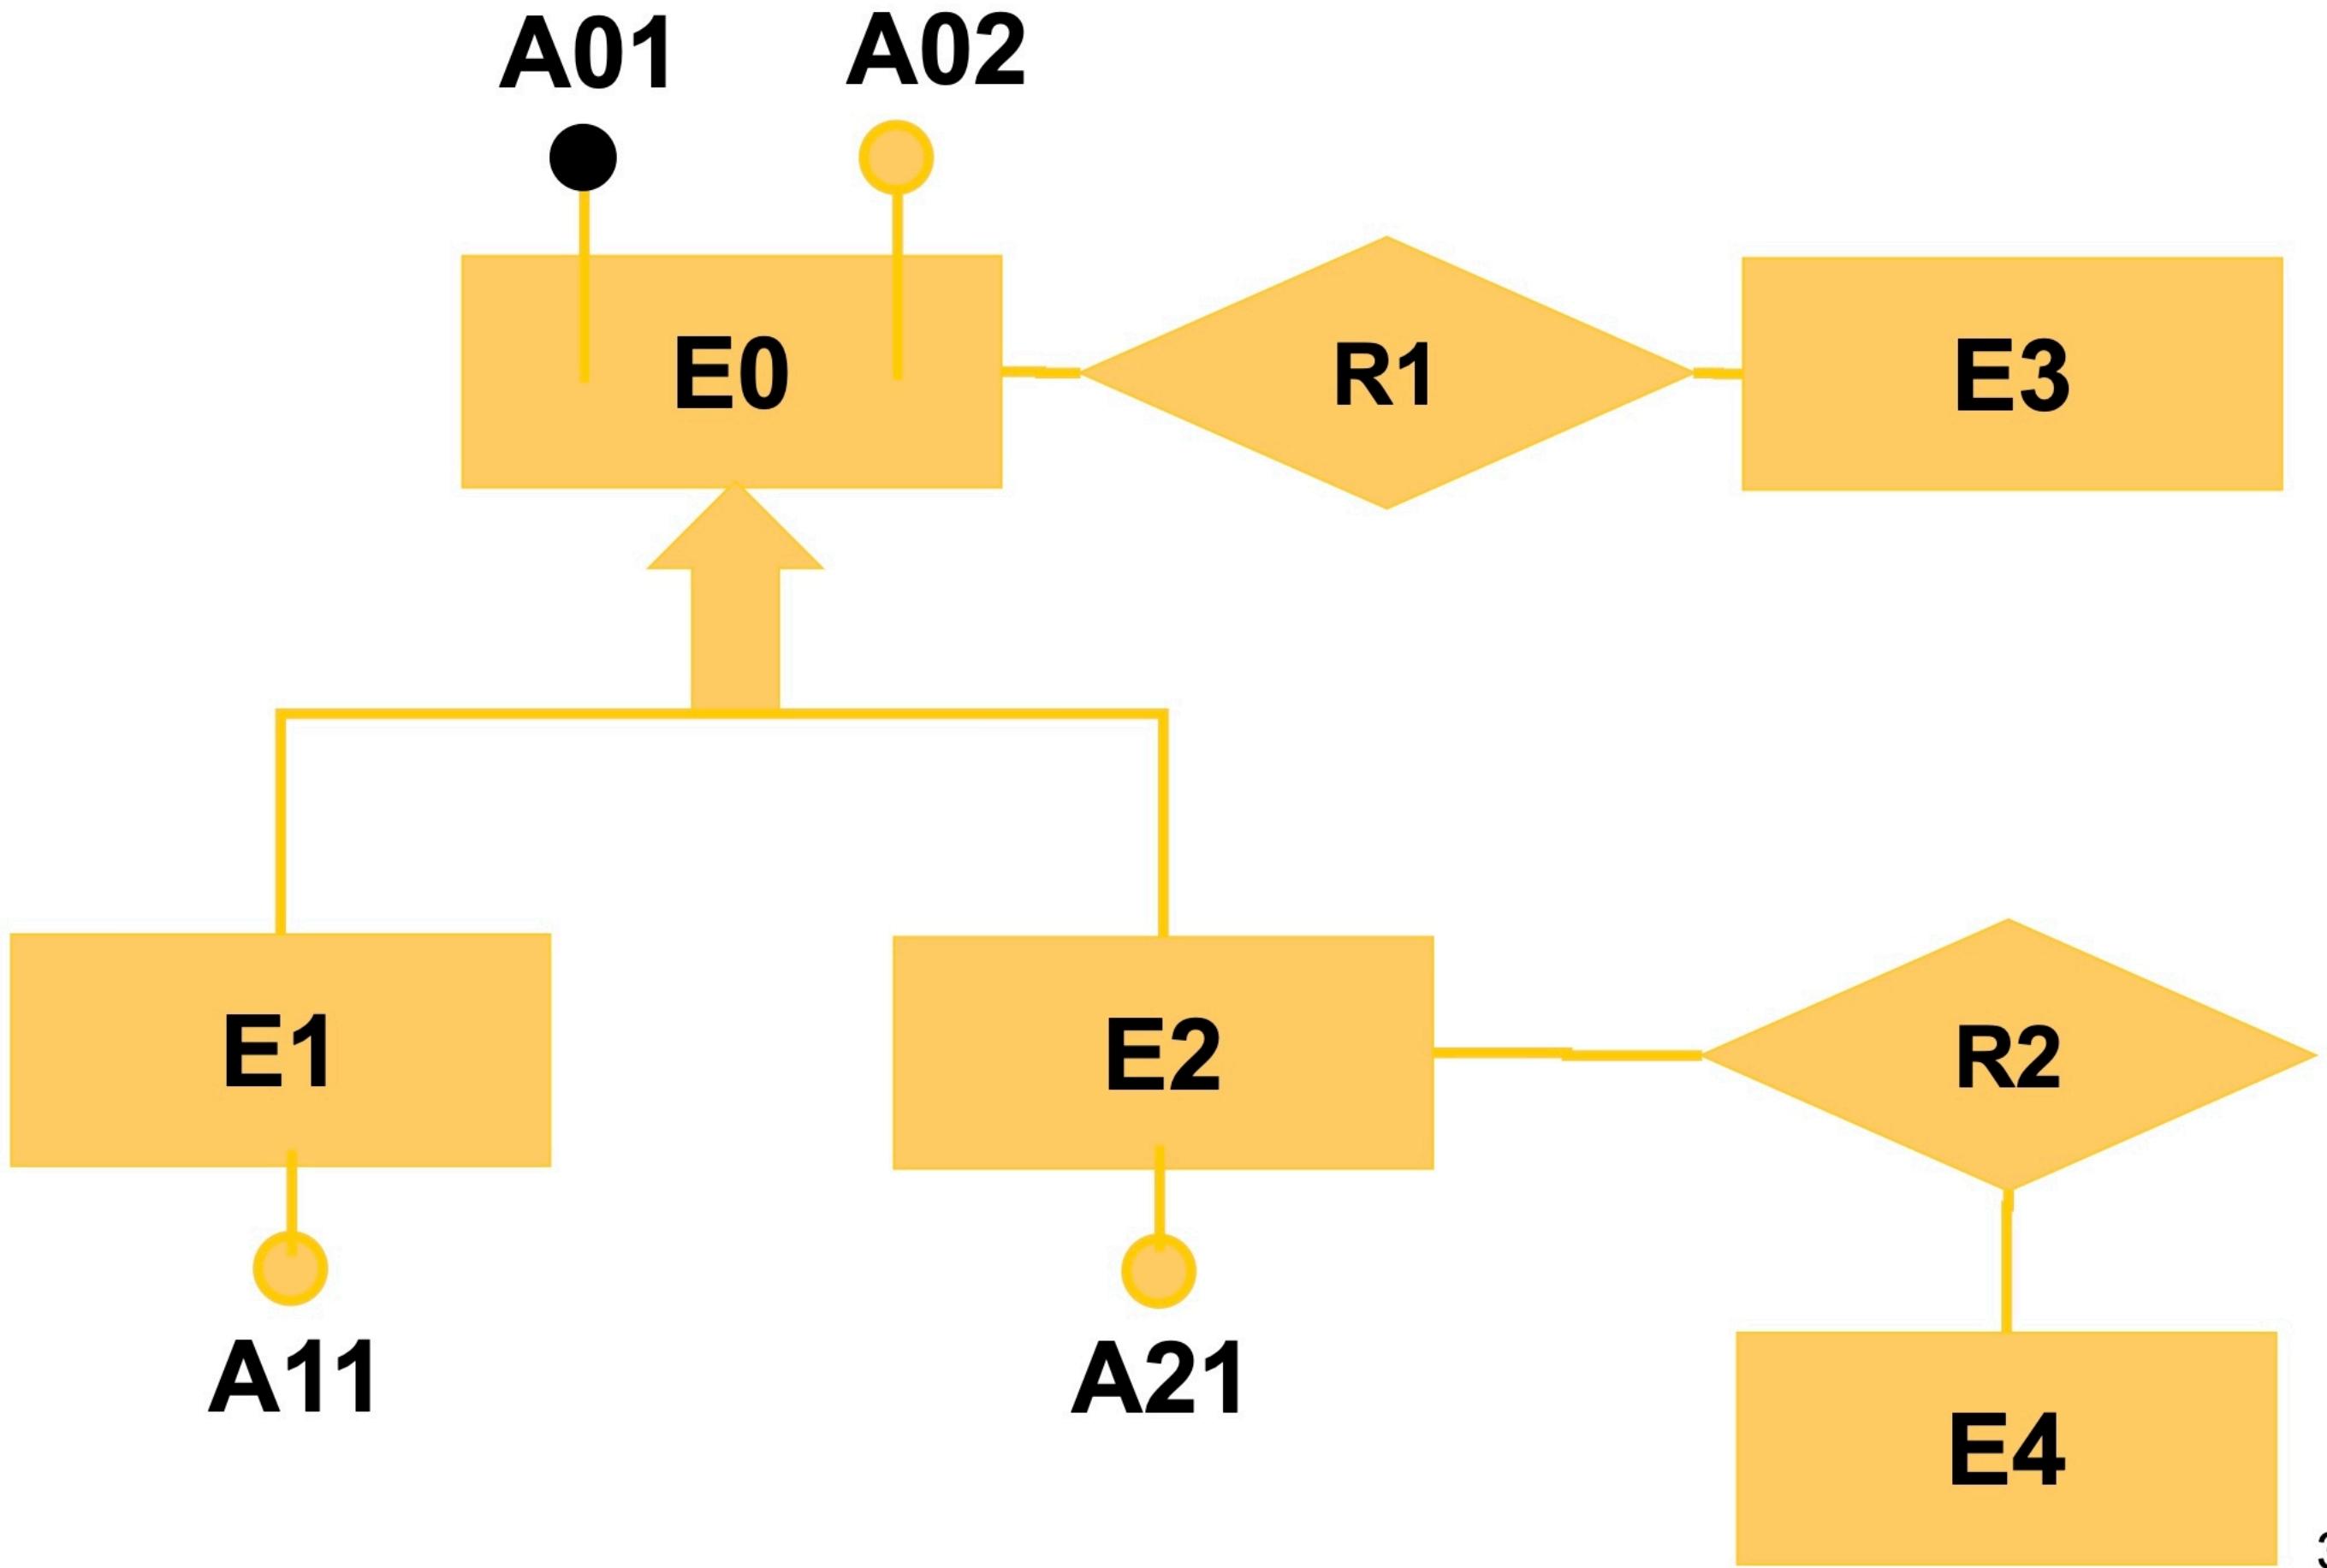
1. Embedding the children of the generalization into the parent.
2. Embedding the parent of the generalization into the children
3. Replacing the generalization with a relationship

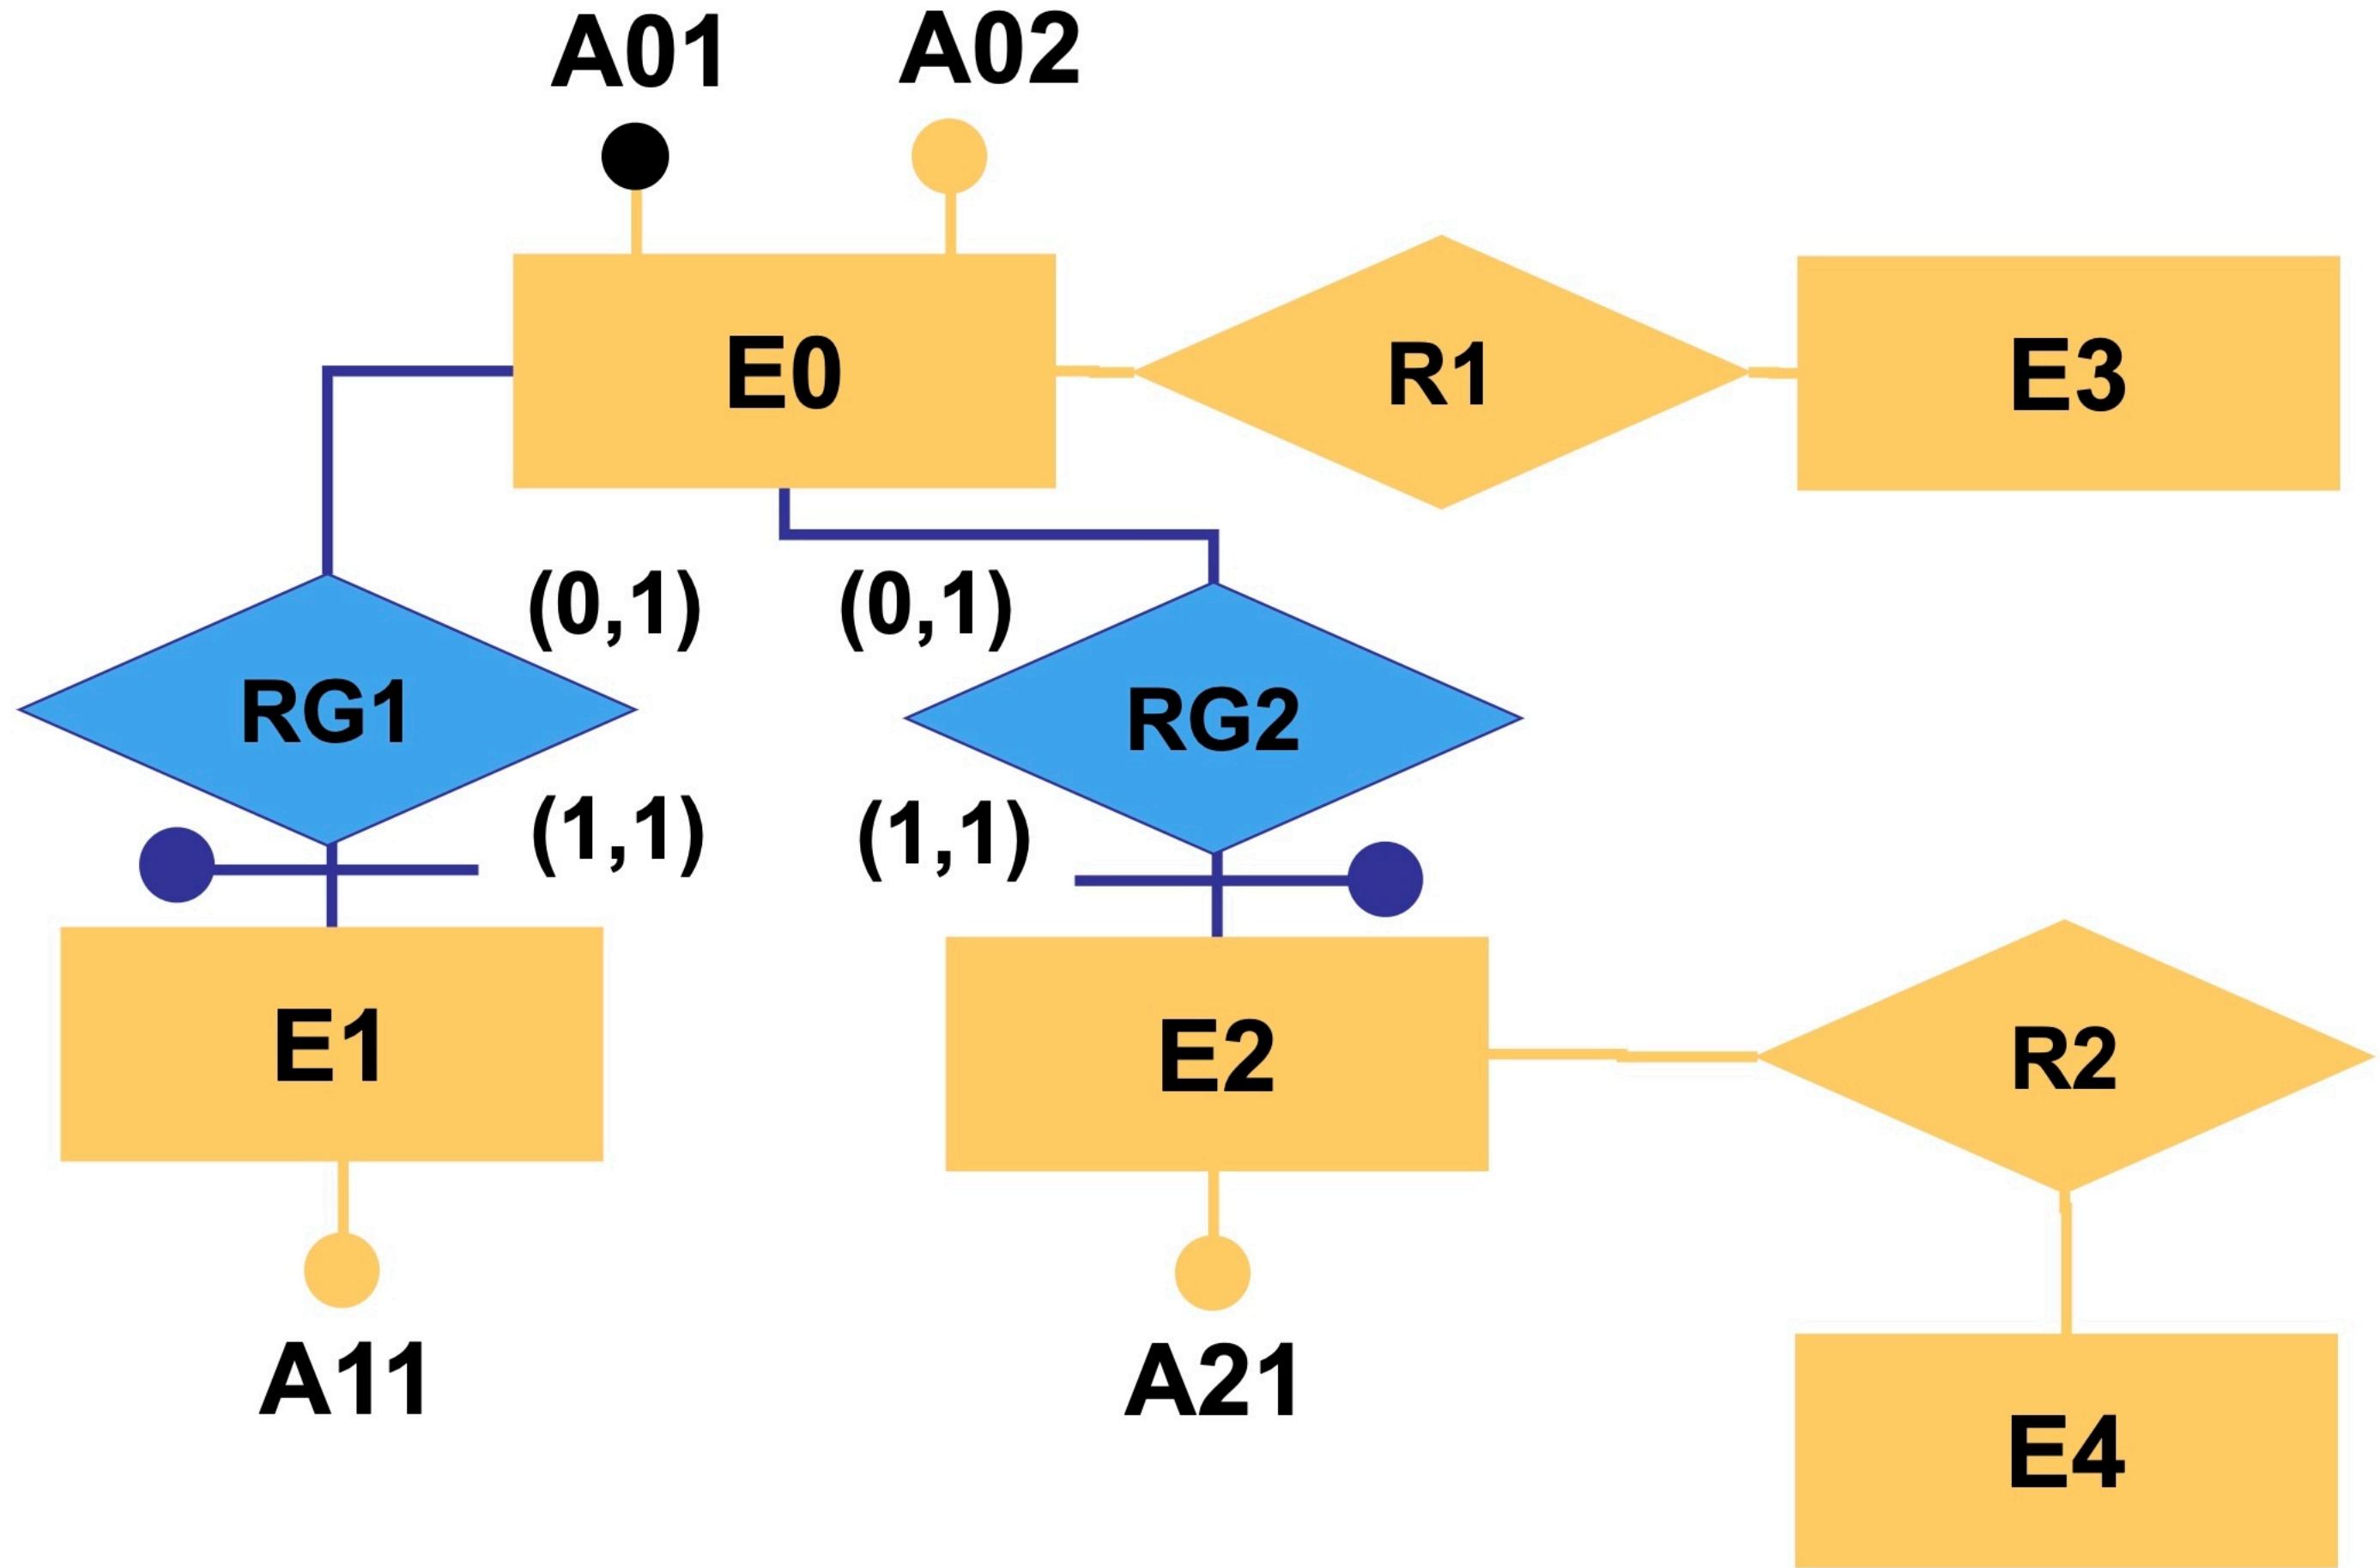








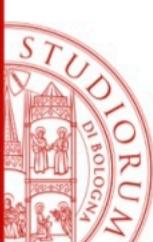






Observations (1)

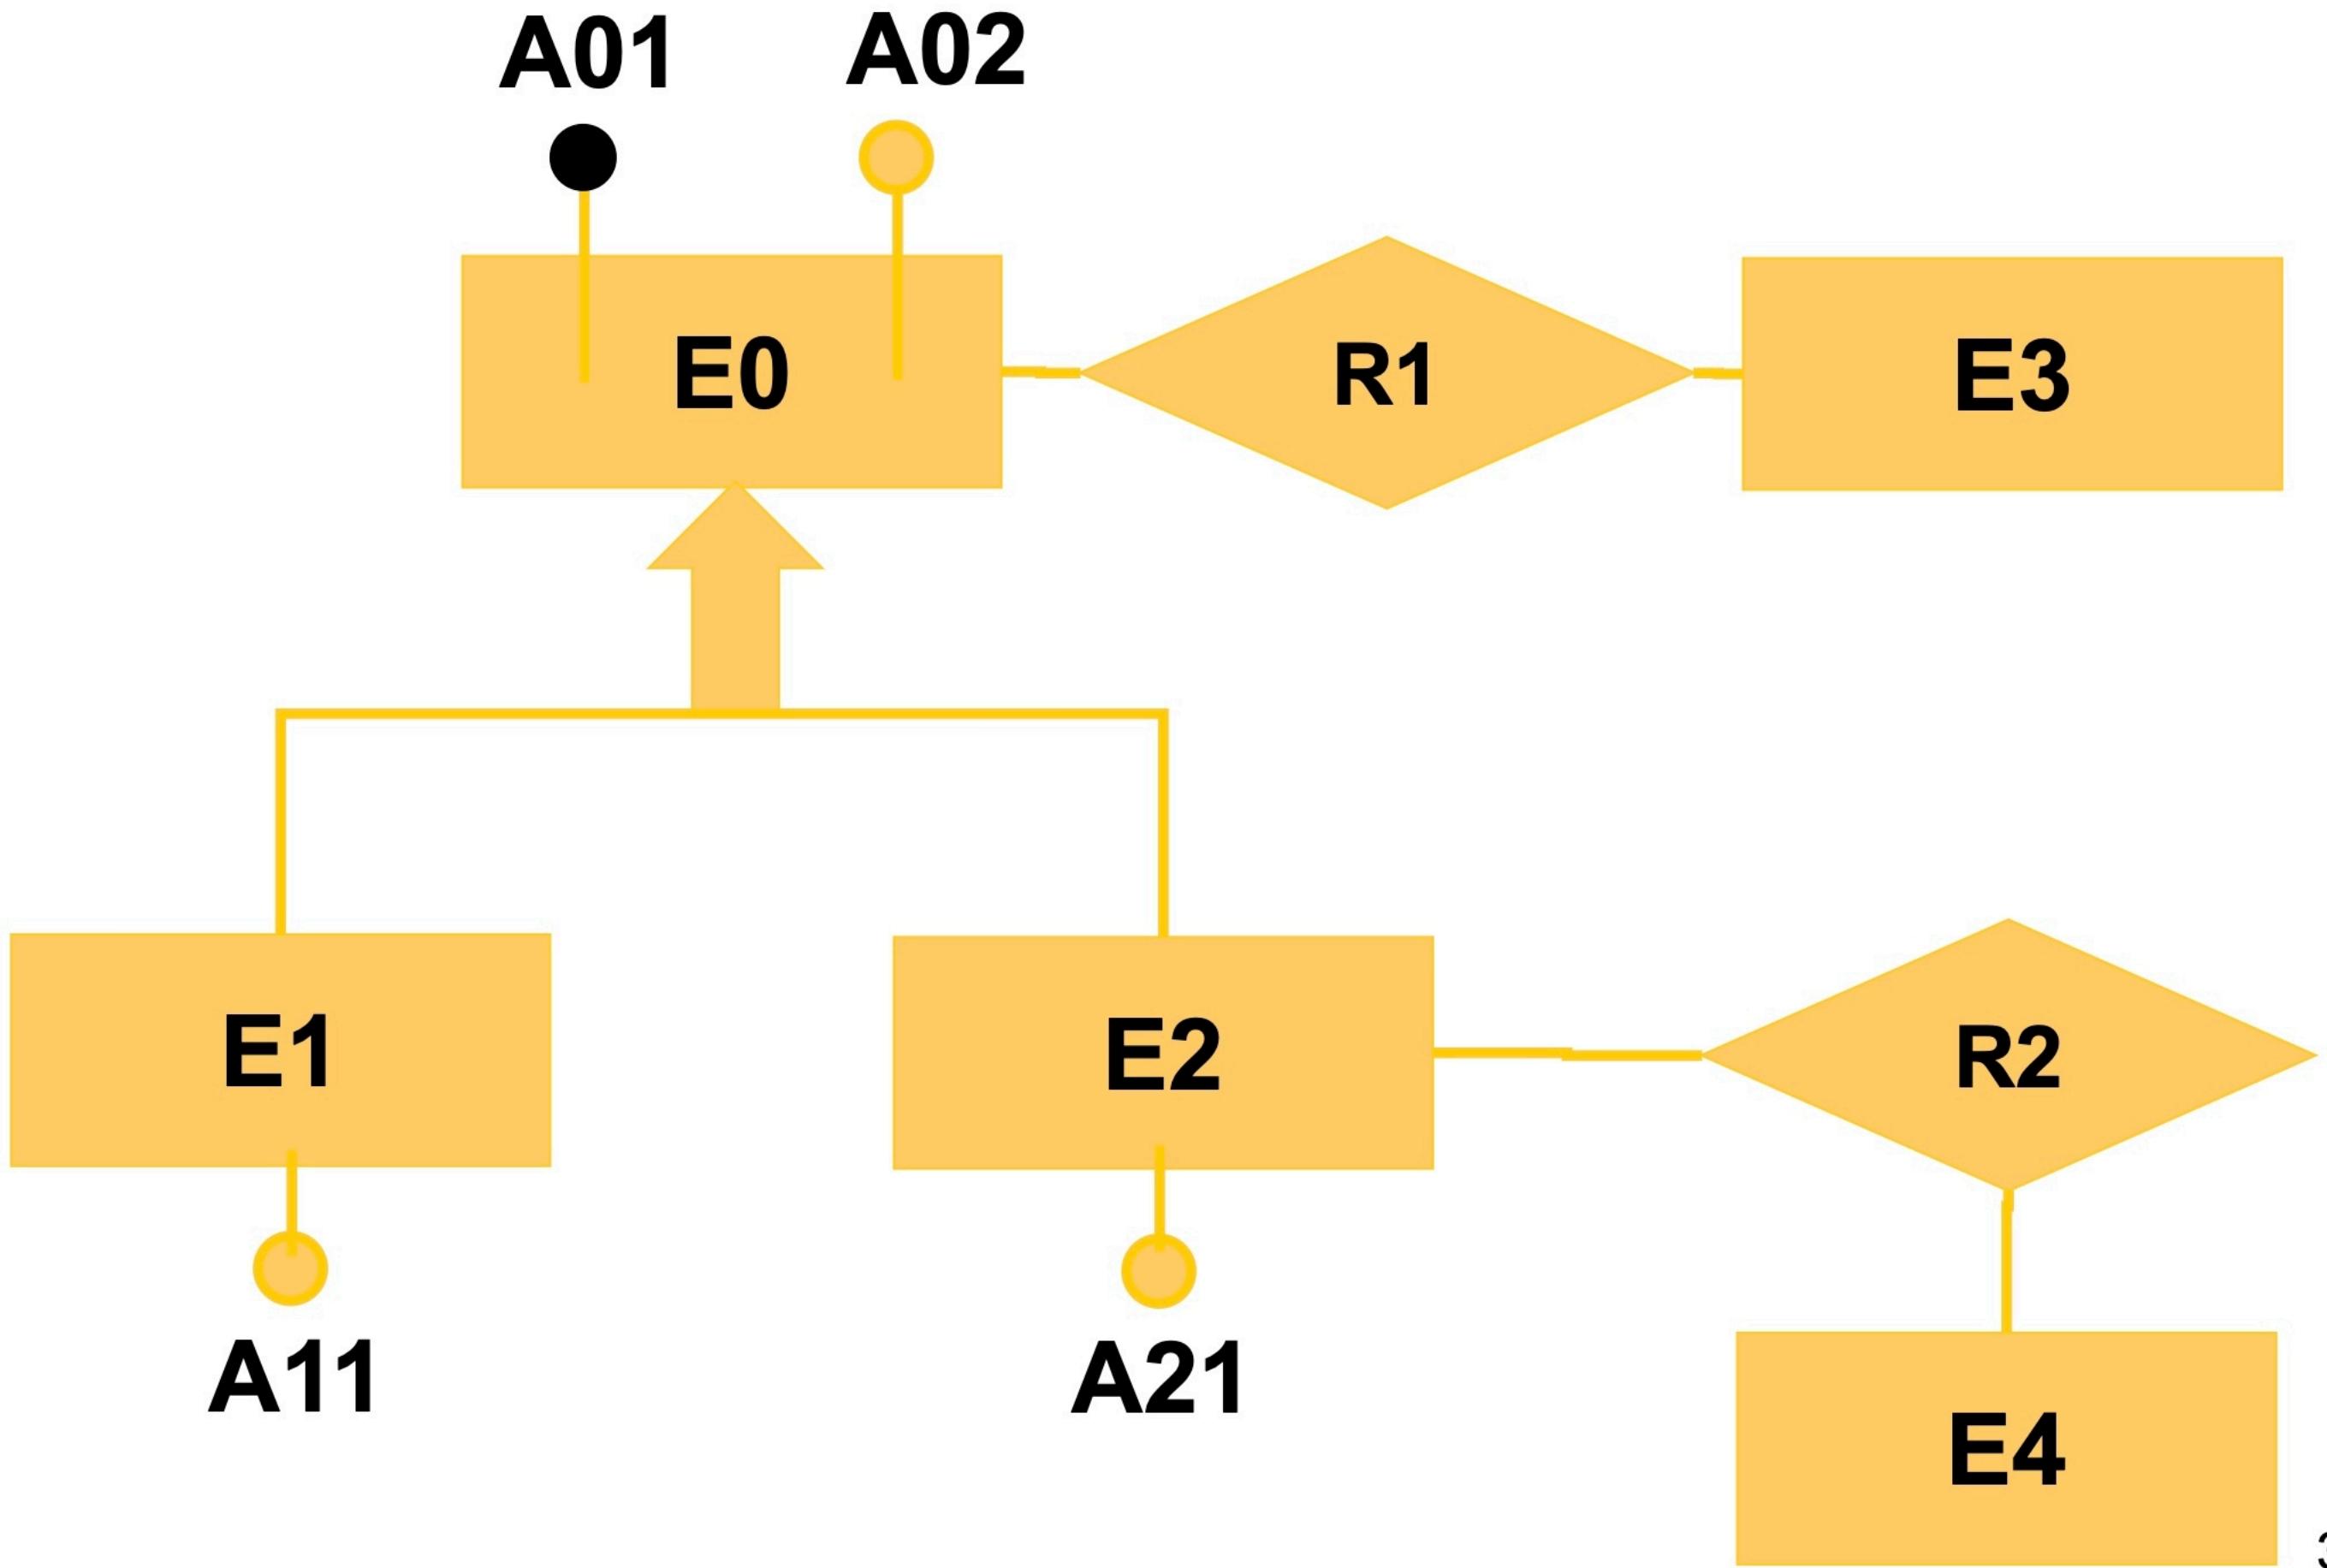
- We can decide between those three alternatives depending on the sizes and access tables (considering only the number of accesses is not enough)
- We now provide some general rules that can be followed

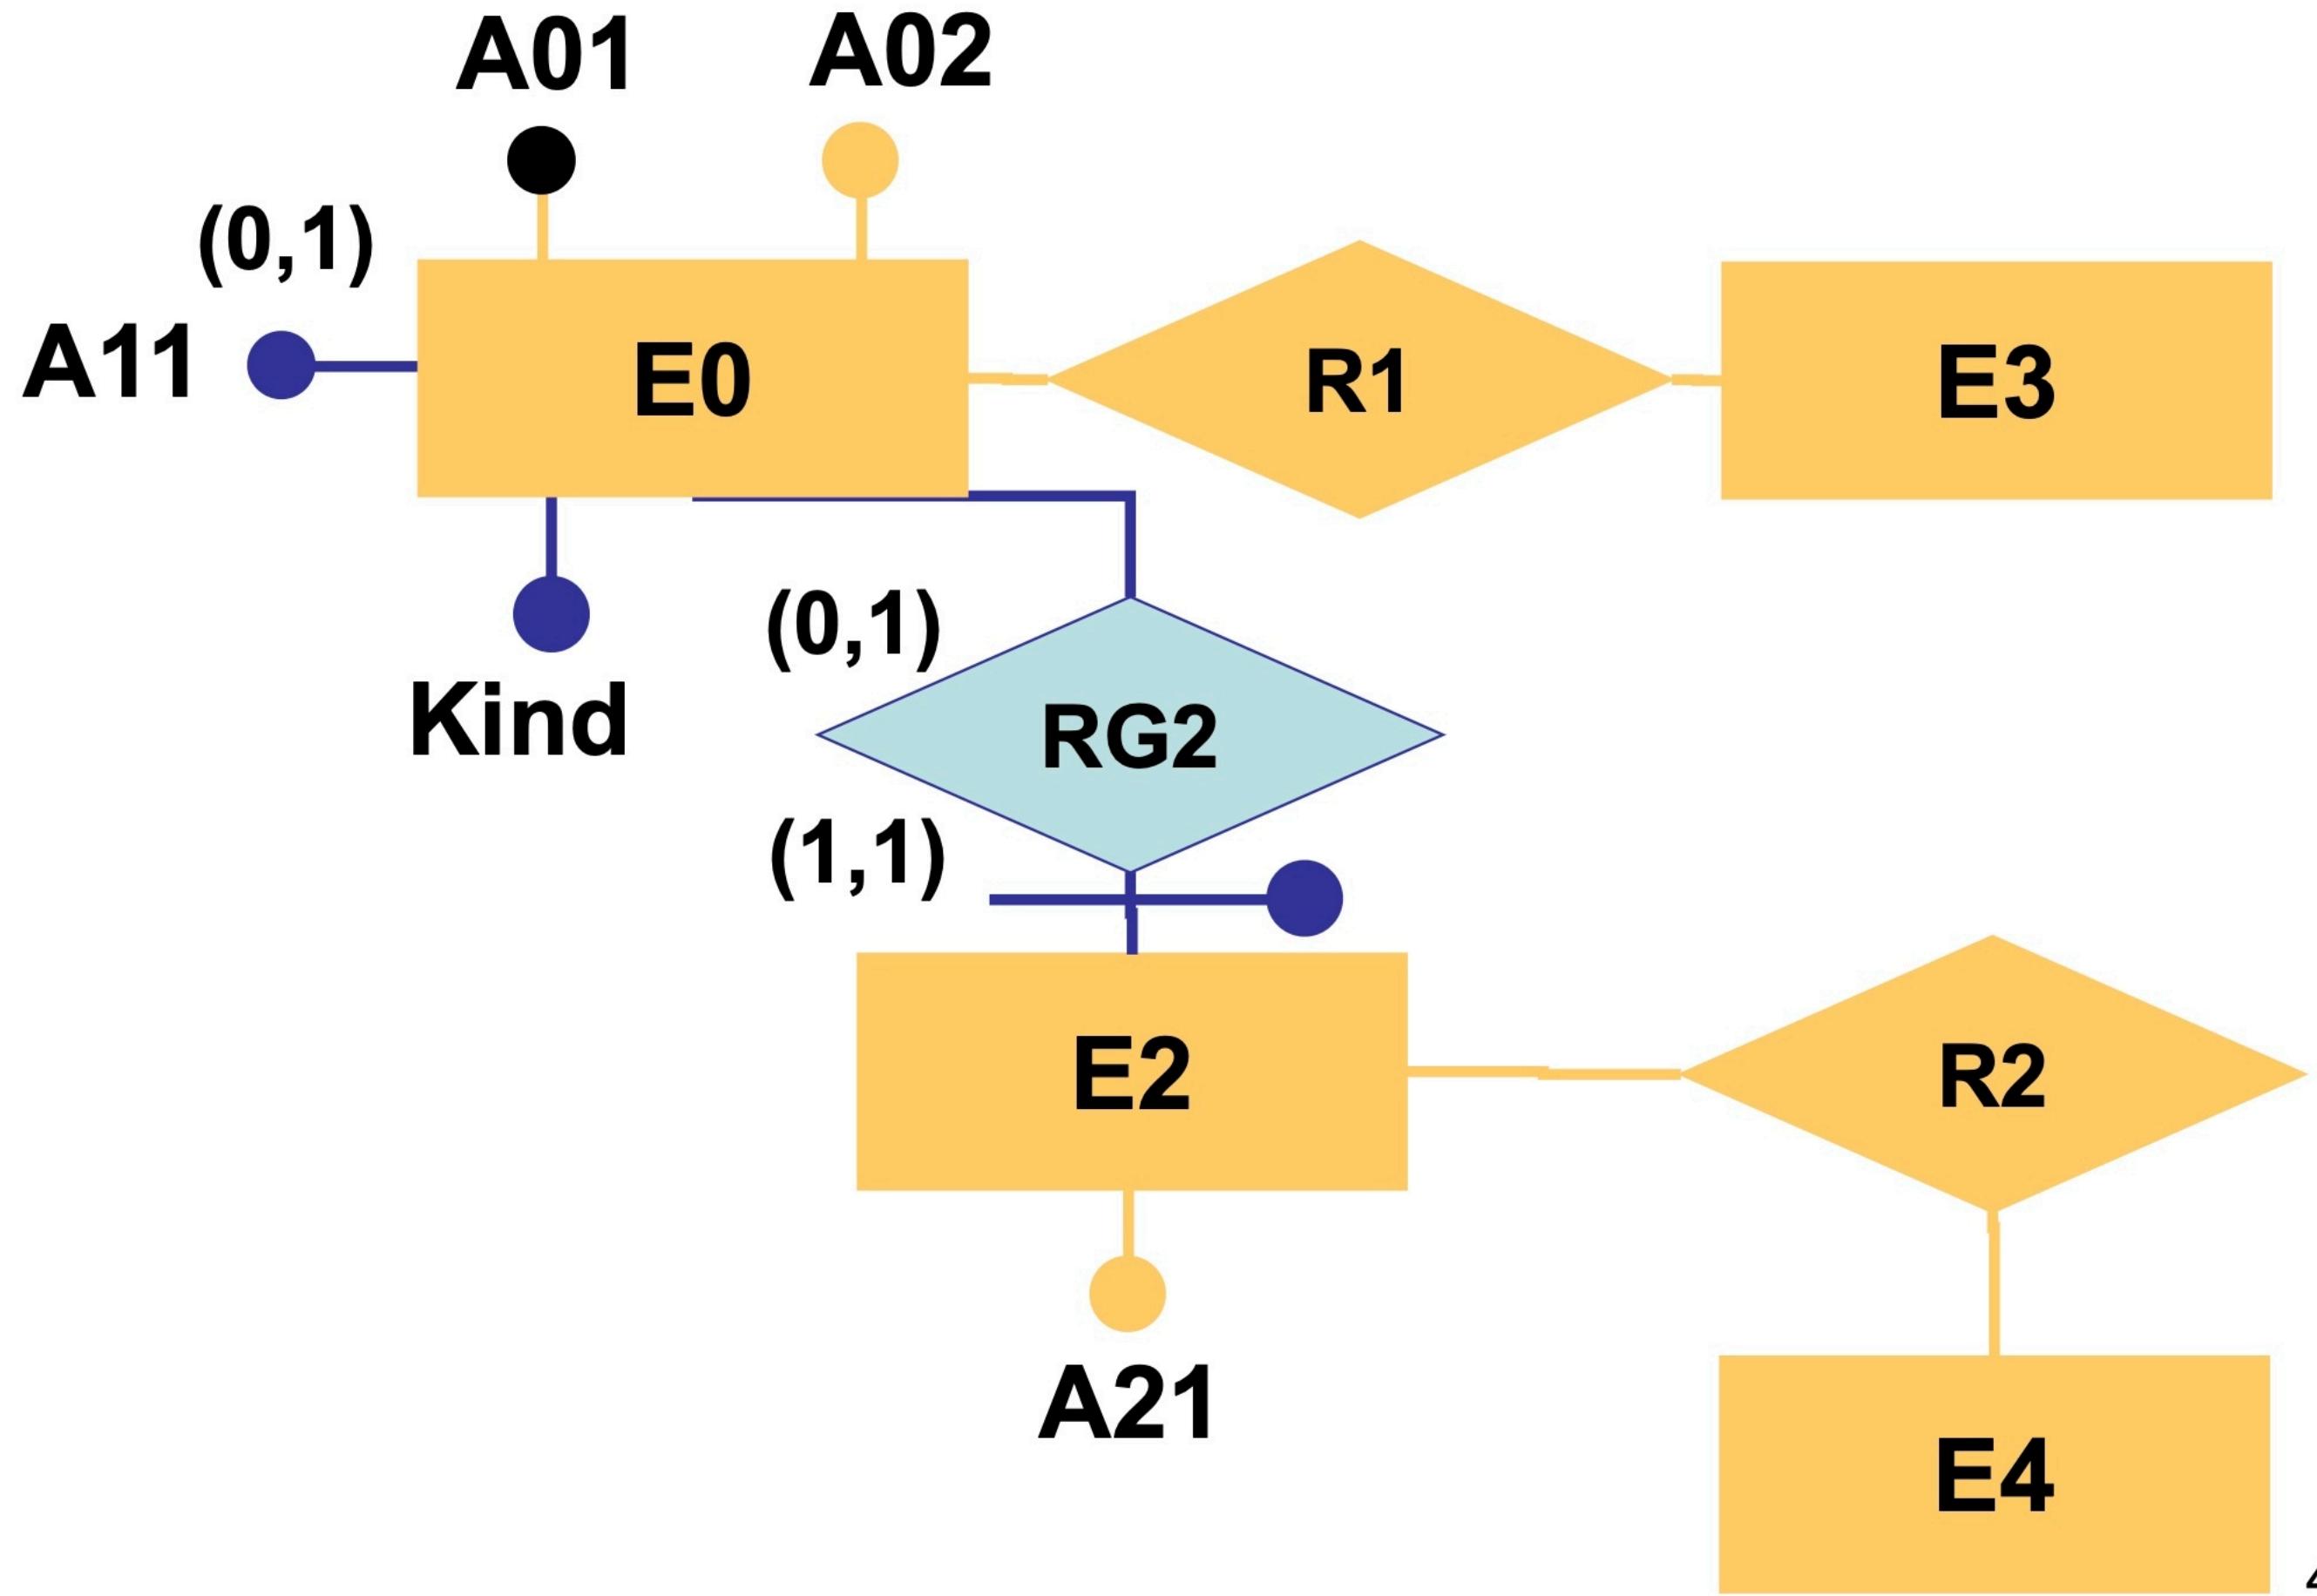


Observations (2)

The three solutions (1, 2 and 3):

1. Should be used when children and father are accessed at the same time
 2. Should be used when children are accessed independently from one another
 3. Should be used when children are accessed independently from the father
-
- We can also apply “hybrid” solutions, especially in hierarchies with more levels.



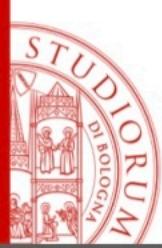




Restructuring activities

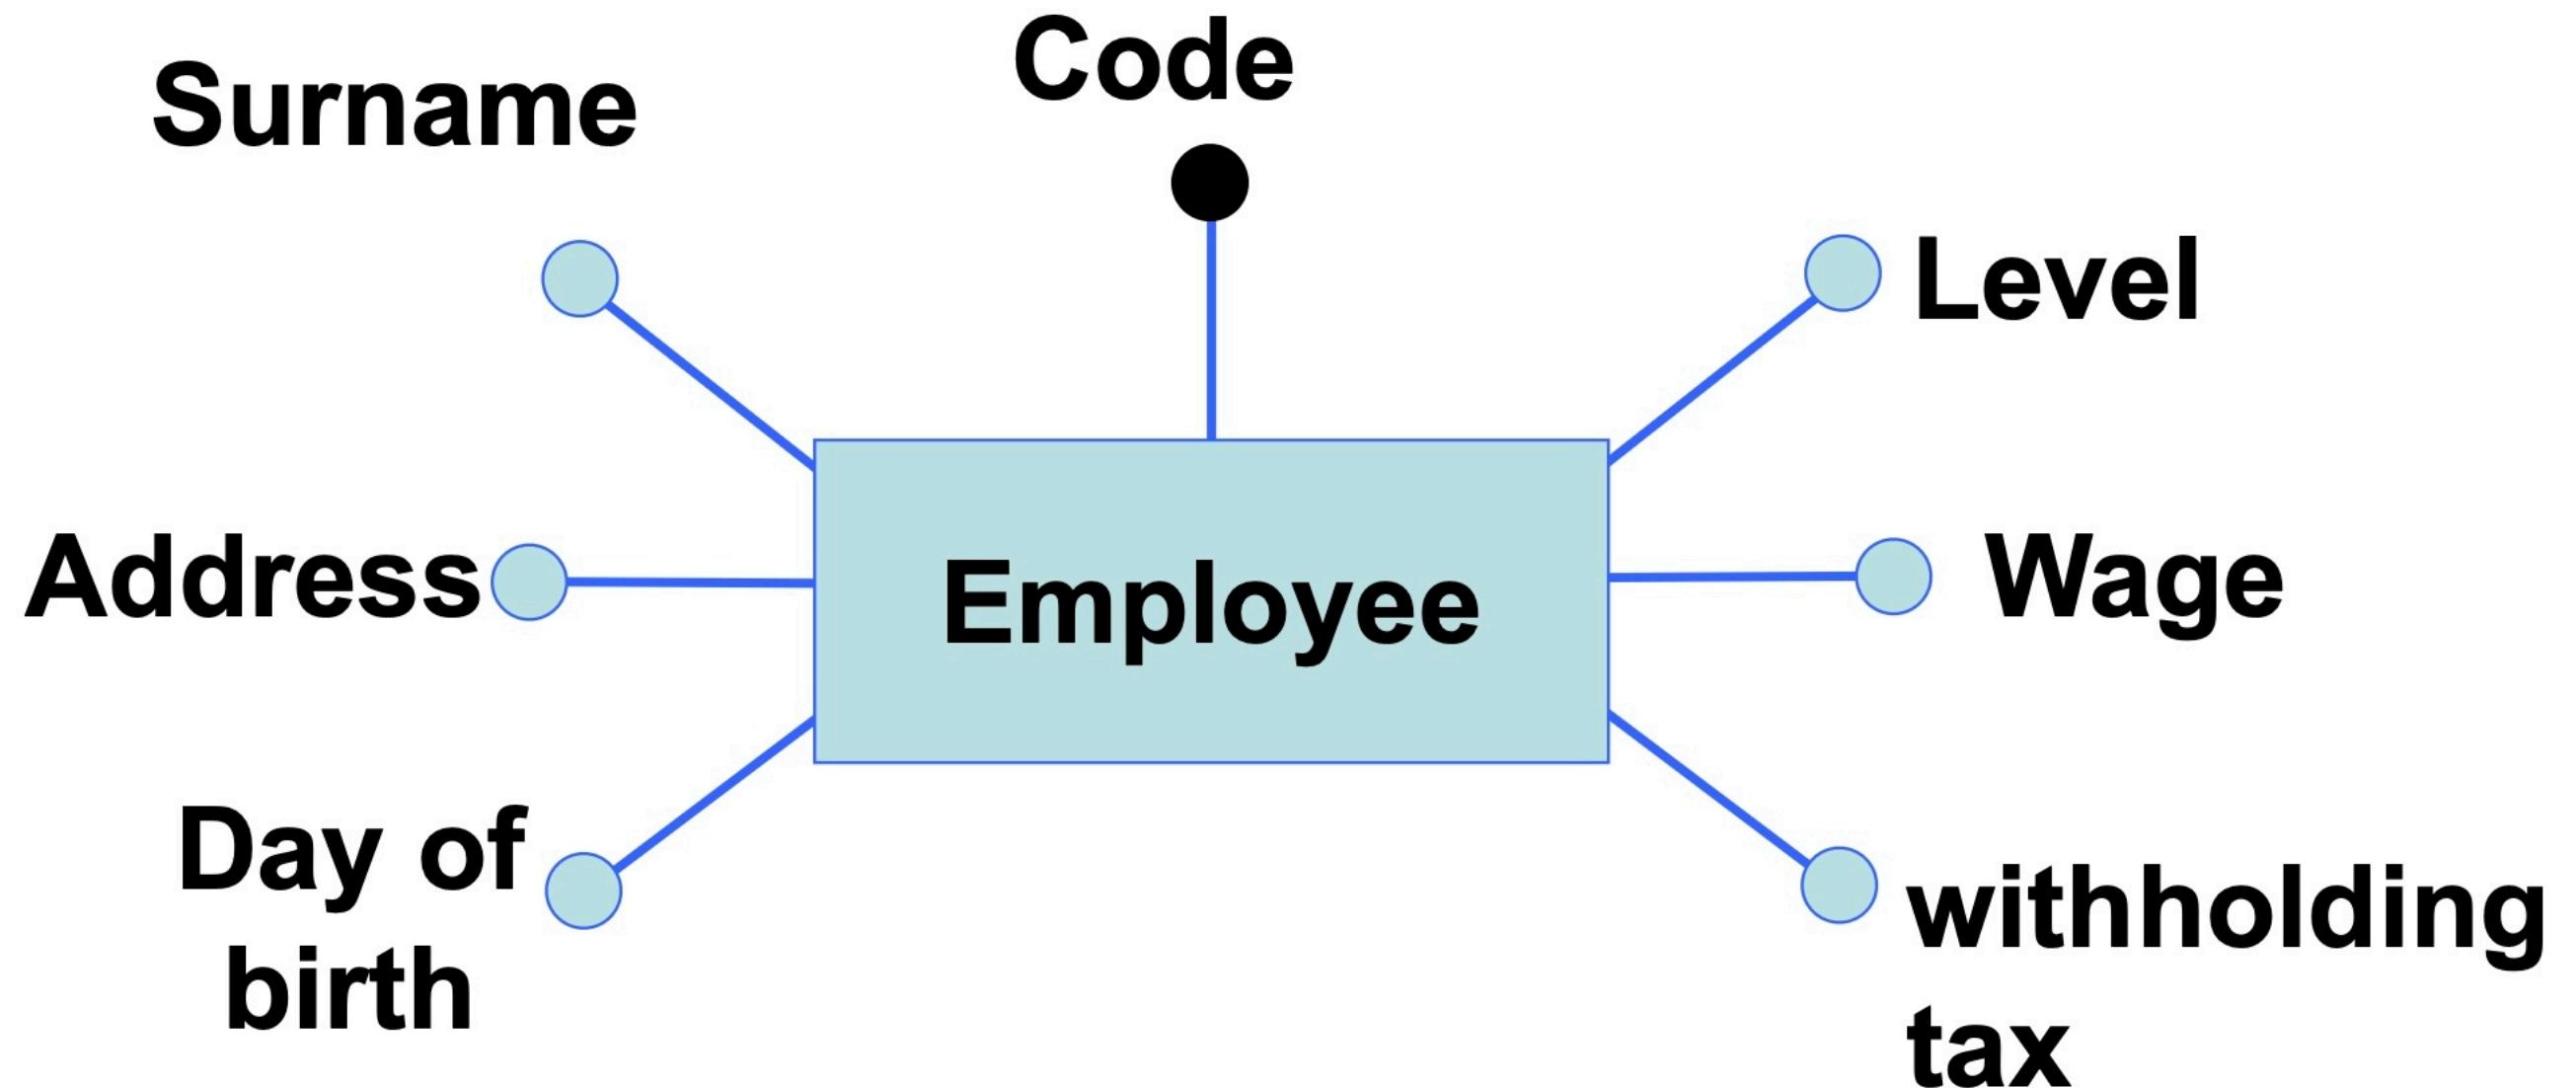
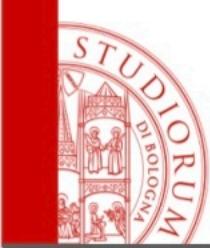
- Redundancies analysis
- Generalizations deletion
- Partitioning/grouping of entities and relationships
- Identifying the primary keys

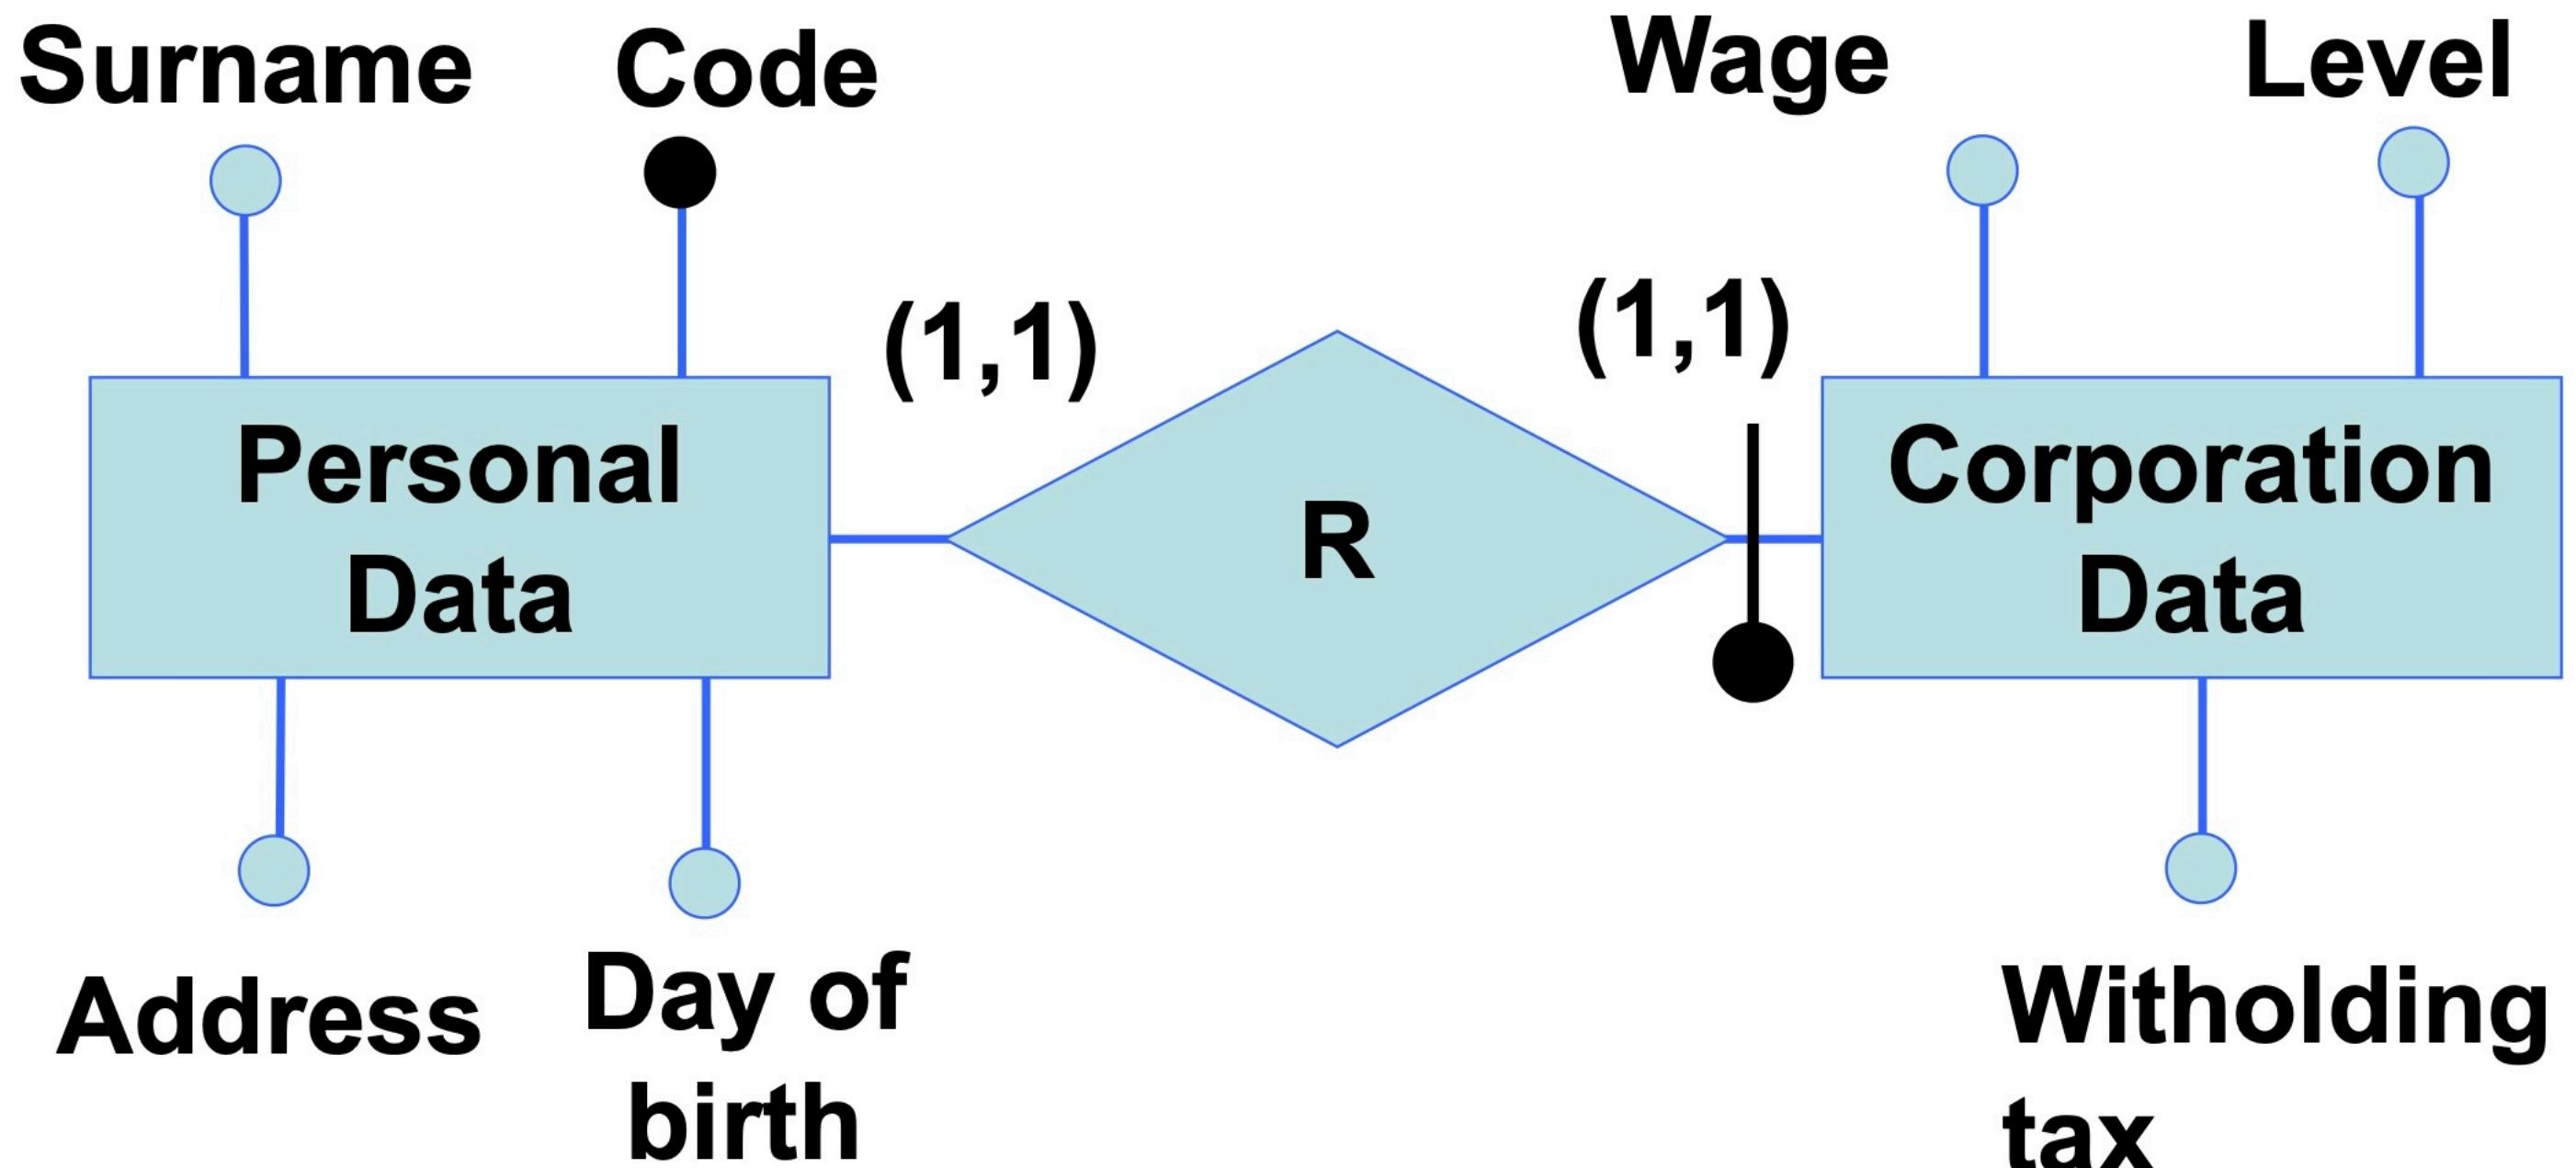
- Restructuring can provide more efficient operations by simply **reducing the number of accesses**:
 - Attributes accessed separately are splitted
 - Attributes accessed on the same time are grouped, even when they belong to different entities/relationships

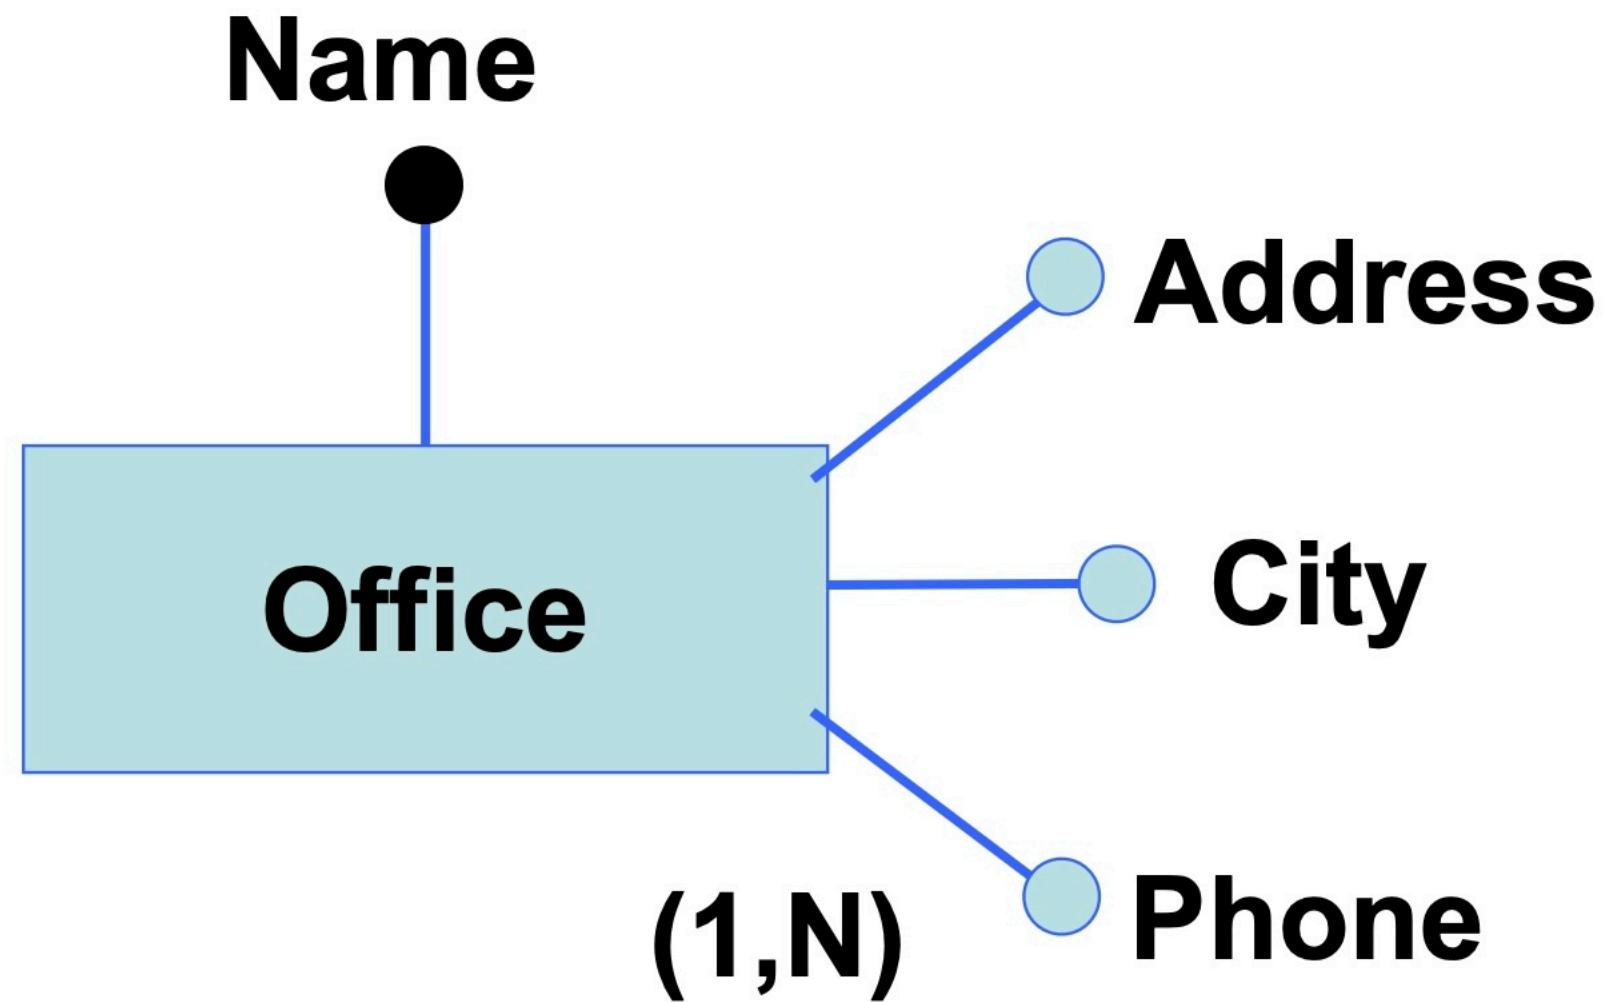


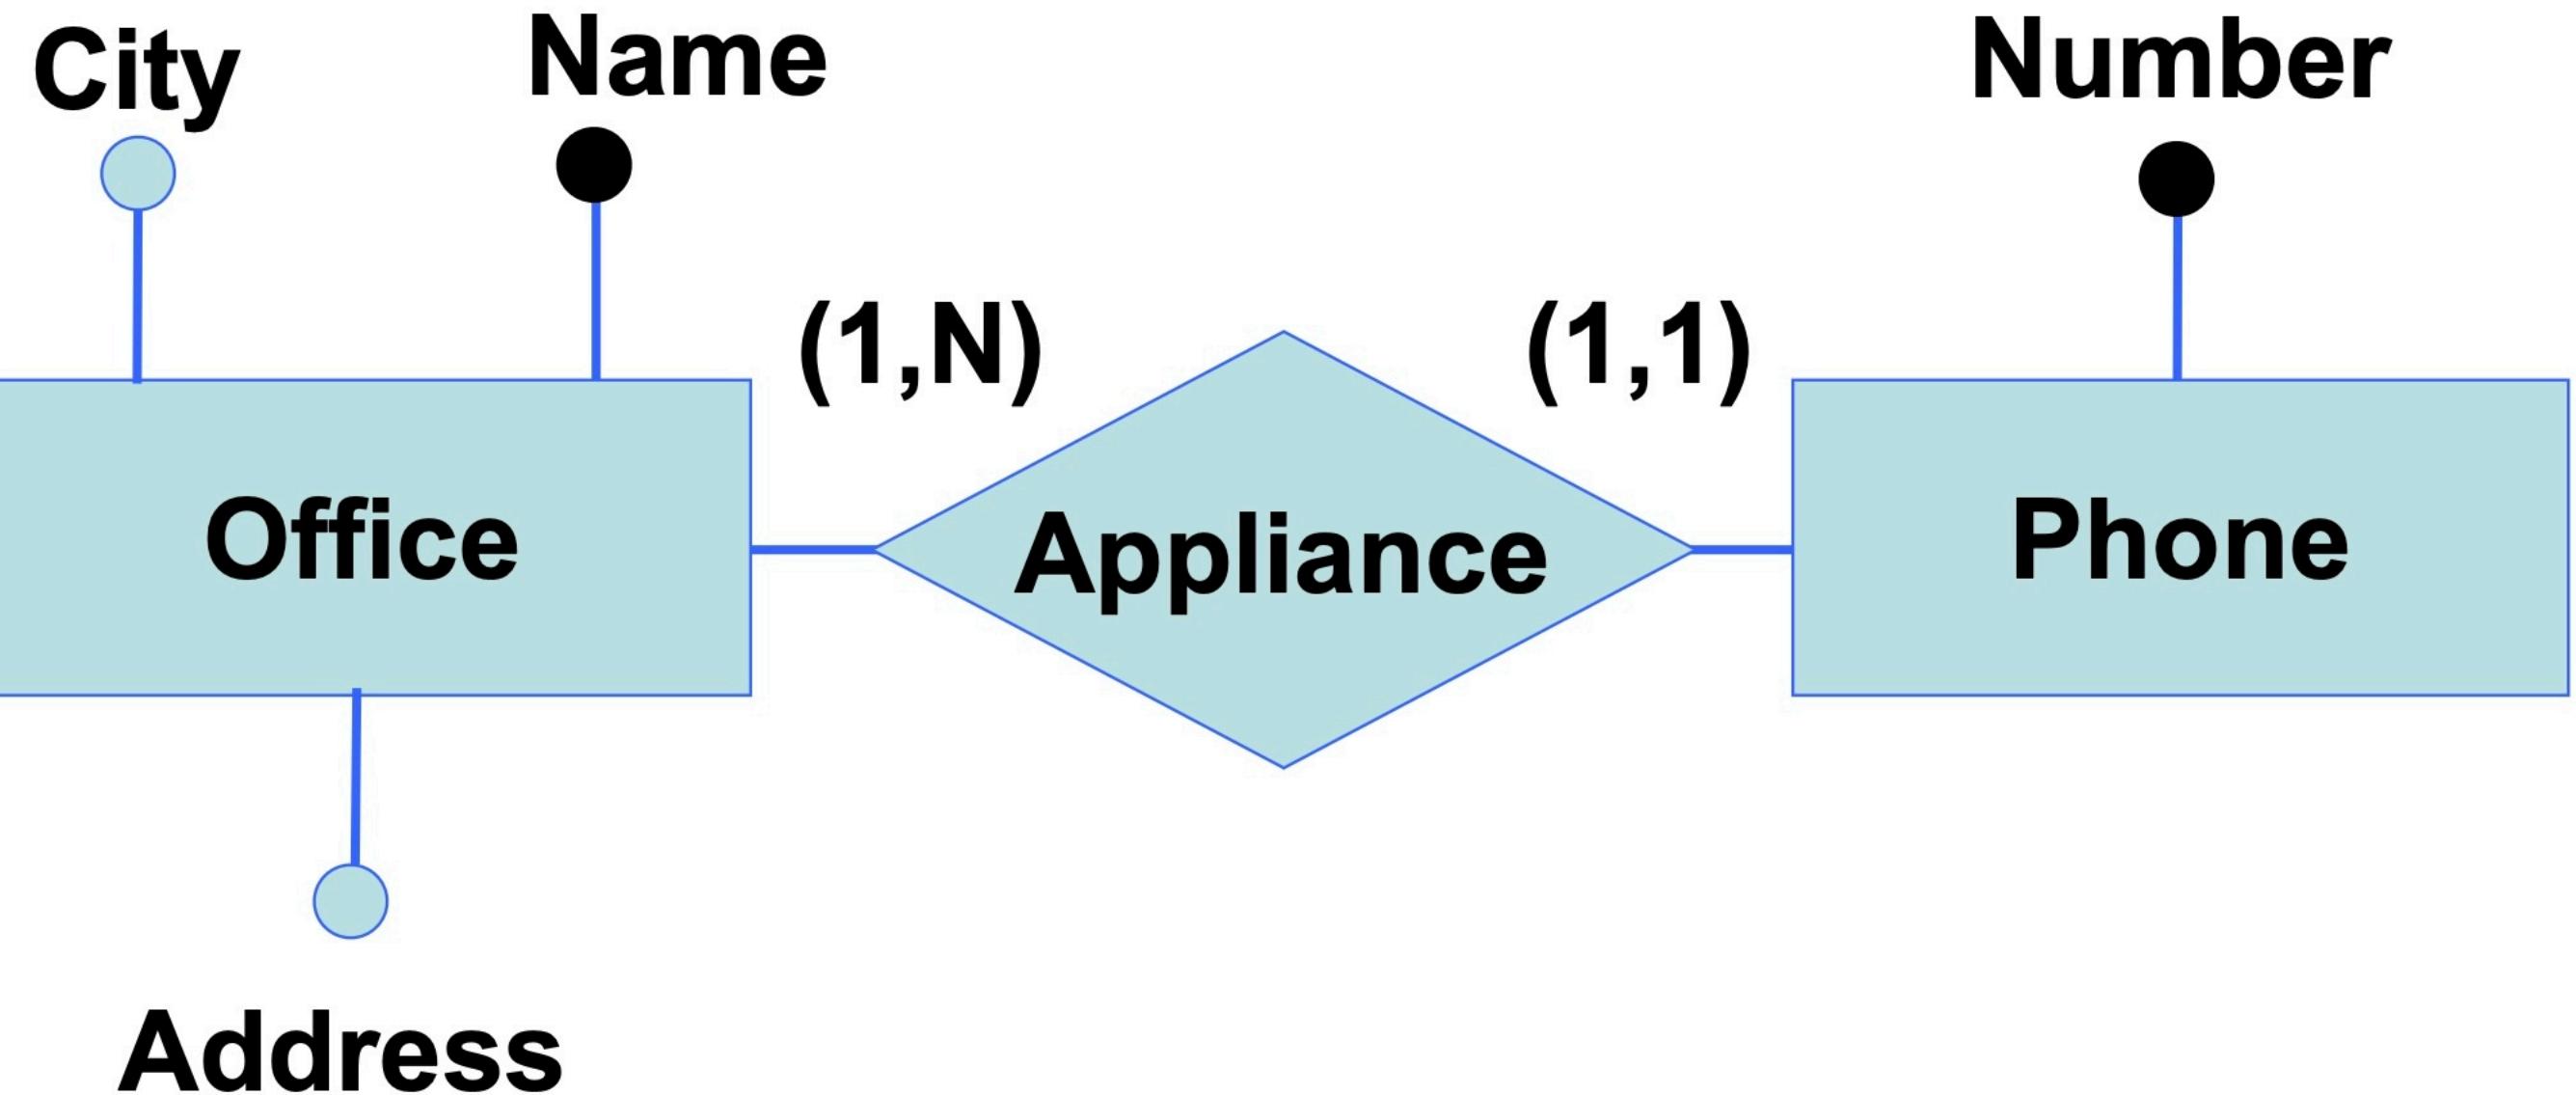
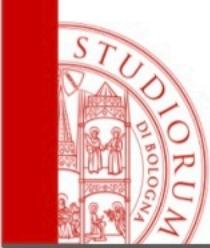
Restructuring, main cases

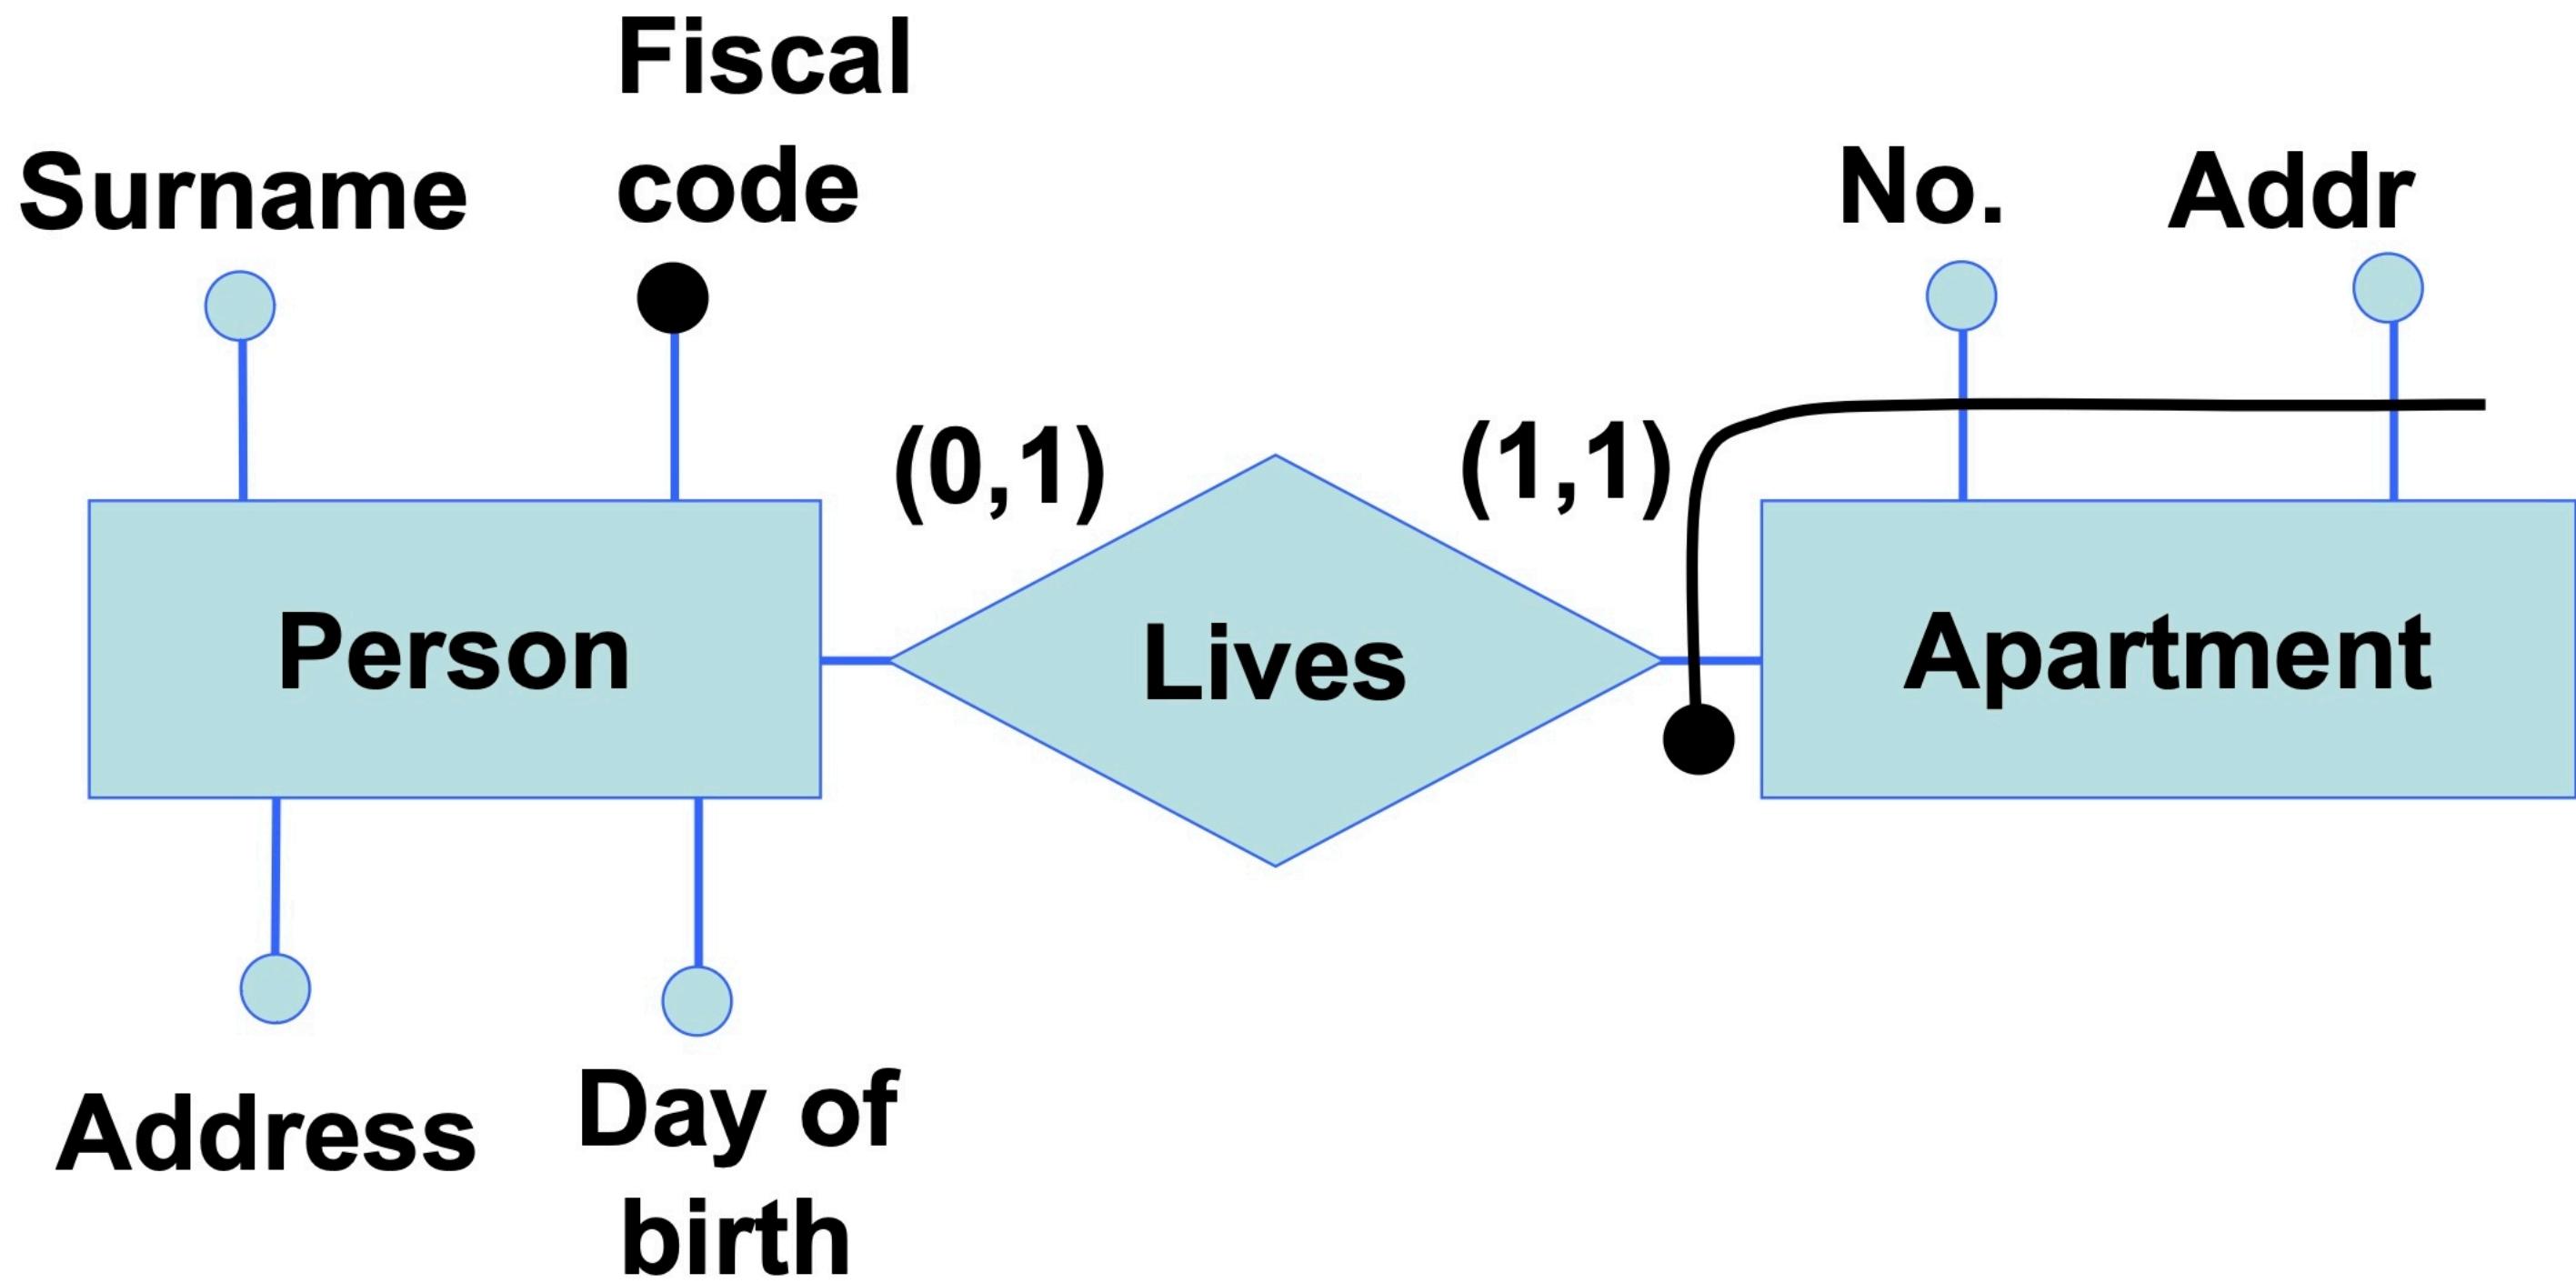
- Entity vertical partitioning
- Relationship horizontal partitioning
- Restructuring multi-valued attributes
- Grouping of entities/relationships

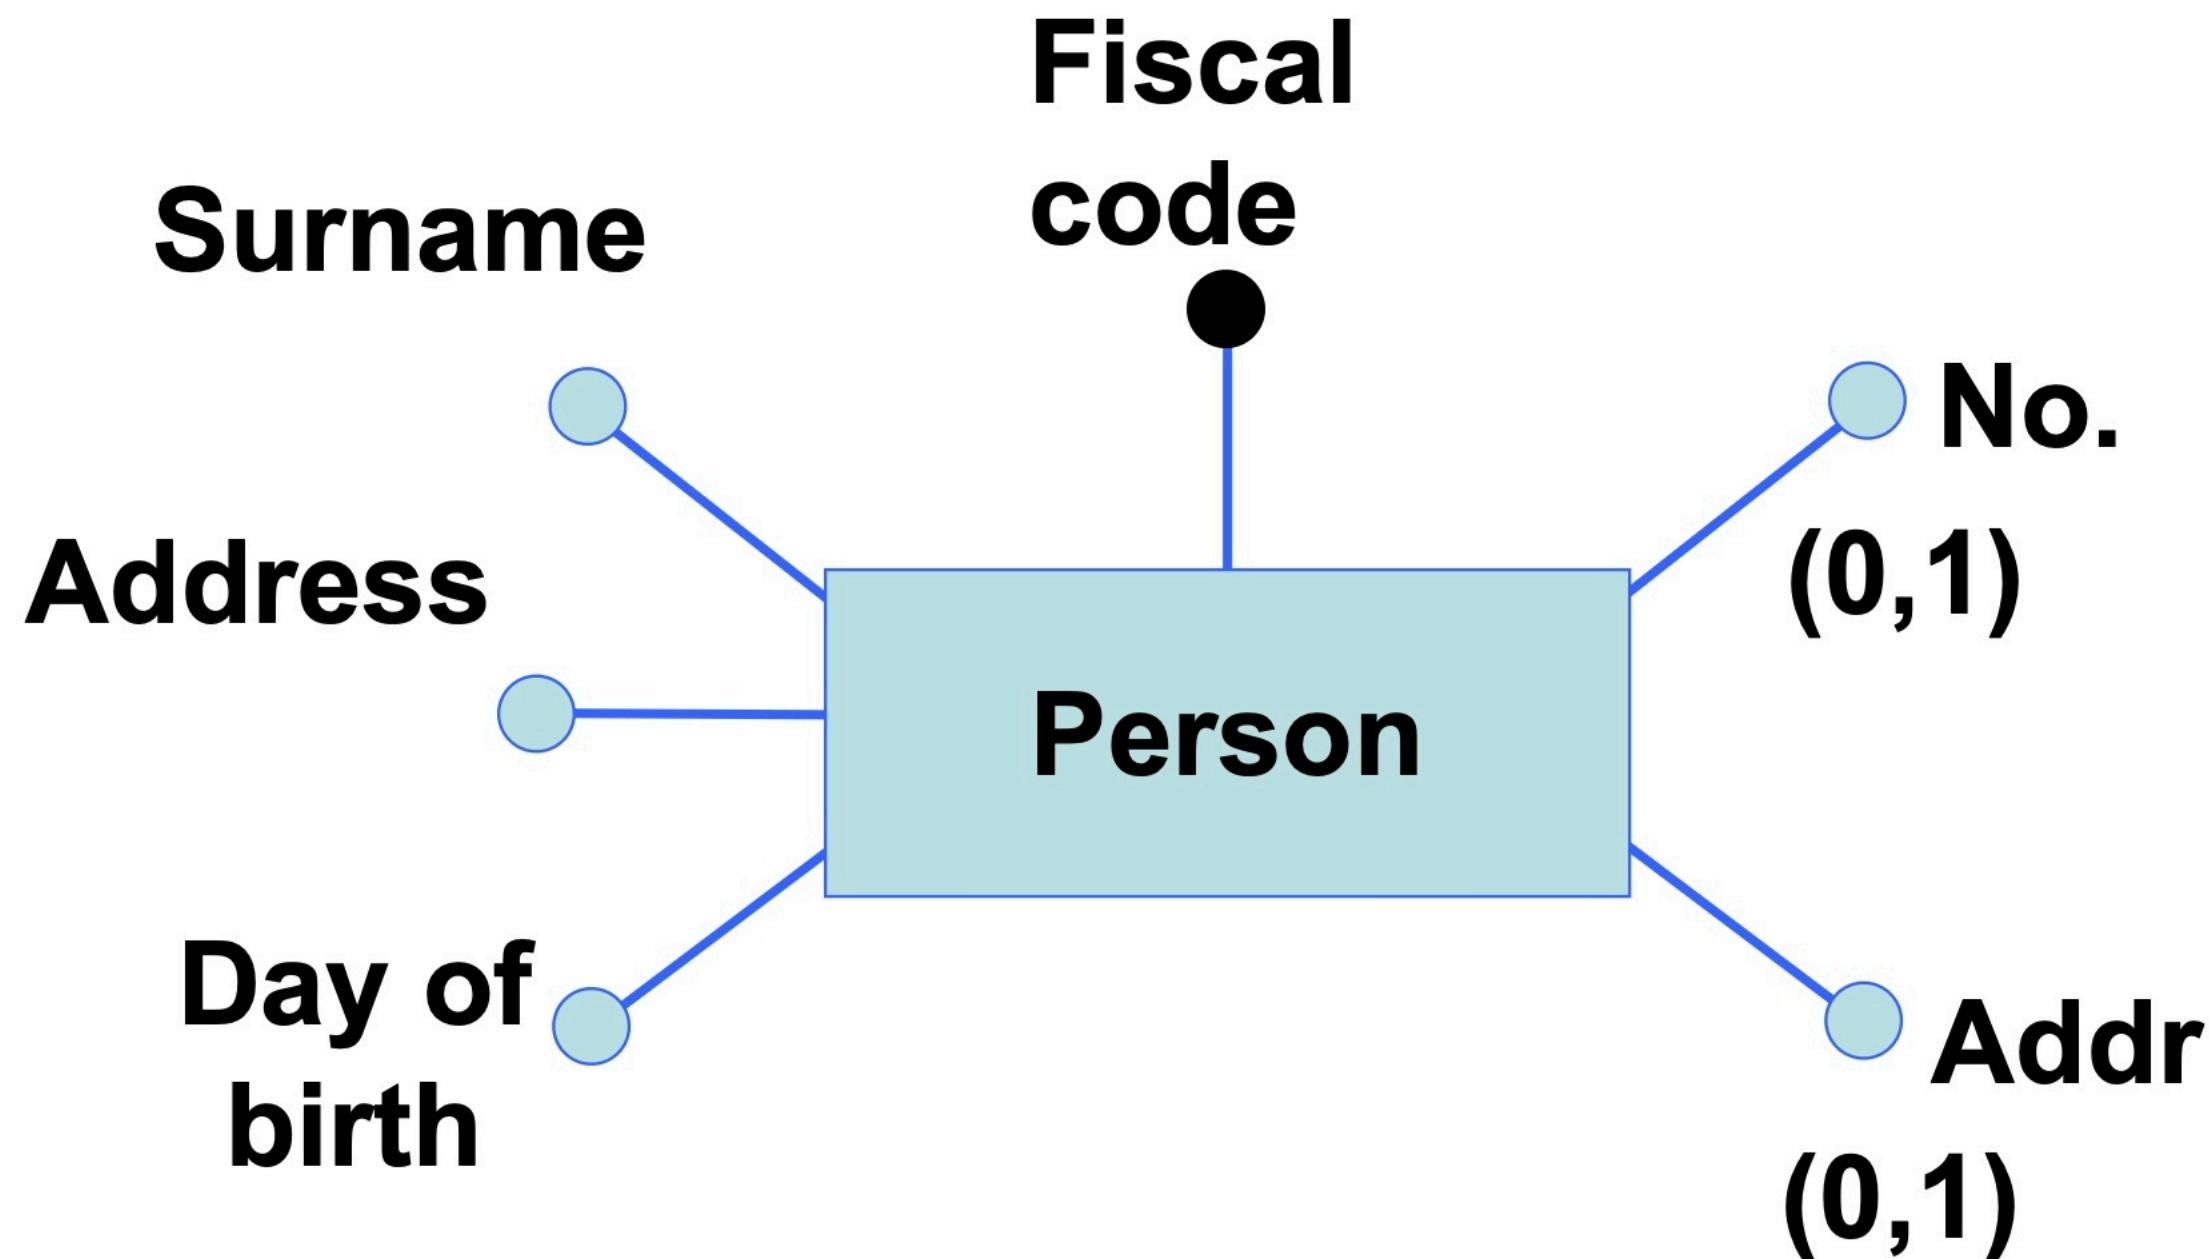


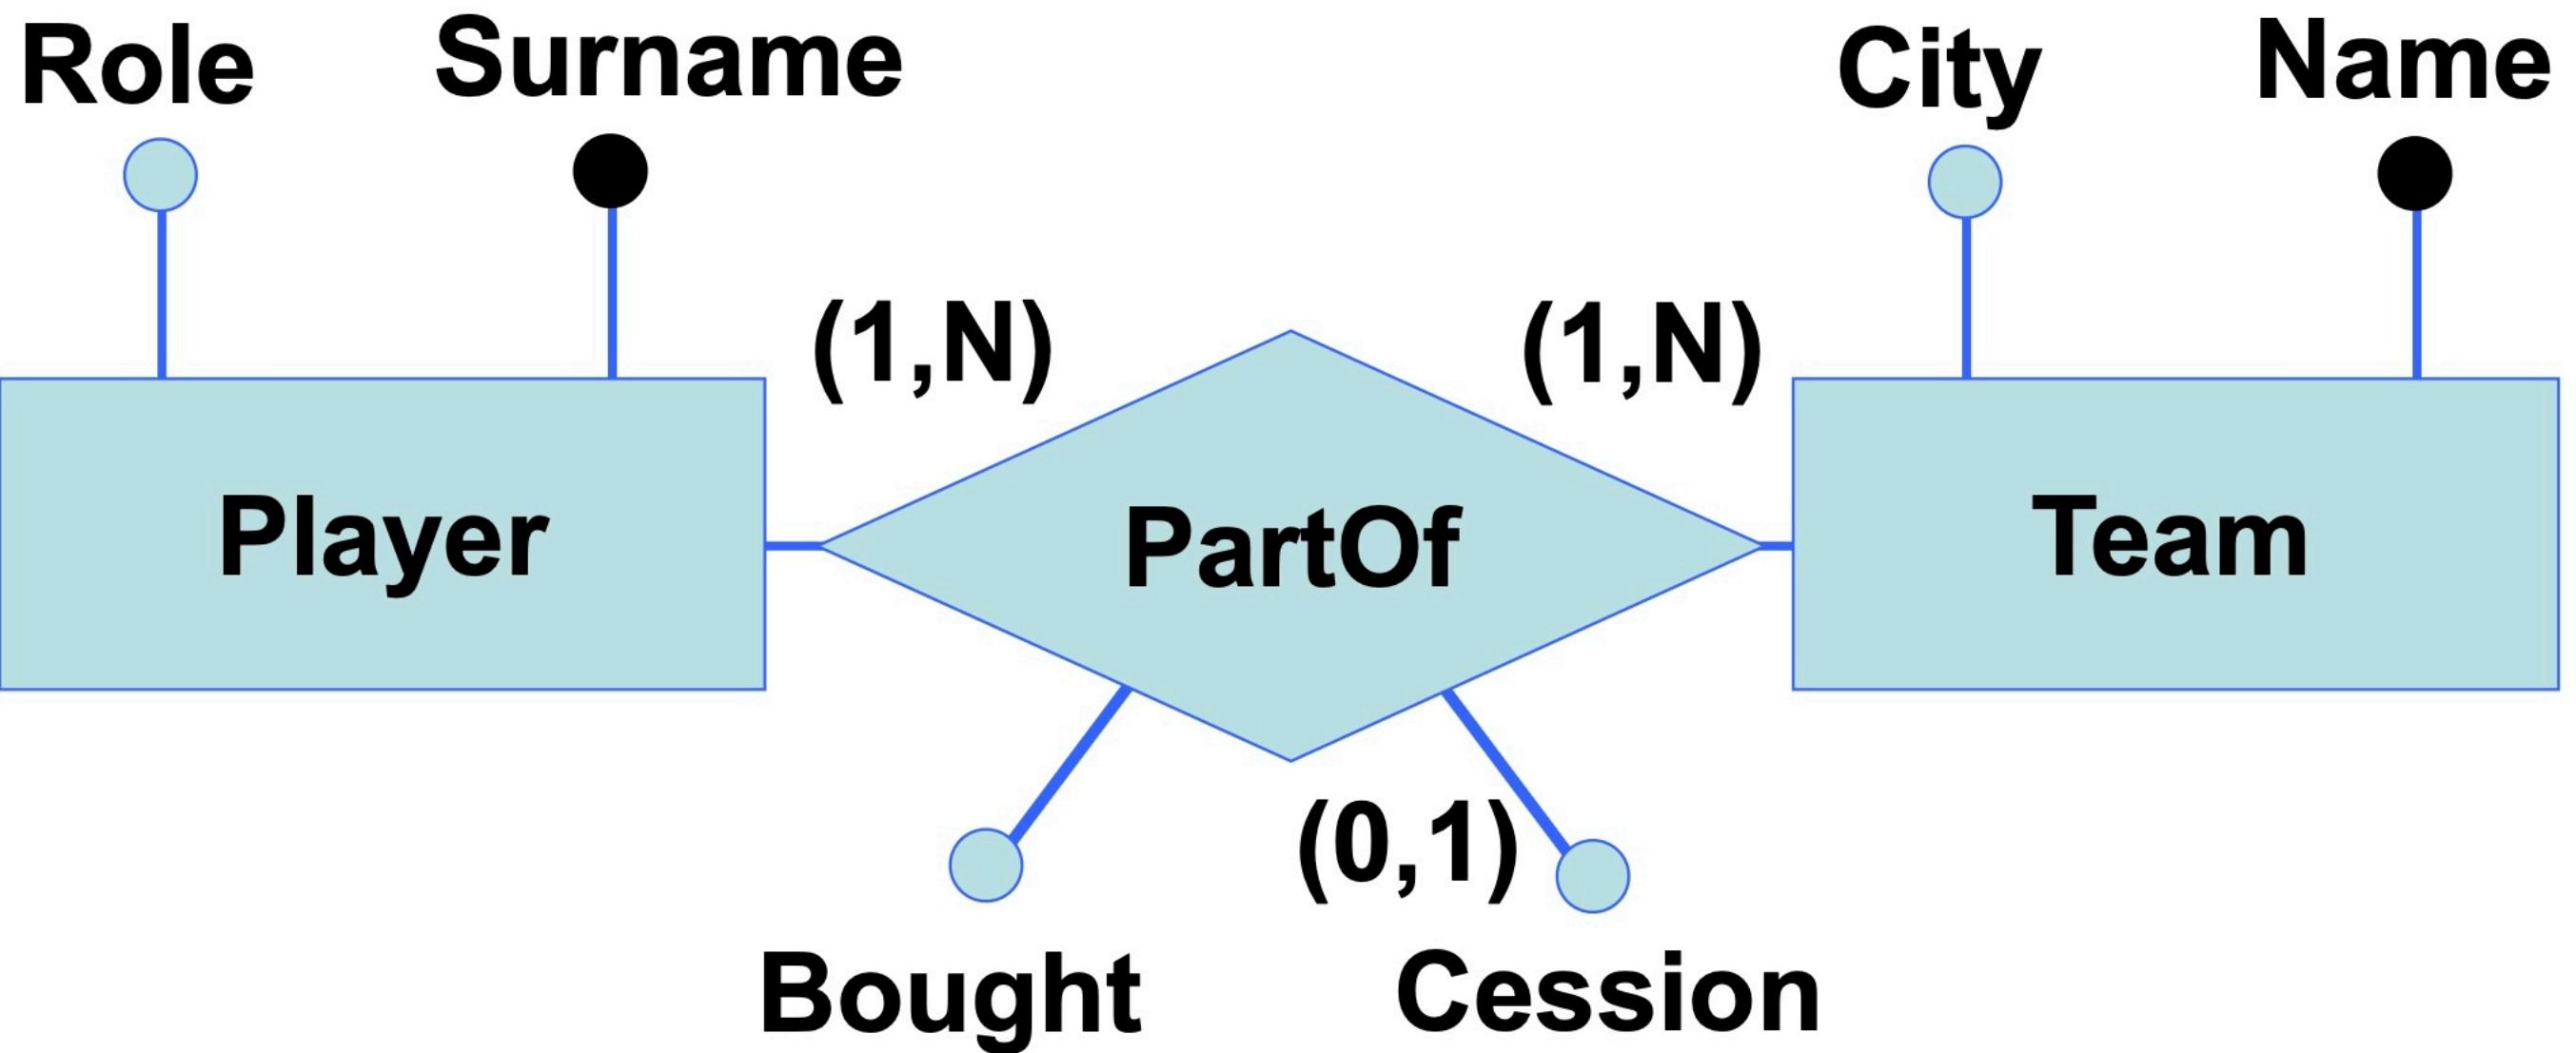
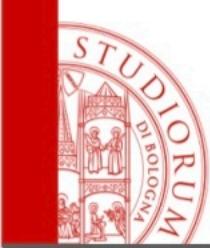


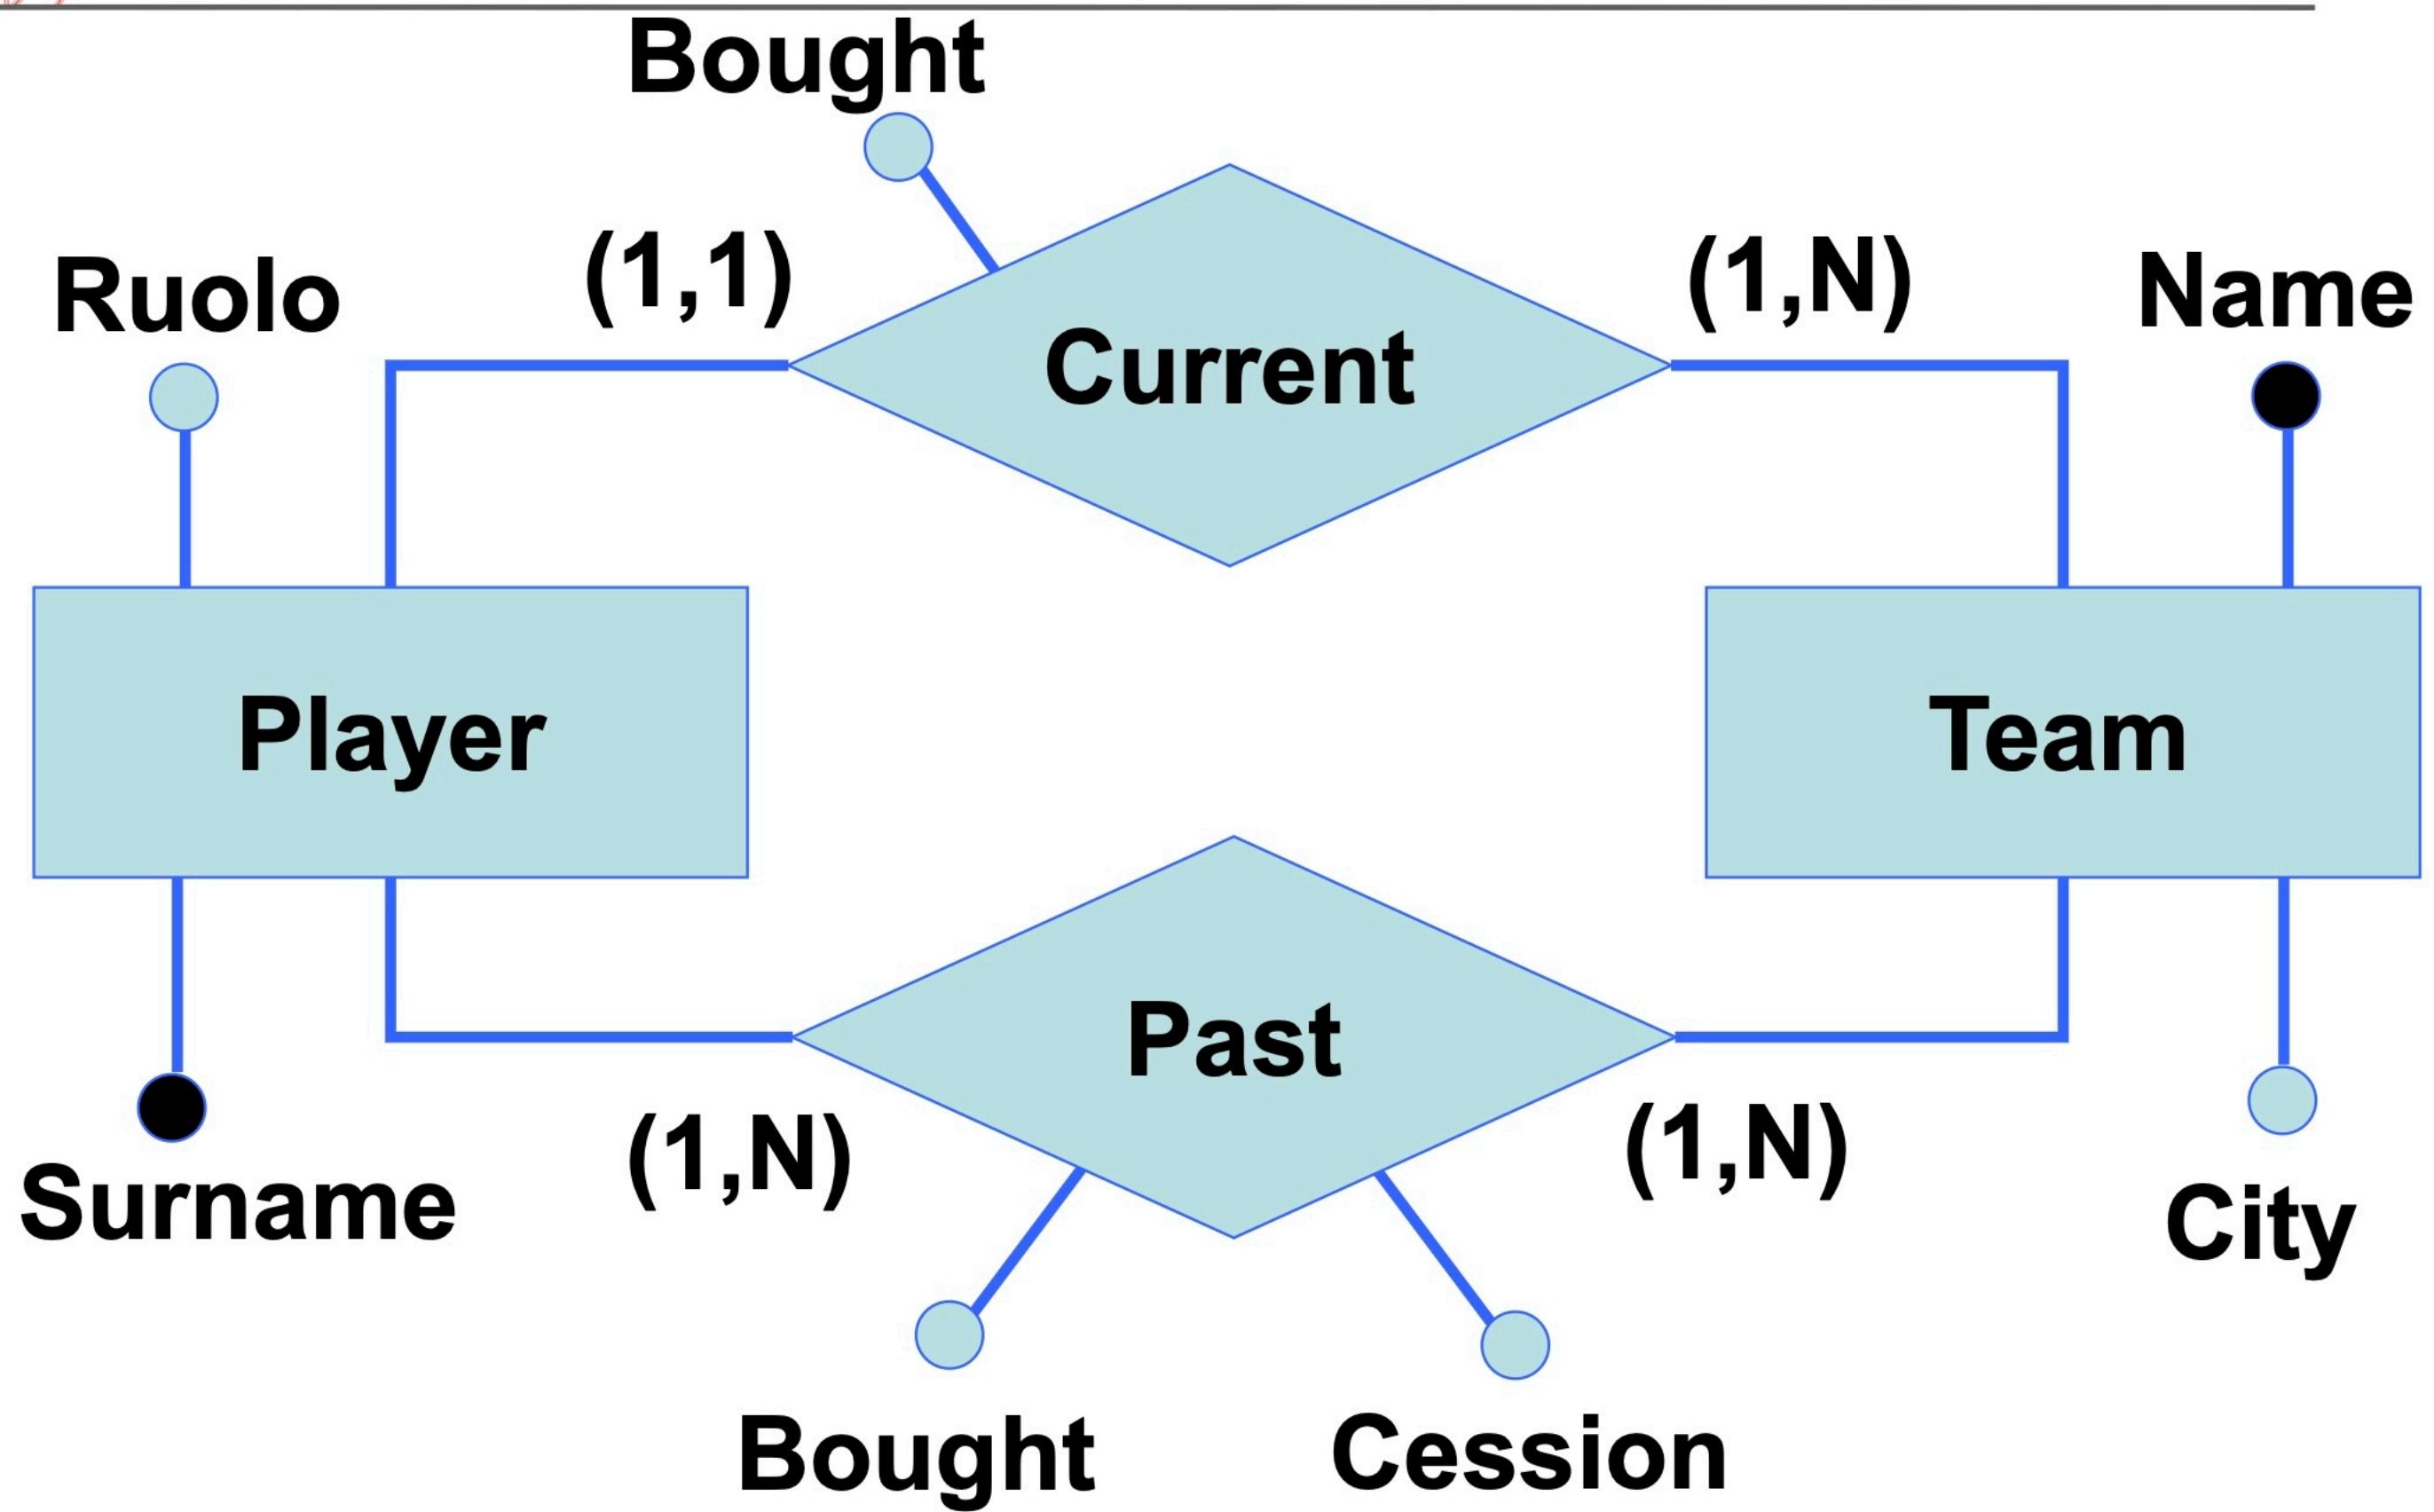
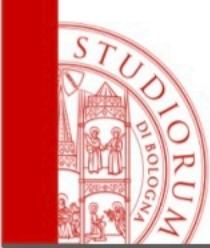








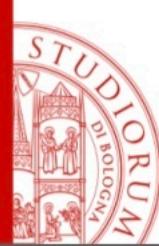






Restructuring activities

- Redundancies analysis
- Generalizations deletion
- Partitioning/grouping of entities and relationships
- Identifying the primary keys



Identifying primary keys

- A mandatory operation for the translation into a relational model
- Criteria
 - Compulsory information
 - Simplicity
 - Used within the most frequent/relevant operation



Identifying primary keys

What if none of the aforementioned conditions are met?

New attributes are introduced using specifically generated codes

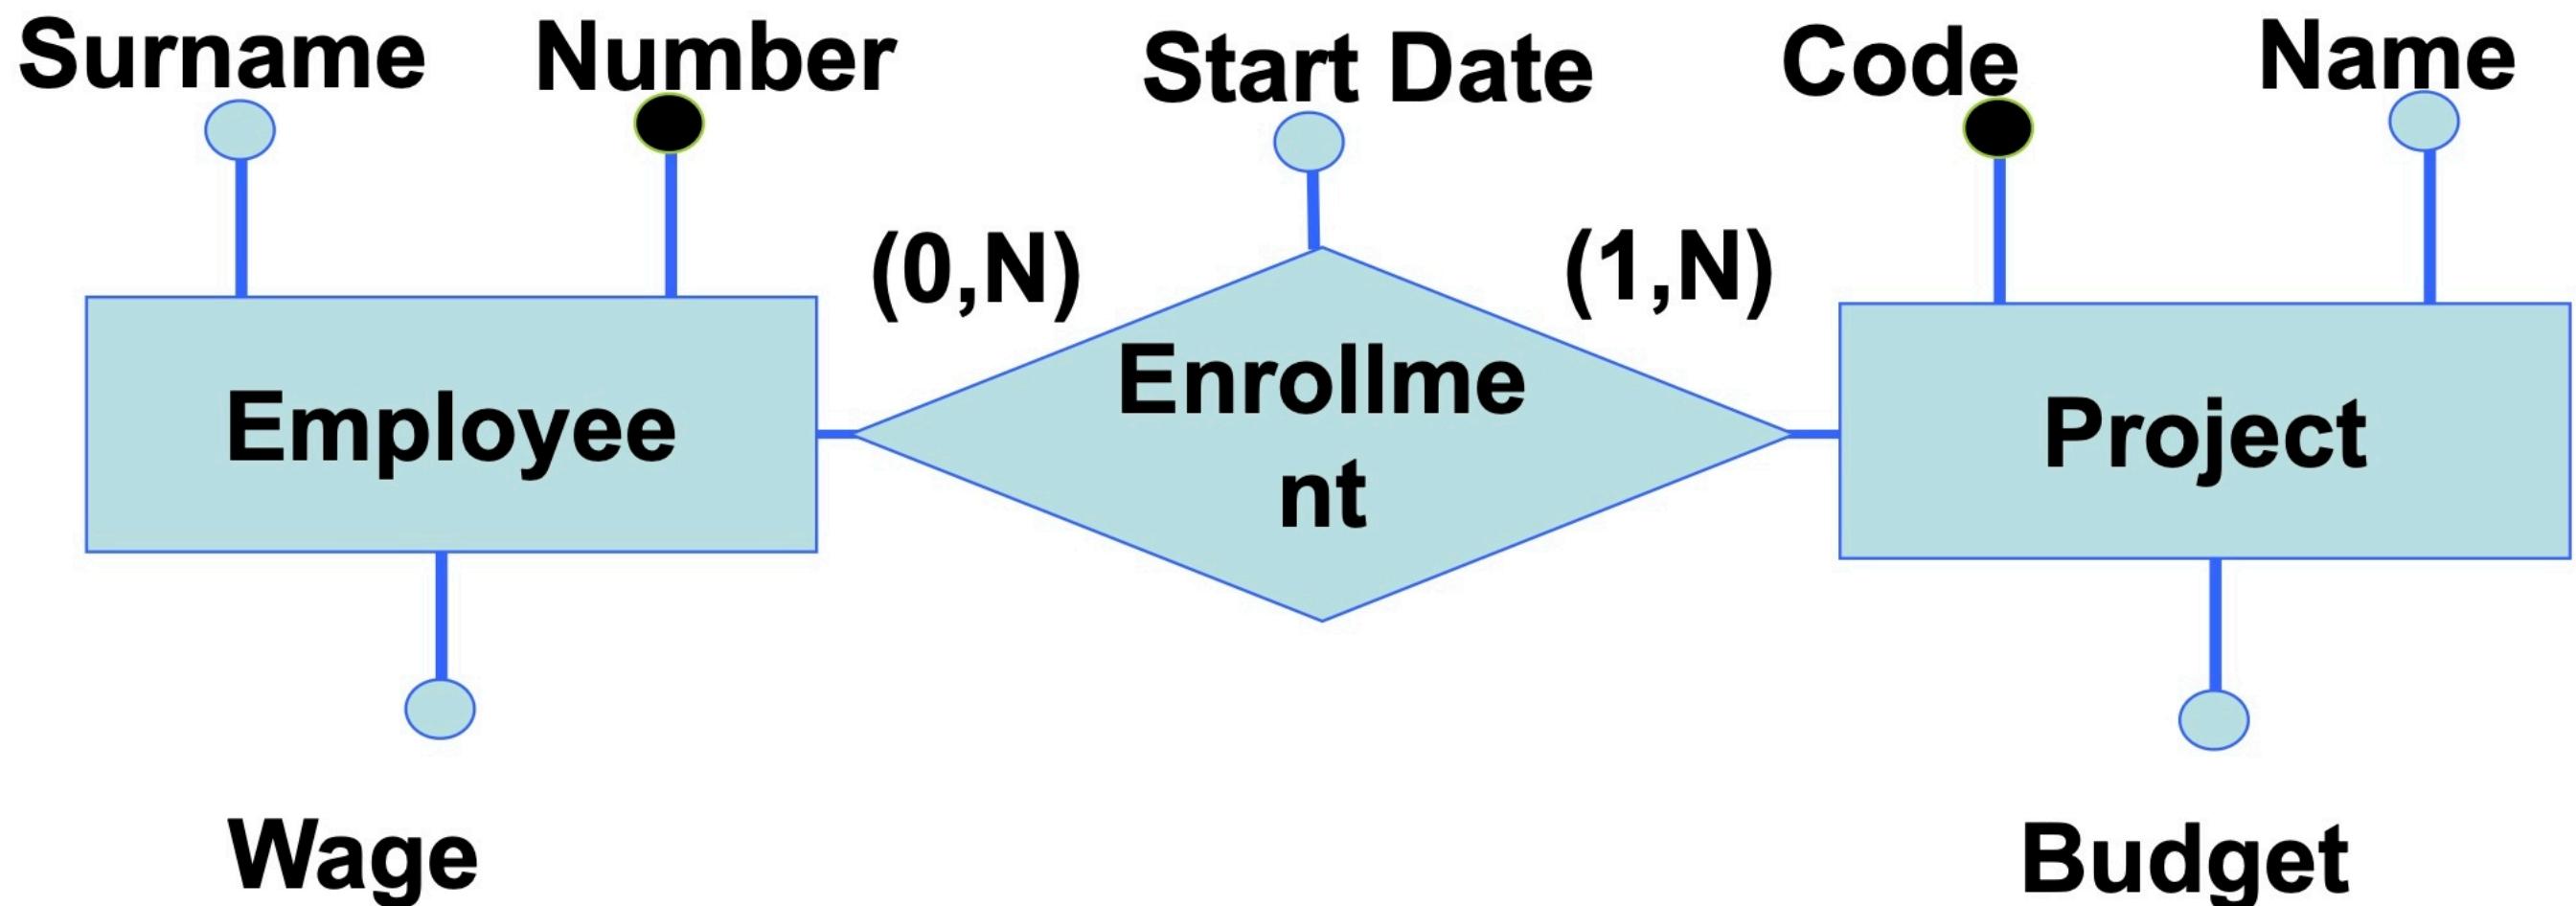


Translation to the relational model

■ Rules of thumb:

- Entities become “tables”, whose schema corresponds to the entities’ attributes
- Relationships become “tables”, their schema correspond to the entities’ attributes, plus the foreign identifiers for the involved entities

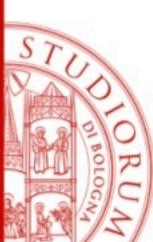
Many-to-Many relationships



Employee(Number, Surname, Wage)

Project(Code, Name, Budget)

Enrollment(Number, Code, StartDate)



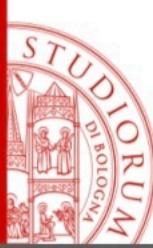
Many-to-Many relationships

Employee(Number, Surname, Wage)

Project(Code, Name, Budget)

Enrollment(Number, Code, StartDate)

- Referential Integrity Constraint between:
 - Number in Enrollment and Employee's key
 - Code in Enrollment and Project's key



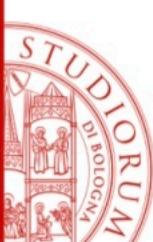
Foreign Keys: use more expressive names

Employee(Number, Surname, Stipendio)

Project(Code, Name, Budget)

Enrollment(Number, Code, StartDate)

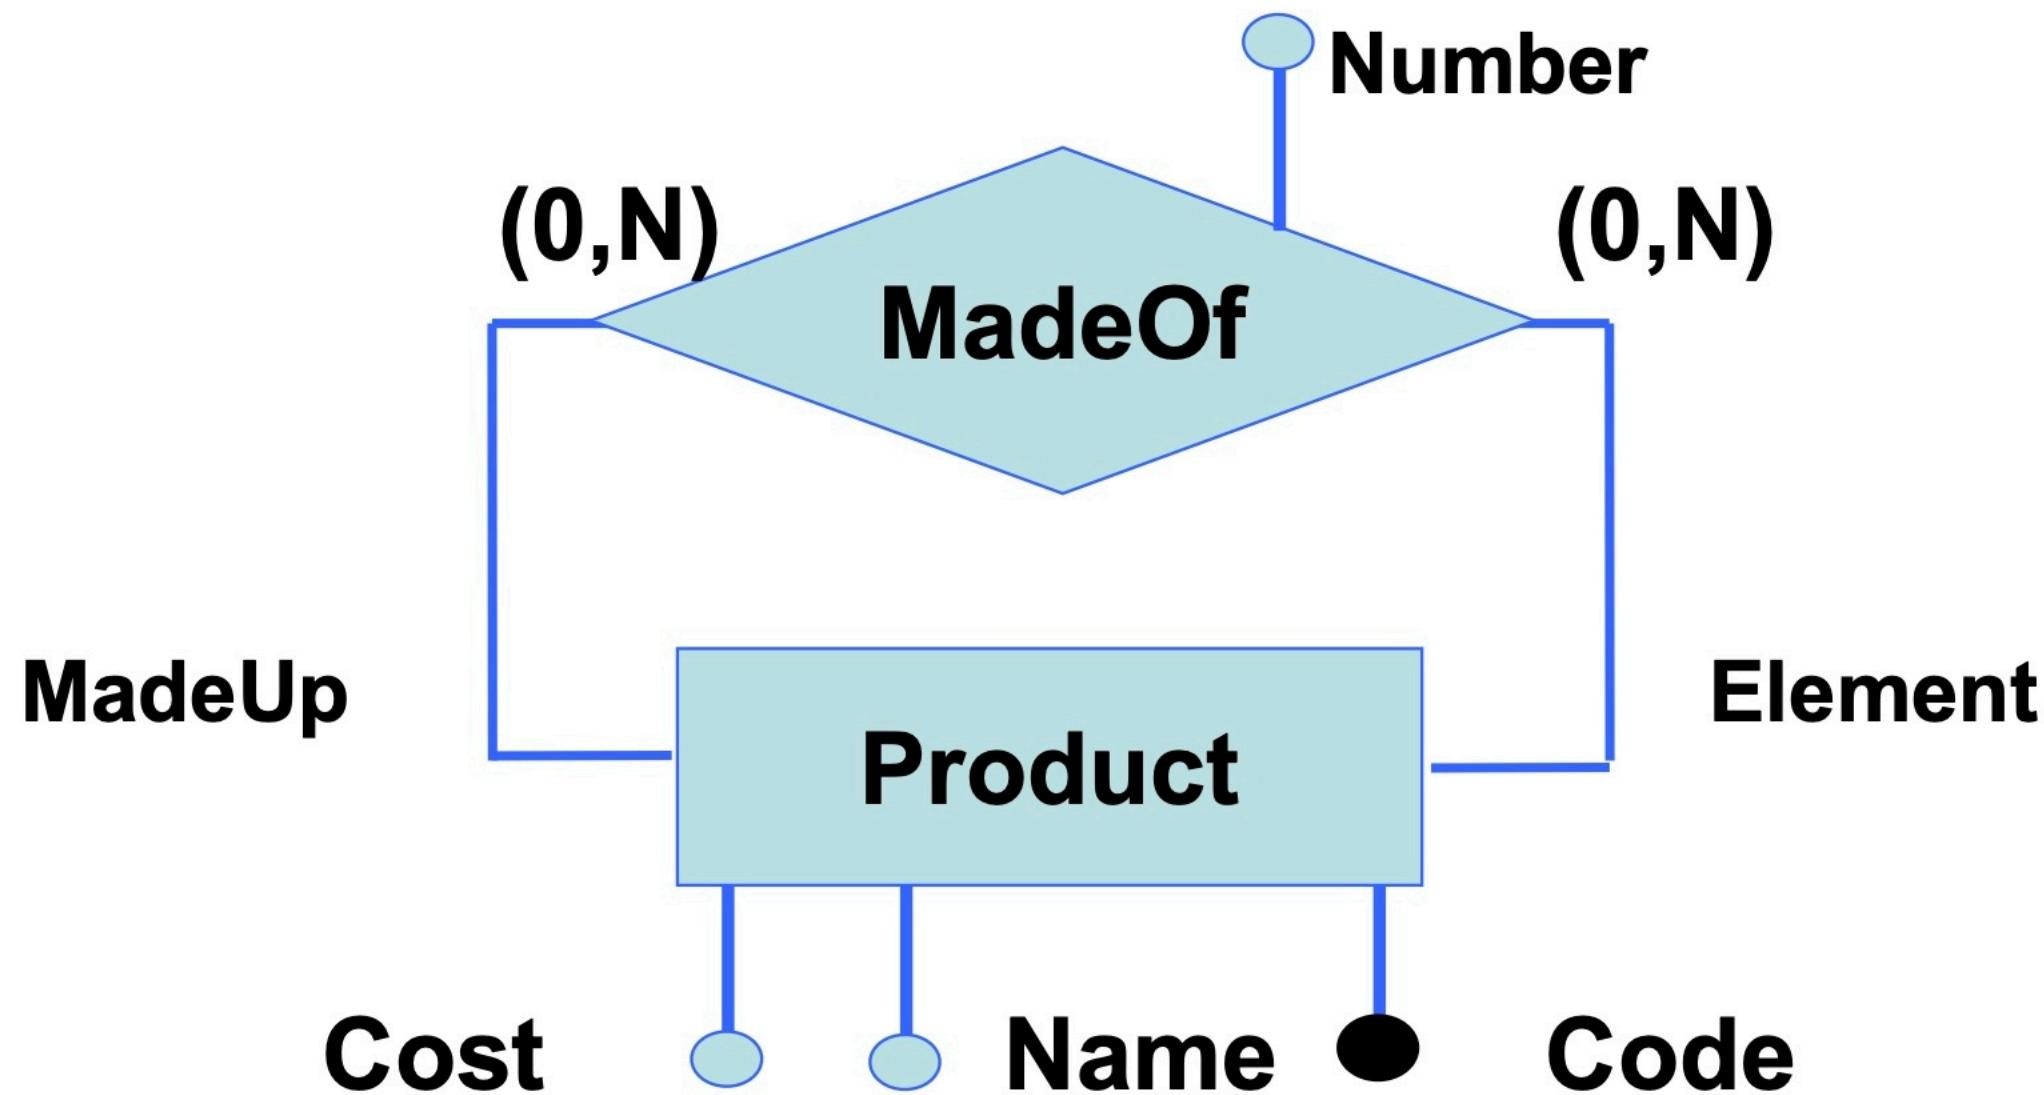
Enrollment(Employee, Project, StartDate)



Please note

- Such translation does not keep into account the many-to-many relationships' minimal cardinality (even if we could use very uncommon and complex CHECKs)

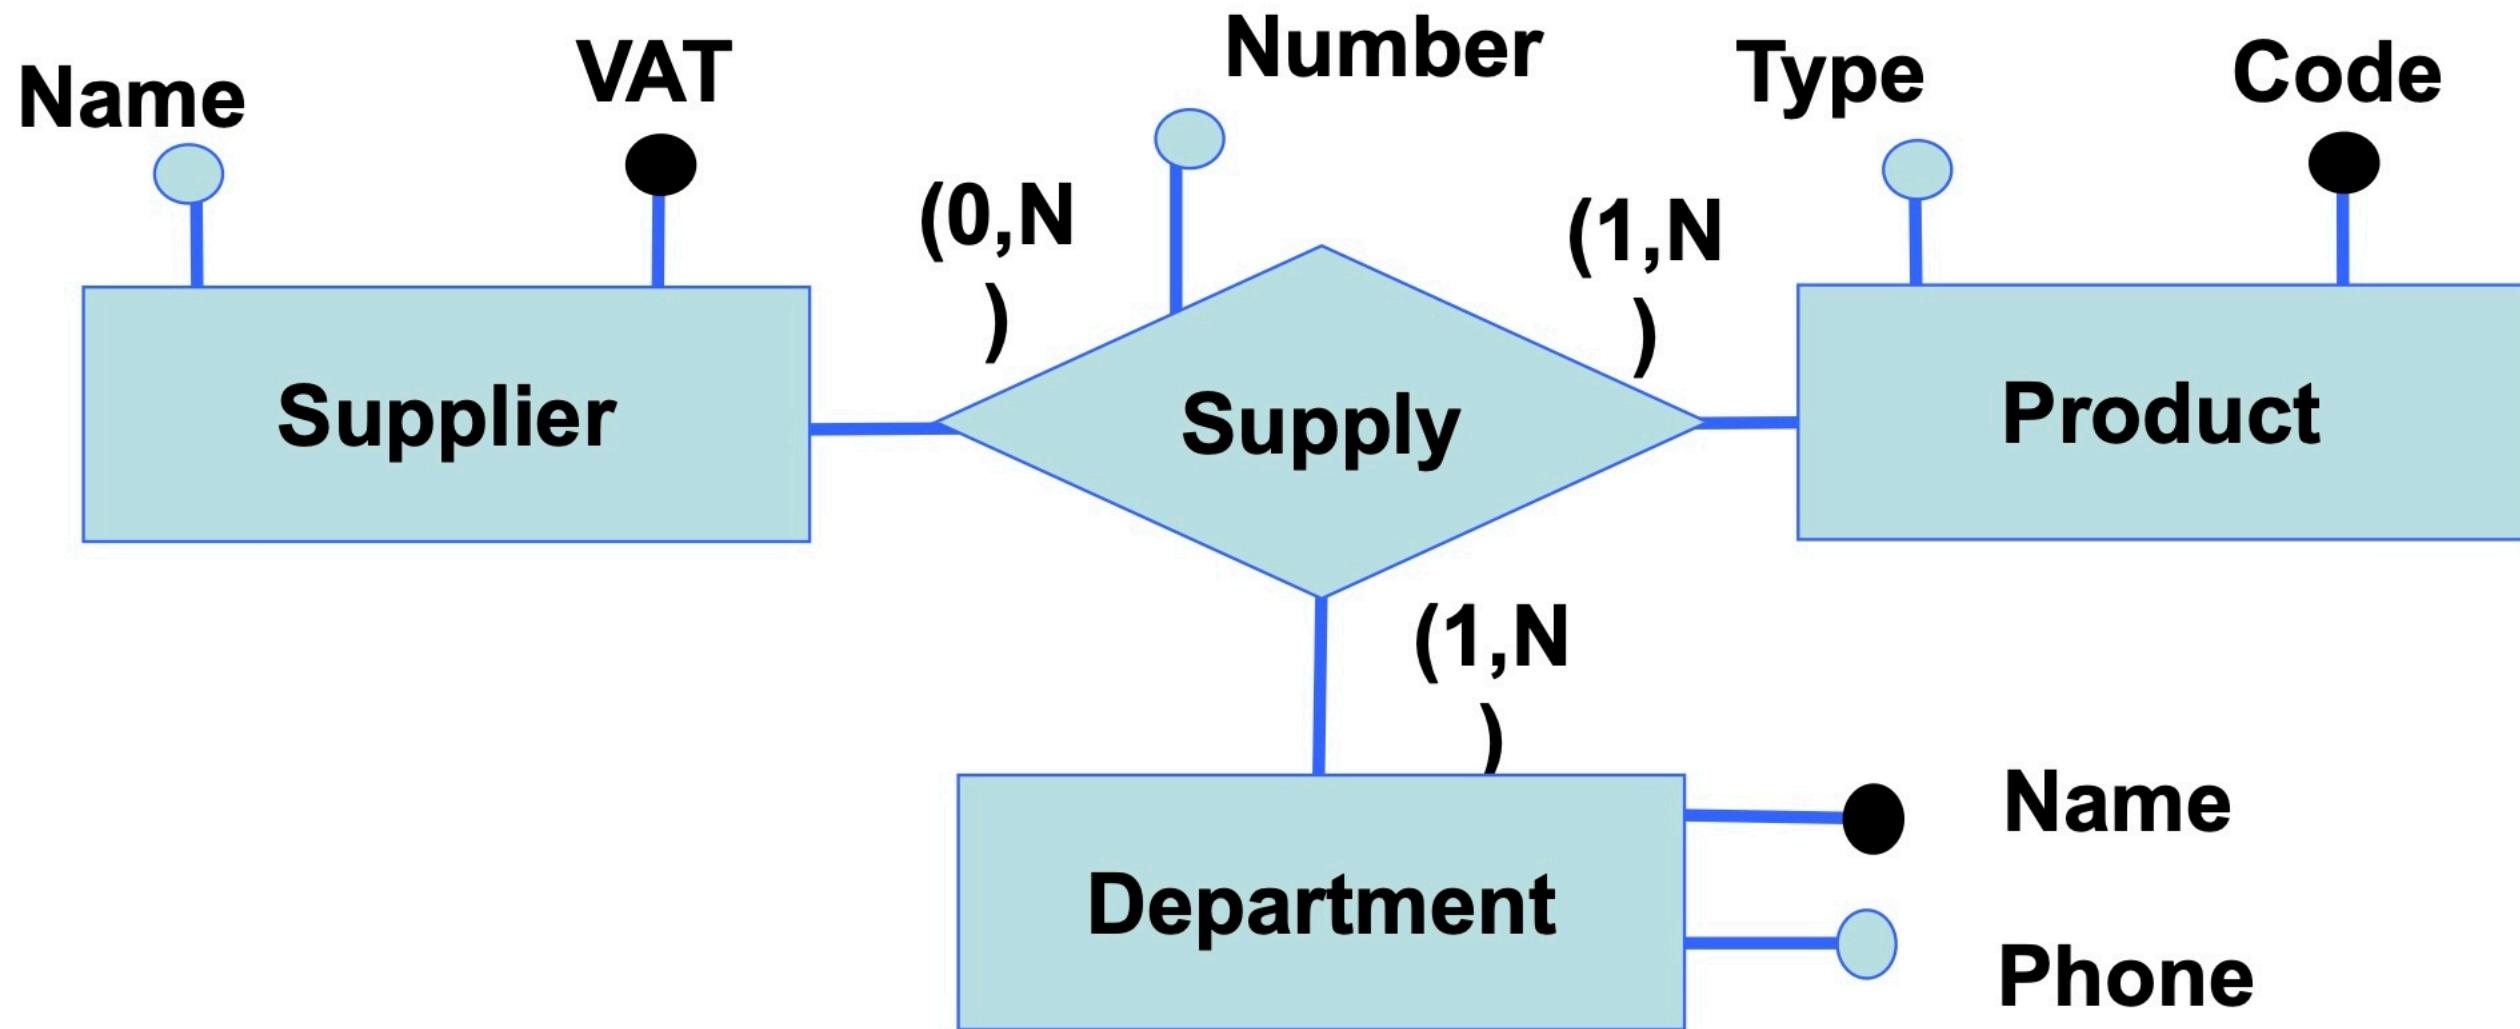
Recursive relationships



Product(Code, Name, Cost)

MadeOf(MadeUp, Element, Number)

N-ary Relationships



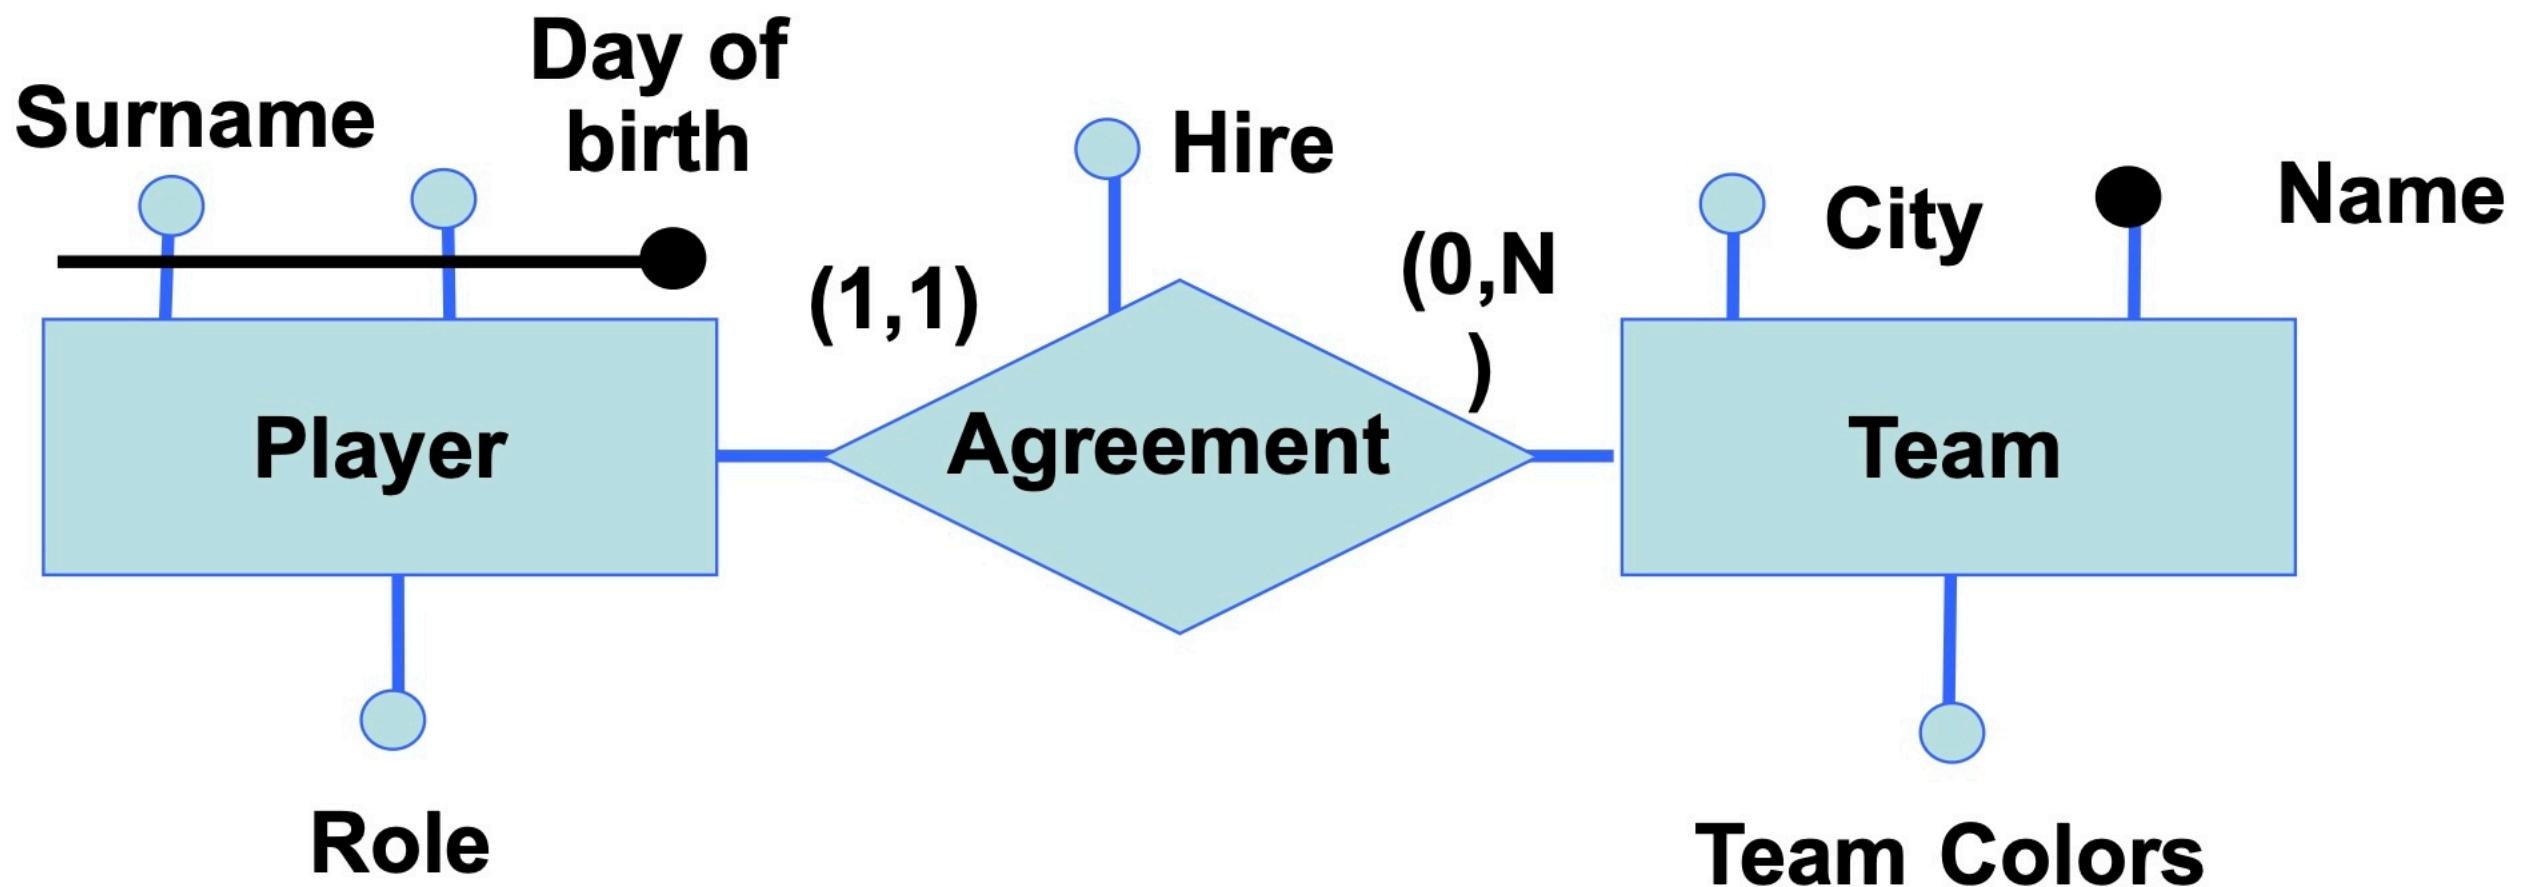
Supplier(VAT, Name)

Product(Code, Type)

Department(Name, Phone)

Supply(Supplier, Product, Department, Number)

One-to-Many Relationship



Player(Surname, DayOfBirth, Role)

Agreement(SurnameP, BDayP, Team, Hire)

Team(Name, City, TeamColors)

- correct?

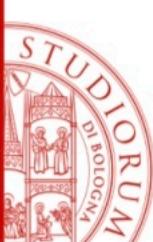


A less redundant solution

Player(Surname, DayOfBirth, Role)
Agreement(SurnameP, BDayP, Team, Hire)
Team(Name, City, TeamColors)

Player(Surname, DayOfBirth, Team, Role, Hire)
Team(Name, City, TeamColors)

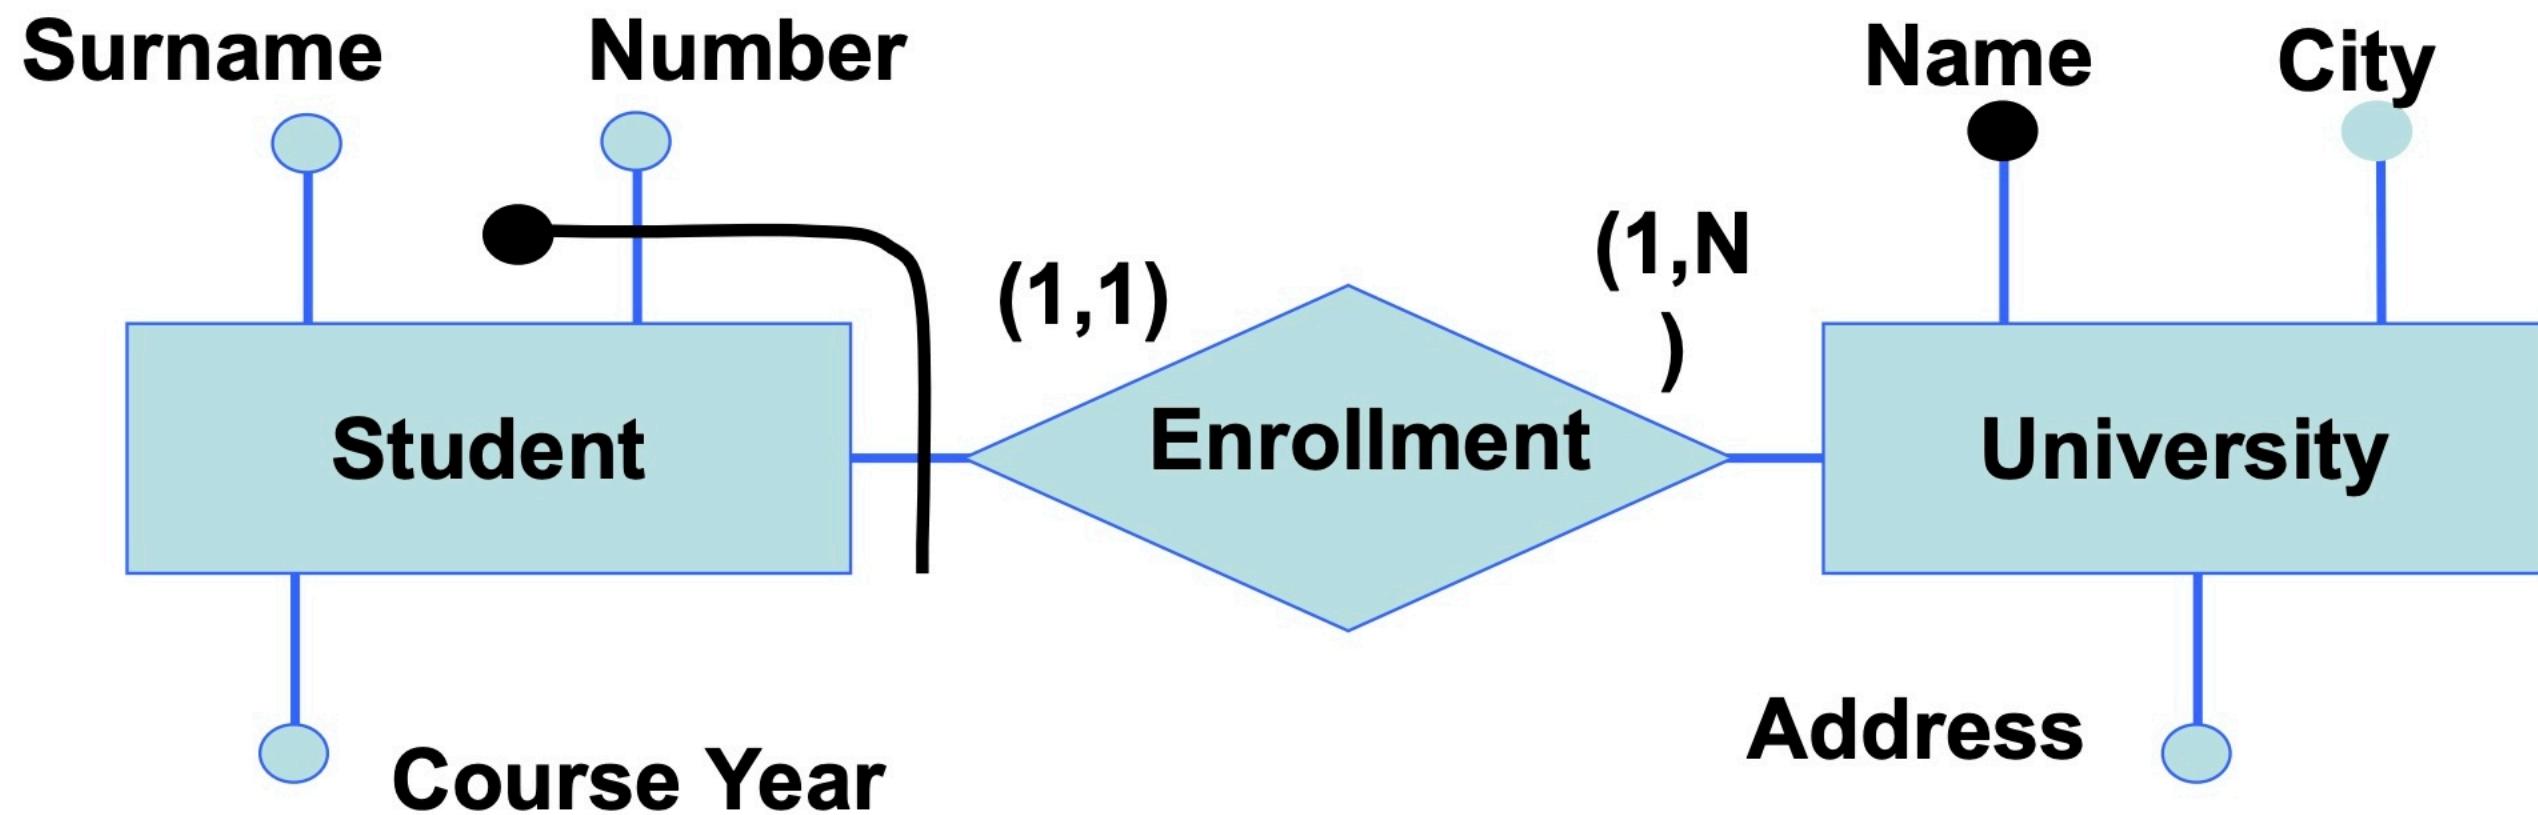
- **Referential Integrity Constraint between Team in Player and Team's key**
- If the relationship's minimal cardinality is 0, then Team in Player must allow NULL values



Please Note

- Such translation could represent the case when 0 is the minimal cardinality and 1 is the maximum one:
 - 0 : NULLs allowed
 - 1 : NULLs **not** allowed

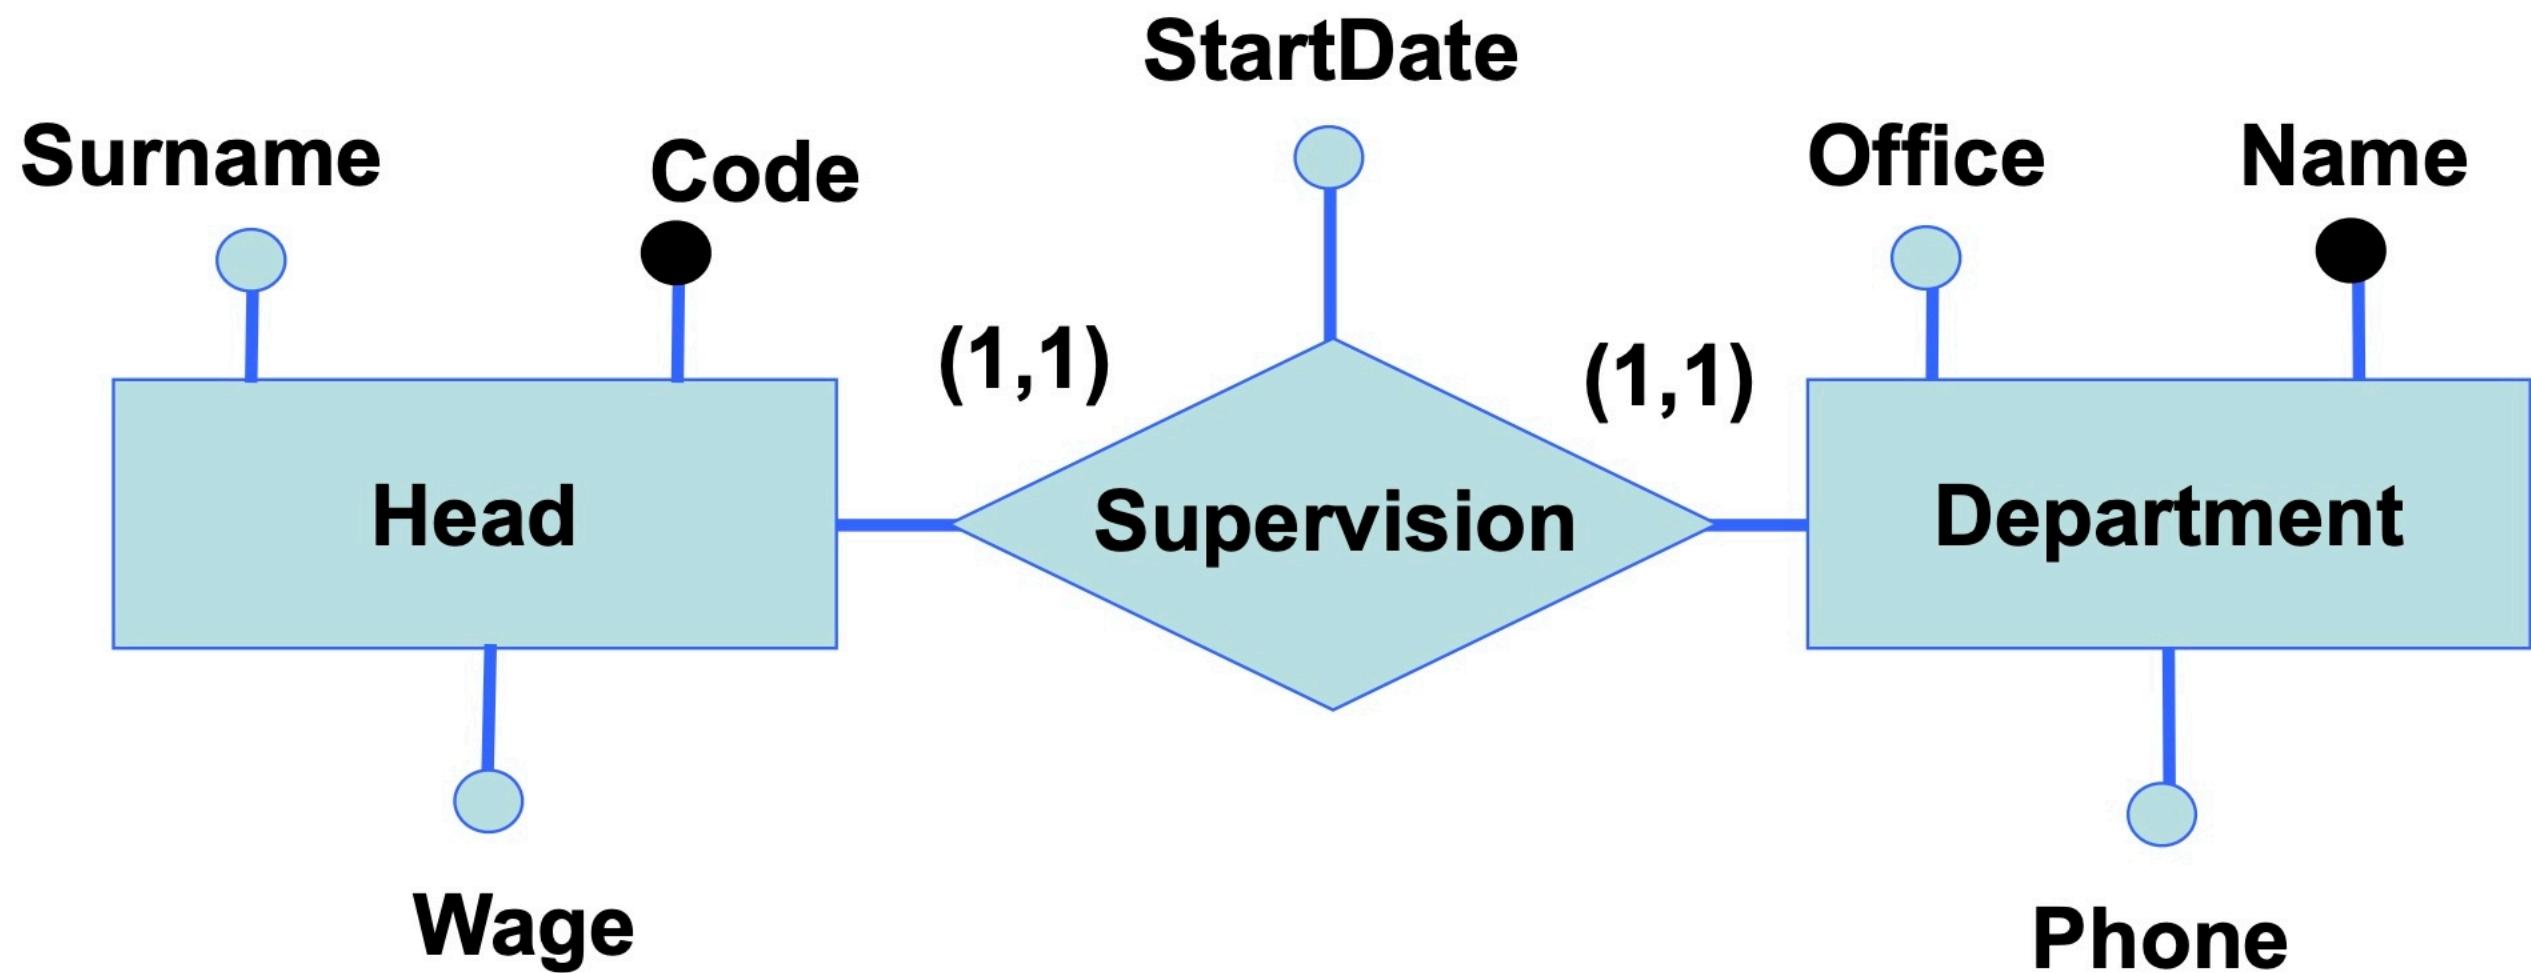
Entity with external identifier



Student(Number, University, Surname, CourseYear)
University(Name, City, Addr)

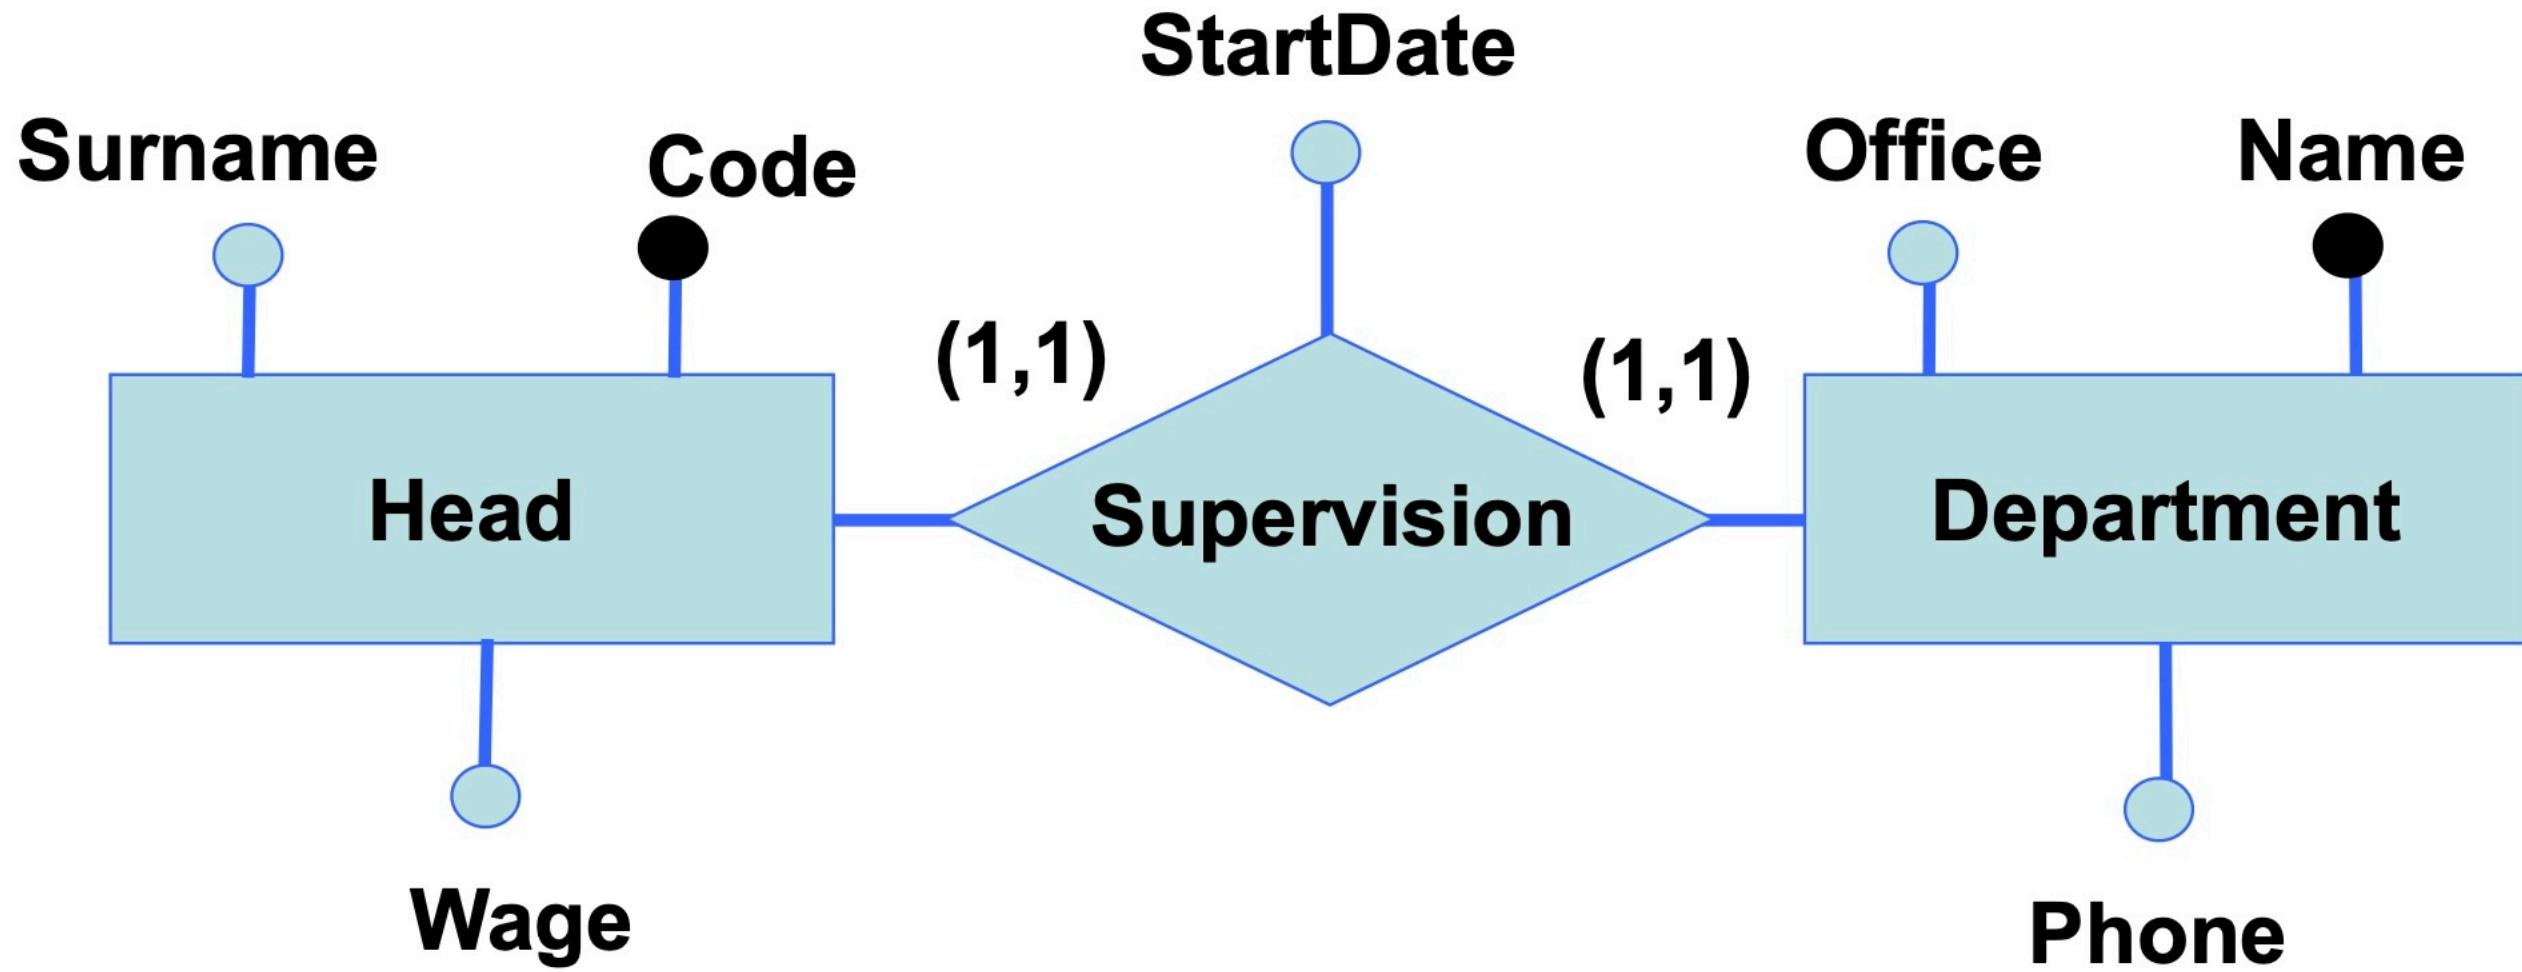
- Constraint: each student is enrolled to only one university

One-to-One Relationship



- Different options:
 - Merge one entity with a relation (on one side or the other)
 - Merge everything together

A preferred solution

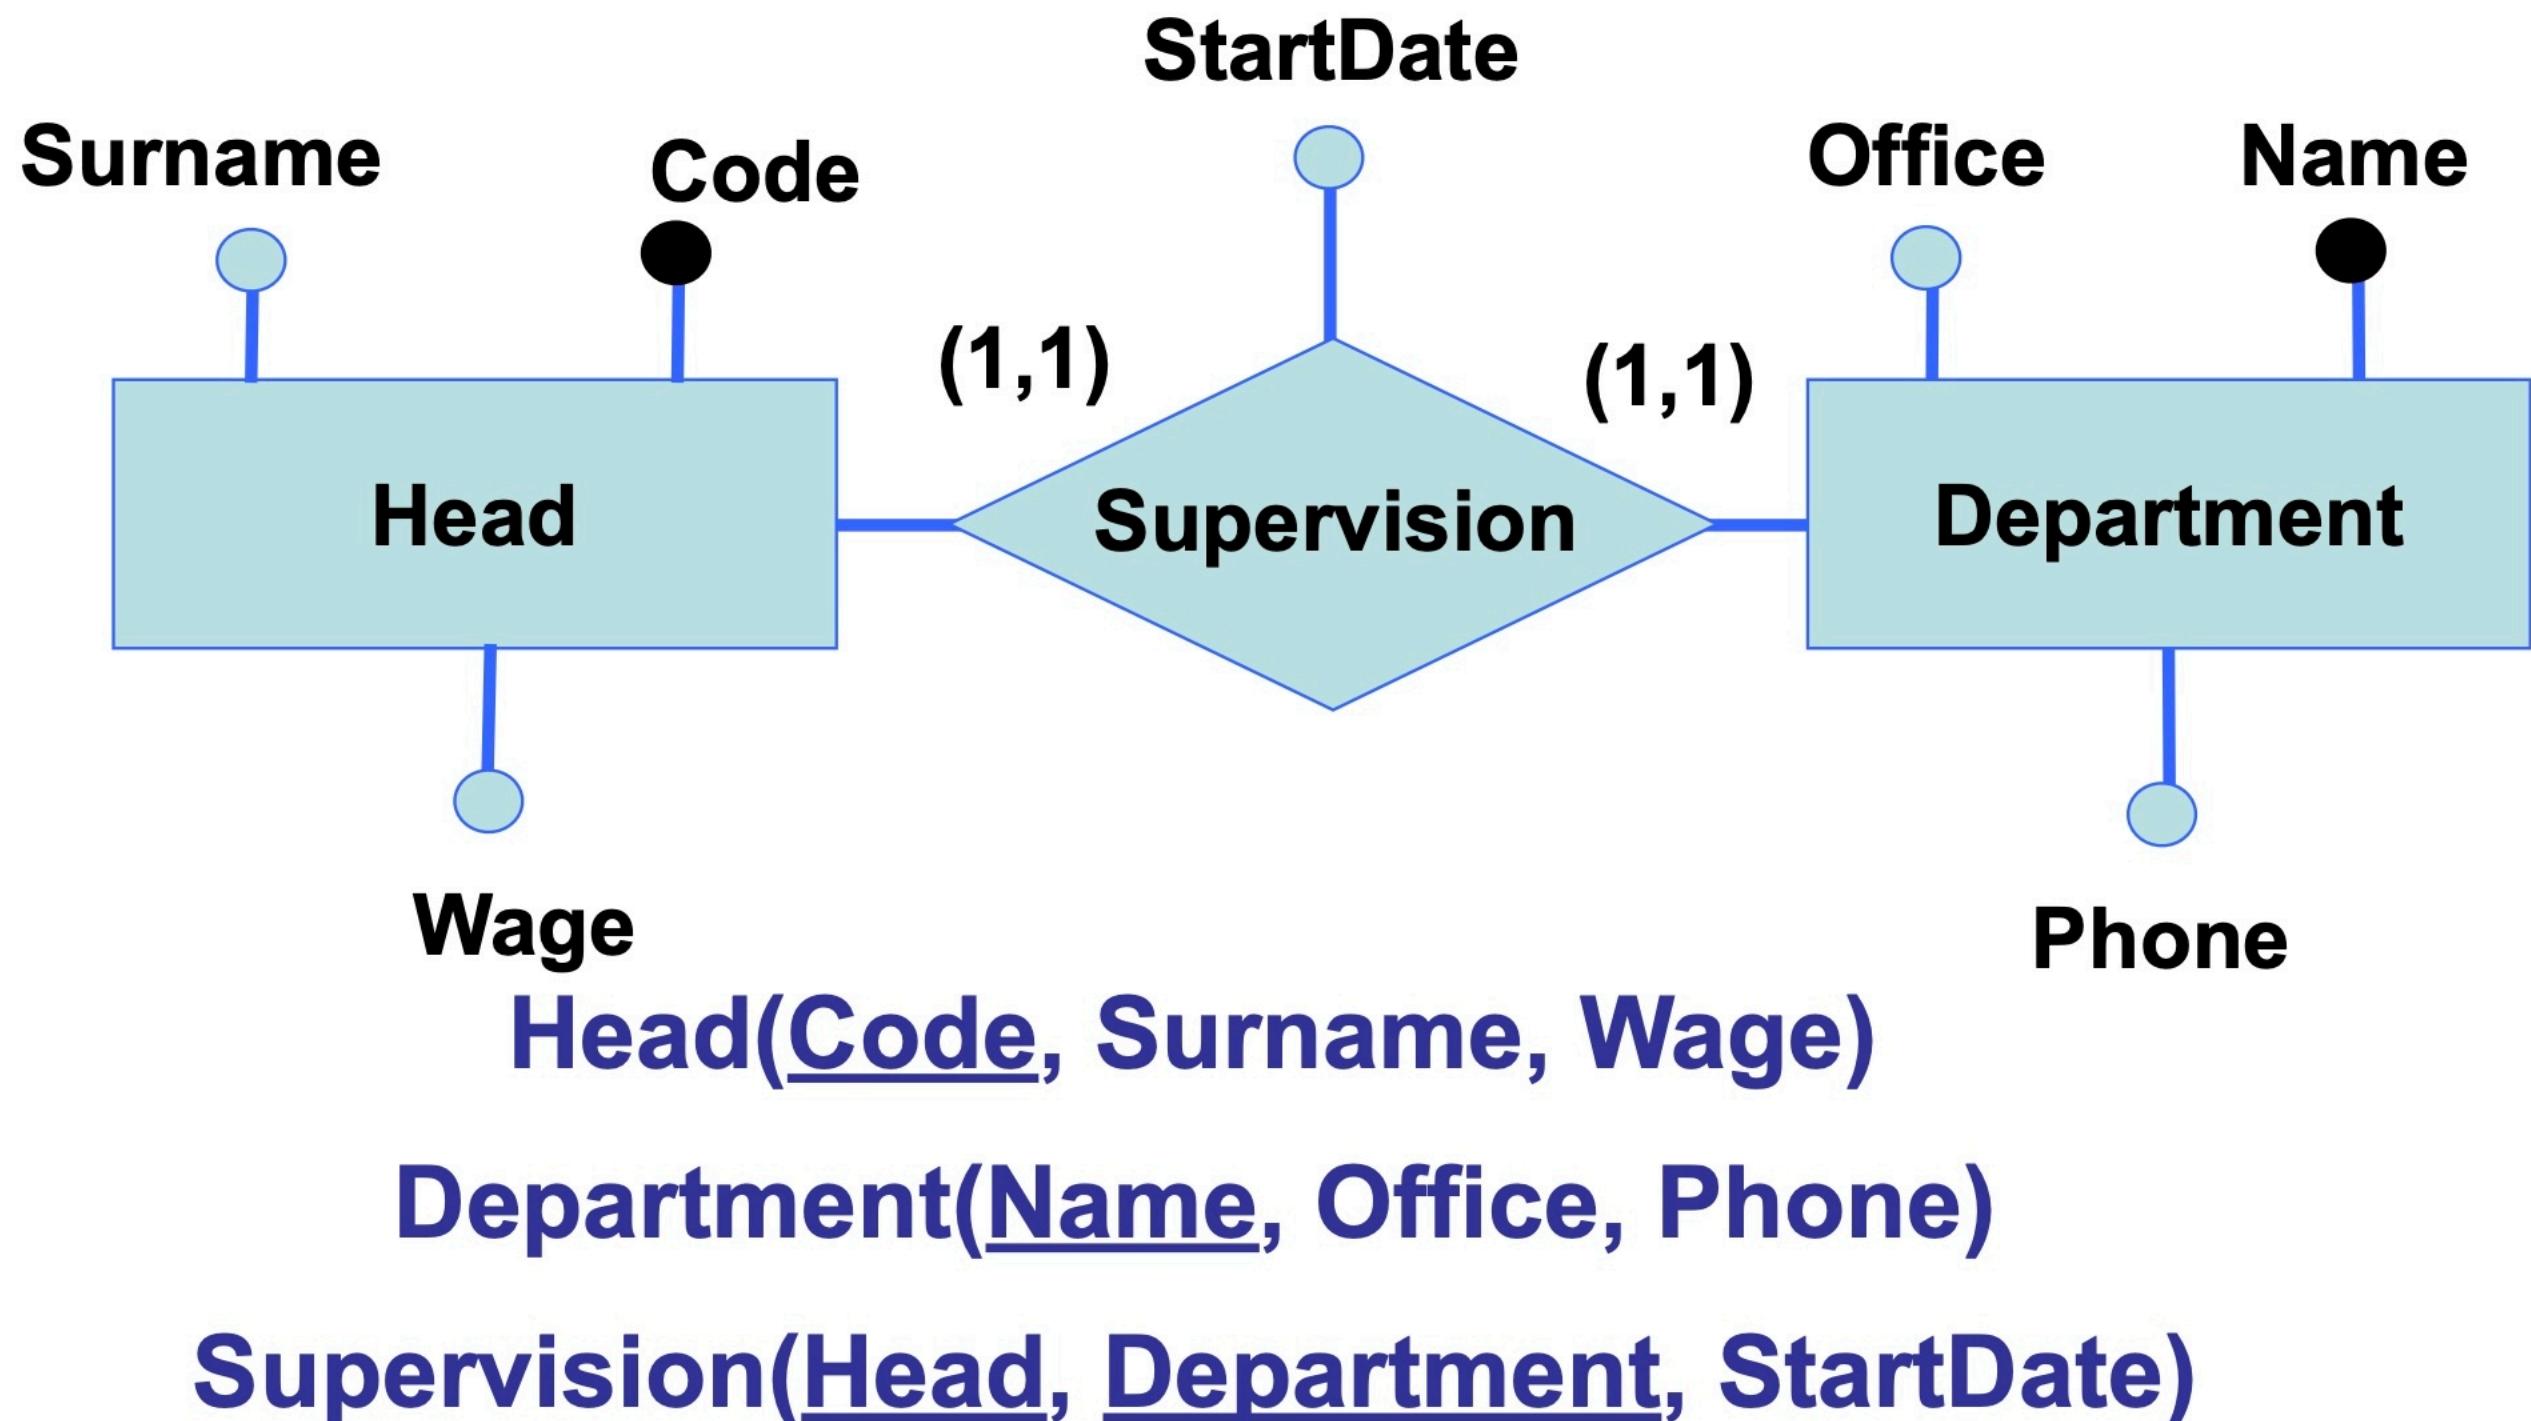


Head(Code, Surname, Wage)

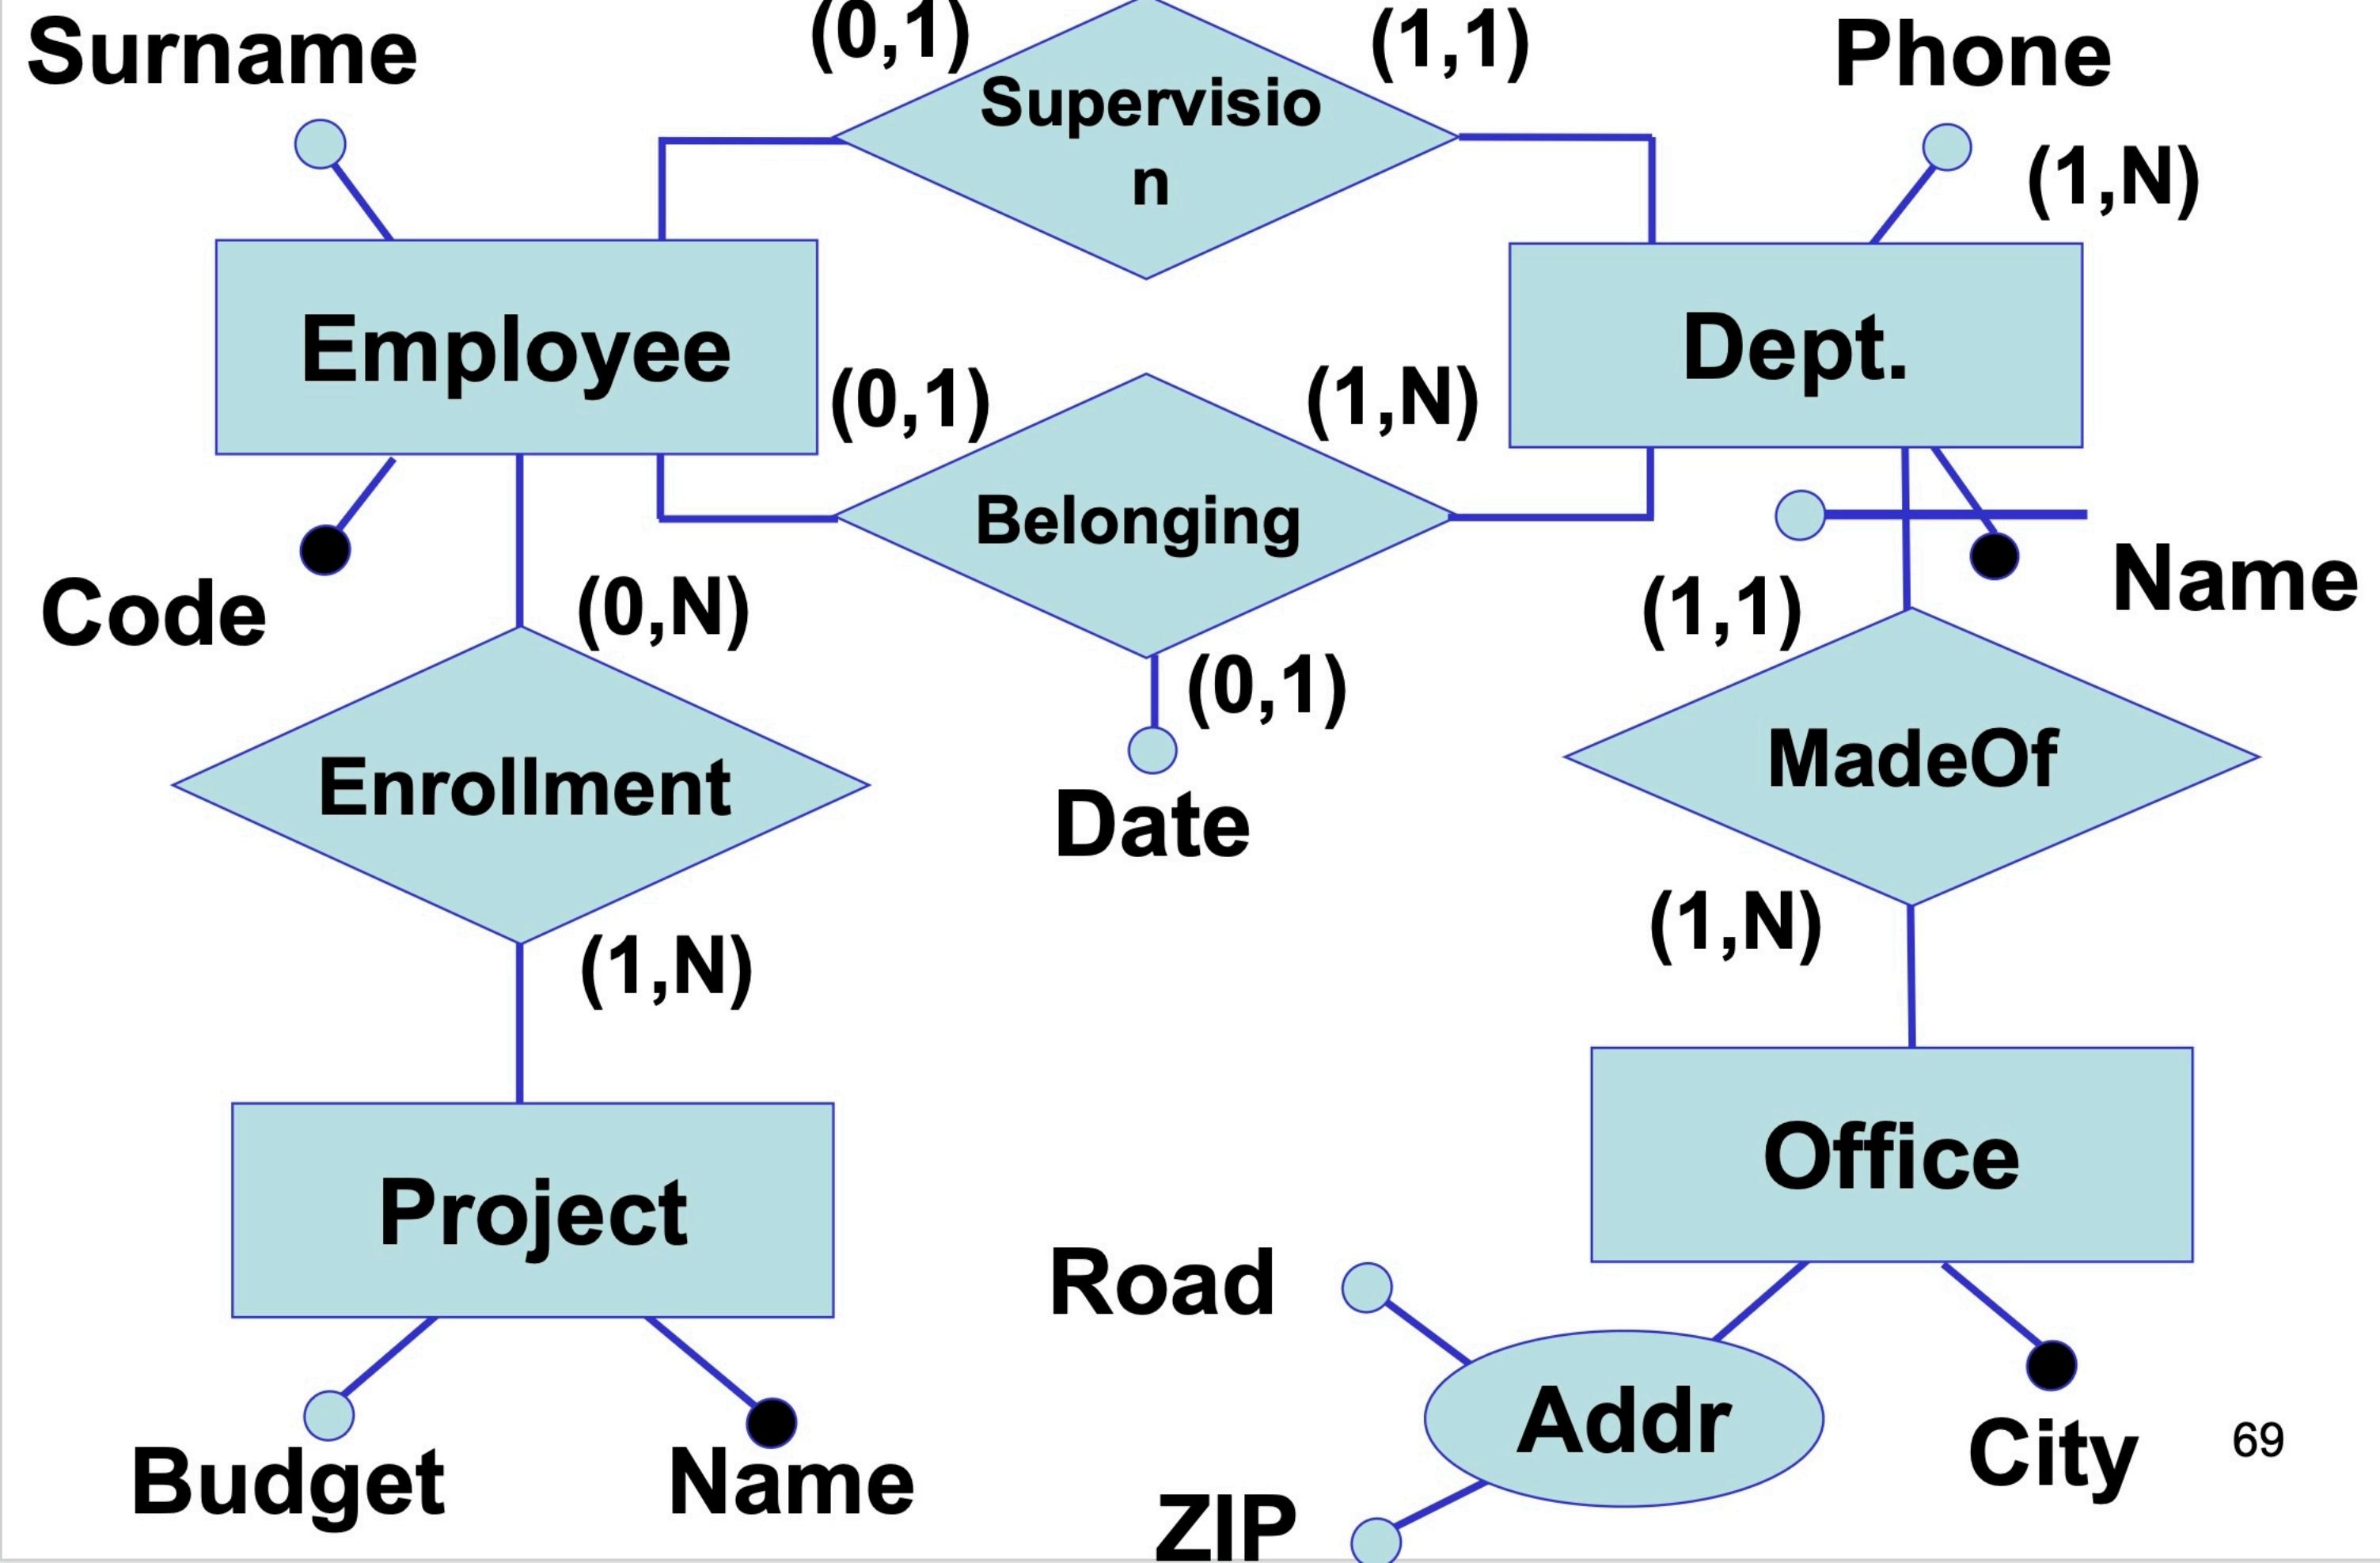
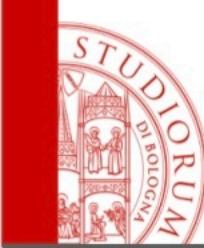
Department(Name, Office, Phone, Head, StartDate)

- **Referential Integrity Constraint**
- **No NULLs allowed**

Another solution



- Two Referential Integrity Constraints
- No NULLs allowed





Final Schema

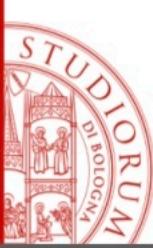
**Employee(Code, Surname,
Department,Office, Date*)**

Dept.(Name, City, Phone, Head)

Office(City, Road, Zip)

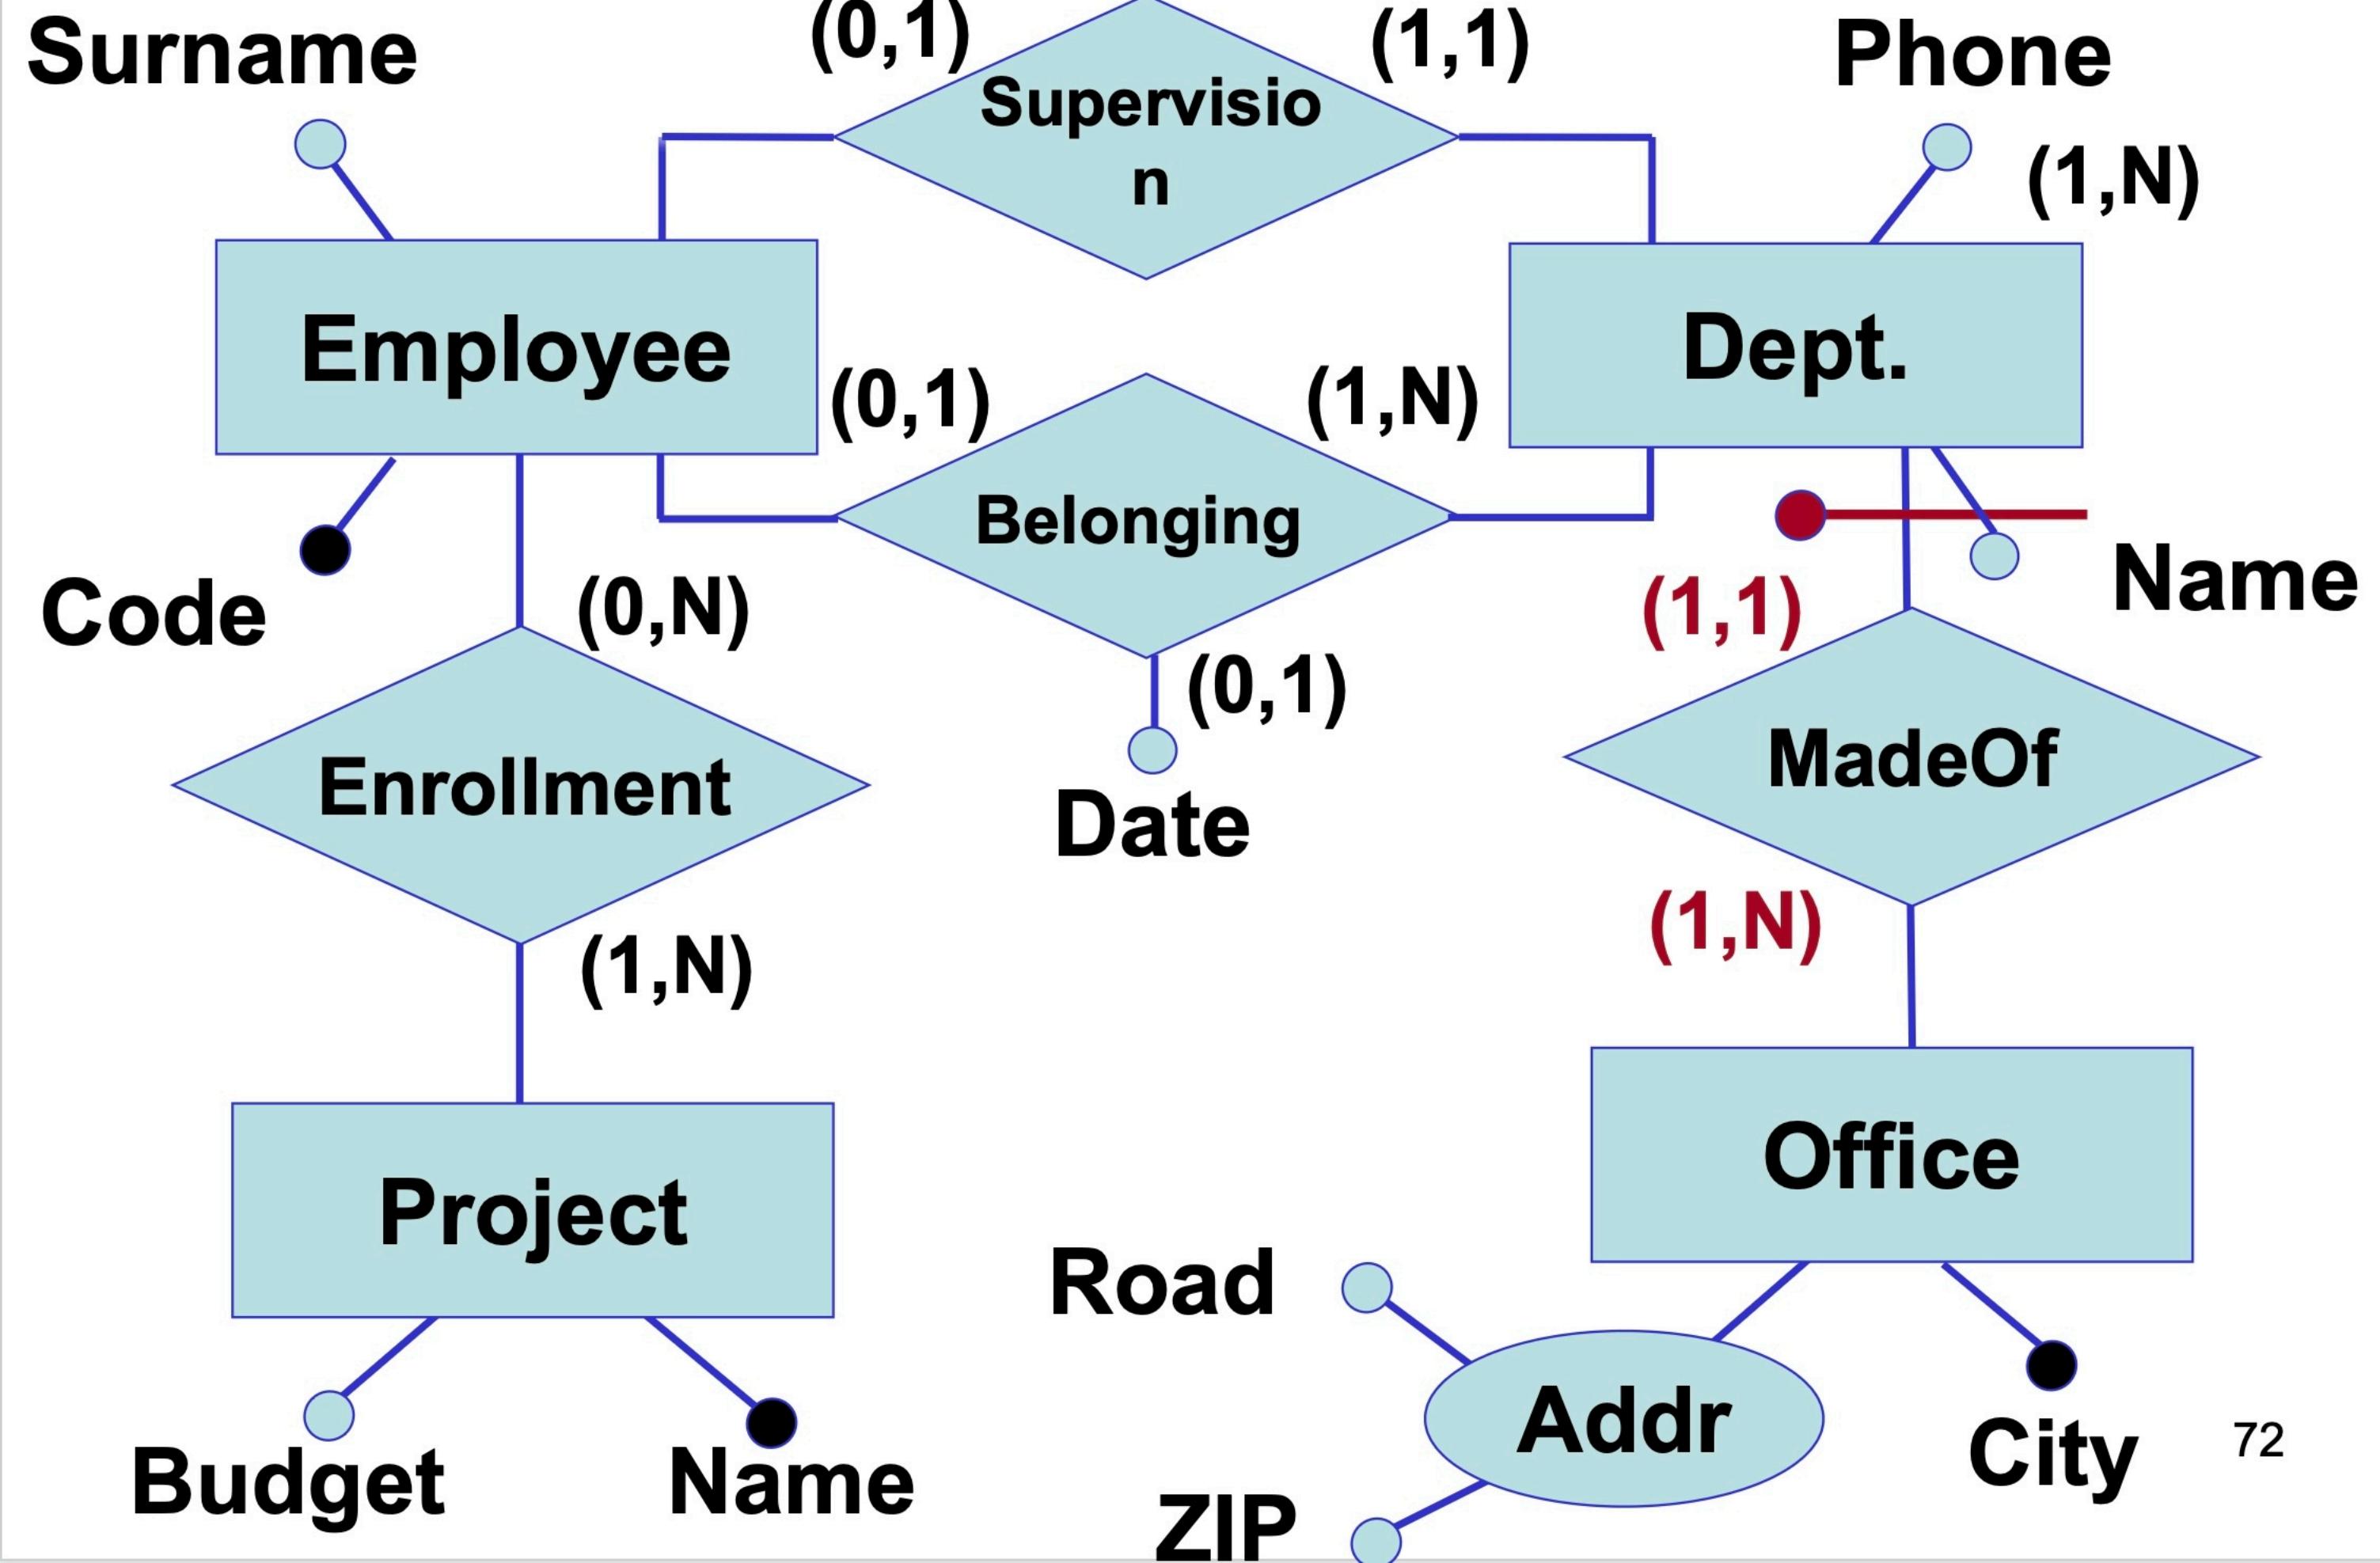
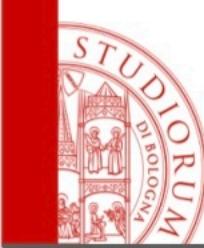
Project(Name, Budget)

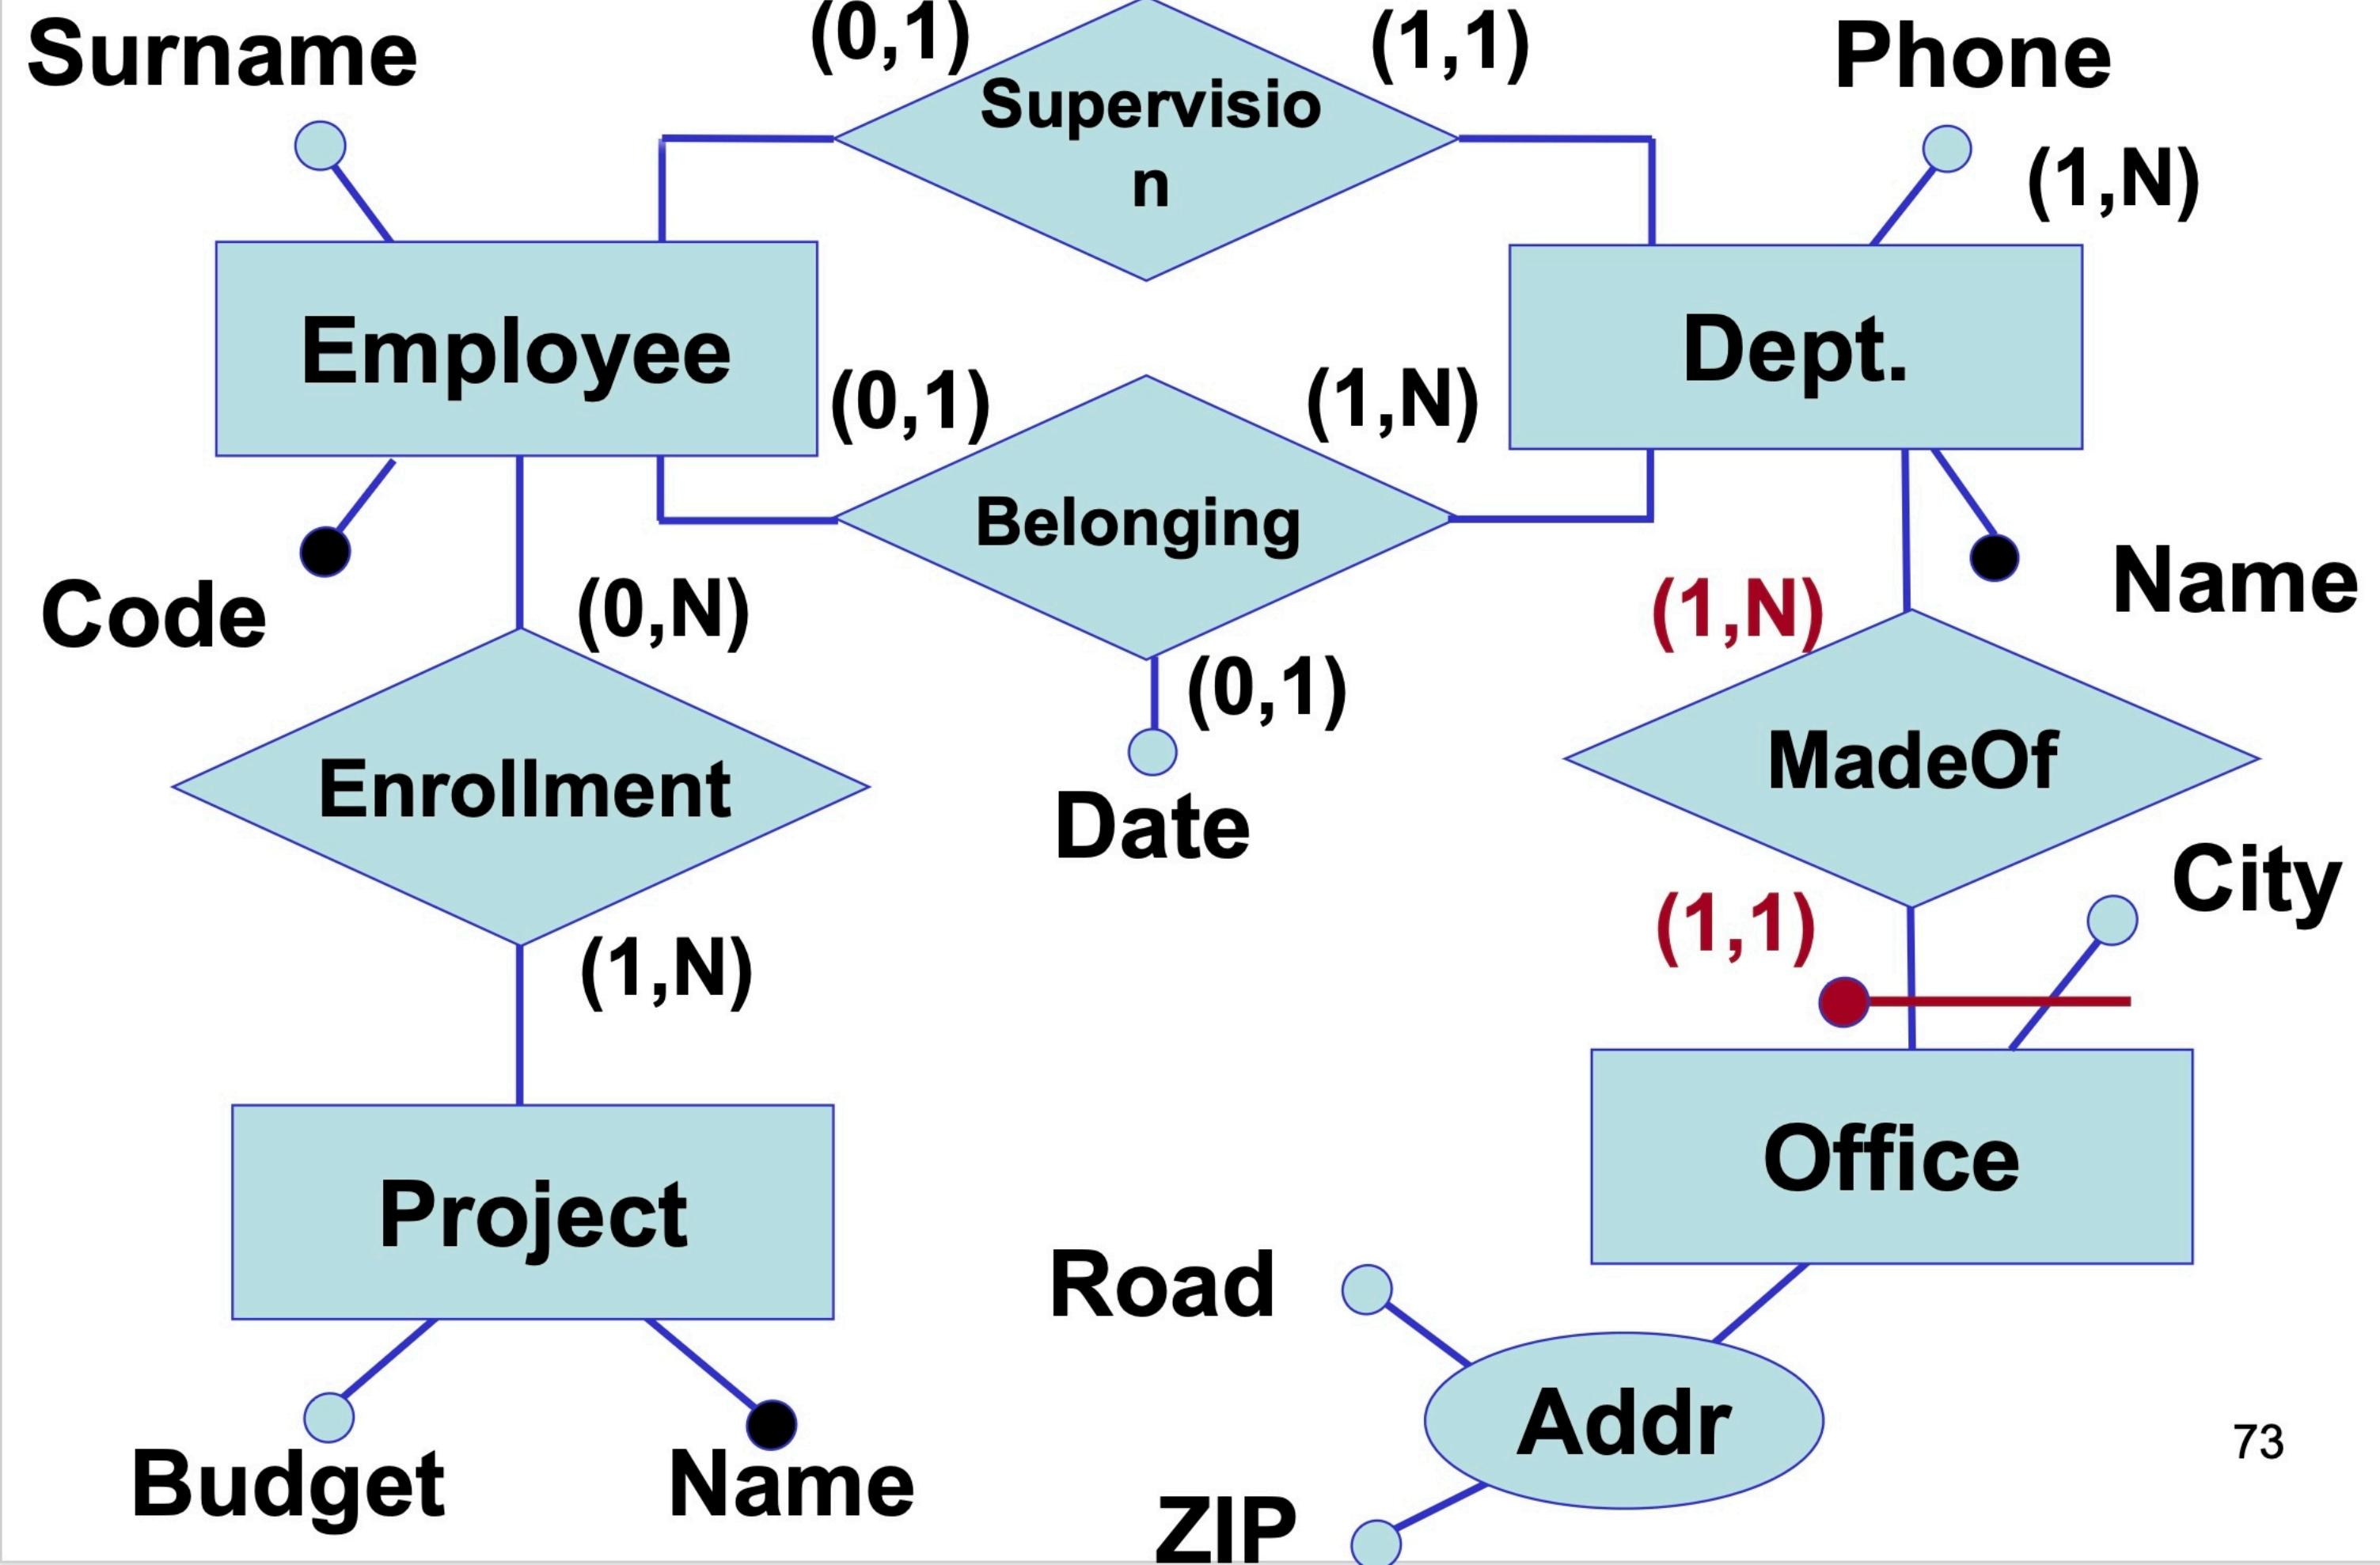
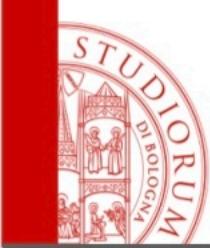
Enrollment(Employee, Project)



Warning

- Apparently small differences in cardinality and identifier choices could lead to very different meanings.







Second final schema

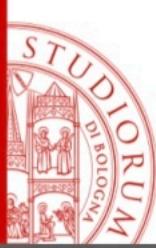
**Employee(Code, Surname,
Dept.,Office, Date*)**

Dept.(Name, City, Phone, Head)

Office(City, Dept., Road, Zip)

Project(Name, Budget)

Enrollment(Employee, Project)



Tools

- There are some CASE off-the-shelf software that provide support throughout all the modelling phases of designing a database

