



Databases

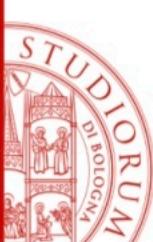
Active Databases



Passive Databases

- Reaction strategies in the integrity constraints are the first example of the need to introduce a **reactive** behaviour in the databases:

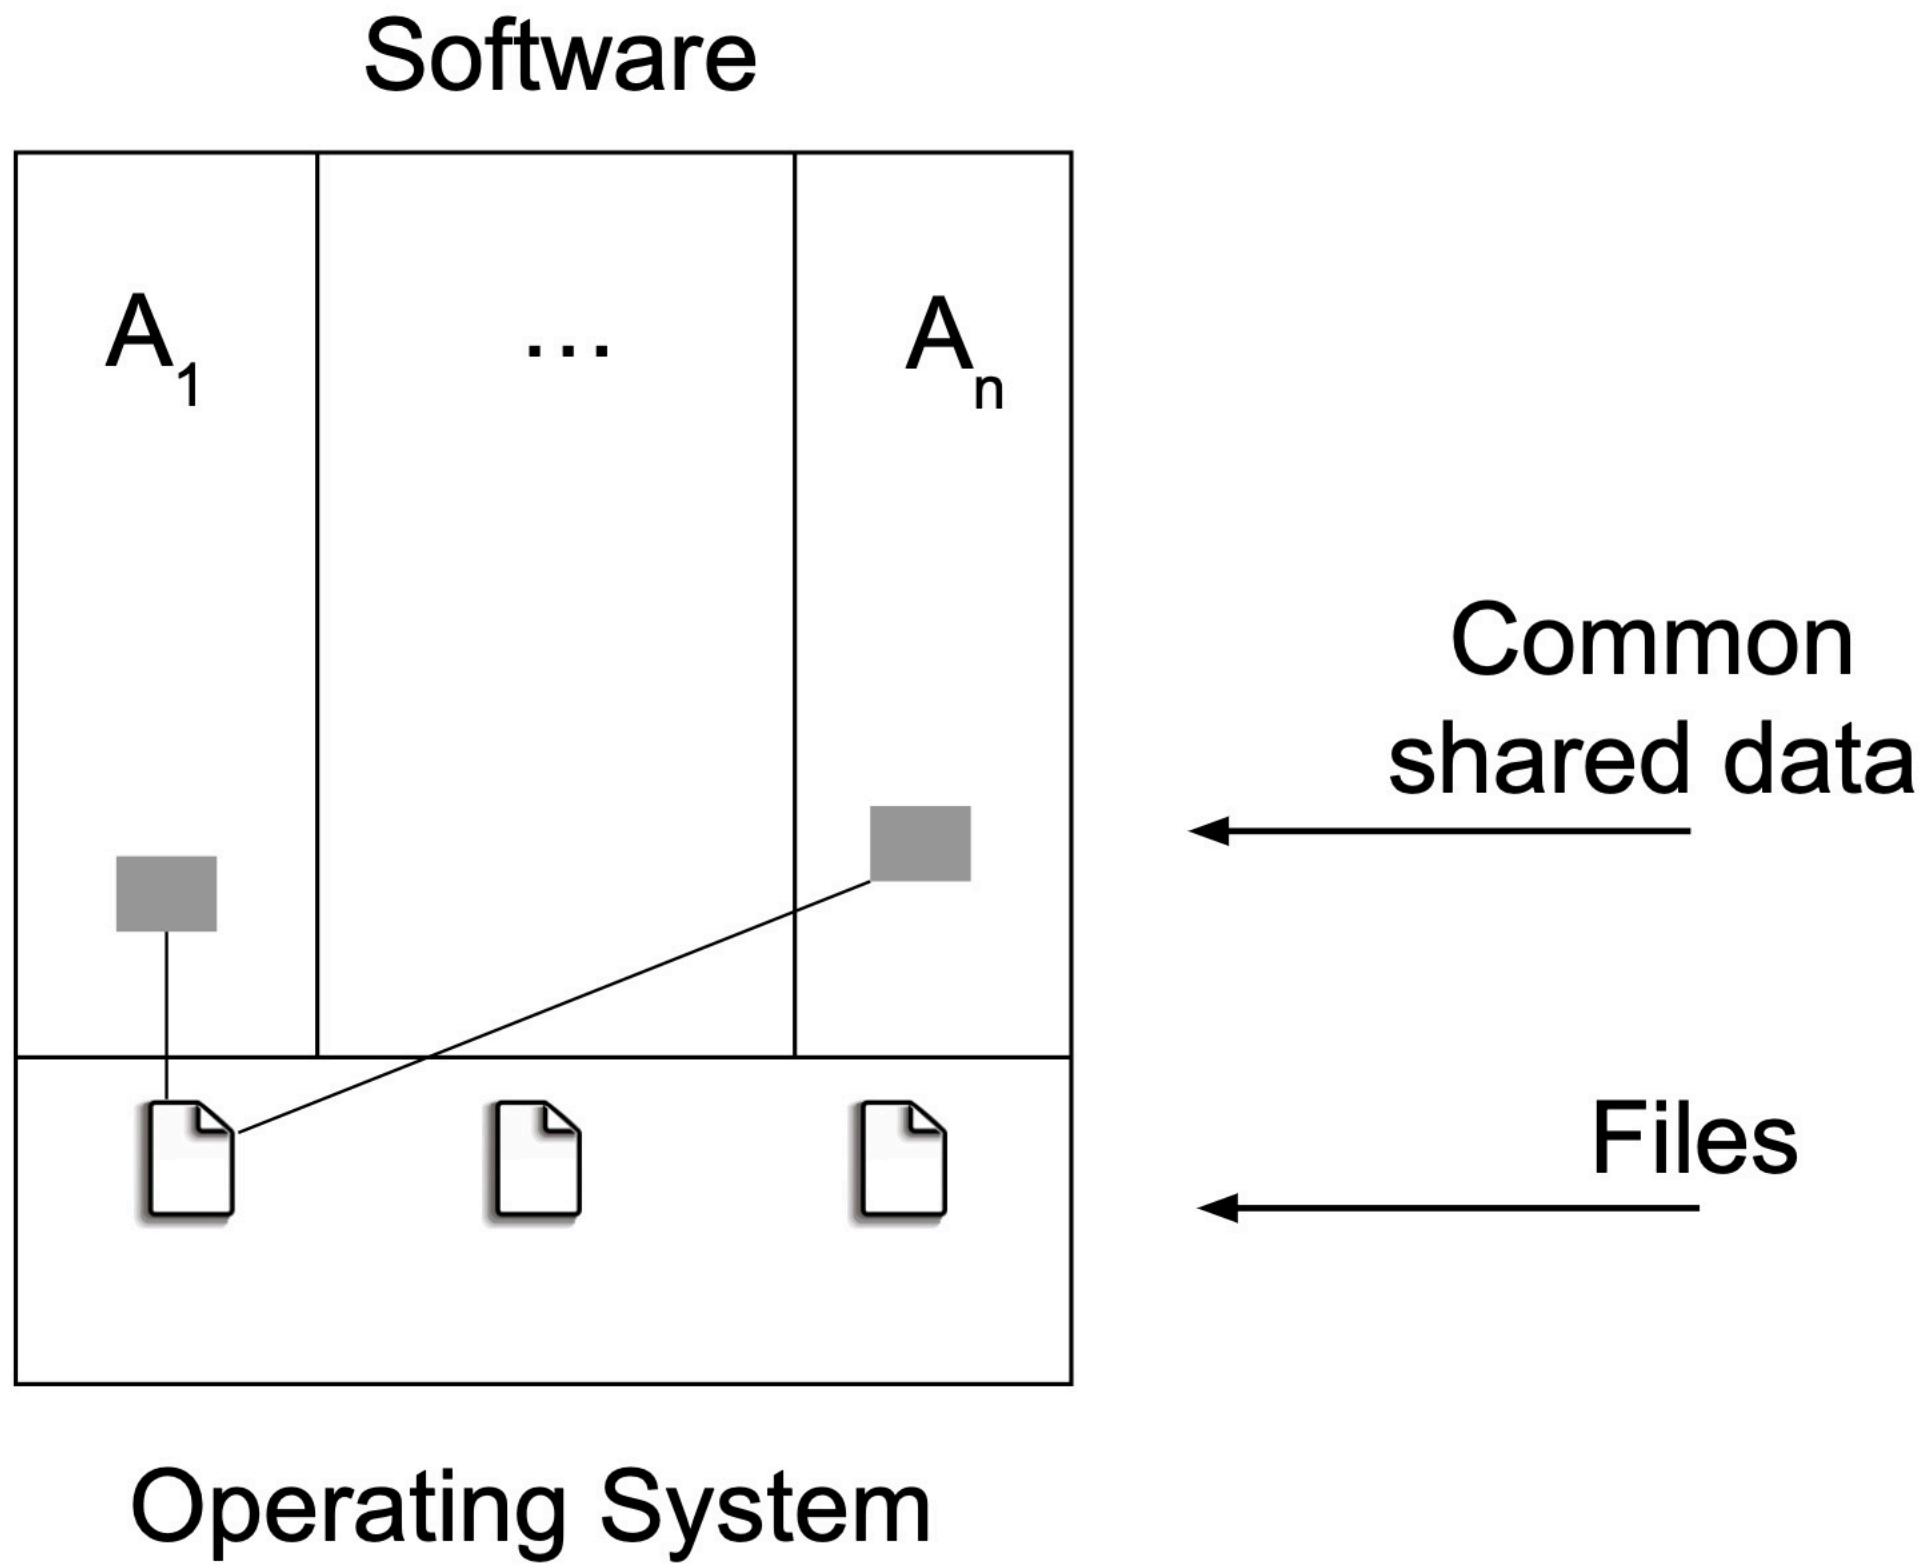
```
on      < delete | updated >
        < cascade | set null | set
              default | no action >
```
- The idea is to introduce **language constructs** specific for this goal
- These constructs are called (active) **rules** to handle a part of the procedural behaviour of an application
- Being at the database level, this behaviour is therefore “shared” among many applications, achieving data independence.



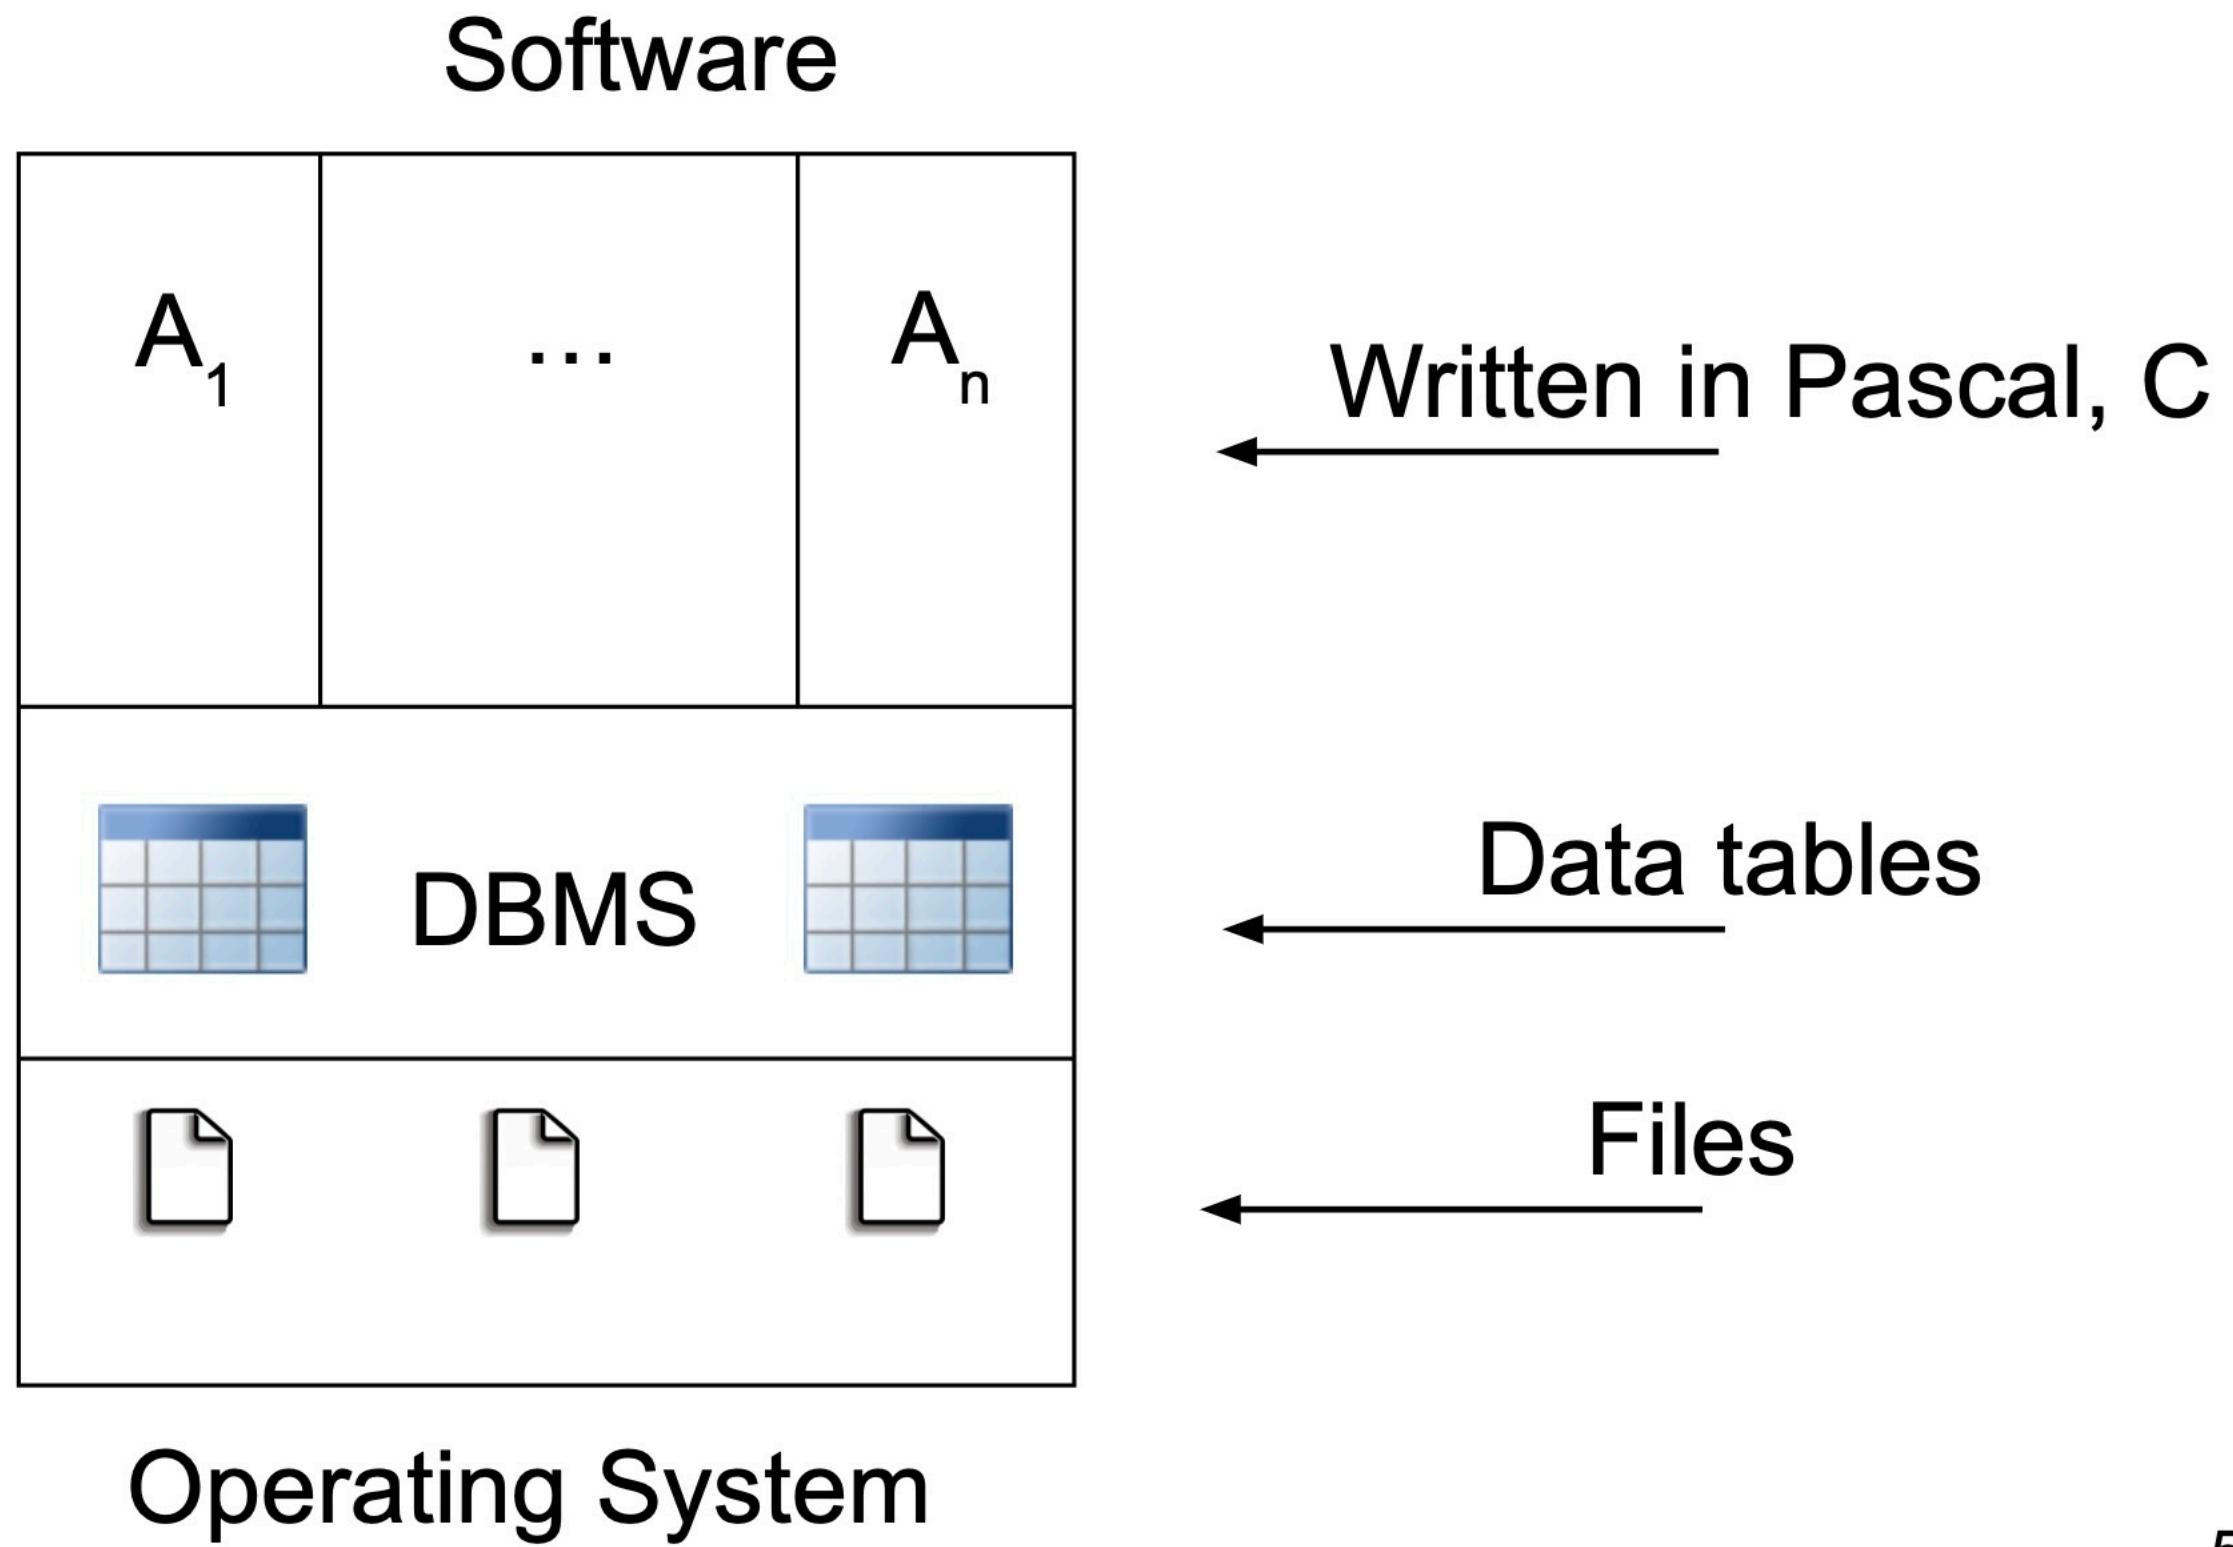
Active Databases

- An active database has a component to handle active rules of the kind **Event-Condition-Action**:
 - events are changes in the database
 - a condition is verified, based on a true/false value
 - one or more actions are executed
- These databases have a **reactive** behaviour (opposite to **passive**): they do not just execute the user transactions, but also the rules
- Commercial DBMSs (SQL3) use **triggers** to specify rules, like the ones in the schema definition

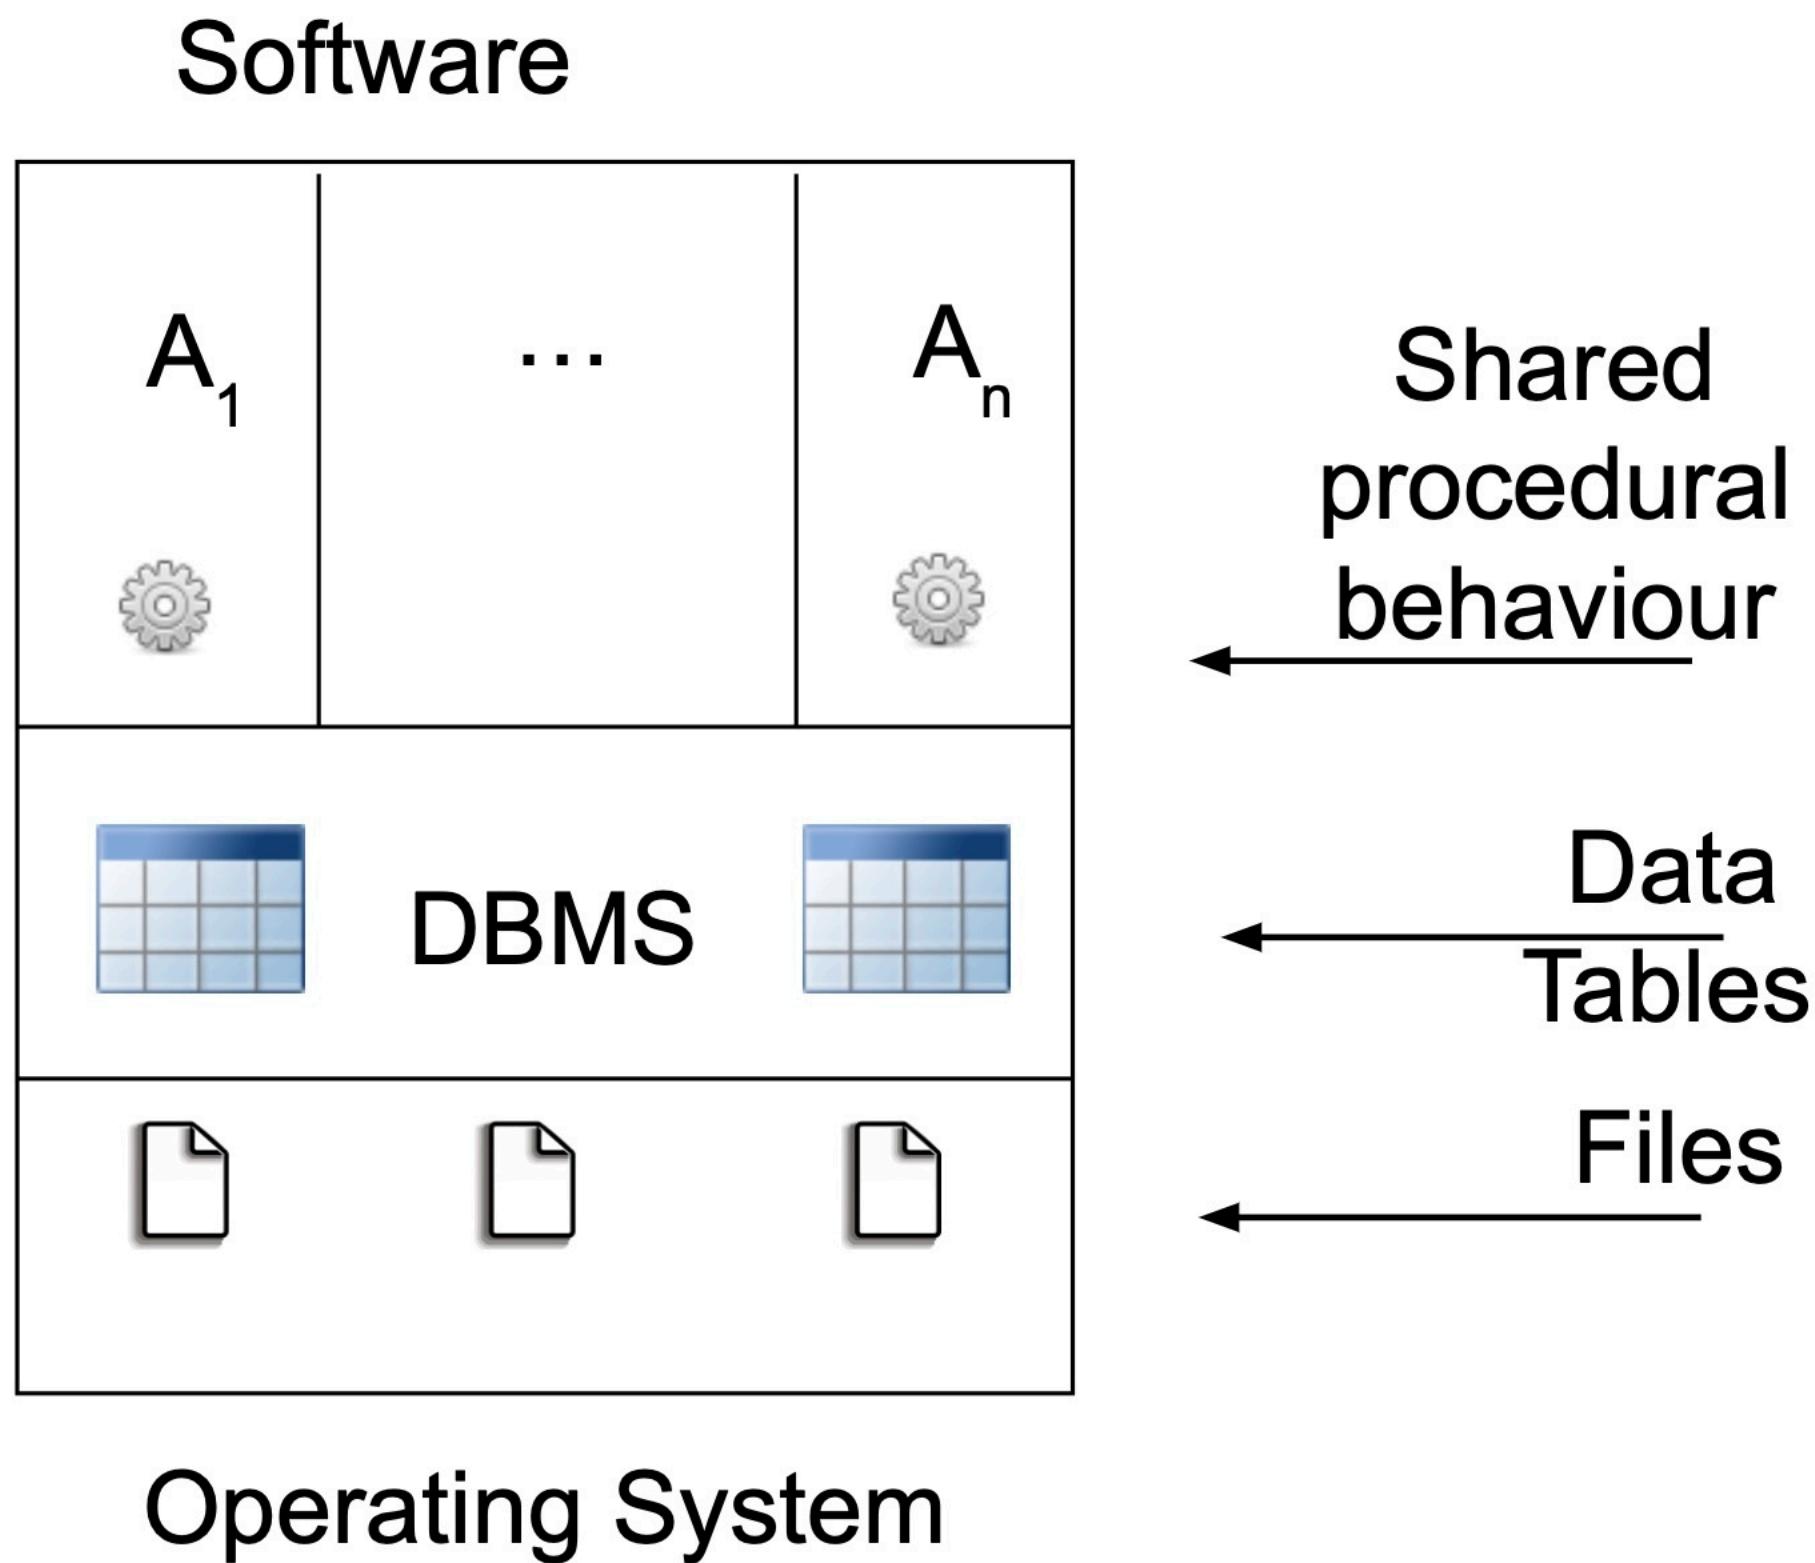
The 70s: no DBMS



The 80s: first DBMSs



The 90s: the procedural behaviour is shared among the software applications



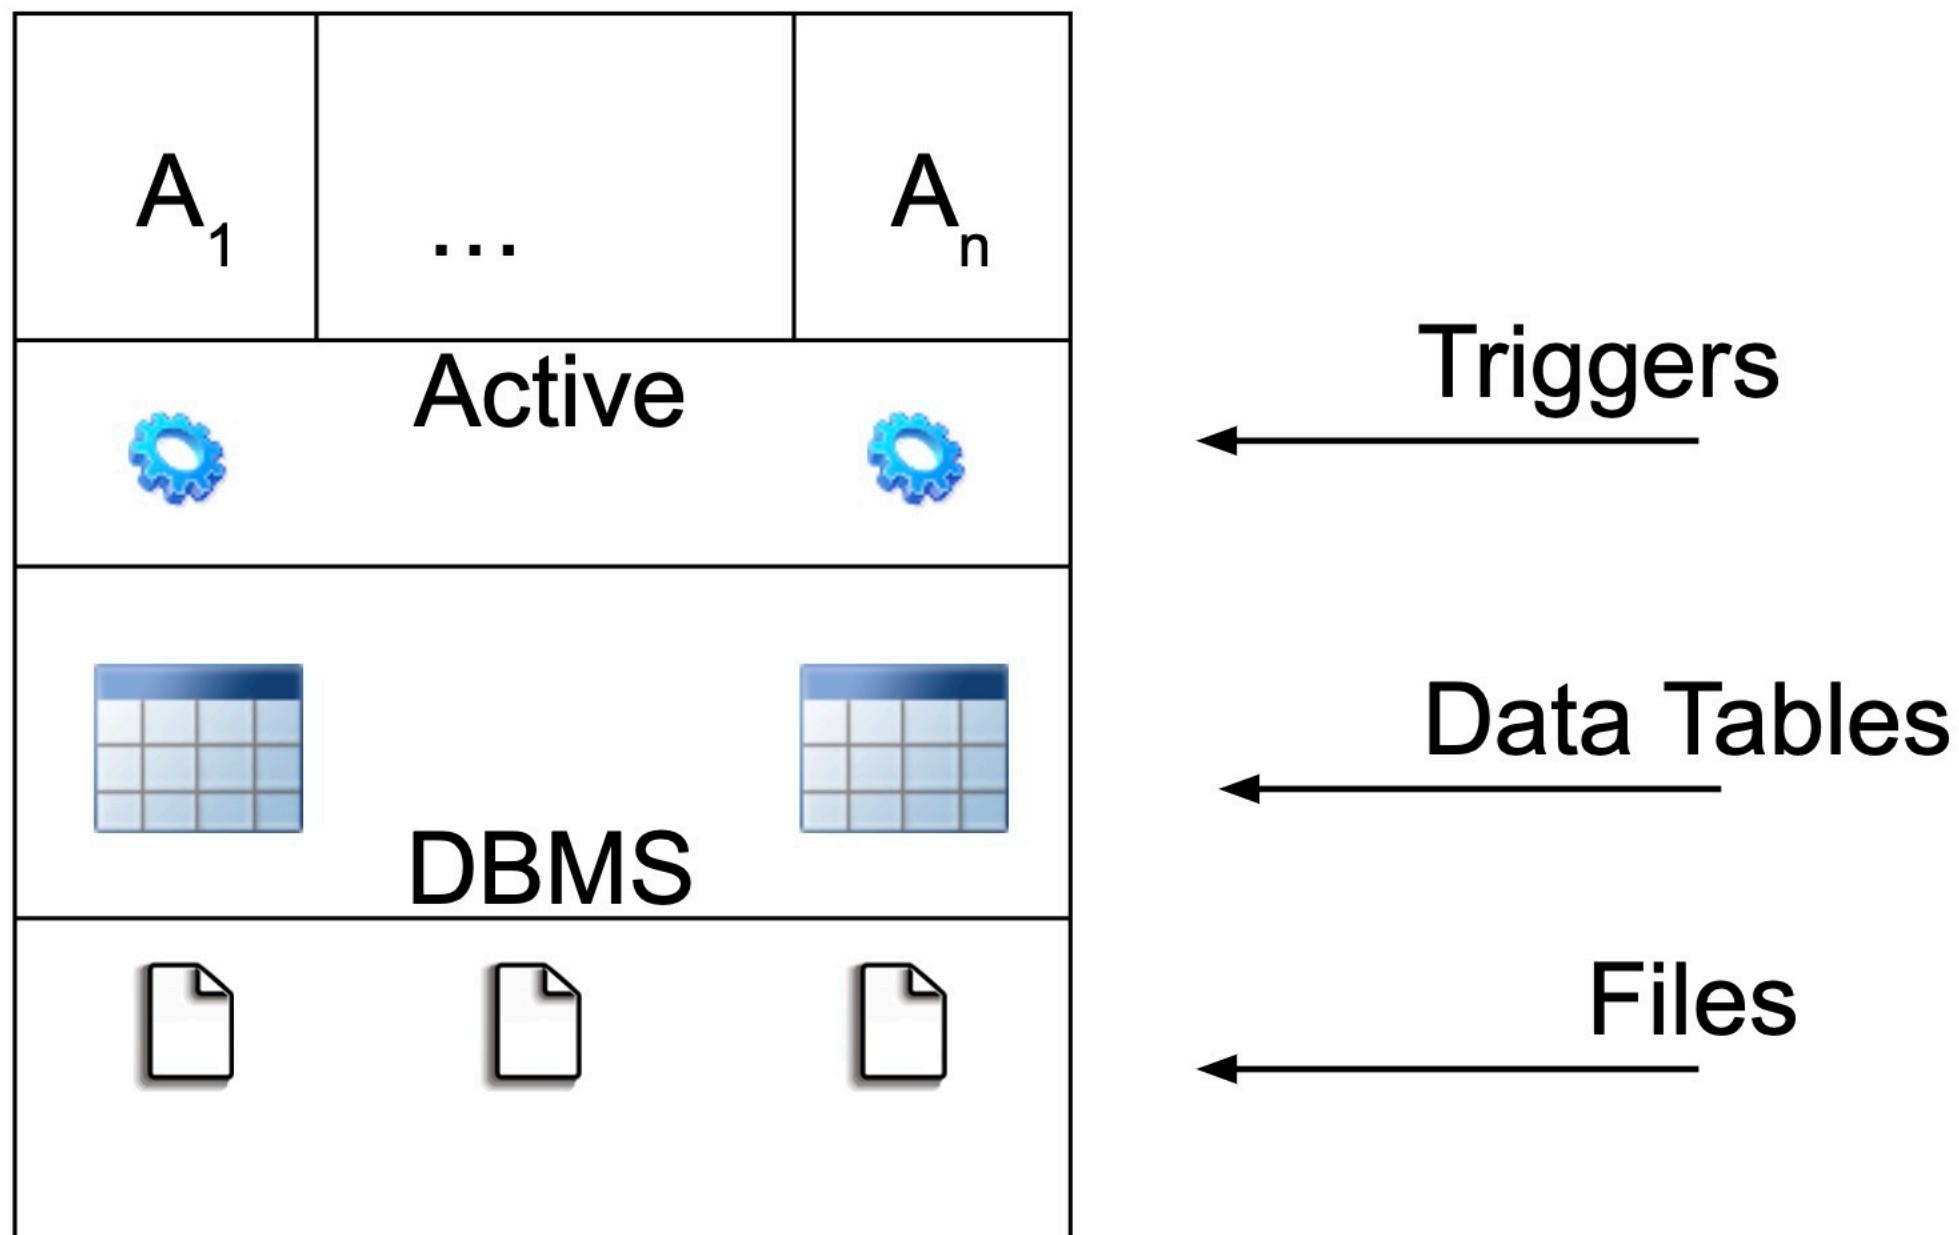


Advances in DBMSs: Stored procedures

- Stored procedures have been introduced to share the common procedural behaviours between different software.
- Stored procedures are not standardized and are affected by the problem of impedance mismatch (we'll see later) with the language used to express such procedures
- As a result specific rules (*triggers*) have been introduced to model the procedural behaviour shared among different software and are handled by the DBMS itself.

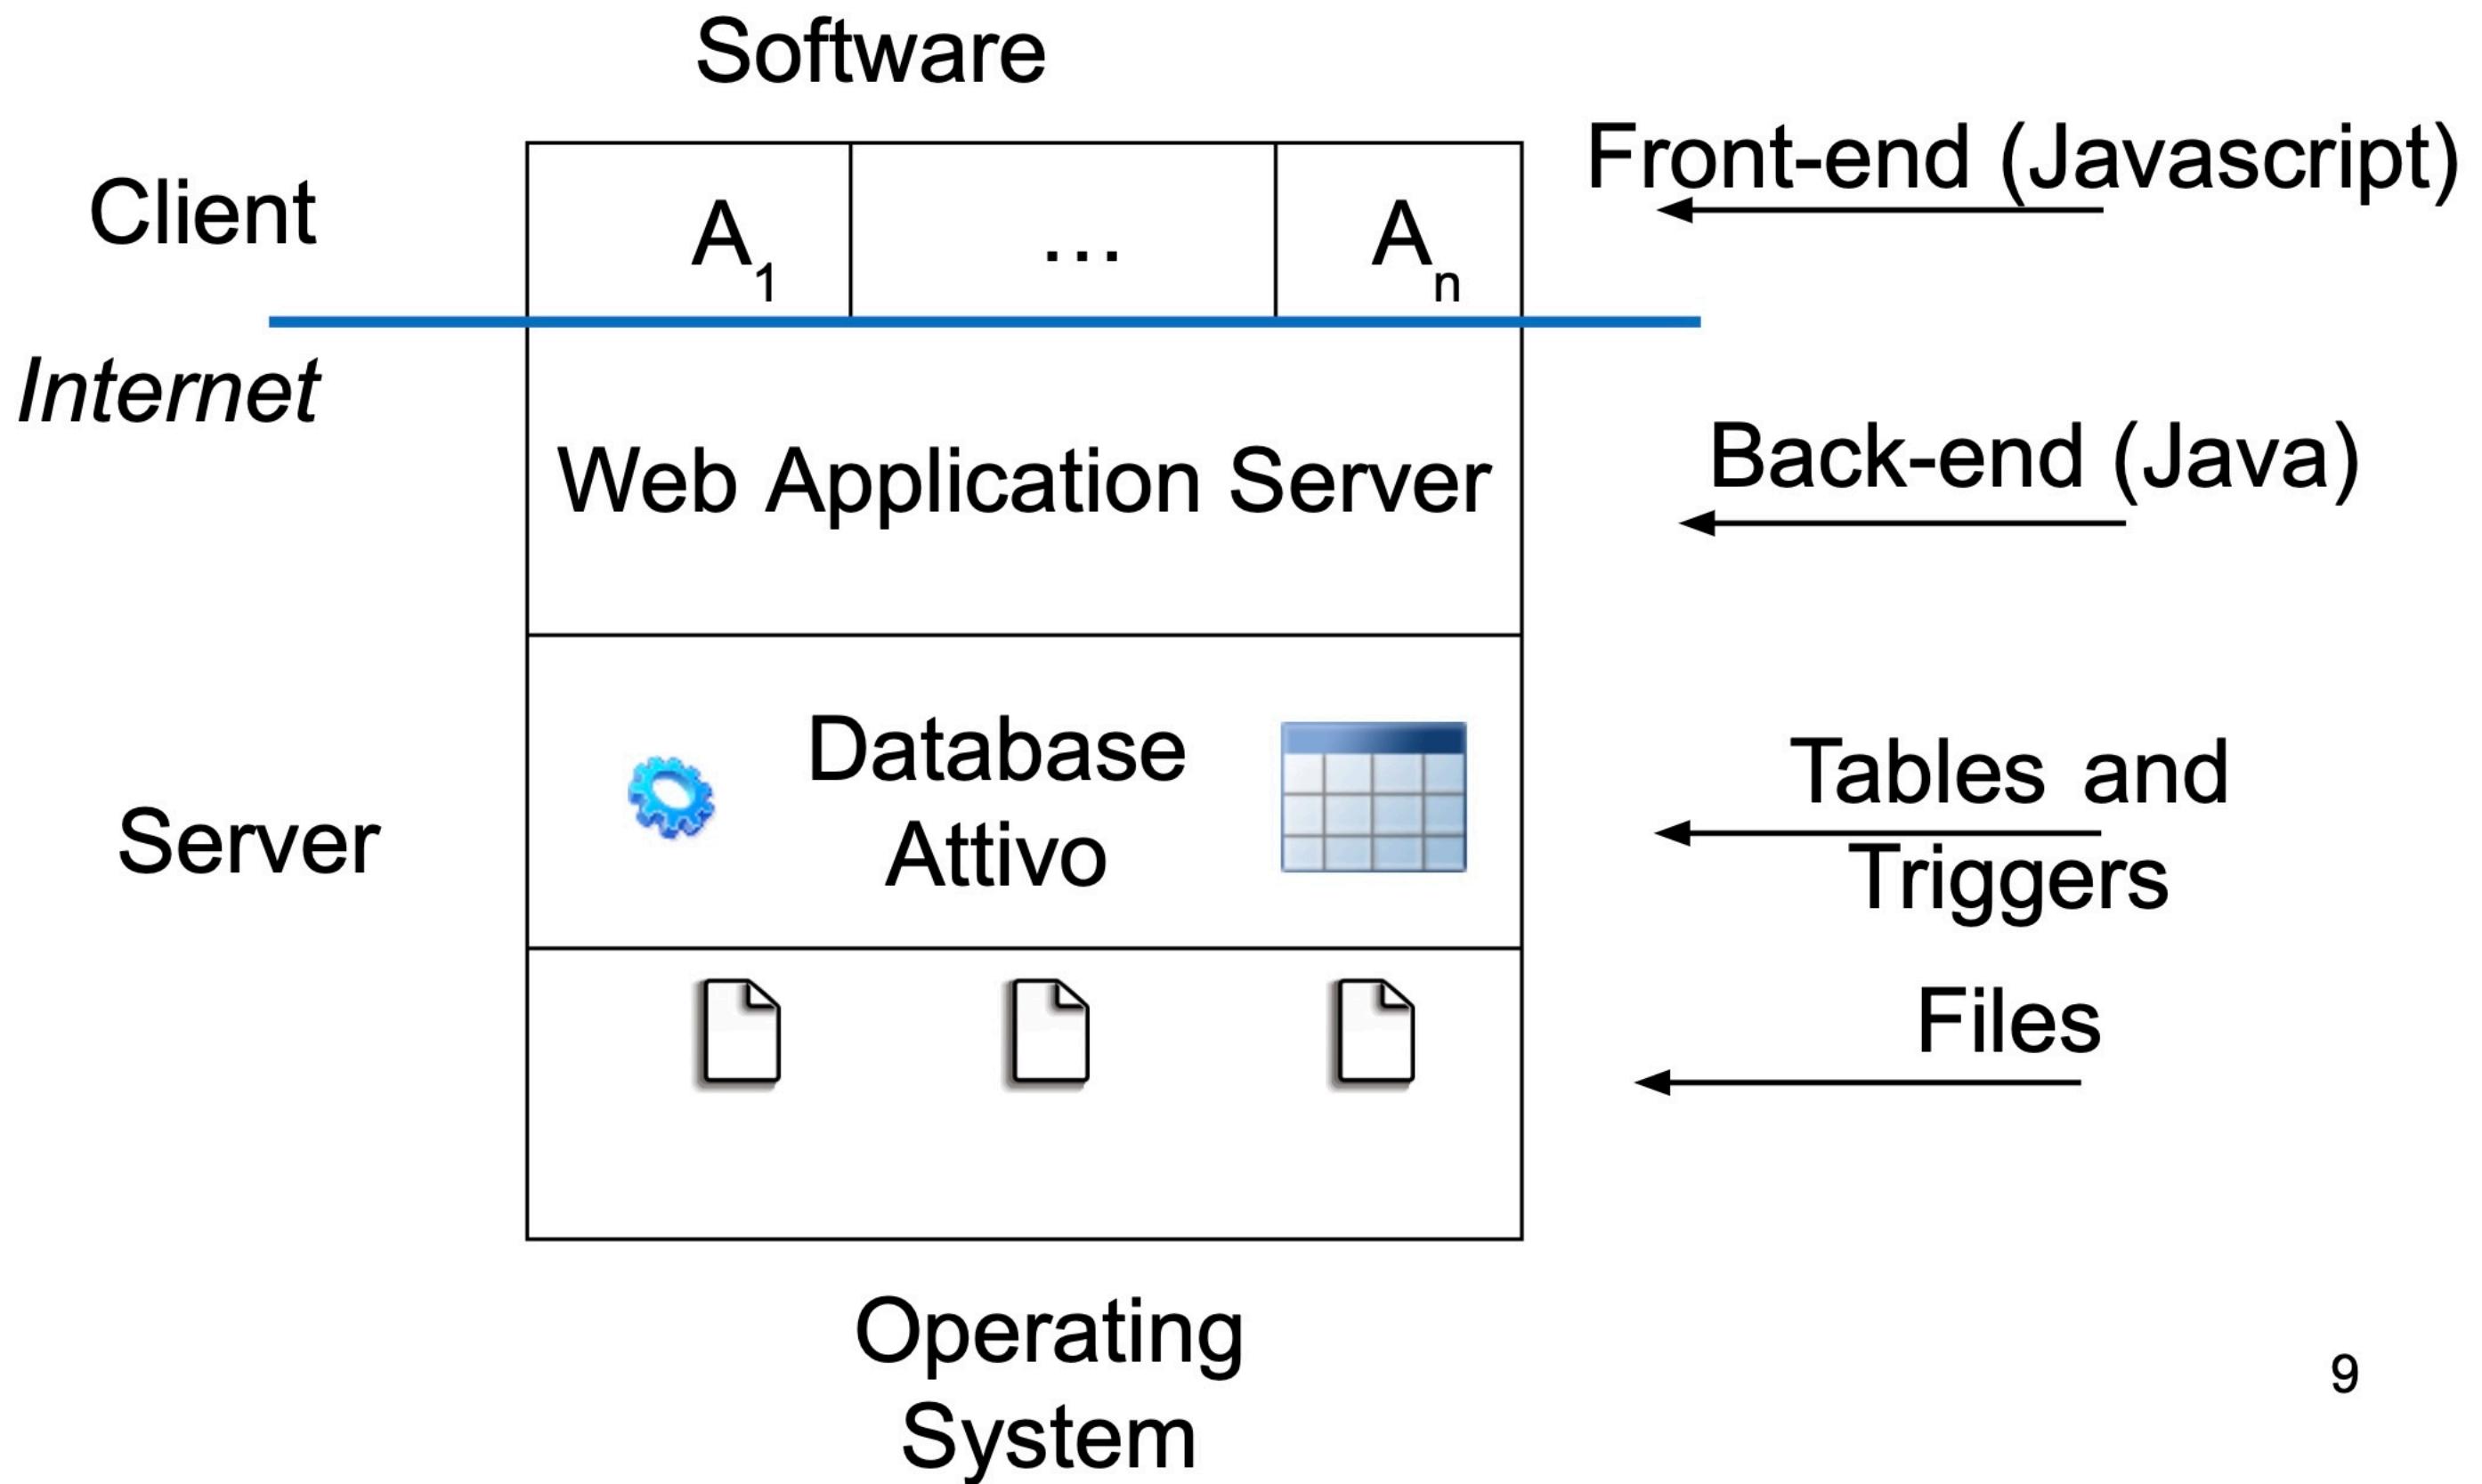
The 90s: Active DBMS

Software

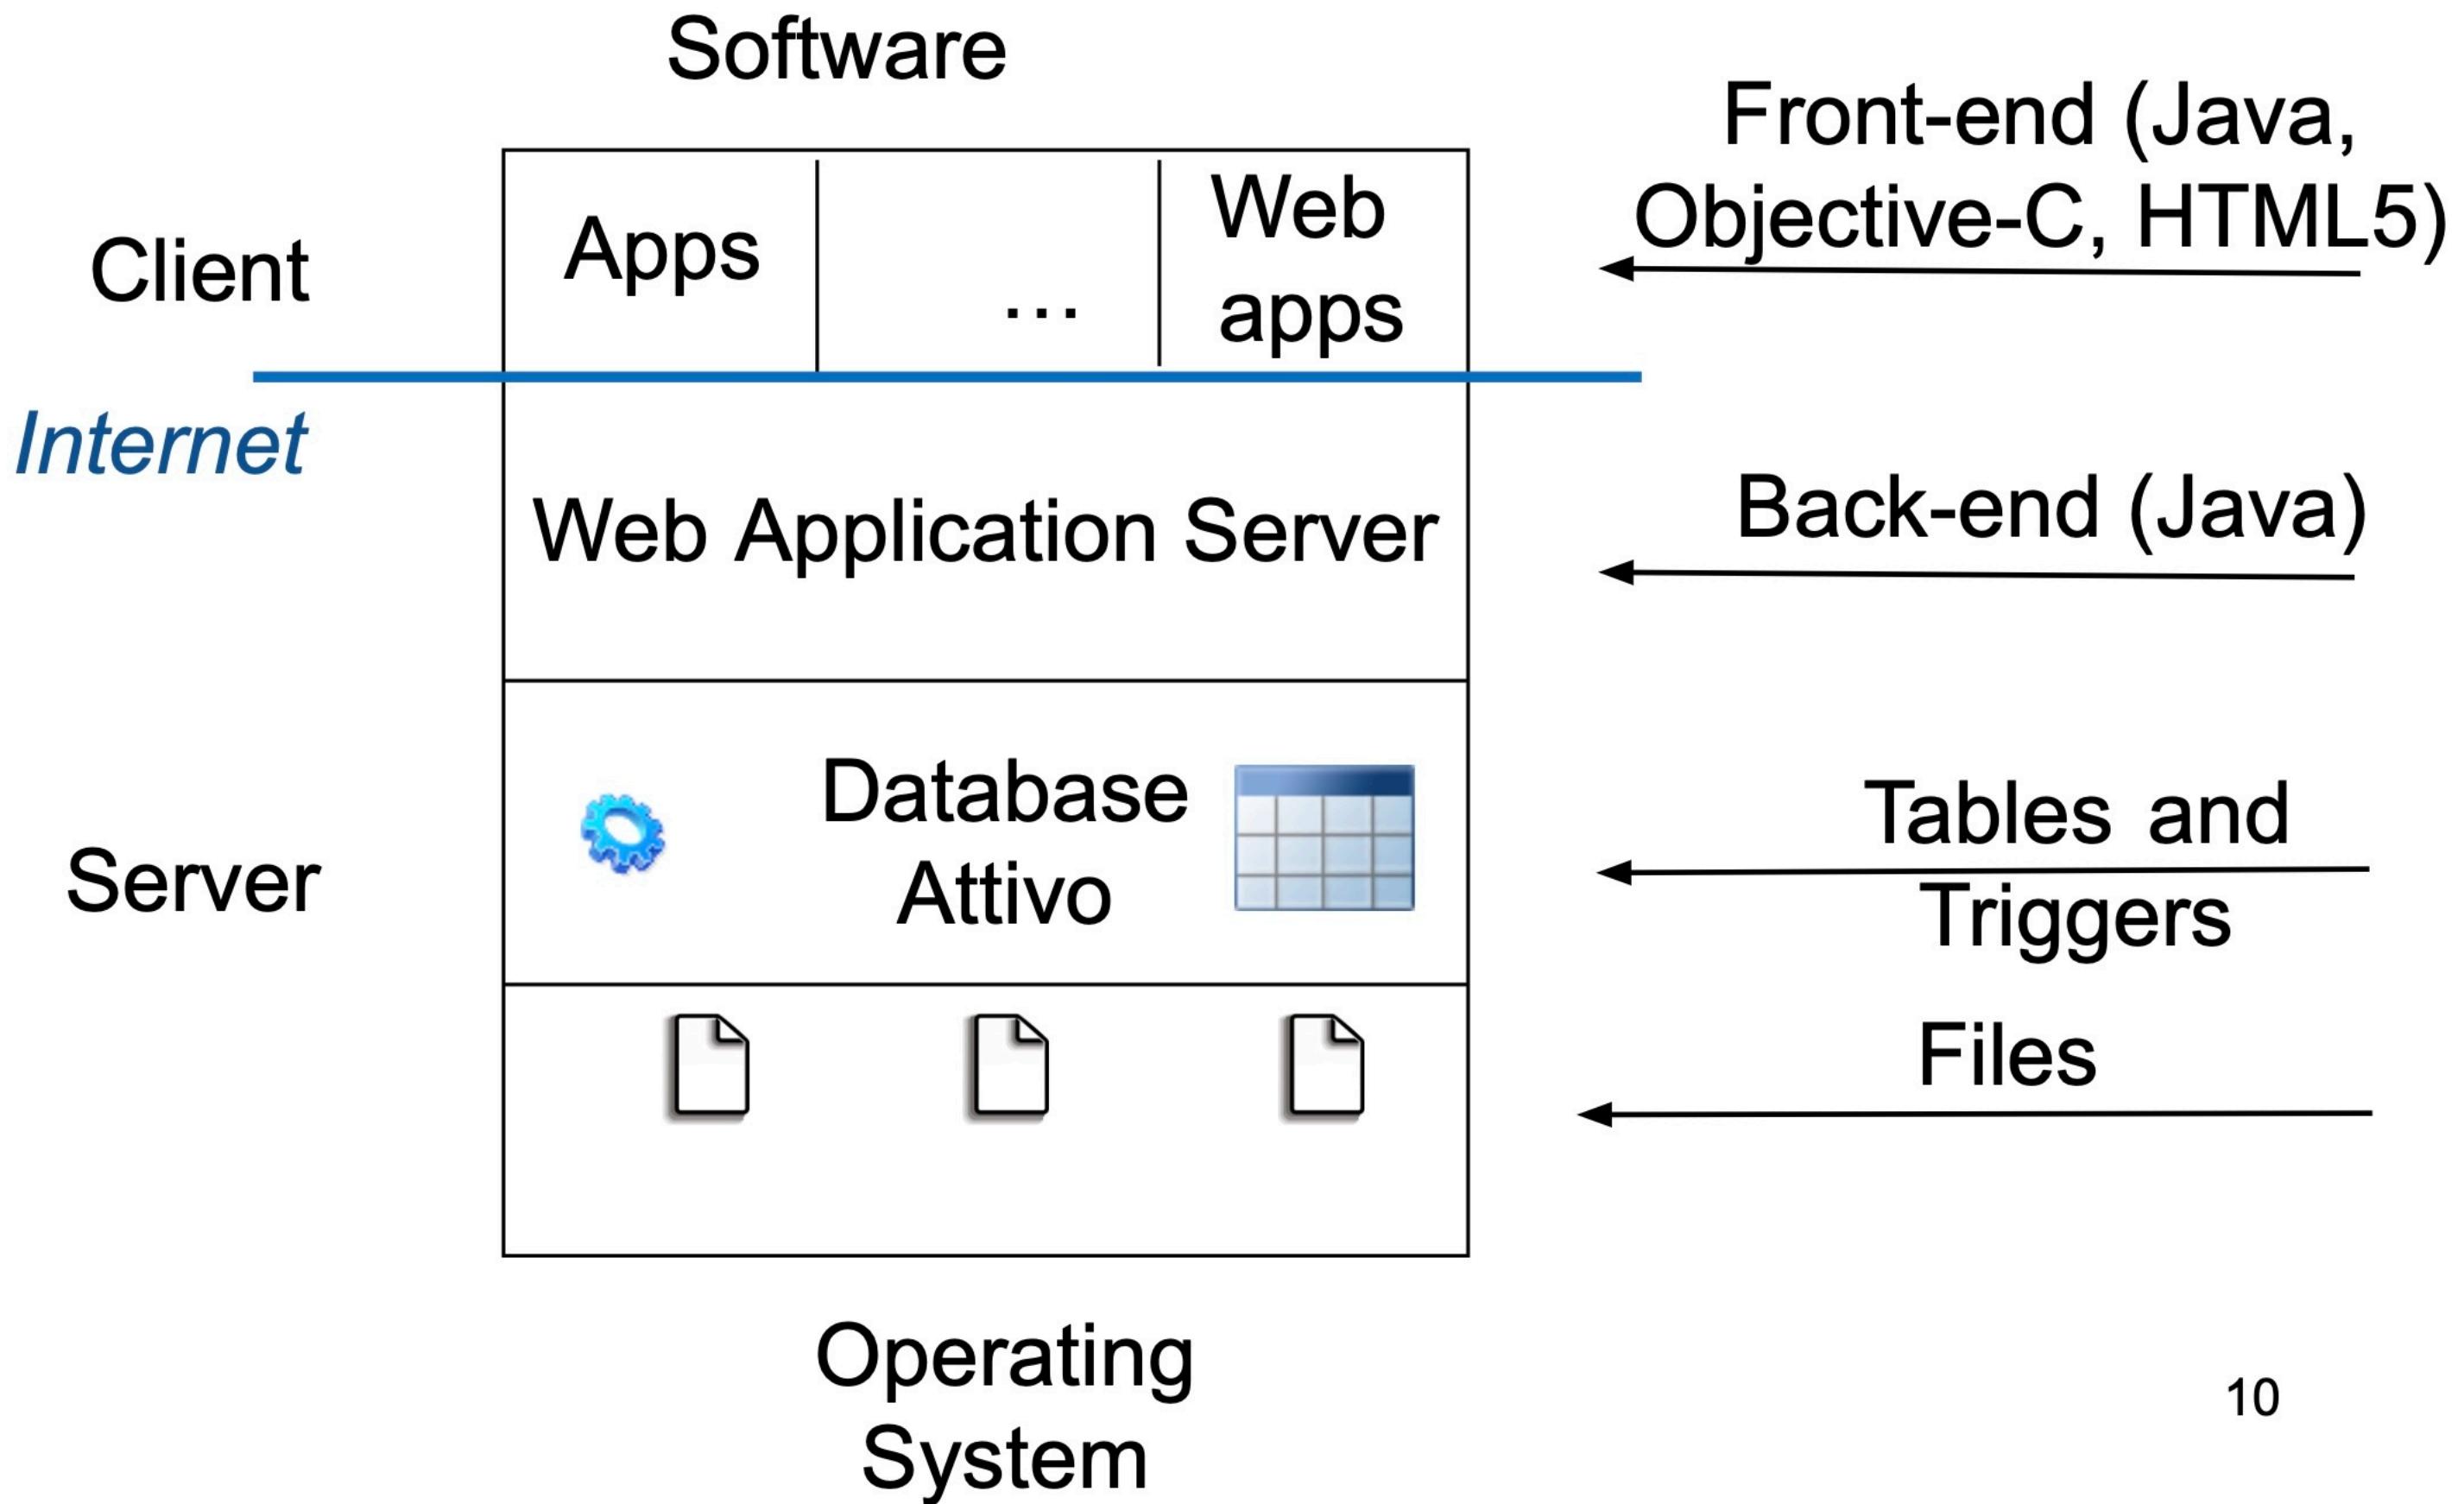


Operating
System

Advance in the 2000s



The 2010s: mobile apps





Trigger

- Defined with DDL instructions (`create trigger`)
 - based on the event-condition-action (ECA) paradigm:
 - event: change in the data, specified using `insert`, `delete`, `update`
 - condition: (optional) SQL predicate
 - action: sequence of SQL instructions (or extensions, for example PL/SQL in Oracle)
 - intuitively:
 - when there is an event (activation)
 - if the condition is satisfied (verification)
 - then execute the action (execution)
 - each trigger refers to a (target) table: it responds to events related to that table



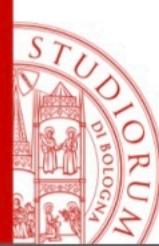
Trigger: granularities and modalities

- Granularity
 - of tuple (row-level): activation for each tuple involved in the operation
 - of operation (statement-level): only one activation for an SQL instruction, with reference to all the tuples involved (“set-oriented”)
- Modality
 - immediate: right after (or right before) of the event
 - deferred: at the time of commit



Computational model

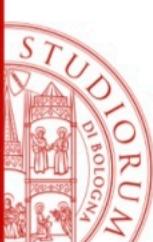
- Let $T^U = U_1; \dots; U_n$ the user transaction
- If the rules of P have form $E, C \rightarrow A$ with E , Event, C , Condition and A , Action, then
- **Immediate** semantic generates $T^I = U_1; \underline{U}_1^P; \dots; U_n; \underline{U}_n^P$
- **Deferred** semantic generates $T^D = U_1; \dots; U_n; \underline{U}_1^P; \dots; \underline{U}_n^P$
- Where \underline{U}_i^P represents the sequence of actions induced by U_i on P
- Problems:
 - Confluence
 - Termination
 - Equivalence



Trigger in Oracle, syntax

```
create trigger TriggerName
  Mode Event {, Event}
  on TargetTable
  [[referencing Reference]]
  for each row
  [when SQLPredicate]
PL/SQLBlock
```

- *Mode*: before o after
- *Event*: insert, update, delete
- for each row specifies the granularity
- *Reference*: it allows to define variable names (usable only for tuple granularity)
 - old as *OldVariable* | new as *NewVariable*



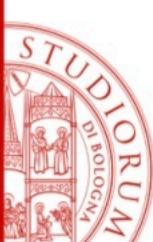
Trigger in Oracle, semantic

- Immediate modality (both after and before)
- execution plan
 - trigger before statement
 - for each tuple involved
 - trigger before row
 - operation
 - trigger after row
 - trigger after row
- in the event of an error, everything is discarded
- priorities between triggers
- max 32 triggers activated in cascade



Trigger in Oracle, example

```
create trigger Reorder
  after update of QtyAvbl on Warehouse
  when (new.QtyAvbl < new.QtyLimit)
  for each row
    declare X number;
begin
  select count(*) into X
  from PendingOrders
  where Part = new.Part;
  if X = 0
  then
    insert into PendingOrders
      values (new.Part, new.QtyReord, sysdate);
  end if;
end;
```



Trigger in Oracle, example 2

Warehouse	Part	QtyAvbl	QtyLimit	QtyReord
	1	200	150	100
	2	780	500	200
	3	450	400	120

T1: update Warehouse

```
set QtyAvbl = QtyAvbl - 70
```

```
where Part = 1
```

T2: update Warehouse

```
set QtyAvbl = QtyAvbl - 60
```

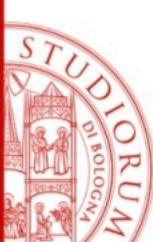
```
where Part <= 3
```



Trigger in DB2, syntax

```
create trigger TriggerName
    Mode Event on TargetTable
    [referencing Reference]
    for each Level
    [when (SQLPredicate)]
    SQLProceduralStatement
```

- *Mode*: before o after
- *Event*: insert, update, delete
- for each *Level* specifies the granularity
- *Reference*: it allows to define variable names (depending on granularity):
 - old as *OldTupleVar* | new as *NewTupleVar*
 - old_table as *OldTableVar* | new_table as *NewTableVar*



Trigger in DB2, semantic

- **Immediate modality** (both after and before)
- before triggers cannot modify the database, apart from variants on the changes caused by the event (therefore they cannot in general activate other triggers)
- In the event of an error, everything is discarded
- No priority between triggers (the order is defined by the system), interaction with compensatory actions on referential integrity constraints
- max 16 triggers activated in cascade



Trigger in DB2, example

```
foreign key (Supplier)
  references Distributor
  on delete set null
```

```
create trigger SoleSupplier
  before update of Supplier on Part
  referencing new as N
  for each row
  when (N.Supplier is not null)
  signal sqlstate '70005' ('Cannot change supplier')
```

```
create trigger AuditPart
  after update on Part
  referencing old_table as OT
  for each statement
  insert into Audit
  values(user, current-date, (select count(*) from OT))
```



Extensions (usually not available)

- temporal events (also periodical) or “user-defined”
- boolean combinations of events
- clause instead of: it does not execute the operation that activated the event, but another action instead
- “detached” execution: an autonomous transaction is activated
- priorities definition
- groups rules, can be activated and deactivated
- rules associated also with queries (not just updates)



Rules properties

- termination (essential)
- confluence
- determinism of observations



Applications

- internal features
 - handling of integrity constraints
 - replication
 - view management
 - materialized: propagation
 - virtual: modification of the queries
- application features: description of the behaviour of the database