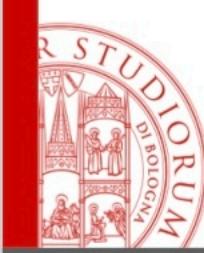


Databases Lab

SQL Exercises

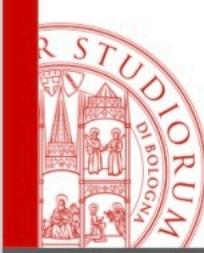
Flavio Bertini

flavio.bertini@smartdata.cs.unibo.it



Different approaches, different styles

- If you are more “theoretic guys”, think first at relational algebra expressions, and then transform them to SQL queries.
- If you are more “geeks”, learning SQL will help you to write down relational algebra expressions.



How to solve the exercises?

Step 0

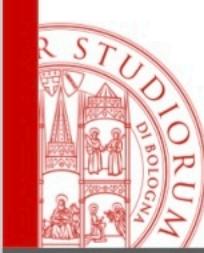
Even before reading the text of the exercise we have to ask ourselves:

0.1. Which **database** are we gonna use?

- Understand the database schema, in order to recognize entities and relations.

0.2. Which are the **primary keys** and the **foreign keys**?

- *Primary keys* are usually underlined, and *foreign keys* uniquely identify a row of another table.



Step 0.1: The database schema

Given the following schema:

LECTURER (Id, Surname, Dept)

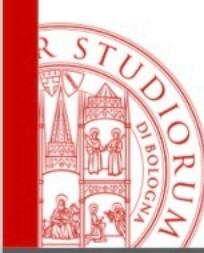
STUDENT (Id, Surname)

COURSE (Code, Name)

EDITION (Course, Year, Lecturer)

EXAM (Student, Course, Year)

Which information can we derive?



Step 0.2: Primary and foreign keys

LECTURER (Id, Surname, Dept)

STUDENT (Id, Surname)

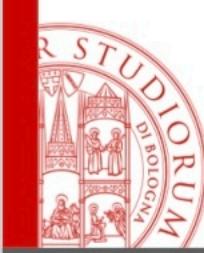
COURSE (Code, Name)

- The underlined attributes are *primary keys*.

EDITION (Course, Year, Lecturer)

EXAM (Student, Course, Year)

- A code identifies a course, but each course could have more than one edition: Course in EDITION is a *foreign key*.
- What about EXAM?



How to solve the exercises?

Step 1

Then, we can read the query, and we have to ask ourselves the following questions:

1.1. Which **relations** are involved?

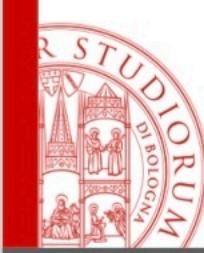
- To list them in the **FROM** clause.

1.2. What should be the **resulting attributes** list?

- To list them in the **SELECT** clause.

1.3. Do we need to **filter the data**?

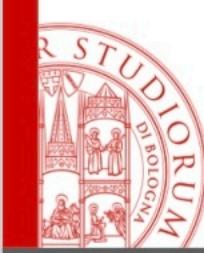
- *Left/Right/Outer join* in **FROM** clause.
- *Theta-join* or value selection in **WHERE** clause.
- Aggregated values in **GROUPBY (HAVING)** clause.



Summary of Syntax

Execution Order

- 5** **SELECT** Attributes + Expressions
- 1** **FROM** Tables + Joins
- 2** [**WHERE** Conditions]
- 3** [**GROUP BY** Attributes]
- 4** [**HAVING** Aggregated conditions]
- 6** [**ORDER BY** Ordering attribute]



Exercise 1

Given the following schema:

LECTURER (Id, Surname, Dept)

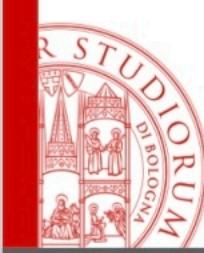
STUDENT (Id, Surname)

COURSE (Code, Name)

EDITION (Course, Year, Lecturer)

EXAM (Student, Course, Year)

Write a SQL query that returns the distinct surnames that belong both to students and to lecturers.



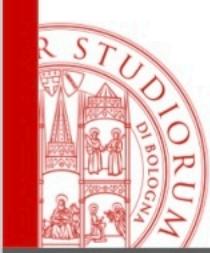
Exercise 1 - step-by-step

Step 1.1 - Let's start by identifying the involved relations, that is LECTURER and STUDENT.

Step 1.2 - Then, to define the resulting attributes list, we have to use the DISTINCT keyword to avoid redundancies (i.e., duplicate Surname).

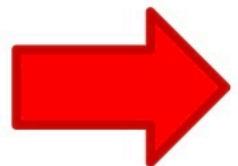
Step 1.3 - Finally, to filter the data we have two options:

1. use the intersection operator (too easy!)
2. perform a join between the two relations.



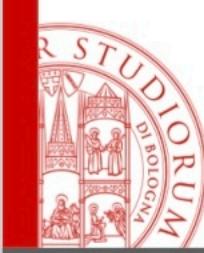
Exercise 1 - easy solution

```
SELECT DISTINCT Surname  
FROM LECTURER
```



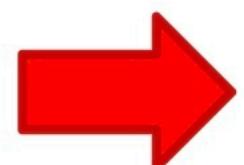
```
INTERSECT
```

```
SELECT DISTINCT Surname  
FROM STUDENT
```

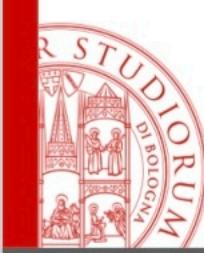


Exercise 1 - tricky solution

```
SELECT DISTINCT LECTURER.Surname  
FROM LECTURER, STUDENT  
WHERE LECTURER.Surname =  
      STUDENT.Surname
```



implicit
equi-join

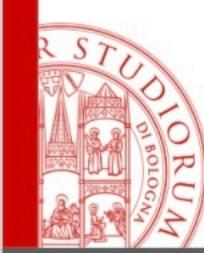


Exercise 1 - good solution

explicit
equi-join



```
SELECT DISTINCT L.Surname  
FROM LECTURER AS L JOIN  
STUDENT AS S ON  
L.Surname = S.Surname
```



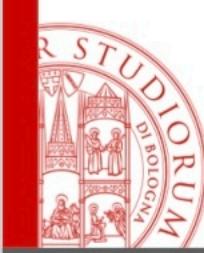
Word of Wisdom

There's no such thing as:

There is only one correct solution!

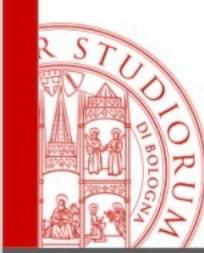
Each of us has a different “style” by following different pragmatisms, as in programming languages.

The important thing is to leave no room for bugs.



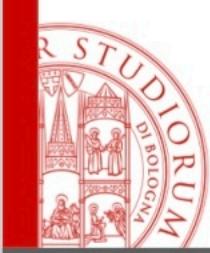
On JOIN Implementation

- The **JOIN operation**, between tables R and S on A and B attributes, is one of the most **time-consuming** operations in query processing.
- These are four of the most common techniques for performing a **JOIN**:
 1. Nested-loop Join
 2. Single-loop Join
 3. Sort-merge Join
 4. Hash-based Join
- Depending on the statistics, **the optimizer chooses** the approach with **the lowest estimated cost**.

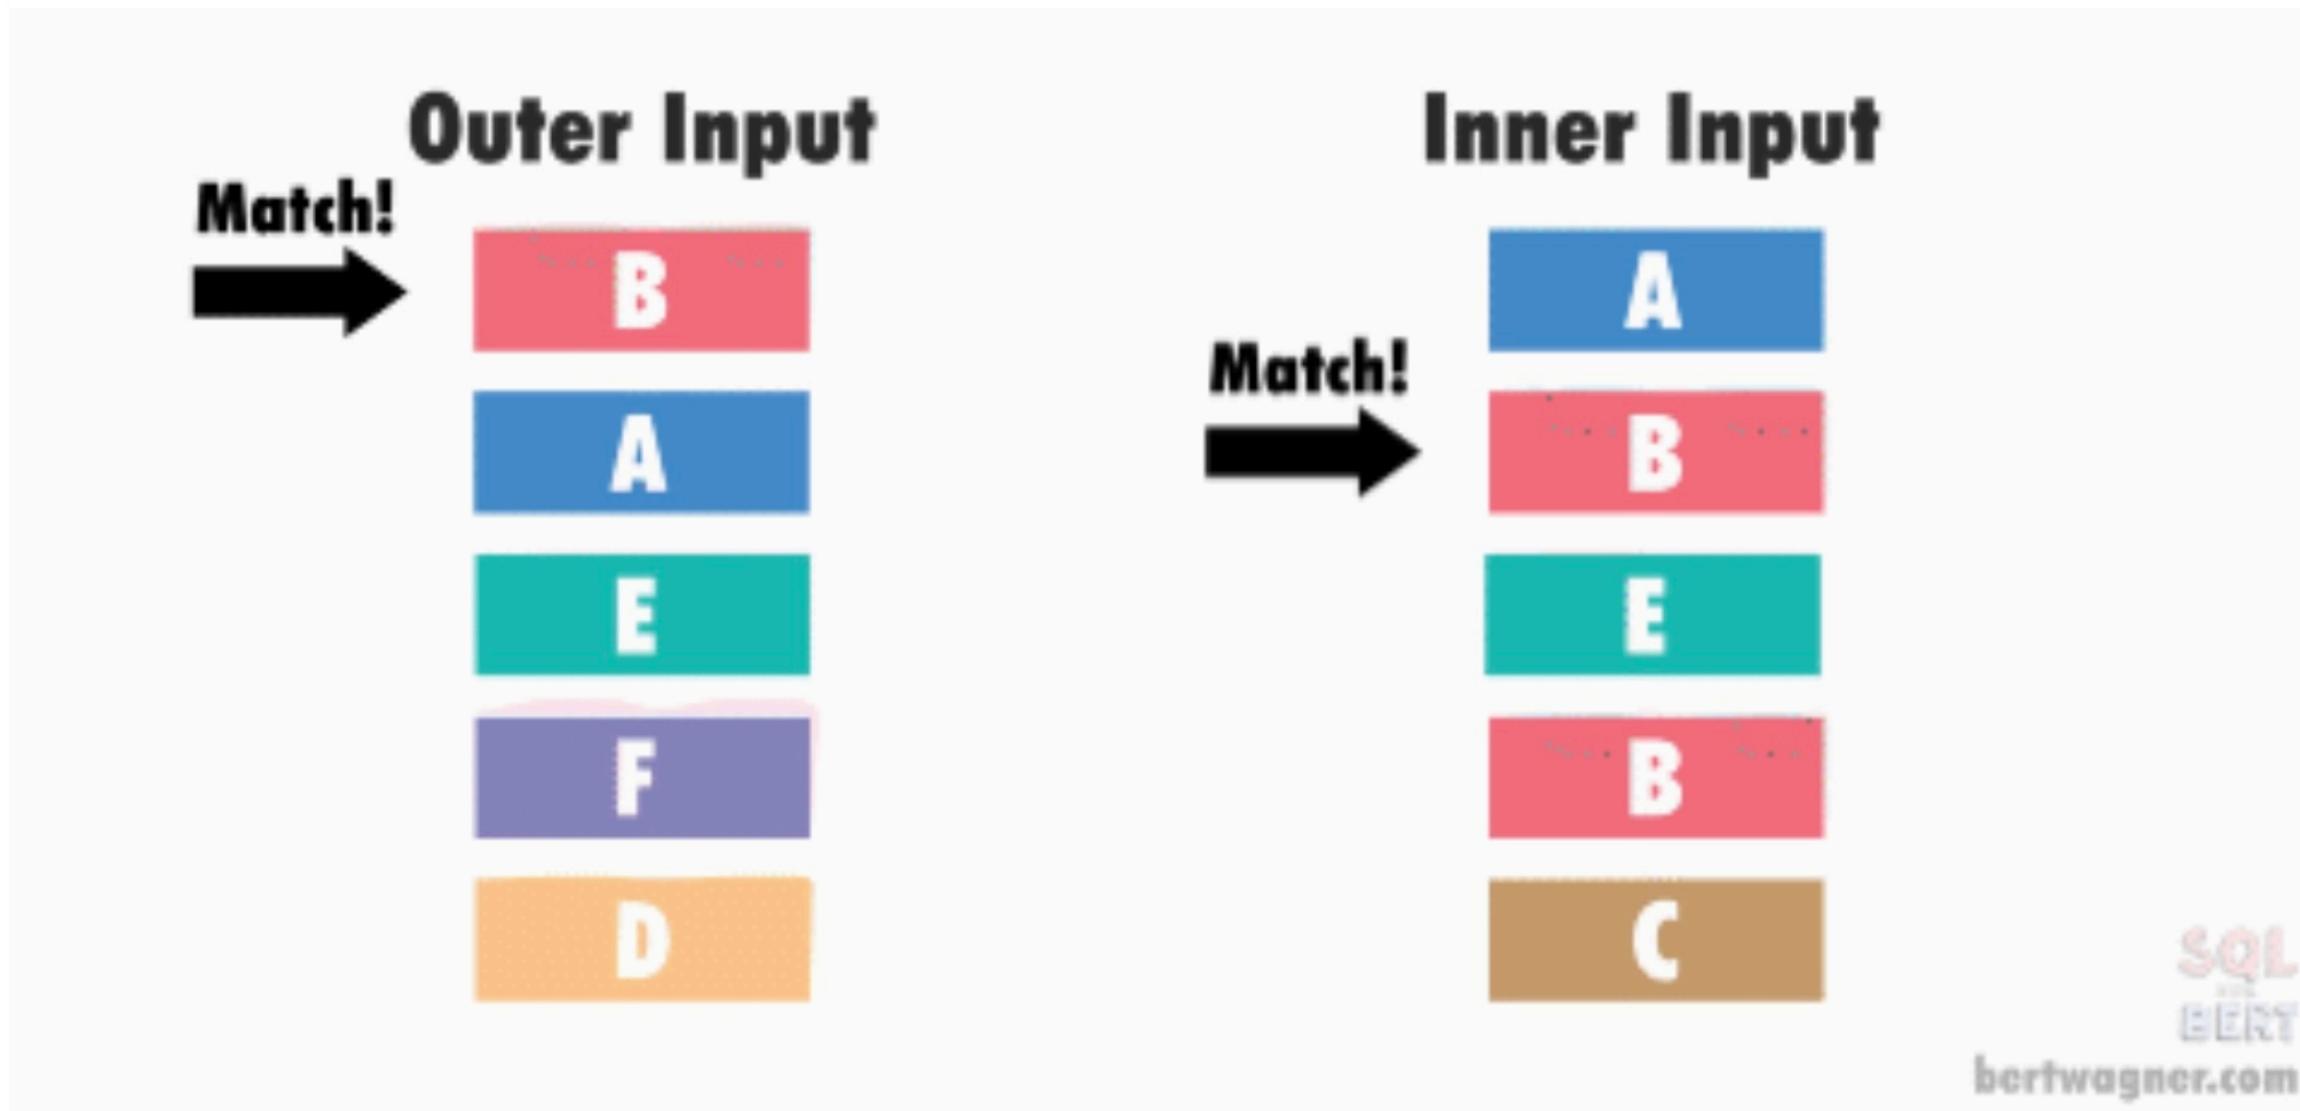


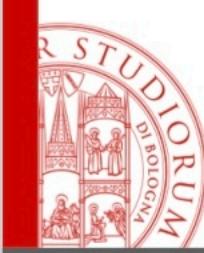
Nested-loop Join

- The table R is set as the *outer table* and the table S is set as the *inner table*.
- For each row in the *outer table* that matches the single-table predicates (the filter F_R), the method retrieves all rows (nested-loop) in the *inner table* that satisfy the join predicate (F_J).
- The cost depends on the assigned role to R and S, the presence of index, and the size of the buffer.



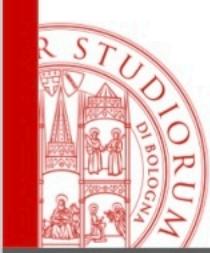
Visualizing Nested-loop Join



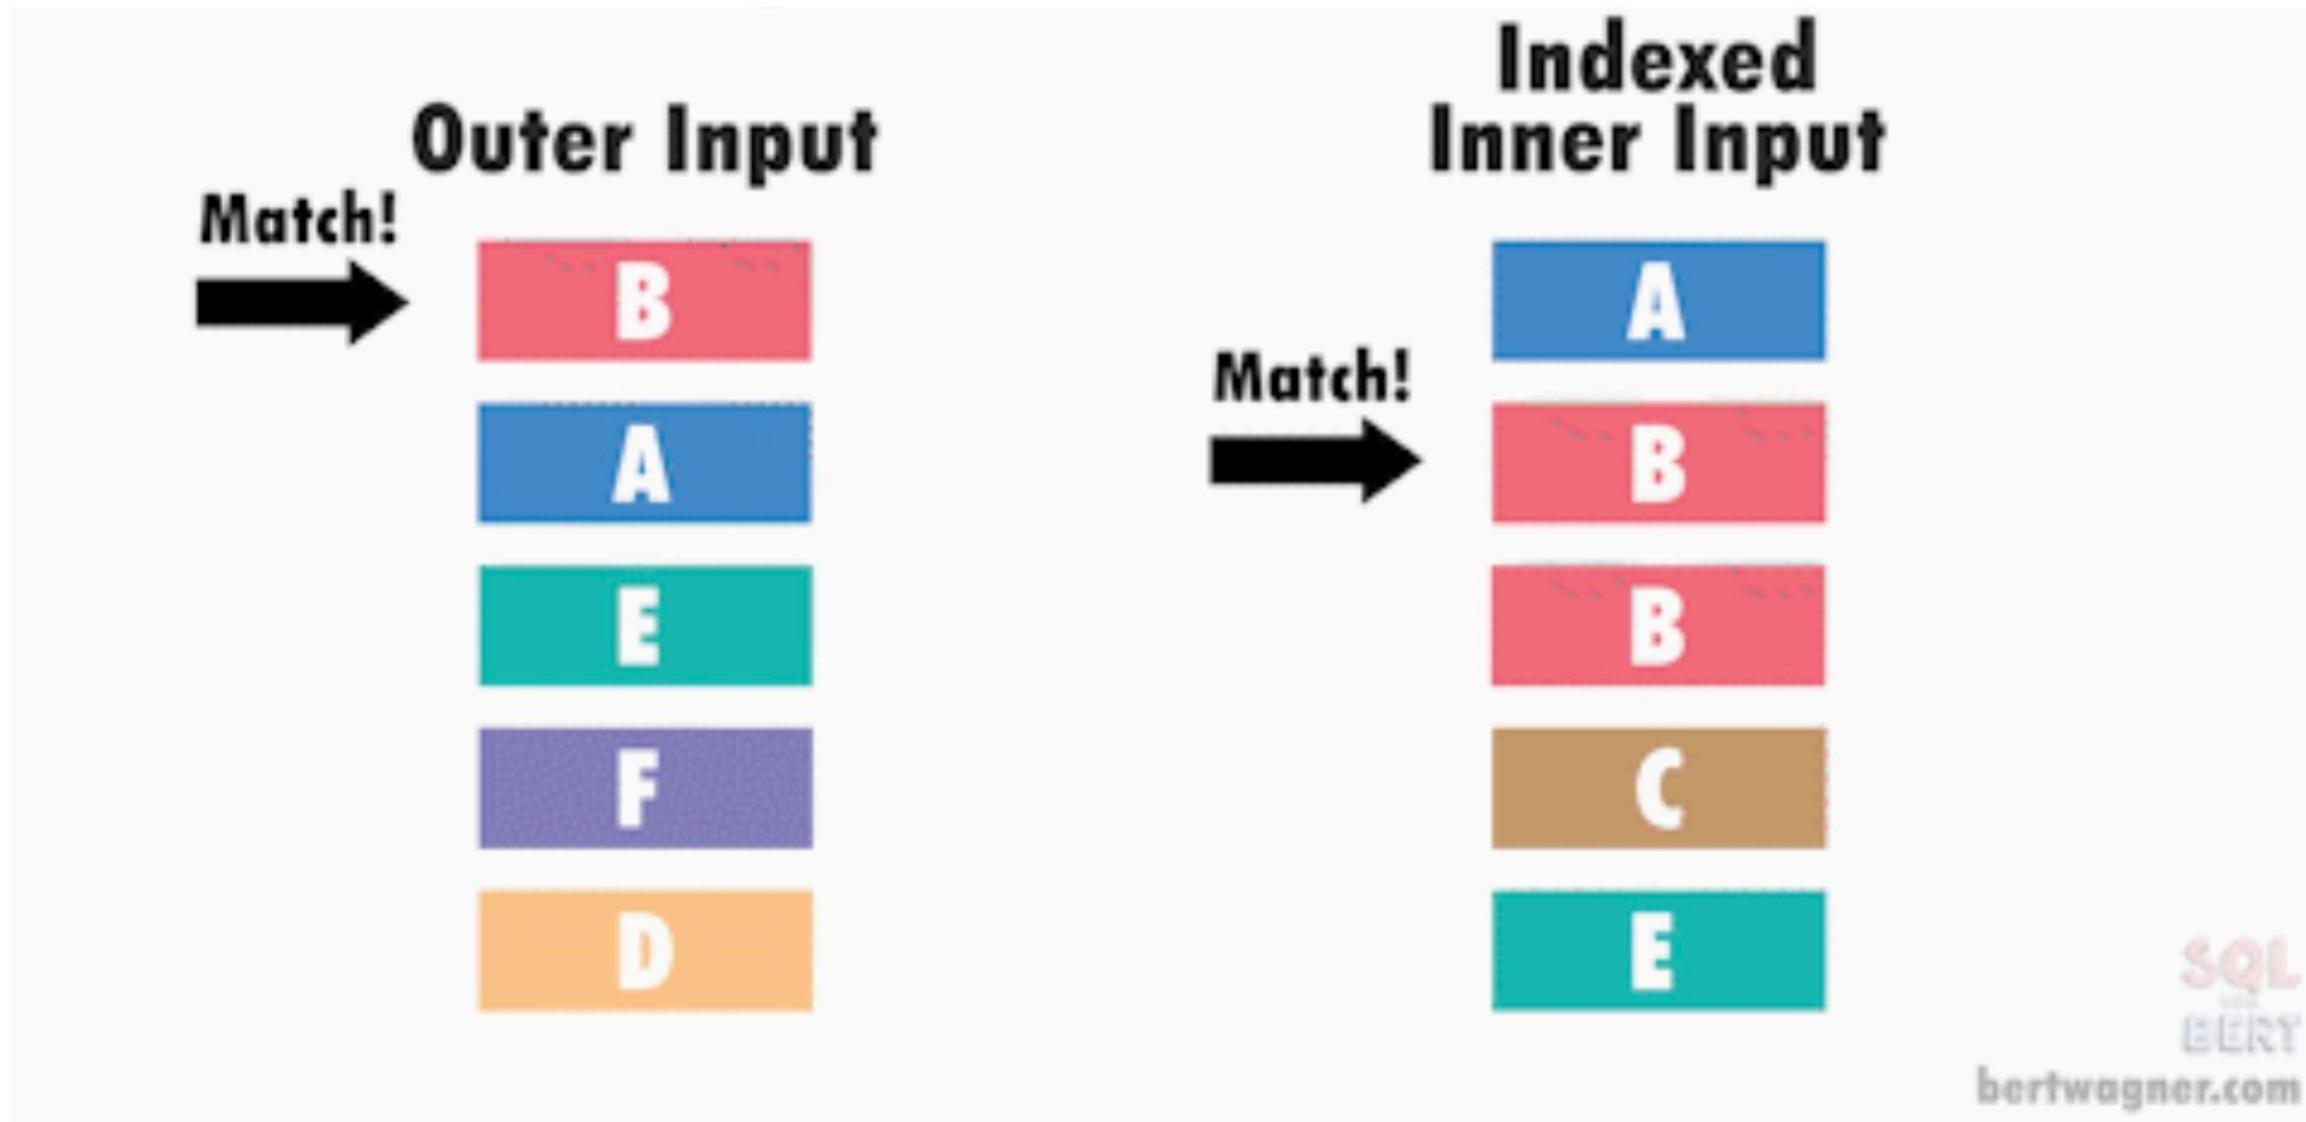


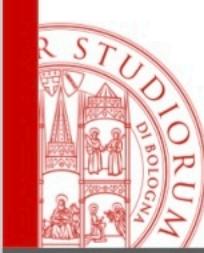
Single-loop Join

- An index (or hash key) exists for one of the two join attributes, for example attribute A of R.
- The method retrieves each row S, one at a time (single-loop), and then uses the access structure to retrieve directly all matching rows from R that satisfy the join predicate.



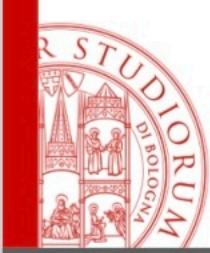
Visualizing Single-loop Join



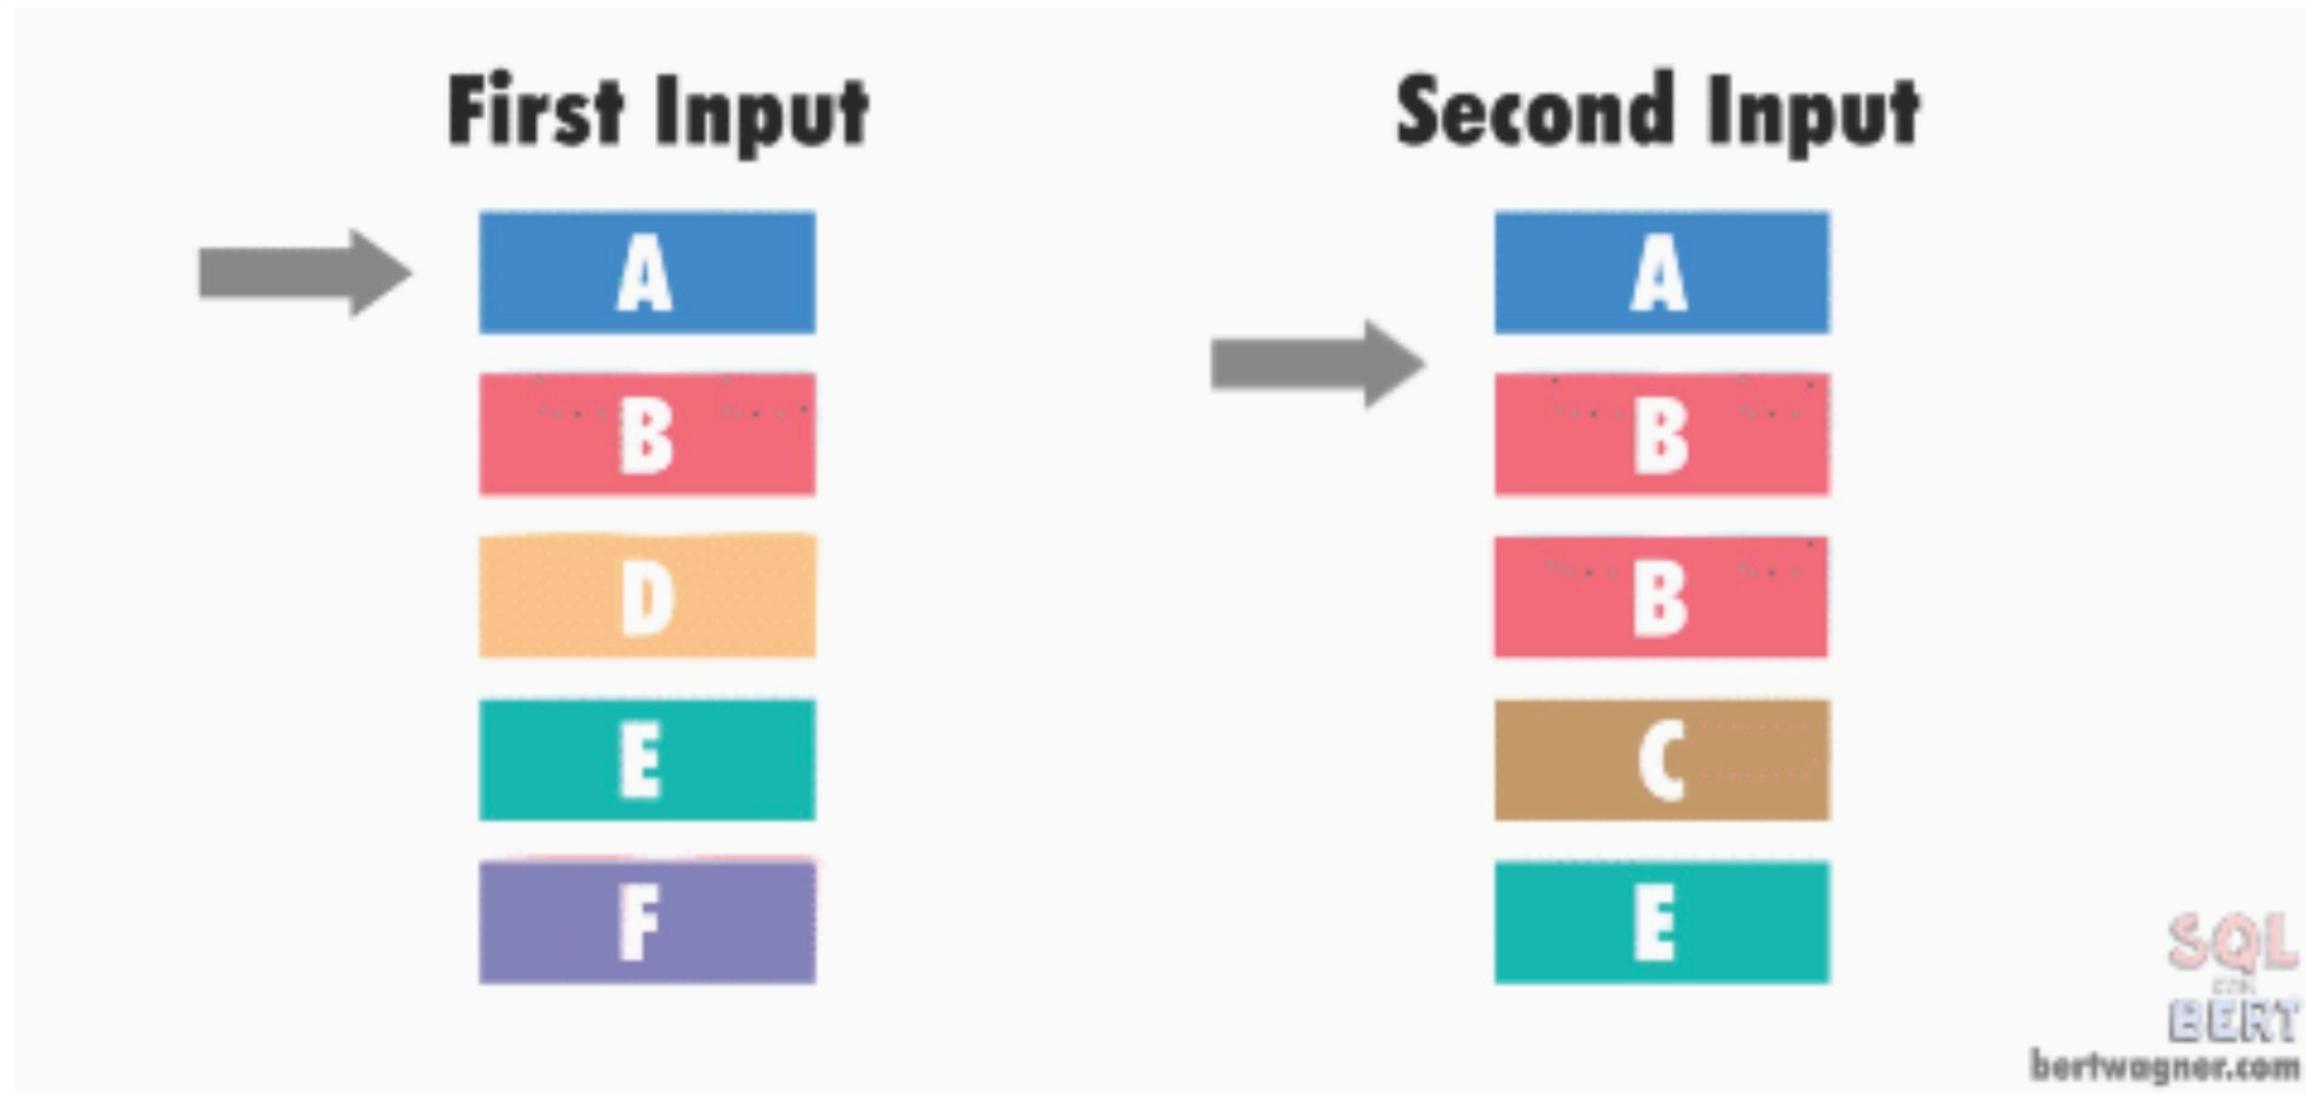


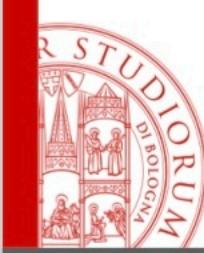
Sort-merge Join

- The R and S tables are physically sorted by value of the join attributes A and B, respectively.
- Both tables are scanned (only once) in order of the join attributes, matching the records that have the same values for A and B.
- The optimizer may choose a sort-merge join for joining large amounts of data when join condition between two tables is not an equi-join (e.g., $<$, \leq , $>$, \geq) or sort is required by other operations.



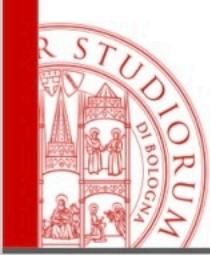
Visualizing Sort-merge Join



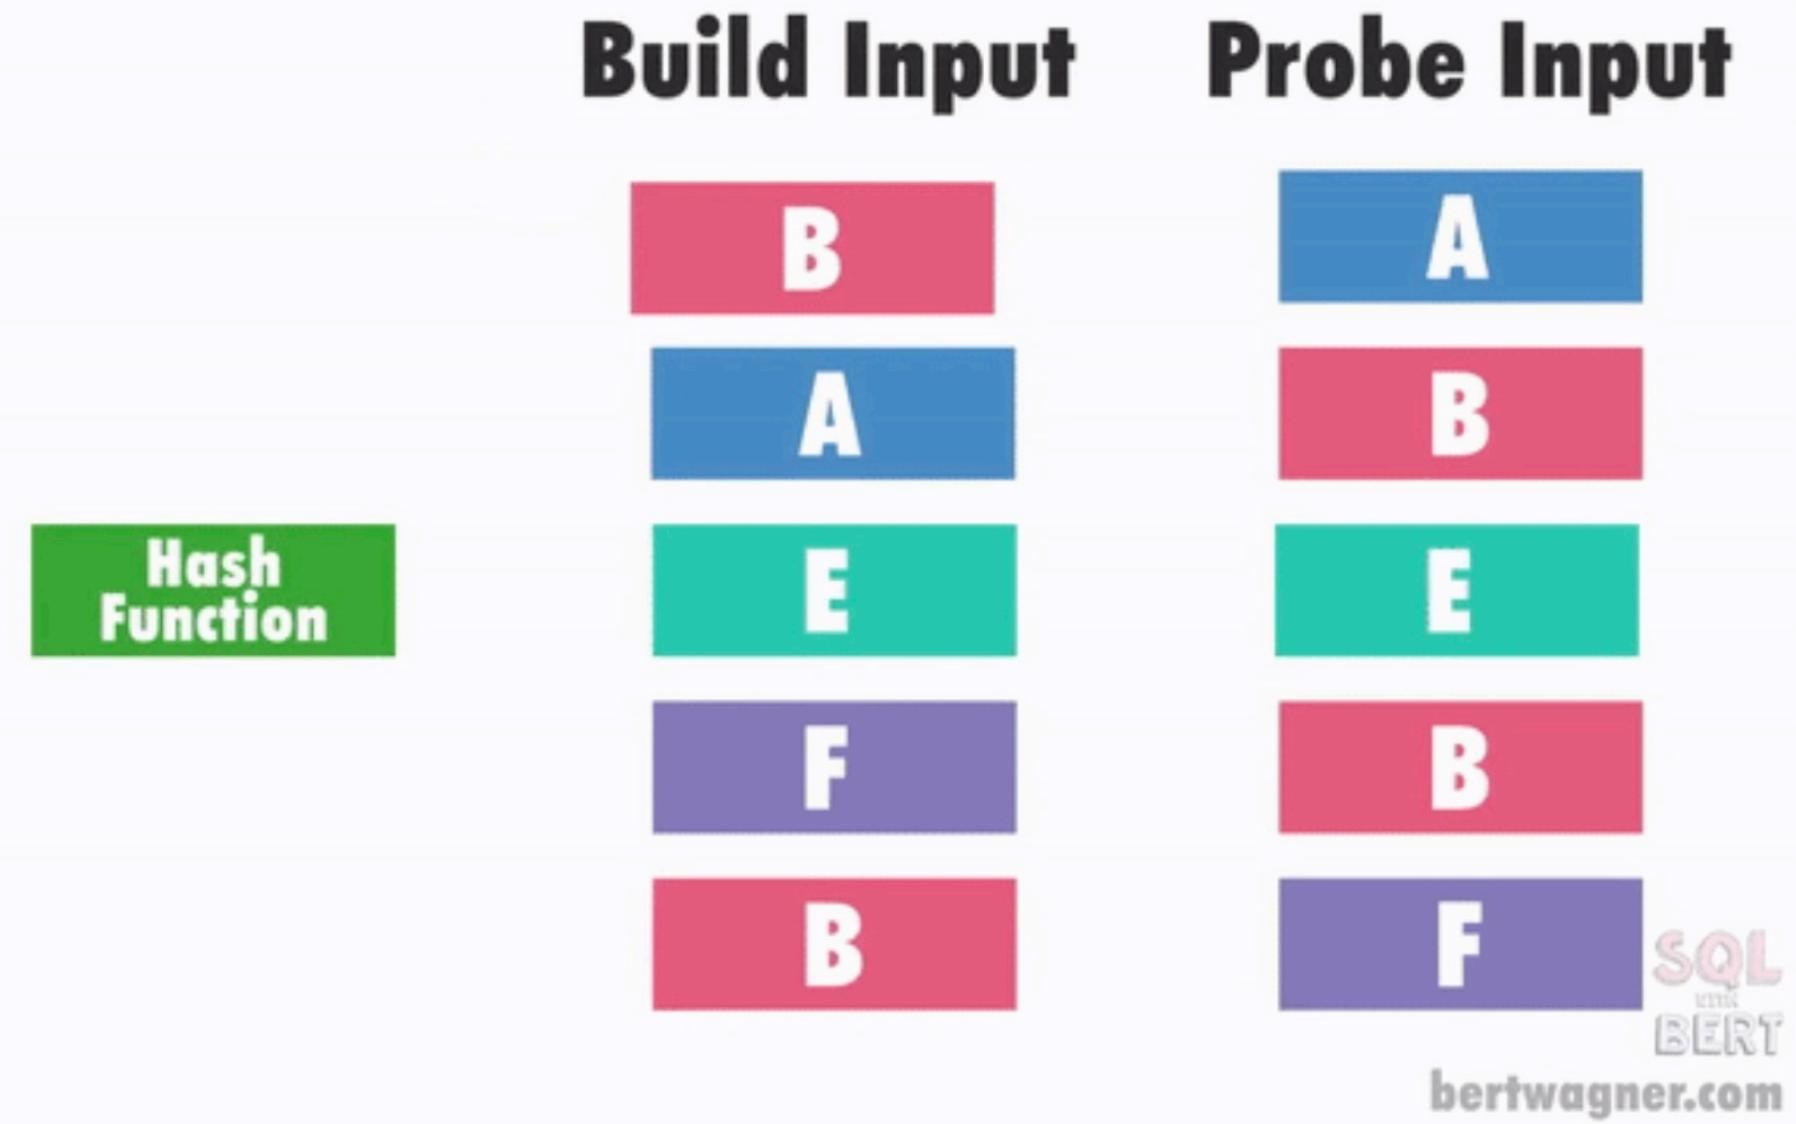


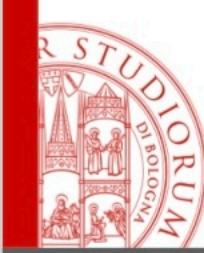
Hash-based Join

- The tables R and S are both hashed using the same hashing function on the join attributes.
- The new hash tables, each with B buckets, allow to find quicker a given join attribute's rows than by scanning the original relation. The final result can be obtained with B simple joins among the B buckets.
- Hash joins are typically more efficient for large tables.



Visualizing Hash-based Join





Exercise 2

Given the following schema:

LECTURER (Id, Surname, Dept)

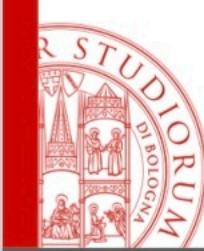
STUDENT (Id, Surname)

COURSE (Code, Name)

EDITION (Course, Year, Lecturer)

EXAM (Student, Course, Year)

Write a SQL query that returns the surnames of those lecturers that have taught in courses for which at least 10 students passed the exam.



Exercise 2 - step 1.1

Let's highlight the involved entities:

LECTURER (Id, Surname, Dept)

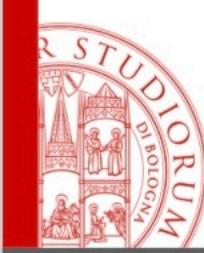
STUDENT (Id, Surname)

? **COURSE** (Code, Name)

? **EDITION** (Course, Year, Lecturer)

EXAM (Student, Course, Year)

Write a SQL query that returns the surnames of those **lecturers** that have taught in **courses** for which at least 10 students passed the **exam**.



Exercise 2 - comments

The following passage within the text is quite ambiguous:

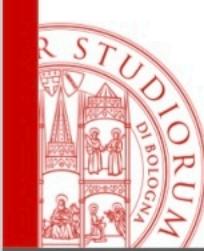
*“... **that have taught in courses** ...”*

Which relation shall we use: COURSE or EDITION?

The EXAM relation has informations on both Course and Year

EXAM (Student, Course, Year)

so we are going to use EDITION.

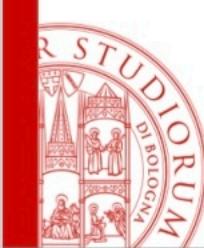


Exercise 2 - step 1.2

Write a SQL query that returns the surnames of those **lecturers** that have taught in “**courses of any edition**” for which at least 10 students passed the **exam**.

- Which is the desired result?

```
SELECT DISTINCT L.Surname  
FROM LECTURER AS L, EDITION AS C,  
EXAM AS E
```



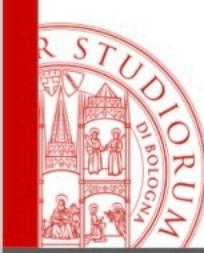
Exercise 2 - step 1.3

Which conditions allow to filter out the data?

- Let's connect the *foreign* and *primary keys*:

WHERE L.Id = C.Lecturer AND
C.Course = E.Course AND
C.Year = E.Year

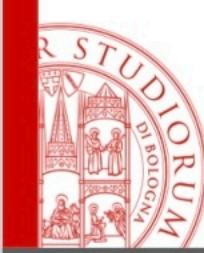
- We aggregate the exams by lecturer, and we filter them by the number of the passed exam:
 - The lecturer's Id is a number, but we have to aggregate Surname too because it must appear in the result.
 - We must aggregate both by lecturer and by the courses because we want the number of exam per course and not per lecturer.



Exercise 2 - solution

```
SELECT DISTINCT L.Surname  
FROM LECTURER AS L, EDITION AS C,  
      EXAM AS E  
WHERE L.Id = C.Lecturer AND  
      C.Course = E.Course AND  
      C.Year = E.Year  
GROUP BY L.Id, L.Surname, C.Course, C.Year  
HAVING COUNT(*) > 10
```

- What happen if we remove C.Year from GROUP BY clause?
- What happen if we remove C.Course and C.Year from GROUP BY clause?

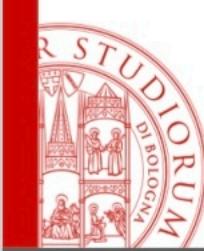


Exercise 3

Given the following schema:

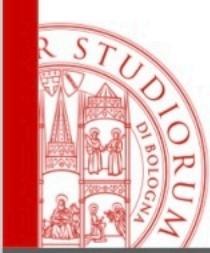
TRAIN (Code, Departure, Arrival, Mile)

Write a SQL query that returns the minimum and the maximum length of the routes between Boston and Chicago without interchanges.



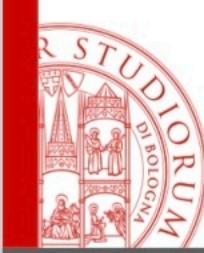
Exercise 3 - comments

- “... without interchanges.” means that we only have to consider direct routes and we have to filter tuples on Departure and Arrival.
- Moreover, we have to use MIN() and MAX() SQL *aggregation functions* for retrieving the tuples with the minimum and the maximum trip length.



Exercise 3 - solution

```
SELECT MIN(T.Mile) AS MinLength,  
       MAX(T.Mile) AS MaxLength  
  FROM TRAIN AS T  
 WHERE T.Departure = "Boston" AND  
       T.Arrival = "Chicago"
```



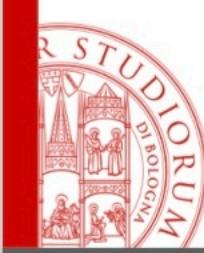
Exercise 4

Given the following schema:

REGION (Name, Population, Surface)

RESIDENCE (Id, Surname, Region)

Write a SQL query that returns the regions having more inhabitants than residents.

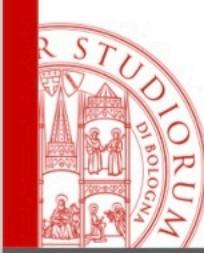


Exercise 4 - a piece of the solution

We can obtain the number of citizen of a region by aggregating RESIDENCE by Region (i.e., the name).

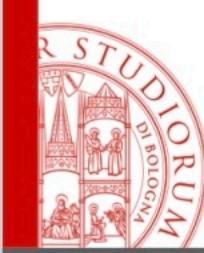
```
SELECT Region, COUNT(*) AS Citizen  
FROM RESIDENCE  
GROUP BY Region
```

Thus, we can define a **view** that could be used in a FROM clause, in order to compare the result with the Population value.



Exercise 4 - final solution

```
SELECT A.Name  
FROM REGION AS A,  
     (SELECT Region, COUNT(*) AS Citizen  
      FROM RESIDENCE  
      GROUP BY Region) AS C  
WHERE A.Name = C.Region AND  
      A.Population > C.Citizen
```



Exercise 4 - optimized solution

CREATE VIEW C AS

SELECT Region, COUNT(*) AS Citizen

FROM RESIDENCE

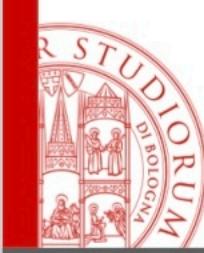
GROUP BY Region

SELECT A.Name

FROM REGION AS A, C

WHERE A.Name = C.Region AND

A.Population > C.Citizen



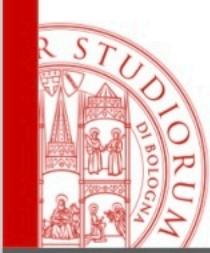
Exercise 5

RATING		
<u>FilmID</u>	<u>UserID</u>	Rating
11234	A984	7
21234	A984	9
31234	B563	8

USER		
<u>UserID</u>	Alias	Age
A984	JamesSmith	20
B563	MaryJohnson	15
F123	JohnBrown	36

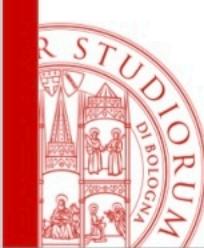
FILM			
<u>FilmID</u>	Title	Year	Director
21234	Blade Runner	1982	Ridley Scott
41234	Pulp Fiction	1994	Quentin Tarantino
11234	Dead Poets Society	1989	Peter Weir

Write a SQL query that returns the titles of those films produced between 1990 and 2000, having an average rating greater than or equal to 7.



Exercise 5 - solution

```
SELECT FILM.Title  
FROM RATING NATURAL JOIN FILM  
WHERE FILM.Year >= 1990 AND  
      FILM.Year <= 2000  
GROUP BY FILM.FilmID, FILM.Title  
HAVING AVG(RATING.Rating) >= 7
```



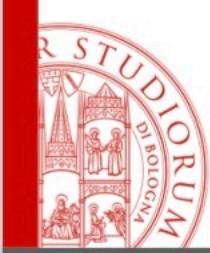
Exercise 6

RATING		
<u>FilmID</u>	<u>UserID</u>	Rating
11234	A984	7
21234	A984	9
31234	B563	8

USER		
<u>UserID</u>	Alias	Age
A984	JamesSmith	20
B563	MaryJohnson	15
F123	JohnBrown	36

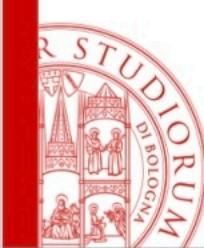
FILM			
<u>FilmID</u>	Title	Year	Director
21234	Blade Runner	1982	Ridley Scott
41234	Pulp Fiction	1994	Quentin Tarantino
11234	Dead Poets Society	1989	Peter Weir

Write a SQL query that returns the minimum age of users who rated “Blade Runner” greater than 8.



Exercise 6 - solution

```
SELECT MIN(USER.Age)
FROM RATING NATURAL JOIN USER
      NATURAL JOIN FILM
WHERE FILM.Title = "Blade Runner" AND
      RATING.Rating > 8
```



Exercise 7

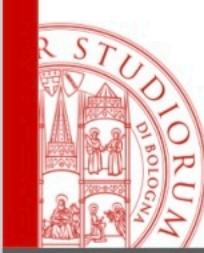
SONG				
<u>ID</u>	Title	Album	Plays	Length
234	Tennessee Whiskey	4444	314	293
112	Purple Rain	3333	111	245
324	Shackled and Drawn	1111	674	226

PLAYLIST		
<u>Name</u>	<u>Position</u>	Song
Summer	1	546
Summer	2	315
Party	1	234

ALBUM			
<u>ID</u>	Title	Year	Artist
1111	Wrecking Ball	2012	234
2222	Sunshine on Leith	1988	657
3333	Purple Rain	1984	756

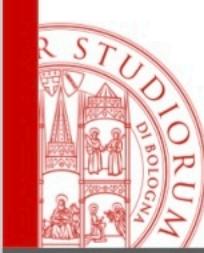
ARTIST		
<u>ID</u>	Name	Labels
234	Bruce Springsteen	Columbia
112	Sting	UMG
324	Chris Stapleton	WCM

Write a SQL query that returns the name of those playlists having at least one song from an album published before 2001 by an UMG artist.



Exercise 7 - solution

```
SELECT DISTINCT PLAYLIST.Name  
FROM PLAYLIST JOIN SONG  
    ON PLAYLIST.Song = SONG.ID  
JOIN ALBUM  
    ON SONG.Album = ALBUM.ID  
JOIN ARTIST  
    ON ALBUM.Artist = ARTIST.ID  
WHERE ALBUM.Year < 2001 AND  
ARTIST.Label = "UMG"
```



Another Schema

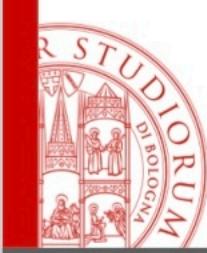
ACTOR(Id, Name, Year, Nationality)

RECITAL(Actor, Film)

FILM(Code, Title, Year, Director,
Country, Genre)

SCREENING(Code, Film, Room, Date,
Profits)

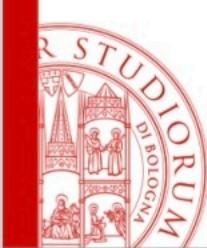
ROOM(Code, Name, Seats, City)



Exercise 8

Write a SQL query that returns the title of Fellini's films produced after 1960.

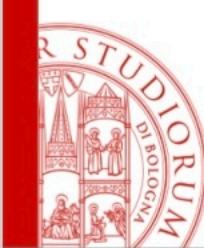
```
SELECT F.Title  
FROM FILM AS F  
WHERE F.Director = "Fellini" AND  
F.Year > 1960
```



Exercise 9

Write a SQL query that returns the title of Japanese science fiction films produced after 1990 or French science fiction films.

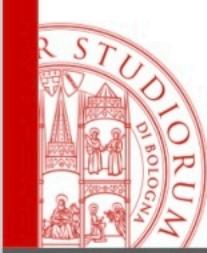
```
SELECT F.Title  
FROM FILM AS F  
WHERE F.Genre = "Sci-fi" AND  
((F.Country = "Japan" AND  
F.Year > 1990) OR  
F.Country = "France")
```



Exercise 10

Write a SQL query that returns the title and the genre of films screened in London on last Christmas Day.

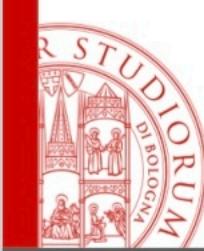
```
SELECT DISTINCT F.Title, F.Genre  
FROM FILM AS F, SCREENING AS S, ROOM AS R  
WHERE S.Date = 2018-12-25 AND  
      R.City = "London" AND  
      F.Code = S.Film AND  
      S.Room = R.Code
```



Exercise 11

Write a SQL query that returns the number of screening room in London with more than 60 seats.

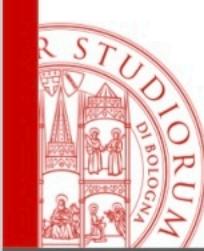
```
SELECT COUNT(*)  
FROM ROOM AS R  
WHERE R.City = "London" AND  
R.Seats > 60
```



Exercise 12

Write a SQL query that returns for each pair of director and actor their names, and the number of films where they have worked together.

```
SELECT F.Director, A.Name,  
       COUNT(*) AS NumMovies  
FROM ACTOR AS A, RECITAL AS R, FILM AS F  
WHERE A.Id = R.Actor AND  
      R.Film = F.Code  
GROUP BY F.Director, A.Id, A.Name
```

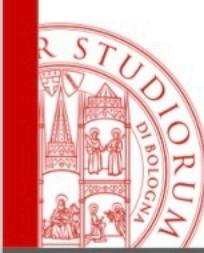


Exercise 13

Write a SQL query that returns the director and the title of films in which less than 6 actors act.

```
SELECT F.Director, F.Title  
FROM FILM AS F, RECITAL AS R  
WHERE F.Code = R.Film  
GROUP BY F.Code, F.Title, F.Director  
HAVING COUNT(*) < 6
```

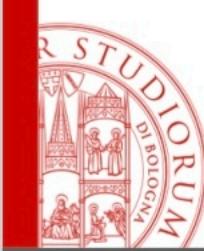
What about of films without actors?



Exercise 13 - a better solution

Write a SQL query that returns the director and the title of films in which less than 6 actors act.

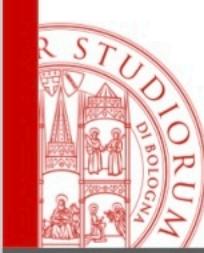
```
SELECT F.Director, F.Title  
FROM FILM AS F  
WHERE 6 > (SELECT COUNT(*)  
            FROM RECITAL AS R  
            WHERE F.Code = R.Film)
```



Exercise 14

Write a SQL query that returns the name and the whole takings of those rooms in Rome that in January 2005 have gained more than \$20'000.

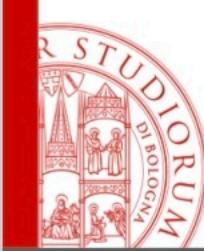
```
SELECT R.Name, SUM(S.Profits)
  FROM ROOM AS R, SCREENING AS S
 WHERE R.Code = S.Room AND
       R.City = "Rome" AND
  2005-01-01 < S.Date AND
  S.Date < 2005-01-31
 GROUP BY R.Code, R.Name
 HAVING SUM(S.Profits) > 20'000
```



Exercise 15

Write a SQL query that returns the title of films that have never been screened in Berlin.

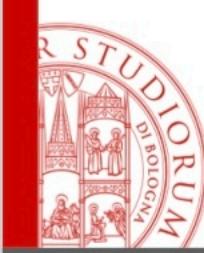
```
SELECT F.Title  
FROM FILM AS F  
WHERE NOT EXISTS (SELECT *  
                   FROM SCREENING AS S, ROOM AS R  
                   WHERE R.City = "Berlin" AND  
                         F.Code = S.Film AND  
                         S.Room = R.Code)
```



Exercise 15 - another solution

Write a SQL query that returns the title of films that have never been screened in Berlin.

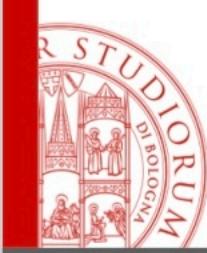
```
SELECT F.Title  
FROM FILM AS F  
WHERE "Berlin" NOT IN (SELECT R.City  
                        FROM SCREENING AS S,  
                        ROOM AS R  
                        WHERE F.Code = S.Film AND  
                              S.Room = R.Code)
```



Exercise 16

Write a SQL query that returns the title of films that have never had a screening with revenues greater than \$500.

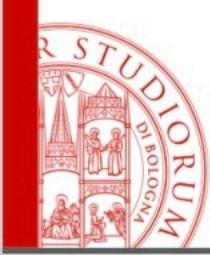
```
SELECT F.Title  
FROM FILM AS F  
WHERE NOT EXISTS (SELECT *  
                   FROM SCREENING AS S  
                   WHERE S.Profits > 500 AND  
                         F.Code = S.Film)
```



Exercise 17

Write a SQL query that returns the title of films that have always earned more than \$500 in all their screening.

```
SELECT F.Title  
FROM FILM AS F  
WHERE 500 <= (SELECT MIN(S.Profits)  
                FROM SCREENING AS S  
                WHERE F.Code = S.Film)
```



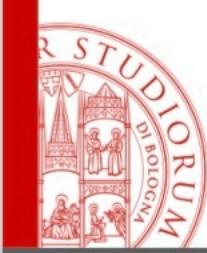
One More Schema

MUSEUM (Name, City)

ARTIST (Name, Nationality)

WORK (Code, Title, NameM, NameA)

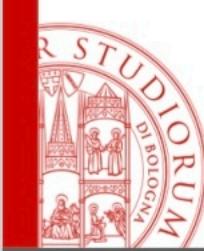
CHARACTER (Name, CodeW)



Exercise 18

Write a SQL query that returns the name of the museums in London that do not have Tiziano's works.

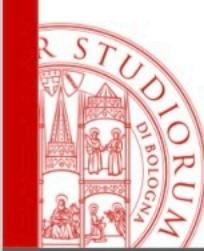
```
SELECT M.Name  
FROM MUSEUM AS M  
WHERE M.City = "London" AND  
NOT EXISTS (SELECT *  
            FROM WORK AS W  
            WHERE W.NameA = "Tiziano"  
              AND M.Name = W.NameM)
```



Exercise 19

Write a SQL query that returns the name of the museums in London that only have Tiziano's works.

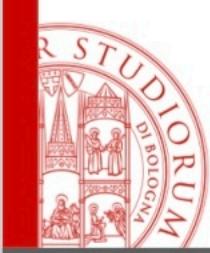
```
SELECT M.Name  
FROM MUSEUM AS M  
WHERE M.City = "London" AND  
      "Tiziano" = ALL (SELECT W.NameA  
                      FROM WORK AS W  
                      WHERE M.Name = W.NameM)
```



Exercise 20

Write a SQL query that returns the name of the museums that have at least 20 works by Italian artists.

```
SELECT W.NameM AS Museum  
FROM WORK AS W, ARTIST AS A  
WHERE W.NameA = A.Nome AND  
A.Nationality = "Italy"  
GROUP BY W.NameM  
HAVING COUNT(*) >= 20
```



Try it Yourself

<https://www.w3schools.com/sql/>

<https://www.w3resource.com/sql-exercises/>